

A Structural Exhaustion Proof of the Erdős–Gyárfás Power-of-Two Cycle Problem

An LLM-assisted minimal-counterexample analysis

Guillem Duran-Ballester
guillem@fragile.tech

July 17, 2026

Abstract

We solve the Erdős–Gyárfás power-of-two cycle problem by proving that every finite simple graph of minimum degree at least three contains a cycle whose length is a power of two. The proof is a minimal-counterexample argument whose main external graph-theoretic input is the P_{13} -free theorem of Hegde, Sandeep, and Shashank. The remaining work is an explicit structural exhaustion: edge-rooted Mersenne-return tests, target-complete replacement, finite P_{13} label algebras, entropy budgets, curvature-rank forcing, surplus ledgers, routed residuals, and local no-go theorems. This paper is the Erdős–Gyárfás instantiation of the structural exhaustion framework developed in the companion paper *Structural Exhaustion* (arXiv link to be provided). The proof was developed with LLM assistance under direct human supervision and guidance; model-generated steps were constrained by named branch states, closure certificates, bookkeeping decreases, and finite enumerations.

Contents

0	Architecture of the proof	2
1	Introduction and overview of the proof	6
2	Target cycles as Mersenne returns	48
3	Minimality and replacement	49
3.1	Boundaried atoms, target-complete quotients, and replacement	94
4	Sparse rank parameters	101
5	Induced P_{13} windows	105
6	The P_{13} curvature algebra	106
7	Entropy bound on independent P_{13} windows	109
8	The P_{13}-free remainder and positive deficiency	132
9	Curvature supply in the atom remainder	136
10	Curvature rank and rank forcing	138

11 Per-vertex remainder entropy and two-budget routing	144
12 The admissible entropy cap	149
13 The large-budget residual	151
14 Local analysis of a negative-net-charge support	155
14.1 Type A: low-deficiency P_{13} -free atom	155
14.2 Type B: high-degree fan-safe surplus obstruction	188
15 The branch-closure theorem	220
15.1 External pressure accounting for the target-defect residual	221
15.1.1 The unified negative Type A collection	222
15.1.2 Every unit of deficit is a silent-excess entry with a nonempty carrier core	222
15.1.3 Unified carrier reduction	224
15.1.4 External pressure accounting for target-defect entries	224
15.1.5 Pressure-defect absorption ledger	227
15.1.6 Window-blocker accounting for the open pressure	230
15.1.7 Large-budget closure from pressure accounting and finite descent	241
A Finite curvature enumeration code	244
B Proof resilience: residual counterexamples under lemma failure	252

0 Architecture of the proof

This preliminary section records the logical architecture of the proof before the local machinery is introduced. Its role is to identify the invariant interfaces through which the later lemmas communicate: the target algebra, the P_{13} window decomposition, the surplus-token routing branch, the curvature-rank budget, and the final net-charge obstruction.

This paper is the Erdős–Gyárfás power-of-two-cycle instantiation of the structural exhaustion framework developed in the companion paper *Structural Exhaustion* (arXiv link to be provided). The proof was developed with LLM assistance under direct human supervision and guidance. In that workflow, model-generated steps were required to enter the record as named branch states, closure certificates, measurable bookkeeping decreases, or finite checks; the human author selected the route, supplied mathematical guidance, and reviewed the resulting claims.

Detailed proof synopsis. The P_{13} -free theorem of Hegde, Sandeep, and Shashank is used as a black box, and auxiliary dense-window quiet-block bookkeeping is recorded in the appendix. After replacing the power-of-two-cycle predicate by the absence of edge-rooted Mersenne returns, minimality and a boundaried replacement principle make every proper atom target-uncompressible; the P_{13} -free theorem then forces a maximal packing of induced P_{13} windows with low density, leaving a large P_{13} -free remainder. The curvature-rank branch routes rank loss to target defect, compression, or delocalization; the whole-graph case is handled by exact closed response profiles, so empty-context data cannot silently identify curvature tests. The final closure is driven by two-budget entropy routing and a large-budget branch governed by a surplus-adjusted positive-deficiency invariant

$$\mathcal{N}(X) = \partial_+(X) - \sigma(X) - \frac{1}{4}|V(X)|$$

on connected admissible supports. The closed part of the local analysis excludes bounded low-deficiency P_{13} -free atoms outside the reduced route-8 silent trace-basin profile and high-degree fan-safe surplus supports whenever the Type B bridge ledger is available. Certificate-marked fan centers have fan degree at most 8 and positive-deficit Type B fan-window data are paid locally by the hybrid incidence ledger. A Type B bridge residual is either a fan-certificate residual center or a failure of the refined support ledger, represented in the latter case by the minimal overlap obstruction of [Definition 14.85](#). Both residual forms are charged to the assigned surplus units at their high-degree fan centers. Before the normalized near-cubic entropy budget is invoked, the sparse surplus-token branch [Proposition 3.87](#) either discharges a sparse surplus exit, routes ordinary decorated Type B fan data to the Type B ledger, or supplies the near-cubic spine estimate of [Definition 4.1](#). In that spine, the assigned surplus consumes the global surplus budget $\sigma(G) = O(\sqrt{n})$; hence all Type B bridge residuals together carry only $o(|R|)$ negative net charge and cannot support the linear large-budget deficit in that spine. The route-8 residual is quantified: if it carries the large-budget charge deficit, it contains linearly many unpaid silent target-complete-minimal trace-basin entries, counted by their routed vertices. The carrier ledger passes each entry to its canonical minimal target-complete carrier core inside the declared u -supported response algebra, closes zero/one-essential-core entries by exits (4)–(7), and reduces any surviving large-budget route-8 branch to a two-carrier route-8 entry. The terminal obstruction is then isolated as a single local no-go theorem, [Theorem 14.61](#), with the carrier-deletion witnesses recorded as part of the residual data. All numerical constants are displayed explicitly, and the finite curvature enumerations are reproduced, with code, in an appendix.

The proof is organized as a routing argument over a fixed lexicographically minimal counterexample G . The load path is

minimal counterexample
 $\Rightarrow$ Mersenne-safe target algebra
 $\Rightarrow$ target-uncompressible atoms
 $\Rightarrow$ low-density induced- P_{13} windows and a large remainder R
 $\Rightarrow$ sparse surplus-token discharge, Type B handoff, or the near-cubic spine
 $\Rightarrow$ full curvature rank or a routed rank defect
 $\Rightarrow$ two-budget entropy closure or large-budget residual
 $\Rightarrow$ negative local net charge
 $\Rightarrow$ Type A or Type B local obstruction
 $\Rightarrow$ terminal two-carrier route-8 obstruction
 $\Rightarrow$ two-carrier route-8 no-go.

Each implication is implemented by a named result in the body of the paper, and the proof-dependency diagram in [Section 1](#) records the same load path at node level.

Phase 1: target algebra and minimality. This phase recasts a global cycle question as local length bookkeeping, and then uses minimality to rule out any smaller graph with the same obstruction. The geometric target is a cycle of length 2^k . The first reduction replaces that target by an edge-rooted arithmetic predicate: by [Lemma 2.2](#), a power-of-two cycle exists precisely when some oriented edge has a simple return of length in

$$\mathcal{M} = \{2^k - 1 : k \geq 2\}.$$

Thus the counterexample condition is expressed as

$$R_e(G) \cap \mathcal{M} = \emptyset \quad \text{for every oriented edge } e.$$

This converts dyadic cycle avoidance into a local embedding constraint. The replacement layer then treats a connected support only through the boundary data visible to outside contexts. The degree-profile fibres, context-universality, and replacement lemma (Lemma 3.103) imply that no proper atom of the minimal counterexample is target-compressible; this is recorded as Corollary 3.105. This uncompressibility statement is the formal minimality lever used whenever a local dependence, quotient, or replacement would otherwise preserve all target hazards.

Phase 2: P_{13} windows, surplus tokens, and the forced remainder. This phase separates the graph into controlled short paths and the part left over, then records the elementary constraints each side imposes on the other. The single load-bearing external theorem at this stage is the Hegde–Sandeep–Shashank P_{13} -free theorem, stated as Theorem 5.1. Since G has no dyadic cycle, G contains an induced P_{13} . A maximal vertex-disjoint family of induced P_{13} subgraphs is removed, producing a packing of windows and a remainder R whose components are P_{13} -free by maximality.

The window family is quantitative rather than cosmetic. Its legal boundary states are governed by the finite P_{13} label algebra and the curvature enumeration. Before any normalized near-cubic budget is used, the surplus-token branch tests whether the surplus above cubic is too large. Proposition 3.87 exhausts that branch: a survivor either realizes a sparse surplus exit, supplies decorated Type B fan data, or satisfies the near-cubic spine estimate of Definition 4.1. On the branch where that estimate is available, the skeleton normalization

$$B_{\text{skel}} = \frac{3}{2}n \log_2 n + o(n \log n),$$

applies. Too many independent windows would then spend more entropy than the skeleton can supply. This gives the window-only cap

$$\theta = \frac{p_{13}}{n} \leq \theta_{\text{win}} + o(1) = 0.01270017798 \dots + o(1),$$

while the high-entropy branch has the sharper $0.01198542083 \dots$ cap in Remark 7.67. Consequently the remainder occupies a fixed linear proportion of the graph, and the external incidences of the windows bound the positive deficiency available in R .

Phase 3: curvature rank and the two-budget split. This phase places the overlapping constraints into one numerical accounting system, so that avoiding one restriction forces a compensating cost elsewhere. The remainder R is large, componentwise P_{13} -free, and boundary-deficient. Its length-two wedges provide the raw curvature tests. The argument first proves a supply lower bound for these tests and then separates raw tests from chargeable entropy by passing through target-rank. A rank drop is not charged as entropy; it is routed through target defect, target-complete proper-support compression, or global delocalization. The target-defect, proper-compression, and proper-delocalization routes are closed by the target algebra, uncompressibility, and proper-support lemmas. The whole-graph delocalization route is closed by the exact-response profile and closed-exactness argument of Definitions 3.95 and 3.98 and Lemma 10.9. Thus the surviving branch has full curvature rank,

$$r_{\Omega}(R) \geq W_2(R) - o(W_2(R)),$$

as stated in Lemma 10.10.

Full curvature rank imposes a forced linear cost. The two-budget routing (Proposition 11.9) uses this cost only after the rank-forcing step has made the coordinates independent. If the remaining non-curvature budget is smaller than the forced curvature cost, the entropy cap closes the branch. If not, the proof enters the large-budget residual, where the same window and remainder estimates yield

$$\Delta_{\text{net}}(R) = \frac{\partial_+(R) - \sigma_R}{|R|} \leq \tau_{\text{win}} + o(1) < \frac{1}{4}.$$

Phase 4: net charge and local exclusion. This phase turns the global accounting deficit into a deficient connected piece, and then proves that the possible local shapes cannot realize it. The large-budget residual is converted into a local obstruction through the net charge

$$\mathcal{N}(X) = \partial_+(X) - \sigma(X) - \frac{1}{4}|V(X)|$$

on connected admissible supports. By Proposition 13.8, the global inequality $\Delta_{\text{net}}(R) < 1/4$ forces a connected admissible support X with

$$\mathcal{N}(X) < 0.$$

The local analysis then proves that, outside the named residual classes, no such X is constructible. If X has no ambient surplus, it is a bounded low-deficiency P_{13} -free atom. Its raw receiver thresholds are first computed as 4, 8, 12; saturated exits (1)–(3), (5), and (6) are closed by the target, compression, and delocalization machinery, exit (4) is routed through the target-defect peeling step, exit (7) is handed to the decorated Type B fan branch, and exit (8) is reduced to target-complete-minimal silent trace-basin entries with the quantitative burden of Lemma 14.47. After these handoffs the remaining unsaturated branch gives the 3, 7, 11 discharging inequality of Lemma 14.45. If X carries high-degree surplus, the high-degree vertices are forced into a fan-safe configuration. A fan center with a certificate labelling is certificate-marked; the structural label-packing cap and its spectral form give the local certificate-closed fan charge. A high-degree center without such a labelling is a fan-certificate residual center and is routed directly to the fan-mass ledger. For certificate-marked centers, the positive-deficit fan branch is paid locally by the hybrid B1 incidence ledger of Lemma 14.79. Summing the certificate-closed and B1 entries across different fan centers requires the refined support ledger B2 of Definition 14.91; once that ledger is available, Lemma 14.98 and Proposition 14.97 bound the support. Failure of B2 gives a minimal Type B overlap obstruction. By Lemma 14.89, that overlap obstruction still carries contextual dyadic safety, high-degree separation, window-stub capacity, and replacement/delocalization routing. The fan-mass closure of Proposition 14.107 charges fan-certificate residual centers and B2 failures to their assigned surplus units. At this point of the main route, the surplus-token branch has either already closed or supplied the near-cubic spine estimate, so those disjoint units have total size $O(\sqrt{n})$. Type B bridge residuals are therefore sublinear in total negative charge and cannot carry the linear net-charge deficit on the surviving route. The route-8 carrier reduction then isolates the terminal two-carrier obstruction class of Definition 14.60, which is ruled out by the local no-go theorem Theorem 14.61.

Redundancy and checkability. The proof is deliberately organized with parallel closures. The entropy cap, the curvature-rank squeeze, and the net-charge collision are arithmetically distinct; a numerical repair in one layer does not erase the formal content of the other layers. The appendices record the finite enumerations and the resilience table for the intermediate branches. Thus each routed branch leaves a specified residual class rather than an unspecified gap, and every residual retains the standing invariants accumulated along the load path above.

1 Introduction and overview of the proof

The Erdős–Gyárfás power-of-two cycle problem asks whether every finite simple graph of minimum degree at least three contains a cycle whose length is a power of two. It belongs to the broader study of which cycle lengths a graph of bounded minimum degree must realize [2, 3], and the spanning-tree and leaf structure of minimum-degree-three graphs invoked below is classical [4]. We prove the following.

Theorem 1.1 (Main closure). *Let G be a finite simple graph with minimum degree at least three. Suppose that G contains no cycle whose length is a power of two, and choose a lexicographically minimal such G . Then no such G exists.*

Proof. Assume, for contradiction, that such a lexicographically minimal counterexample G exists. By Lemma 2.2, a Mersenne edge-rooted return is equivalent to a power-of-two cycle, so no such return occurs in G . Since G has no power-of-two cycle, the Hegde–Sandeep–Shashank theorem Theorem 5.1 and Corollary 5.2 place G on the P_{13} -packing spine.

Before the normalized near-cubic budget is used, the sparse surplus-token branch is exhausted by Proposition 3.87. If a sparse surplus exit occurs, that sparse exit closes in the sparse branch. If a role-homogeneous same-token bottleneck produces decorated Type B fan data, that data is routed to the Type B fan ledger and is not a terminal non-near-cubic leaf. The remaining possibility is that G satisfies the near-cubic spine estimate of Definition 4.1. Thus every surviving minimal counterexample reaches the derived near-cubic entropy/net-charge spine.

The rank branch is exhausted by Lemma 10.4, Lemmas 10.6 and 10.9, and Lemma 10.10: a rank drop is routed to the target-defect, compression, or delocalization closures, while the complementary branch has full curvature rank. On that branch, Proposition 12.2 closes the small-budget entropy alternative and Proposition 11.9 routes the remaining case to the large-budget net-charge analysis. At that point every survivor has the near-cubic estimate, so the rank/full-budget analysis sends it to the residual families treated by Theorem 15.1.

Theorem 15.1 excludes all large-budget residuals except the Type A target-defect/route-8 ledger and the Type B bridge residual. The Type B bridge residual has sublinear total negative mass by Proposition 14.107, so it cannot carry the linear large-budget deficit. The remaining Type A ledger is closed by Theorem 15.48: target-defect entries are canonical exit-(4) peel data and finite descent removes them, while the terminal route-8 two-carrier entry is excluded by Proposition 14.62 and Theorem 14.61. This contradiction proves that no large-budget residual survives, and therefore no lexicographically minimal counterexample exists. $\square$

Remark 1.2 (Status of the routed leaves). *The normalized entropy and fan-mass arguments are invoked only after the sparse surplus-token branch has been discharged. Definition 4.1 names the estimate available on the surviving branch, and Lemma 4.4 computes the labelled skeleton budget once that estimate is available. The estimate is obtained, or the branch is already discharged, before the near-cubic budget is invoked: Proposition 3.87 proves that every survivor of the sparse surplus exits uses the full active-demand family, an $O(n)$ -deficit baseline demand, and the capacity-token ledger, and then has only three possible outcomes: the near-cubic spine estimate, a sparse surplus exit, or decorated Type B fan data routed to the Type B fan ledger. The free/blocked pair split of Definition 3.37 feeds the total route Definition 3.42 and Lemma 3.43: blocked pairs retain their canonical capacity-token assignment, while free pairs use their first selected port in the primitive class. Together with the token supply identity Lemma 3.44, this gives the exact total-pair high-load alternative of Theorem 3.57. The token-fibre graph accounting of Definitions 3.46 and 3.47 and Lemma 3.53 converts any large token load into either a same-token matching or a same-token*

star, and [Lemma 3.49](#) refines it to a role-homogeneous same-token pattern in the window-incidence, remainder-surplus, or primitive-carrier token class. Bounded role-fibre load routes to the near-cubic spine by [Theorem 3.63](#). If bounded load fails, the exact numerical frontier is [Theorem 3.57](#): it gives the complete pair demand, splits it over the three token classes and their subtypes, supplies

$$e(R, W), \quad 2e_{\times}(W), \quad \sigma_R, \quad n, \quad 2m, \quad \sigma(G),$$

and gives the lower bound

$$\psi\left(\frac{N_*(G)}{36(4n + 15p_{13} + 3\sigma(G))}\right)$$

for the largest role-homogeneous pattern in the overloaded subtype. The exact same-token bottleneck routing of [Definition 3.74](#), [Lemma 3.84](#), and [Theorem 3.85](#) then discharges that pattern: it gives a sparse surplus exit, ordinary Type B fan data, or fixed homogeneous caps and hence the near-cubic estimate. The exact window-join pressure of [Corollary 13.6](#) records the global budget available for window-incidence load: without a negative admissible support, the graph must put a linear surplus excess on packed P_{13} -windows. The rank-dependence part is routed by [Lemma 3.67](#): non-independent pair responses force a boundary-profile blocker, target-response blocker, target-defective quotient, target-complete compression, or global exact-profile contradiction. This is the surplus-token branch used before the near-cubic entropy spine. The route-8 Type A residual is reduced by [Proposition 14.59](#) to the terminal two-carrier package of [Definition 14.60](#), and that package is excluded by [Lemmas 14.54](#) and [14.55](#) and [Theorem 14.61](#). The Type B bridge residual recorded in [Definition 14.91](#) is not a terminal leaf: [Proposition 14.107](#) bounds its total negative net charge by the global surplus budget. The Type A target-defect/route-8 ledger is closed by [Theorem 15.48](#), which combines the pressure ledger, finite exit-(4) descent, and the terminal two-carrier route-8 no-go. All other branches used in [Theorem 15.1](#) are closed inside the paper or are explicitly routed to one of the displayed residual ledgers. In particular, the nonrepetitive low-entropy case in [Proposition 11.9](#) is not a terminal leaf; it is routed to the large-budget net-charge analysis and therefore enters the same Type A route-8 and Type B bridge closure.

The proof assumes such a counterexample exists and chooses one, G , lexicographically minimal by

$$(|V(G)|, |E(G)|),$$

and shows that any such graph is forced into the route-8 terminal package, which is ruled out by [Theorem 14.61](#); this gives the closure stated in [Theorem 1.1](#). The main route uses the P_{13} -free theorem of Hegde, Sandeep, and Shashank ([Theorem 5.1](#)) as a black box. The dense-window and route-8 quiet-block bookkeeping is recorded in the appendix as auxiliary resilience data.

Conventions. Throughout, G is a finite simple graph with $n = |V(G)|$ and $m = |E(G)|$. We write

$$\mathcal{D} = \{2^k : k \geq 2\} = \{4, 8, 16, \dots\} \quad \text{and} \quad \mathcal{M} = \{2^k - 1 : k \geq 2\} = \{3, 7, 15, \dots\}$$

for the dyadic lengths and their predecessors. The *target* is a power-of-two (dyadic) cycle; a *target- X* notion means X taken with respect to the predicate “contains a power-of-two cycle”. All asymptotic statements ($o(\cdot)$, $O(\cdot)$) are as $n \rightarrow \infty$, and “for large n ” abbreviates “for all sufficiently large n ”.

Standing notation and constants

The recurring graph parameters are

$$\begin{aligned}\beta &= m - n + 1 && \text{(cycle rank),} \\ \lambda &= 2n - 3 - m && \text{(sparsity rank),} \\ \sigma(G) &= 2m - 3n && \text{(surplus above cubic).}\end{aligned}$$

together with the maximum disjoint induced- P_{13} packing p_{13} , its density $\theta = p_{13}/n$, and the P_{13} -free remainder R . For a subgraph X we use the positive deficiency $\partial_+(X)$, ambient surplus $\sigma(X)$, length-two wedge count $W_2(X)$, curvature target-rank $r_\Omega(X)$, per-vertex remainder entropy $\eta(R)$, and net charge $\mathcal{N}(X)$, each defined where it first appears. The explicit numerical constants are collected here for reference; each is established at the cited place and reproduced by the code in [Section A](#).

Constant	Value	Meaning	Established in
c_Ω	2.28922315244...	flatness cost of one two-step curvature test	Lemma 6.3
c_{13}	118.108581006...	multi-scale P_{13} -window constant	Remark 7.1, Section A
θ_{win}	0.01270017798...	window-only P_{13} density cap on the derived spine branch	Remark 7.67
θ_{high}	0.01198542083...	sharper high-entropy P_{13} density cap	Remark 7.67
ω_{win}	2.54365026308...	wedge density of the remainder on the derived spine branch	Lemma 9.1
ω	2.57407357888...	sharper high-entropy wedge density	Lemma 9.1
$K_{\text{win}} = c_\Omega \omega_{\text{win}}$	5.82298307395...	window-only forced curvature cost per remainder vertex, under the rank closure	Corollary 10.11
$K = c_\Omega \omega$	5.89262883286...	sharper high-entropy forced curvature cost per remainder vertex	Corollary 10.11
τ_{win}	0.22817486846...	net-deficiency cap on the derived spine branch	Section 13
$\Theta(n)$	finite- n threshold	packing threshold closing the entropy cap	Definition 12.1

The proof is organized as a case analysis. Raw curvature tests are not charged directly as independent entropy. Instead, the rank-forcing argument first proves that any surviving residual has full curvature target-rank. This forced curvature cost is then used only in the range where the remaining non-curvature budget is strictly smaller than that cost. If the remaining budget is not small, the analysis reduces to an explicitly described large-budget residual and then to a local net-charge obstruction.

The proof proceeds in the following steps.

- (1) Absence of power-of-two cycles is equivalent to absence of edge-rooted Mersenne returns.
- (2) Minimality gives no proper internal 3-core and forbids target-complete replacements of proper atoms.
- (3) The P_{13} -free theorem forces an induced P_{13} .

- (4) The finite P_{13} curvature algebra gives the constants

$$c_\Omega = \log_2(543958/111286) = 2.28922315244\dots$$

and

$$c_{13} = 118.108581006\dots$$

- (5) The sparse surplus-token branch is exhausted before the normalized skeleton budget is used: it either closes a sparse surplus exit, routes decorated Type B fan data to the Type B ledger, or supplies the near-cubic spine estimate.
- (6) On the surviving spine branch, the independent- P_{13} window accounting gives a first upper bound on the packing density $\theta = p_{13}/n$.
- (7) Removing a maximal disjoint induced- P_{13} packing leaves a large componentwise P_{13} -free remainder R .
- (8) The rank-forcing argument proves that any rank drop among raw atom-curvature tests gives target defect, replacement, target-complete proper-support compression, or whole-graph de-localization. Therefore a surviving residual has full curvature rank,

$$r_\Omega(R) \geq W_2(R) - o(W_2(R)).$$

- (9) Full curvature rank imposes the forced linear cost

$$c_\Omega W_2(R) \geq 5.82298307395\dots |R| - o(|R|)$$

from the window-only constants, with the sharper high-entropy value $5.89262883286\dots |R| - o(|R|)$ when the sharper deficiency cap is available.

- (10) If the remaining non-curvature budget is smaller than this forced cost, the entropy cap closes the branch.
- (11) If the remaining budget is not small, the derived spine window-only cap still gives a large P_{13} -free remainder and the surplus-adjusted positive deficiency satisfies

$$\frac{\partial_+(R) - \sigma_R}{|R|} \leq \tau_{\text{win}} + o(1) < \frac{1}{4}.$$

- (12) Thus any surviving large-budget residual must contain a connected admissible support X with negative net charge

$$\mathcal{N}(X) = \partial_+(X) - \sigma(X) - \frac{1}{4}|V(X)| < 0.$$

Such an X is either a bounded low-deficiency P_{13} -free atom or a high-degree fan-safe surplus obstruction. The local exclusions close Type A outside route 8 and Type B outside the bridge residual. The final large-budget squeeze then extracts all route-8 Type A deficit, charges the remaining Type B bridge mass to the global surplus budget, applies the carrier ledger to reduce the surviving route-8 branch to a two-carrier entry, and invokes [Theorem 14.61](#).

The proof-dependency diagram below records the branch structure in full detail. It is intentionally redundant: each displayed node corresponds to a definition, lemma, proposition, corollary, theorem, or local obstruction class in the body of the paper. One principal case split is

$$\text{rank drop} \Rightarrow \text{target defect/compression/delocalization},$$

while the complementary branch forces full curvature rank. The final spine is the large-budget net-charge branch up to the terminal two-carrier route-8 no-go; Type B fan-certificate residual centers and B2 failures are controlled there by the fan-mass estimate.

Proof-dependency diagram

The proof-dependency diagram is split across eleven numbered parts. Each diagram node has a unique stable bracketed integer. Retired nodes leave gaps, so later verified identifiers are never renumbered. Continuation arrows refer to later nodes by number rather than repeating the same node number. Round nodes are closed branches.




Figure 1: Proof-dependency diagram, Part I. Rectangles are assertions or residual states, diamonds are yes/no branch tests, ellipses are terminal outcomes, and edge labels such as yes or no state which branch is followed. The bracketed numbers are unique diagram-node numbers used by the dependency table. Nodes [19] and [20] record the non-near-cubic surplus accounting branch used before the normalized spine budget; node [20] is the continuation expanded in Fig. 10, and node [25] continues in Fig. 2.




Figure 2: Proof-dependency diagram, Part II. Rectangles, diamonds, ellipses, bracketed node numbers, and edge labels use the convention of Fig. 1. At node [29], $\sigma_R = \sigma(R)$ denotes the ambient surplus on the remainder; the positive-deficiency bound $\partial_+(R) \leq 15p_{13} + o(n)$ is the incidence-supply input, and the $\partial_+(R) - \sigma_R$ bound is its surplus-adjusted consequence. The diagram produces two continuations: the rank-drop branch, closed in Part III, and the full-rank residual, continued in Part IV.




Figure 3: Proof-dependency diagram, Part III. Rectangles, diamonds, ellipses, bracketed node numbers, and edge labels use the convention of Fig. 1. The node-[43] payload contains an already-admitted quotient. Its certified-reduction clause and minimality force injectivity on the declared raw labels, so the retained distinct-coordinate rank-drop witness closes at node [46]; the complementary CT15 execution feeds the full-rank join at node [47].




Figure 4: Proof-dependency diagram, Part IV. Rectangles, diamonds, ellipses, bracketed node numbers, and edge labels use the convention of Fig. 1. The surviving continuation node is the large-budget net-deficiency residual.




Figure 5: Proof-dependency diagram, Part V. Conventions are as in Fig. 1. This panel performs the net-charge split [57]–[64]. Node [63] continues in Fig. 8, and node [64] continues in Fig. 6.




Figure 6: Proof-dependency diagram, Part VI. The Type B continuation starts from node [64], receives the incoming Type A handoff through the dashed input [66], splits at [68], and applies [69]–[76] before the route-8 continuation [77]. At node [74], B2 gives the exact local alternative proved in [Proposition 14.97](#): nonnegative charge, a saturated Type A receiver, or a boundary-transfer overload bounded by assigned center surplus. The later nodes consume the two residual alternatives. The degree-4 no-branch of [68] is expanded in [Fig. 7](#).




Figure 7: Proof-dependency diagram, Part VII. This panel expands the degree-4 no-branch of [68]. Node [80] fixes the certificate attached to each center. Node [81] then makes an exhaustive local decision on those exact certificates: an unresolved local entry routes directly to [84]; otherwise CT12 returns either a full disjoint B2 choice or a minimal overlap obstruction. A full choice gives the exact post-ledger alternative at [82], while the obstruction is retained at [83] and routed to [84]. Node [84] carries the selected-center and grouped fan-mass accounting needed on the original edge to [85].




Figure 8: Proof-dependency diagram, Part VIII. Rectangles, diamonds, ellipses, bracketed node numbers, and edge labels use the convention of Fig. 1. Type A expands from node [63]. The 3/7/11 calculation applies only after the saturated branch [89] has been exhausted. If [89] is yes, exits 1–7 are tested explicitly: exits 1–3, 5, and 6 close at [96], [98], [100], [104], and [106], exit 4 peels one target-defective routed load at [102] and returns to [89] with the residual load L_4 , while exit 7 hands off to Type B [65] through node [108]. On the no-branch of [93], the visible-first excess calculation gives $S_{\text{sil}}^{\text{exc}}(X) \geq 4D_A(X)$. Exit (8), the route-8 residual [109], is expanded in Fig. 9.




Figure 9: Proof-dependency diagram, Part IX. Rectangles, diamonds, ellipses, bracketed node numbers, and edge labels use the convention of Fig. 1. The figure expands exit (8) from Fig. 8 and records how it feeds the full pressure closure. On the derived spine branch, the route-8 carrier reduction closes the no-two-carrier branch: the global squeeze extracts $\mathcal{X}_A$ carrying D_A , the route-8 burden gives $N_{\text{basin}} \geq 4D_A$, zero/one-essential-core entries trigger exits (4)–(7), and absence of a two-carrier entry gives $3N_{\text{basin}} \leq \tau_{\text{win}}|R| + o(|R|)$, contradicting $3N_{\text{basin}} \geq 12(1/4 - \tau_{\text{win}})|R| - o(|R|)$ and $\tau_{\text{win}} < 3/13$. A two-carrier entry enters node [123], the large-budget pressure descent: target-defect entries are canonical exit-(4) peel data and decrease Λ_4 , while a terminal route-8 entry is sent to node [124], where the local no-go theorem excludes it.




Figure 10: Proof-dependency diagram, Part X. This panel expands the sparse non-near-cubic branch [20]. Nodes [126]–[128] build the full active surplus family, and node [129] records the separate baseline-demand definition with its exact deficit. Node [130] performs the canonical blocker split. At [131], every blocker-free pair is assigned to its first selected surplus port, which is already a primitive token, with the distinct role `freeAnchor`. Blocked pairs pass through [132]–[134] and receive their canonical capacity token. The two routes join at [136], where the 25 token–role labels partition the complete scheduled-pair family without double counting. Node [137] is the exact coupled CT9 decision of Proposition 3.59; it consumes the literal pair count $\binom{\sigma(G)}{2}$, not a Boolean-realization estimate. If $D_{25} = 0$, the explicit quadratic surplus bound routes to [138]. If $D_{25} > 0$, CT9 returns an actual overloaded token–role fibre. Nodes [139] and [141] inspect the literal constructor of that token and route it to the window-incidence, remainder-surplus, or primitive-carrier geometric audit [140], [142], [143]. A free-anchor fibre lies in the primitive class and is a literal same-first-port star. Node [144] consumes the finite attachment and local-separator support internally and returns exactly the original sparse-exit, Type B, or near-cubic-spine alternative.




Figure 11: Proof-dependency diagram, Part XI. This panel preserves the original hot/cold P_{13} -window topology behind nodes [22]–[24]. If $\theta < 1/78$, route 8 closes at [147]. Otherwise the hot entropy cap yields [149] or a linear unpaid cold family [150], followed by the near-cubic and stub-excess stages [151]–[152]. Node [153] contains the graph-owned local attachment, return, component, and D1–D7 first-failure computations; only exhaustive F1–F4 absence enters the germ trichotomy [154] as F5. The three original consumers remain [155], [156], and [157].

Detailed dependency table

The table records every named object in the proof-dependency diagram. The *Diagram node(s)* column uses the bracketed node numbers from the eleven numbered parts. The Type A continuation node [63] is expanded in Fig. 8, the Type B continuation node [64] is expanded in Fig. 6, the route-8 residual node [109] is expanded in Fig. 9, the surplus-pair accounting branch [20] is expanded in Fig. 10, and the hot/cold window interface is expanded in Fig. 11. The item numbers remain theorem-level bookkeeping and are not diagram-node identifiers.

Item	Diagram node(s)	Node / theorem	Formal content	Failure route	Label
0	[1]–[3]	Counterexample setup	$\delta(G) \geq 3$ and no C_{2j}	If false, theorem holds for G	Definition 2.4
1	[4]	Minimality	choose minimal (V , E) counterexample	smaller counterexample impossible	Assumption
2	[5]–[7]	Edge-rooted target equivalence	no C_{2j} iff all $R_e(G) \cap \mathcal{M} = \emptyset$	Mersenne return gives target cycle	Lemma 2.2
3	[8]	No proper core	every proper subgraph has $\delta \leq 2$	proper smaller counterexample	Lemma 3.1
4	[9], [10], [67]	High-degree independence	$V_{\geq 4}$ independent	high-high edge deletion	Lemma 3.2
5	[11]	Boundaried graphs	atoms are T -boundaried supports	replacement formalism	Definition 3.88
6	[11], [36], [37]	Boundary degree profiles	quotients fibrewise over $\mathbf{d}_\partial$	otherwise target-defective	Lemma 3.101
7	[12], [36]	Context universality	target-complete identities work against every context	otherwise quotient defect	Lemma 3.102
8	[13]	Replacement	smaller, at most as obstructing representative gives smaller counterexample	minimality contradiction	Lemma 3.103
9	[14], [38], [39]	Hereditary uncompressibility	no proper target-complete compression	replacement contradiction	Corollary 3.105
10	[31]	Target rank	independent coordinates realize 2^k states	entropy only counts independent coordinates	Definition 3.106
11	—	Skeleton budget	labelled edge sets dominate all canonical data	no extra auxiliary degrees of freedom	Lemma 4.2
12	—	Near-cubic budget	$B_{\text{skel}} = \frac{3}{2}n \log_2 n + o(n \log n)$	budget scale	Lemma 4.4
13	[15], [16]	P_{13} -free theorem	P_{13} -free $+\delta \geq 3$ gives target cycle	forces induced P_{13}	Theorem 5.1
14	[18]	P_{13} label algebra	399 labels and relations C_s	finite algebra basis	Lemma 6.1
15	[21]	Curvature enumeration	543958, 432672, 111286 and c_Ω	verified finite computation	Lemma 6.3
16	[21], [145]	Weighted live/cold window interface	in packing order, extend the recoverable joint hot aggregate by one complete safe/flat package at $\log_2 n - O(1)$ separated scales, or retain the identical window with failure of that exact extension	bounded dependent hot product or cold geometric route	Remark 7.1 and Definition 7.51; Section A
16a–r	[153]	Unnumbered local support for first-failure extraction	same-window attachment, return normalization, local component schedules, D1–D7 projections, and bounded CT10/CT8 obligations are retained as internal producers and residuals of node [153]	no auxiliary diagram vertices or edges; each computation scans only its actual finite support	Remarks 7.34, 7.49 and 7.50
17	[22]–[24]	P_{13} density bound	after surplus-token routing supplies Definition 4.1, $\theta \leq \theta_{\text{win}} + o(1)$; high entropy gives $\theta \leq 0.01198542083 \dots + o(1)$	window entropy overflow / window-only cap	Remark 7.1, Remark 7.67

18	[145]–[157]	Hot/cold window branch	if $\theta < 1/78$, route 8 closes; otherwise hot failure forces $C \geq (\theta - \theta_{\text{win}})n - o(n)$ and $b(\mathfrak{S}_{\text{cold}}) \geq 13C - o(n)$; node [153] reads prefixes stage-major, with a literal attachment scan for F1 followed at the same stage by F2, F3 and F4, and only exhaustive F5 enters [154]	the graph-owned F2–F5 producer and the resulting germ bound $N_{\text{germ}} \geq 13C/D_{\text{cold}} - o(n)$ remain open	Definitions 7.51, 7.52, 7.54, 7.57, 7.58 and 7.63, Remarks 7.53 and 7.66, and Lemmas 7.55, 7.56, 7.60 to 7.62, 7.64 and 7.65
19	[17], [25]–[27]	Remainder R	exact partition, certified floor $L_R \leq R $, and componentwise P_{13} -freeness	packing maximality and the incoming coverage certificate	Definition 8.1 and Section 8
20	[28]	Positive deficiency	$\partial_+(X) = \sum_v \max(0, 3 - d_X(v))$	high degree cannot cancel deficient vertices	Definition 8.2
21	[29]	External-incidence supply	$\partial_+(R) \leq 15p_{13} + o(n)$ and hence $\partial_+(R) - \sigma_R \leq 15p_{13} + o(n)$	few windows give small supply	Lemma 8.4
22	[30]	Wedge lower bound	componentwise $W_2(C) \geq 3 V(C) - 2\partial_+(C)$; with $\partial_+(R) \leq (\tau_{\text{win}} + o(1)) R $, this gives $W_2(R) \geq \omega_{\text{win}} R - o(R)$; high entropy gives $\omega R - o(R)$	degree-count proof	Lemma 9.1
23	[31]	Curvature target-rank	$r_\Omega(R)$ counts independent curvature tests via functional admissible rank quotients of exact response profiles	raw tests are not charged directly; closed empty-boundary data stay label-injective	Definitions 3.95, 3.98, 3.99 and 10.1
24	[32], [33], [35]–[37]	Localization of curvature dependence	rank drop gives a finite functional target-dependence and routes to target defect/admissible compression/support	dependence cannot be silent locally	Lemmas 10.3 and 10.4
25	[35]–[37]	Separated wedge testers	identical separated wedge balls are context-universal or defective	supports full-rank forcing	Lemma 10.5
26	[40]–[42]	Proper delocalization support	enlarged $Z \subsetneq G$ gives target defect or target-complete proper-support compression	context defect or replacement contradiction	Lemma 10.6
27	[44]	Repair identity	$s = p - 2 + 2\beta_Z - \sigma_Z$	global delocalization rank barrier	Lemma 10.8
28	[43], [45]	No silent global delocalization	$Z = G$ gives target-defect, smaller closed counterexample, or exact label-injective global data	no global silent rank reduction	Definitions 3.95 and 3.98 and Lemma 10.9
29	[34], [46], [47]	Full curvature rank	$r_\Omega(R) \geq W_2(R) - o(W_2)$	if rank drops, contradiction	Lemma 10.10
30	[48]	Forced curvature cost	$c_\Omega W_2(R) \geq K_{\text{win}} R - o(R)$; high entropy gives $K R - o(R)$	finite-size cost	Corollary 10.11
31	[49]	Finite realized-state entropy	$\eta(R) = \log_2 S_R / R $ on the node-[48] realized family	separates high-/low-state-count regimes	Definition 11.1
32	[49], [50]	Dominant local type	the low-power branch plus a separately proved subexponential radius-2 type coordinate and relabelling closure gives a dominant rooted type	otherwise too many labelled type vectors	Lemma 11.5
33	[51]	High remainder bit contribution	$n^{ R } \leq N_R^{10}$ gives $(R /10) \log_2 n \leq \log_2 N_R$	one denominator-free power comparison	Lemma 11.2
34	[50]–[54]	Two-budget routing	high-entropy branch uses window entropy; repetitive low-entropy branch gives positive curvature rank; nonrepetitive low-entropy branch is routed forward	prevents entropy overcounting	Proposition 11.9

35	[54]	Finite threshold $\Theta(n)$	high- θ branch closed by entropy cap	small-budget closure	Definition 12.1, Proposition 12.2
36	[55]–[57]	Net-deficiency cap	large-budget branch has $\Delta_{\text{net}} \leq \tau_{\text{win}} + o(1) < 1/4$	very low net deficiency required	Section 13
37	[58]	Net charge	$\mathcal{N}(X) = \partial_+(X) - \sigma(X) - X /4$	failure localizes to negative support	Definition 13.4
38	[59]–[62]	Net charge localization	negative global net charge gives connected negative support	Type A/Type B split	Lemma 13.5, Proposition 13.8
39	[63], [86]–[109]	Type A obstruction	bounded low-deficiency P_{13} -free atom split into closed saturated exits, target-defect peeling, Type B handoff, unsaturated discharging, and the route-8 trace-basin branch	raw 4/8/12 thresholds, exits 1–3, 5, and 6 close, exit 4 peels exactly one target-defective routed load, exit 7 hands off to Type B through connector-germ routing, cubic-switch absorption, and high-degree handoff, unsaturated 3/7/11 discharging, route-8 burden $S_{\text{sil}}^{\text{exc}}(X) \geq 4D_A(X)$, declared u -supported response algebra, canonical minimal target-complete carrier core, cut parity, zero/one essential-carrier collapse, private-carrier reduction to a two-carrier entry, carrier-deletion witnesses, and the terminal no-go theorem	Definition 14.3, Definition 14.11, Definition 14.16, Definition 14.33, Lemma 14.23, Lemma 14.24, Lemma 14.25, Lemma 14.34, Lemma 14.35, Lemma 14.36, Lemma 14.37, Lemma 14.39, Lemma 14.27, Lemma 14.40, Lemma 14.43, Lemma 14.45, Definition 14.15, Definition 14.46, Definition 14.48, Lemma 14.47, Lemma 14.56, Lemma 14.58, Definition 14.51, Lemma 14.52, Lemma 14.53, Lemma 14.54, Lemma 14.55, Definition 14.60, Theorem 14.61, Proposition 14.62

40	[64]–[85], [108]	Type B obstruction	degree-4 fan-safe surplus support, degree-greater-than-four heavy-center side branch, and decorated handoff fan data	degree-four centers give either a fan-compatible open pair or two triangular ports; heavy centers reduce to a fan-compatible open pair or three triangular ports; in either case two fan-closed ports give a positive-deficit Type B fan; triangular shoulders have only central/cross/outside completion landings; certificate-marked fans have the local cap and local certificate/B1 ledger; fan-certificate residual centers go to fan-mass; B2 failures form minimal overlap obstructions; on the derived spine branch all Type B bridge residual mass is bounded by the global surplus budget	Definition 14.19, Lemma 14.26, Definition 14.18, Definition 3.7, Definition 3.15, Lemma 3.9, Definition 3.10, Lemma 3.11, Corollary 3.14, Lemma 3.12, Corollary 3.13, Lemma 3.16, Lemma 3.17, Lemma 3.18, Lemma 3.19, Definition 14.69, Lemma 14.70, Proposition 14.80, Corollary 14.81, Corollary 14.82, Definition 14.68, Definition 14.72, Definition 14.73, Definition 14.74, Lemma 14.75, Lemma 14.76, Lemma 14.77, Lemma 14.78, Lemma 14.79, Definition 14.83, Definition 14.84, Definition 14.85, Lemma 14.87, Lemma 14.89, Proposition 14.90, Definition 14.91, Proposition 14.97, Definition 14.71, Lemma 14.98, Definition 14.101, Lemma 14.104, Lemma 14.105, Lemma 14.106, Proposition 14.107
----	------------------	--------------------	--	--	--

41	[76], [85]	Branch-closure theorem	After surplus-token routing has supplied the near-cubic spine estimate, pure Type A outside route 8 closes locally and Type B bridge residual centers and B2 failures are sublinear globally	outside explicit Type A route 8 and Type B bridge residuals, the entropy-admissible branch has no survivor	Theorem 15.1
42	[123]	Large-budget pressure descent	The remaining linear Type A deficit is placed in the unified target-defect/route-8 ledger; every target-defect two-carrier entry is canonical exit-(4) peel data, and finite descent removes it	the terminal large-budget branch can survive only through a two-carrier route-8 entry, which is then routed to node [124]	Definition 15.2, Definition 15.12, Definition 15.45, Lemma 15.46, Lemma 15.44, Lemma 15.47, Theorem 15.48
43	[110]–[124]	Route-8 exit closure	after target-defect descent, the route-8 burden gives $N_{\text{basin}} \geq 4D_A$; canonical carrier cores close zero/one-essential-core entries; absence of a two-carrier entry forces incompatible private-carrier upper and lower bounds; a two-carrier entry supplies a terminal obstruction with declared deletion witnesses	node [124] rules out the terminal two-carrier obstruction because the two-carrier condition places each declared deletion quotient in the canonical exit-(4) family, while exit (4) is absent in a true route-8 residual	Definition 14.46, Definition 14.48, Definition 14.50, Lemma 14.47, Lemma 14.56, Lemma 14.58, Proposition 14.59, Definition 14.51, Lemma 14.52, Lemma 14.53, Lemma 14.54, Lemma 14.55, Definition 14.60, Theorem 14.61, Proposition 14.62, Theorem 15.48
44	—	Enumeration appendix	reproduces all finite constants including c_{13}	constants are checkable	Section A

45	[19], [20]	Non-near-cubic surplus branch	every survivor of the sparse surplus exits has the full active-demand family and, after an $O(n)$ -deficit baseline demand is fixed, satisfies the exact capacity-token budget; the geometric same-token bottleneck lemma gives fixed homogeneous caps unless the load realizes a sparse surplus exit or ordinary Type B fan data, and the caps give $\sigma(G) = O(\sqrt{n})$	derives the near-cubic spine before the normalized entropy budget is used, except for branches that are already routed to sparse exits or to the Type B fan ledger; blocker-free pairs use their first selected surplus port in the primitive token class, while blocked pairs use the canonical capacity-token map; the exact complete-pair token-fibre budget gives the load forced by any attempted violation, and Lemma 3.84 turns a large homogeneous fibre into a geometric first-separator/fan route	Lemma 3.3 , Lemma 3.26 , Definition 3.27 , Definition 3.31 , Lemma 3.32 , Definition 3.34 , Definition 3.37 , Lemma 3.38 , Definition 3.39 , Lemma 3.40 , Definition 3.41 , Lemma 3.44 , Lemma 3.45 , Definition 3.42 , Lemma 3.43 , Definition 3.46 , Definition 3.47 , Definition 3.74 , Lemma 3.53 , Lemma 3.49 , Lemma 3.84 , Theorem 3.85 , Definition 3.69 , Lemma 3.70 , Lemma 3.71 , Theorem 3.57 , Definition 8.3 , Lemma 8.5 , Corollary 13.6 , Proposition 3.87
46	[125]–[128]	Ordered surplus activation	a sparse-pressure survivor yields the full excess-slot ledger; the ordered CT6 audit scans high centres and either returns the first non-cubic incident endpoint or certifies every centre active	deletion criticality rules out the failure terminal; the ledger weight is exactly $\sum_v (d(v) - 3) = \sigma(G)$	Lemma 3.28 , Lemma 3.30 , Lemma 3.32
46a	[129]	Baseline spine demand	after the active family has been established, record a declared baseline demand and its exact deficit	a linear deficit bound is a separate downstream estimate, not a premise of the CT6 activation audit or a consequence of the definition	Definition 3.69
46b	[67]–[69], [78]–[81], [130]	Finite local-port audit	four-cycle avoidance, selected-port classification, center-fibre overload, overload-only adjacency-response comparison, the exact two-shoulder ledger, the nonadjacent same-center fan-compatibility implication, the all-incident-port high-center dichotomy, and triangular shoulder-completion bookkeeping	scans only declared vertex and neighbour schedules; it does not materialize port pairs, paths, or subsets, and does not establish the later Type B/free-blocked routing	Remark 3.8 , Lemmas 3.9 , 3.11 , 3.12 and 3.16 , Definition 3.7 , and Corollary 3.13
47	[130]–[132]	Free/blocked pair split and free-anchor route	pair-response coordinates split into blocker-free and blocked pairs; each free pair is assigned to its first selected surplus port with role <code>freeAnchor</code>	blocked pairs retain the canonical first-hit payload for [134], while free-anchor fibres enter the primitive audit [143]	Definition 3.33 , Definition 3.65 , Lemma 3.67 , Definition 3.42 , Lemma 3.43

48	[132]–[136]	Blocker and total token ledger	canonical blockers, primitive carriers, and capacity tokens count blocked pairs exactly once; the free-anchor route uses selected ports already in the primitive universe	[136] is the exact complete-pair token–role partition and no-overcounting closure	Definition 3.34, Lemma 3.35, Definition 3.36, Definition 3.37, Lemma 3.38, Definition 3.39, Lemma 3.40, Definition 3.41, Lemma 3.44, Lemma 3.45, Definition 3.42, Lemma 3.43
49	[135], [136]	Window-join pressure	the exact packed-window join identity supplies $15p_{13} + \sigma_W$ and the surplus-aware window-stub pressure	supplies the window/remainder capacity side of the token ledger	Definition 8.3, Lemma 8.5, Corollary 13.6, Remark 3.86
50	[137]–[143]	Coupled high-load split	the single-graph high-load test splits the forced homogeneous pattern into window-incidence, remainder-surplus, and primitive-carrier token classes	[139]–[143] are the three token-class audits	Definition 3.46, Definition 3.47, Definition 3.48, Lemma 3.56, Lemma 3.53, Theorem 3.57, Theorem 3.63
51	[140], [142]–[144]	Geometric bottleneck classification and semantic frontier	a large same-token fibre gives a homogeneous matching or star; internal node-[144] support computes the coarse collision, classifies the rooted germs, and exposes the local separator split	the finite support stages are verified locally; the sparse-exit, CT3, Type B, or fixed-cap consumer remains open inside [144]	Lemmas 3.76, 3.78 and 3.84

The constraint ledger

The reduction tracks a fixed list of *invariants*: numbered structural constraints that a minimal counterexample G must satisfy, together with the entropy or deficiency *budget* each one spends and the labeled results that introduce or consume it. Throughout, G is the lexicographically minimal counterexample, so $\delta(G) \geq 3$ and $R_e(G) \cap \mathcal{M} = \emptyset$ for every oriented edge. Budgets quoted in the skeleton normalization $B_{\text{skel}} = \frac{3}{2}n \log_2 n + o(n \log n)$ are used only after [Proposition 3.87](#) has either routed the non-near-cubic sparse-surplus branch to a closed ledger or supplied the near-cubic spine estimate of [Definition 4.1](#). We write $\theta = p_{13}/n$, with densities taken per $|R|$ or per n as noted. We recall the recurring parameters $\beta = m - n + 1$ (cycle rank), $\lambda = 2n - 3 - m$ (sparsity rank), $\sigma = 2m - 3n$ (surplus above cubic), and, for an atom C , its deficiency $d(C) = \sum_v (3 - d_C(v))$. The ledger is organized into blocks; the reverse (result $\rightarrow$ invariant) view is given in the next subsection.

External inputs (black boxes). The argument imports four external results; only the first enters the final proof as a load-bearing black box. The non-near-cubic surplus branch is internal: it uses the free/blocked pair split, the capacity-token ledger, and the same-token bottleneck routing of [Lemma 3.84](#).

Input	Statement used	Constraint it imposes on G	Used by
HSS P_{13} -free theorem (Theorem 5.1)	P_{13} -free + $\delta \geq 3 \Rightarrow$ a power-of-two cycle	G must contain an induced P_{13} ; after removing a maximal disjoint P_{13} family, every remainder component is P_{13} -free and boundary-deficient	Corollary 5.2 , Lemma 8.4 , invariants 20–24, entire remainder analysis

Input	Statement used	Constraint it imposes on G	Used by
Bondy–Vince / Gao–Ma	Few sub-cubic vertices force two cycle lengths differing by 1 or 2; quiet-block cap $ X < 5b(X)^2$	Large almost-cubic atoms must export close cycle lengths or pay quadratic boundary	Appendix resilience bookkeeping
Kleitman–West max-leaf	$\delta \geq 3 \Rightarrow$ spanning tree with $\geq n/4 + 2$ leaves	Cotree fundamental cycles must avoid Mersenne tree distances	Appendix tree–cotree profile bookkeeping
Degree-3-critical / 1–3 tree literature	Extremal graphs with no proper $\delta \geq 3$ subgraph have 1–3 tree / leaf-path structure	Guides the repair-network identity	Appendix marked length-rank bookkeeping

Block I: minimality and target algebra (invariants 1–8).

#	Invariant	Node (first tracked)	Constraint on G	Budget	Lemmas/theorems using it
1	Mersenne-return target algebra	Target [5]	No dyadic cycle $\iff R_e(G) \cap \mathcal{M} = \emptyset$ for every directed edge	— (defines the target predicate)	Lemma 2.2 (states it); Definition 2.5 ; consumed by <i>every</i> quotient/atom/profile result: Lemma 3.101 , Lemma 3.102 , Lemma 3.103 , Corollary 3.105 , invariants 30–38
2	Minimal proper-subgraph condition	Core [8]	Every proper subgraph has $\delta \leq 2$	—	Lemma 3.1 (states it); underlies Lemma 3.5 , Lemma 3.103
3	Edge-deletion criticality	Delete [9]	Every edge touches a degree-3 vertex	—	Lemma 3.2 (states it)
4	High-degree independence	High-degree [10]	$V_{\geq 4}$ is independent	—	Lemma 3.2 (corollary); used by Lemma 14.98 (fan aggregation), invariant 16
5	No passive degree-2 storage	Counterex. [2]	Final G has no degree-2 vertex; frontier degree-2 must be compensated	—	Motivates 1–3 networks; Lemma 8.4 , repair-network auxiliaries
6	Replacement dominance	Replace [13]	If X' is smaller with $H(X') \subseteq H(X)$, replacing X by X' gives a smaller counterexample	—	Lemma 3.103 (states it); Lemma 3.101 , Lemma 3.102 feed it; Corollary 3.105 derives from it
7	Lumpability defect	Univ. [12]	A quotient is valid only if it preserves target-hitting under every outside context	—	Lemma 3.102 (states it); used in final spine step 7; Lemma 6.1 (defect at path-lengths 1, 2, 3)
8	Hereditary target-uncompressibility	Uncomp. [14]	No proper atom admits a nontrivial target-complete compression	—	Corollary 3.105 (states it); the contradiction lever in final spine step 7; consumed by Lemma 14.45 , Lemma 14.98 , invariant 29

Block II: degree, rank, sparsity (invariants 9–19).

#	Invariant	Node (first tracked)	Constraint on G	Budget	Lemmas/theorems using it
9	Minimum edge count	Counterex. [2]	$m \geq \lceil 3n/2 \rceil$	—	Lemma 3.5 ; near-cubic specialization Lemma 4.4
10	Sparse upper envelope	Core [8]	$m \leq 2n - 2$	—	Lemmas 3.3 and 4.3 ; Laman identity (invariant 12)
11	Cycle-rank pressure	Minimality [4]	$\beta = m - n + 1 \geq n/2 + 1$ (up to parity)	—	Lemma 3.5 (states it); feeds 12, 13
12	Laman slack identity	Rank ids [4]	$\beta + \lambda = n - 2$	—	Definition 11.1 , two-budget split; standing notation; Lemma 4.4
13	Surplus–rank relation	Rank ids [4]	$\sigma = 2\beta - n - 2$, $\lambda = n/2 - 3 - \sigma/2$	—	Near-cubic accounting; Lemma 13.5 (via σ)

#	Invariant	Node (first tracked)	Constraint on G	Budget	Lemmas/theorems using it
14	High-degree surplus bound	Surplus branch [125]–[144] supplies spine estimate	$\sigma = O(\sqrt{n})$ for survivors after sparse exits and Type B handoff data are routed, because Theorem 3.85 supplies fixed role-homogeneous same-token caps for all three token classes	σ -budget: $O(\sqrt{n})$	Definition 4.1 , Definition 3.46 , Definition 3.47 , Definition 3.74 , Lemma 3.53 , Lemma 3.49 , Lemma 3.84 , Theorem 3.85 , Corollary 3.50 , Proposition 3.87 ; Proposition 13.8
15	Near-cubic edge count	Surplus branch [125]–[144] supplies spine estimate	$m = \frac{3}{2}n + O(\sqrt{n})$, $\bar{d} = 3 + O(n^{-1/2})$ for the same survivors	sets B_{skel} coefficient $\frac{3}{2}$	Definition 4.1 , Theorem 3.85 , Proposition 3.87 , Lemma 4.4 , Lemma 4.2 , all entropy inequalities
16	Certificate-marked fan degree	High-degree [67]	certificate-marked Type B fans satisfy $d_G(h) \leq 8$; fan-certificate residual centers are charged by $\sum_h (d_G(h) - 3)$	label-packing cap plus σ -budget	Lemmas 14.66 and 14.98 and Proposition 14.107 ; related to invariant 4
17	Spectral feasibility mask	Spectral [70]	$A(H) \leq M$, $\delta(H) \geq 3 \Rightarrow \rho(M) \geq 3$	—	Auxiliary (spectral masks, backup)
18	Spectral equality branch	Spectral [70]	$\rho(M) = 3$ closing $\Rightarrow H$ cubic	—	Auxiliary
19	Spectral surplus branch	Spectral [70]	$\rho(M) > 3 \Rightarrow$ surplus used as dense/compensation/critical/entropy	—	Auxiliary (organizes repair recursion)

Block III: induced P_{13} packing and forced remainder (invariants 20–24).

#	Invariant	Node (first tracked)	Constraint on G	Budget	Lemmas/theorems using it
20	Low P_{13} density	Density [24]	after Proposition 3.87 supplies Definition 4.1 , $p_{13} \leq \theta_{\text{win}}n + o(n)$; high entropy gives $p_{13} \leq 0.01198542083 \cdot n + o(n)$	window-only entropy vs B_{skel} ; sharper high-entropy cap	Remark 7.1 , Remark 7.67 ; Eq. (1) ; Proposition 12.2 ; final spine steps 1–2
21	Forced remainder	Residual A [25]	$ R \geq 0.83489768623 \cdot n - o(n)$	$= (1 - 13\theta)n$	Lemma 8.4 , Lemma 9.1 , invariant 28; all atom analysis
22	P_{13} -free atom condition	Residual A [25]	Components of R are P_{13} -free, boundary-deficient	—	Direct from Theorem 5.1 ; Lemma 8.4 ; Lemma 14.45 , Lemma 14.98
23	Window stub capacity	Stub accounting [29]	Exact split $e(R, W) + 2e_{\times}(W) = 15p_{13} + \sigma_W$; on the derived spine branch this gives $e(R, W) \leq 15p_{13} + o(n)$	15 per window plus explicit window surplus	Lemma 8.5 , Lemma 8.6 , Lemma 8.4 ; invariant 24
24	Remainder deficiency density	Stub accounting [29]	$\partial_+(R) \leq 15p_{13} + O(\sqrt{n})$ on the derived spine branch; before that estimate is supplied, $\partial_+(R) - \sigma_R \leq 15p_{13} + \sigma_W - \sigma_R$ and no-negative-support forces $\sigma_W - \sigma_R \geq (n - 73p_{13})/4$	15θ aggregate plus window-join pressure	Lemma 8.4 , Lemma 8.6 , Corollary 13.6 ; final spine step 3; Proposition 13.8 ; net-charge collision (vs Type A/B output)

Block IV: curvature constraints (invariants 25–29).

#	Invariant	Node (first tracked)	Constraint on G	Budget	Lemmas/theorems using it
25	Legal P_{13} labels	Label algebra [18]	399 legal labels; the zero-defect quotient through path-lengths 1, 2, 3 is the identity	finite (399)	Lemma 6.1 (states it); Lemma 6.3; Remark 14.67; invariant 7; Lemma 14.45, Lemma 14.98
26	Two-step curvature density	Curvature enumeration [21]	$432672/543958 = 0.795414 \dots$ of locally safe wedges are curvature-positive	finite enumeration	Lemma 6.3 (states it); invariant 28; Section A
27	Flatness entropy cost	Cost [48]	$c_\Omega = 2.28922315244 \dots$ bits per independent curvature coordinate	bits/coord	Corollary 10.11, Remark 10.12, Proposition 11.9, Eq. (2)
28	Raw atom curvature supply	Wedge lower bound [30]	after surplus-token routing supplies Definition 4.1, $W_2(R) \geq 2.54365026308 \cdot R $; high entropy gives $2.57407357888 \cdot R $	lower bound (demand)	Lemma 9.1 (states it); Lemma 10.10, Corollary 10.11; final spine step 4
29	Curvature compression ratio	Full-rank squeeze [47]	$r_\Omega/W_2 \leq 0.23758 + o(1)$ forces a rank-defect route unless the full-rank squeeze applies	budget cap $0.611561 \cdot R $ independent tests	Lemma 10.10, Proposition 11.9, Theorem 15.1; final spine steps 5–7; collides with invariant 8 outside explicit residuals

Block V: carrier / cycle-space constraints (invariants 30–38). These refine the target predicate (invariant 1) on the cycle space; all are consumed by the target-algebra layer and the replacement/uncompressibility machinery.

#	Invariant	Node (first tracked)	Constraint on G	Used by
30	Two-path criterion	Target algebra [5]	Two disjoint paths summing to $2^k \Rightarrow$ forbidden dyadic cycle	Lemma 2.2; Lemma 14.45 (return spectrum)
31	Theta closure	Target algebra [5]	All branch-pair sums in a theta avoid powers of two	Atom hazard profiles; Lemma 3.103
32	Ear closure	Target algebra [5]	Ear + base-path length avoids powers of two	Hazard profiles; Lemma 3.102
33	Symmetric difference	Context universality [12]	Overlapping cycles: primitive dyadic output kills	Lemma 3.102
34	Rank-one dyadic support	Delocalization support [40]	All dyadic signals must have support rank ≥ 2	Lemma 10.9, Lemma 10.6
35	Composite cancellation	Repair identity [44]	Dyadic cancellation must stay nonprimitive, non-realized	Lemma 10.8
36	Odd carrier nonempty	Target algebra [5]	Every cycle length has an odd prime divisor (empty carrier = power of two)	Target predicate refinement
37	No common carrier	Carrier transport [35]	No single odd prime divides all cycle lengths	Carrier-transport analysis
38	Polarized carrier transport	Separated testers [35]	Overlap formula $q_p(E) = q_p(C) + q_p(D) - 2t$ flat and non-killing	Lemma 10.5, delocalization lemmas

Block VI: curvature-rank routing and two-budget machinery. These manuscript results operate the curvature invariants (25–29) without being separate ledger invariants; the result $\rightarrow$ invariant direction is displayed for completeness.

Result	Role	Invariants consumed
Lemma 10.4	Localizes where curvature dependence can occur (3 cases)	26, 28
Lemma 10.5	Separated wedges tested by disjoint contexts	28, 38

Result	Role	Invariants consumed
Lemma 10.6	Proper delocalization supports forbidden	34
Lemma 10.8	Enlarged delocalization = repair networks	35
Lemma 10.9	No silent global delocalization by closed exact profiles	34, 35; Definitions 3.95 and 3.98
Lemma 10.10	Forces full curvature rank $r_\Omega \geq W_2 - o(\cdot)$	28, 29
Corollary 10.11	Curvature rank $\rightarrow$ entropy tax on the derived spine branch	27, 28, Definition 4.1
Lemma 11.5	Dominant local type in low-entropy branch	15, 26
Lemma 11.7	Dominant type $\Rightarrow$ independent curvature translates	26, 29
Proposition 11.9	Two-budget routing (entropy, repetitive low entropy, nonrepetitive low entropy) on the derived spine branch	20, 27, 29, Definition 4.1
Proposition 12.2	Closes high- θ branch via entropy cap on the derived spine branch	20, 27, Definition 4.1
Remark 12.4	Closure survives at $K = 0$ (curvature dispensable)	8, 24, 29

Block VII: large-budget residual and local exclusion (the spine's terminus).

Result	Role	Invariants consumed
Definition 13.2	Canonical support decomposition	21, 22
Definition 13.3	Admissible support (carries cumulative constraints)	1, 6, 8, 22, 23, 24
Lemma 13.5	Exact net-charge localization over the connected components of R	13, 14, 24
Proposition 13.8	Large-budget residual is negative net charge	14, 24, Lemma 13.5
Definition 14.3, Lemma 14.4	Define receivers, missing-port counts $q(w)$, and routed load $L(w)$	23, 24
Definition 14.5	Defines anchored returns, receiver-entry returns, visible loads, and silent loads at a completion port	1, 23, 30
Lemma 14.7, Lemma 14.8	First entries of anchored returns are receivers; degree-two receiver entries satisfy the stated port budget	23, 24
Lemma 14.9	Every completion port has an anchored return	1, 3, 30; Lemma 3.4
Definition 14.10, Lemma 14.12, Definition 14.15	Define the reduced route-8 trace-basin data, the declared u -supported response algebra, and connector-band constraints	1, 22, 23, 25, 30, 31
Definition 14.16	Enumerates the eight saturated receiver exits	1, 6, 8, 22, 25, 30, 31, 34, 35
Definition 14.33, Lemma 14.34, Lemma 14.38	Defines exit-(4) peeling and proves the exact one-load, 1/4-charge update	6, 8; Definition 14.16
Lemma 14.35	Routes four visible unpeeled receiver-entry returns to one of exits 1–7	1, 8, 25, 30, 31; Definition 14.33

Result	Role	Invariants consumed
Lemma 14.36	Routes the no-four-visible-return case to exits 4–8, with exit 4 carrying an unpeeled excess load	6, 8, 34, 35; Definition 14.33, Definition 14.15
Lemma 14.37	Combines the visible and silent residual routing lemmas after prior exit-(4) peeling	Lemma 14.35, Lemma 14.36
Definition 14.11	Defines connector germs, first separators, and finite continuation classes	6, 8, 30, 31
Lemma 14.23	Routes nonabsorbed four-coordinate continuation data to exits 4–6 or a surviving first separator	6, 8, 30, 31; Definition 14.11
Lemma 14.24, Lemma 14.25	Turns surviving first separators into decorated Type B handoff envelopes	4, 6, 8; Definition 14.11
Lemma 14.39	Closes exits 1–3, 5, and 6, peels exit 4, and transfers exit 7 to Type B	1, 6, 8, 25, 30, 31, 34, 35; Lemma 14.38, Lemma 14.25
Lemma 14.27	Computes raw receiver thresholds $H_0 \leq 4$, $H_1 \leq 8$, $H_2 \leq 12$	4, 24
Lemma 14.28, Lemma 14.32, Lemma 14.30, Lemma 14.31	Split the saturated 4/8/12 receiver branch into routes 1–8 and quantify the reduced silent trace-basin branch	1, 8, 22, 25, 30, 31, 34, 35
Lemma 14.40	Removes closed exits 1–3, 5, and 6 and all exit-4 peeled target-defective loads before Type A discharging, hands exit 7 to Type B, and leaves reduced route 8 as a residual	Lemma 14.28, Lemma 14.32, Lemma 14.35, Lemma 14.36, Lemma 14.37, Lemma 14.39, Definition 14.15
Definition 14.46	Defines the Type A route-8 deficit D_A	14, 24; Definition 14.15, Lemma 14.31
Lemma 14.47	Gives the indexed basin lower bound $N_{\text{basin}} \geq 4D_A$ and the large-budget lower bound for D_A	14, 24; Definition 14.46
Definition 14.48, Definition 14.50	Defines route-8 entries, essential carrier cores, private carriers, and true route-8 residual entries	8, 23, 30, 31
Lemma 14.56	Forces mixed target events in a carrier core to cross at least two carriers	1, 8, 30, 31; Definition 14.48
Lemma 14.58	Uses the cut-parity constraint to close zero/one essential-carrier entries by exits 4–7	1, 8, 30, 31; Lemma 14.56
Proposition 14.59	Spends the route-8 basin burden against the boundary-incidence ledger and reduces every surviving large-budget route-8 branch to a two-carrier entry	14, 23, 24; Corollary 8.7, Lemma 14.47, Lemma 14.58
Definition 14.51, Lemma 14.52, Lemma 14.53	Records deletion witnesses forced by inclusion-minimality of the essential carrier core and records their declared support and carrier support	6, 8, 23, 24; Definition 14.48

Result	Role	Invariants consumed
Lemma 14.54, Lemma 14.55	Proves that a two-carrier declared deletion quotient is canonical exit (4), and that it realizes exit (4)	1, 6, 8, 30, 31; Lemma 14.53
Definition 14.60, Theorem 14.61, Proposition 14.62	Packages the terminal two-carrier route-8 obstruction, proves the local no-go theorem, and derives full route-8 closure	Proposition 14.59, Lemma 14.54, Lemma 14.55
Lemma 14.43	Unsaturated receivers give the 3/7/11 charge inequality	4, 24
Lemma 14.45	Dense $\sigma = 0$ supports outside target-defect peeling, route 8, and the Type B handoff satisfy $\partial_+ \geq X /4$; a negative one has target-defect, route-8, or Type B handoff data	1, 4, 8, 22, 23, 24, 25, 30, 31; Lemma 14.40, Lemma 14.43
Definition 14.68, Definition 14.69, Definition 14.72	Define fan-window incidences, fan-closed ports, and the hybrid window/non-window credit ledger for one Type B fan	23, 24, 25
Definition 14.73, Definition 14.74, Lemma 14.75, Lemma 14.76	Direct fan-window and two-window cycles are eliminated structurally before incidence credit is counted	1, 23, 30, 31
Lemma 14.77, Lemma 14.78, Lemma 14.79	Every certificate-marked positive-deficit survivor has deficit $D_B = c - \frac{11-k}{4}$ paid locally by half-credit on disjoint window and non-window incidences before the global B2 ledger is invoked	23, 24, 25
Proposition 14.80, Corollary 14.81, Corollary 14.82	Compatible-pair, degree-four triangular, and heavy triangular local routes enter the local Type B fan ledger	4, 8, 25, 30, 31
Definition 14.95 and Proposition 14.96	Deleting every actual high center gives the choice-free bound $\mathcal{N}_-(X) \leq \frac{21}{4} \sum_h (d_G(h) - 3) + \frac{1}{4} \Omega_A(X)$, where Ω_A is the exact Type A receiver overload	4, 14, 24; Hegde–Sandeep–Shashank internal-three-core exclusion
Definition 14.83, Definition 14.84, Definition 14.85, Lemma 14.87, Lemma 14.89, Proposition 14.90, Definition 14.91, Proposition 14.97	For supports with no fan-certificate residual center, local Type B closure is reduced to the B2 refined support ledger after local B1 is proved; disjoint-carrier failure is a minimal overlap obstruction inheriting global constraints	1, 4, 8, 22, 23, 24, 25, 34, 35
Definition 14.19, Lemma 14.26	Records Type A exit-(7) data as an admissible decorated Type B handoff envelope	1, 4, 8, 22, 25; Lemma 14.25
Definition 14.18, Remark 14.67	Defines certificate-marked Type B fans and records the fan-degree label-packing cap	16, 25
Definition 3.7, Definition 3.15, Lemma 3.9, Definition 3.10, Lemma 3.11, Corollary 3.14, Lemma 3.12, Corollary 3.13, Lemma 3.16, Lemma 3.17	Splits degree-four and degree-greater-than-four centers into fan-compatible open pairs or triangular-port alternatives	4, 8, 16, 25
	Shows triangular shoulders have completion edges and that relevant returns enter through shoulders	1, 4, 8, 25, 30, 31

Result	Role	Invariants consumed
Lemma 3.18, Lemma 3.19	Exhausts first landings for triangular shoulders and supplies the second port in the cross-shoulder case	1, 4, 8, 25, 30, 31
Lemma 14.70	Routes compatible open pairs into two fan-closed ports	4, 8, 25; Definition 3.10
Definition 14.71, Lemma 14.98	Excludes certificate-marked high-degree supports with no fan-certificate residual center under B2; positive-deficit fans are paid locally and B2 ledger failures are isolated	4, 8, 16, 24, 25; Lemma 14.78, Lemma 14.79
Definition 14.101, Lemma 14.104, Lemma 14.105, Lemma 14.106, Proposition 14.107	Fan-certificate residual centers, grouped decorated envelopes, and B2 failures are charged to high-degree fan-center surplus units, with any route-8 non-window core extracted into the Type A deficit ledger; on the derived spine branch the remaining bridge mass is $O(\sigma(G)) = o(R)$	4, 14, 16, 24, 35, Definition 4.1, Lemma 14.79, Definition 14.91
Theorem 15.1	Near-cubic entropy-admissible branch closure outside explicit Type A route 8 and Type B bridge residuals	20, 27, 28, 29, 35, 39, Definition 4.1
Theorem 15.48	Closes the remaining Type A target-defect/route-8 ledger: target-defect two-carrier entries are peeled by finite exit-(4) descent, and the terminal two-carrier route-8 entry is impossible	20, 27, 28, 29, 35, 39, Definition 4.1, Proposition 14.107, Definition 15.45, Lemma 15.46, Lemma 15.44, Lemma 15.47, Proposition 14.59
Theorem 1.1	Main closure through sparse-surplus routing, the near-cubic entropy spine, and the terminal two-carrier route-8 no-go (using HSS)	all spine invariants: 1–8, 15, 20–24, 25–29, 35–41, 43–44; Proposition 3.87, Definition 4.1

Spine dependency summary (the load path through route 8).

Theorem 5.1 plus *Definition 4.1* $\Rightarrow$ inv 20 (low P_{13} density, via *Remarks 7.1* and *7.67*)
 $\Rightarrow$ inv 21 (forced remainder $\geq 84.4\%$)
 $\Rightarrow$ inv 23, 24 (stub capacity 15/window $\Rightarrow$ deficiency $\leq \tau_{\text{win}} = 0.2282\dots$)
 $\Rightarrow$ inv 28 (raw curvature supply $\geq 2.543|R|$; high entropy gives $\geq 2.574|R|$) [DEMAND]
 $\Rightarrow$ inv 29 (budget caps independent tests $\leq 0.611|R|$) [SUPPLY]
 $\Rightarrow > 76\%$ correlation forced
 $\Rightarrow$ inv 7/8 (lumpability defect *or* target-complete compression)
 $\Rightarrow$ *Corollary 3.105*: contradiction outside route 8, with Type B fan-certificate residual centers and B2 failures made sublinear by *Proposition 14.107*

Auxiliary path. The auxiliary invariants (spectral masks 17–19, tree-cotree / Kleitman–West, marked length-rank, HIST branch) are not used as load-bearing inputs in *Theorem 1.1*. They record possible repair directions if one later replaces a quantitative estimate in invariants 27–29.

Likewise, [Corollaries 3.52, 3.54 and 3.58](#) are recorded for the surplus-token repair path; the proof of [Theorem 1.1](#) uses the sharper routed statements [Theorems 3.63, 3.73 and 3.85](#).

Per-lemma invariant requirements

The following per-result checklist lists, for each labeled object, the invariants it *requires as input*, a plain-language description, its role, and its position in the proof flow. A lemma may be verified in isolation by confirming that every invariant in its **Requires** cell is already established at an earlier-or-equal stage. Two position columns are given so the reader may cross-check against either presentation: a coarse **Stage** label and the fine-grained **diagram-node** identifier of the proof-dependency diagram.

Stage legend. **R** root/setup · **MIN** minimality · **REP** replacement layer · **ENT** entropy/skeleton budget · **REM** P_{13} -free remainder · **CURV** curvature supply/rank · **2BUD** two-budget routing · **LB** large-budget residual · **EXC** local exclusion (terminus).

Remark 1.3 (Two distinct integer scales). *Ledger-invariant numbers and diagram-node numbers are separate namespaces. The tables below say “inv” for ledger invariants and use bracketed integers only for unique diagram nodes.*

Setup and target algebra.

Result	Stage	Diagram node(s)	Requires	What it does	Role
Lemma 2.2	R	[5]–[7] (target/ret/cycle)	inv 1	A graph has a 2^k cycle iff some edge has a return path of length $2^k - 1$ (a Mersenne number).	Converts the geometric target into an arithmetic one; the foundation.
Definition 2.5	R	[5] (target)	inv 1	Defines “dyadic-safe” <i>contextually</i> (a cycle may exit and re-enter).	Makes safety a property of embedding — the cumulative constraint.

Minimality and replacement.

Result	Stage	Diagram node(s)	Requires	What it does	Role
Lemma 3.1	MIN	[8] (core)	inv 2, minimality	Every proper subgraph has a vertex of degree ≤ 2 .	Forces 2-degeneracy; no smaller $\delta \geq 3$ piece.
Lemma 3.2	MIN	[9] (delete), [10] (hiind)	inv 3, 4	Every edge touches a degree-3 vertex; $V_{\geq 4}$ independent.	Buffers high-degree surplus through cubic vertices.
Lemma 3.5	MIN	[4] (min) $\rightarrow$ rank identities	inv 9, 11	Cycle rank $\beta \geq n/2 + 1$ from $\delta \geq 3$.	Quantifies forced cycle structure.
Lemma 3.101	REP	[11] (bd)	inv 1, 6	Distinct boundary degree profiles cannot be quotient-merged.	Protects boundary bookkeeping under replacement.
Lemma 3.102	REP	[12] (univ)	inv 1, 7	A target-complete identification stays valid in every outside context.	Makes hazard profiles context-independent.
Lemma 3.103	REP	[13] (replace)	inv 1, 6, 7; Lemma 3.101 , Lemma 3.102	Swap a sub-atom for a smaller no-more-hazardous one $\Rightarrow$ smaller graph.	Engine of minimality.
Corollary 3.105	REP	[14] (uncomp)	inv 8; Lemma 3.103	No proper atom can be target-compressed.	The contradiction lever.

Entropy / skeleton budget.

Result	Stage	Diagram node(s)	Requires	What it does	Role
Lemma 3.107	ENT	[18] (p13alg) → [22]	inv 1	Independent target-testable coordinates each consume skeleton entropy.	Sets the bit “currency.”
Lemma 4.2	ENT	[21] (enum)	inv 15	The labelled skeleton count dominates auxiliary data.	Justifies one common budget B_{skel} once the edge-count branch is fixed.
Lemma 4.3	ENT	[21] (enum)	inv 10	Budget valid across the edge-count range.	Robustness to edge-count variation.
Lemma 4.4	ENT	[21] (enum)	Definition 4.1	Near-cubic budget $B_{\text{skel}} = \frac{3}{2}n \log_2 n + o(n \log n)$.	Pins the $\frac{3}{2}$ constant on the branch where the surplus-token routing has supplied the spine estimate.
Lemma 3.32	ENT/ SURP	[125]–[128] (active setup)	inv 10, 14, 15; sparse surplus exits absent	A sparse-pressure survivor has the full active surplus family.	Sets up the active-demand side of the surplus branch.
Definition 3.69	ENT/ SURP	[129] (baseline)	inv 14, 15; Lemma 3.32	Defines a declared baseline demand and records its exact deficit.	Retains the exact baseline datum; the surplus proof proceeds through literal token accounting.
Lemma 3.70 , Lemma 3.71	ENT/ SURP	auxiliary arithmetic audit	inv 10, 14, 15; Definition 3.69	Computes the cubic baseline budget and the incremental edge-count room above cubic.	Records exact standalone skeleton arithmetic without feeding the surplus-token route.
The pre-retokenization clause of Definition 3.42 and the free-anchor clause of Lemma 3.43 Definition 3.37	ENT/ SURP	[131] (free-anchor dispatch)	Lemma 3.32 and Definitions 3.34 and 3.65	Assigns every blocker-free pair to its first selected surplus port with role <code>freeAnchor</code> ; every blocked pair retains its canonical blocker kind and first-hit record for the downstream capacity-token map. The five pre-retokenization role fibres partition the complete pair schedule.	Closes node [131] locally and supplies the exact free and blocked inputs for [134]–[136].
Definition 3.37	ENT/ SURP	[132]–[134] (canonical blockers)	inv 10, 14, 15; the blocked output of node [130]	Defines the canonical blocker assigned to each blocked response.	Makes the blocked-pair assignment single-valued.
Lemma 3.38	ENT/ SURP	[134] (blocker ledger)	inv 10, 14, 15; Definition 3.37	Proves the canonical blocker ledger counts each blocked response once.	Prevents overcounting before capacity tokens are introduced.
Definition 3.39 Lemma 3.40	ENT/ SURP	[134] (primitive carrier)	inv 14, 15; Definition 3.37	Primitive sparse blocker carriers supply the primitive-carrier token class.	Supplies the primitive-carrier side of the token ledger.
Definition 3.41	ENT/ SURP	[134]–[136] (capacity tokens)	inv 10, 14, 15; Lemma 3.38	Defines the capacity-token fibres attached to the blocked-pair ledger.	Names the token side of the blocked-pair accounting.
Lemma 3.44	ENT/ SURP	[136] (token supply)	inv 10, 14, 15; Definition 3.41	Records the window, remainder-surplus, and primitive-carrier token supplies.	Supplies the three capacity classes for the coupled high-load test.
Lemma 3.45	ENT/ SURP	[136] (token partition)	inv 10, 14, 15; Definition 3.41	Proves the token fibres partition the blocked-pair set exactly.	Prevents double counting on the token side.
Definition 8.3	ENT/ SURP	[135] (join)	inv 14, 23, 24; Definition 3.41	Defines the window/remainder surplus split.	Sets the notation for the packed-window join identity.
Lemma 8.5 , Lemma 8.6	ENT/ SURP	[135] (join)	inv 14, 23, 24; Definition 8.3	Gives the exact window-incidence supply $15p_{13} + \sigma_W$ and the surplus-aware stub pressure.	Supplies the window side of the token capacities.
Corollary 13.6	ENT/ SURP	[136] (supply)	inv 14, 23, 24; Lemma 8.5 , Lemma 8.6	Globalizes the join pressure into the coupled token budget.	Supplies the window/remainder side of the token capacities.

Result	Stage	Diagram node(s)	Requires	What it does	Role
Definition 3.46	ENT/ SURP	[136] (token fibres)	inv 14, 15; Definition 3.41	Defines same-token fibres.	Names the fibres used by the high-load split.
Definition 3.47	ENT/ SURP	[139]–[143] (token roles)	inv 14, 15; Definition 3.46	Defines the window-incidence, remainder-surplus, and primitive-carrier roles. Executes CT9 on the literal complete pair schedule and the exact 25-role alphabet.	Names the three token classes used in the high-load split.
Proposition 3.49	ENT/ SURP	[137]–[139], [141] (exact coupled decision)	green nodes [126], [130]–[136]; Lemmas 3.43 and 3.44		A positive excess returns an actual overloaded fibre and its unique token-class route; the other side proves the explicit node-[138] quadratic surplus bound.
Theorem 3.63	ENT/ SURP	auxiliary uniform-load consequence	inv 10, 14, 15; Lemmas 3.43 and 3.44	Converts a separately supplied fixed bound on every token load into a near-cubic estimate.	This consequence is not used to execute the node-[137] CT9 decision.
Definition 3.74	ENT/ SURP	[140], [142], [143] (role audits)	inv 4, 8, 14, 15; Definition 3.47	Defines the routing germs used in the three role audits.	Names the geometric data audited by the bottleneck lemma.
Lemma 3.53	ENT/ SURP	[140], [142], [143] (role audits)	inv 4, 8, 14, 15; Definition 3.74	A large same-token load contains a matching or star in the active role class.	Produces the combinatorial pattern audited by homogeneous extraction.
Lemma 3.49	ENT/ SURP	[140], [142], [143] (role audits)	inv 4, 8, 14, 15; Lemma 3.53	The matching or star contains a homogeneous role subpattern.	Produces the geometric pattern audited by the bottleneck lemma.
Lemma 3.75	ENT/ SURP	[144] (coarse collision)	green nodes [140], [142], [143]	Executes the exact 49-item/48-code scan and retains two distinct equal-code pairs, literal attachment maps, and two canonical rooted germs.	Routes the unchanged residual to the verified finite classifier [144] without asserting a semantic conclusion.
Lemma 3.76 and Remark 3.77	ENT/ SURP	[144] (finite classifier)	green node [144] and its exact typed trigger	Scans the literal $78p_{13}$ attachment coordinates, returns the first mismatch or full alignment, and then records one of four exact rooted-germ shapes in an exhaustive CT10 table.	Routes the unchanged five-way residual to [144]; it asserts no sparse-exit, CT3, Type B, fixed-cap, or spine conclusion.
Lemma 3.84, Theorem 3.85	ENT/ SURP	[144] (semantic bottleneck)	inv 4, 8, 14, 15; node [144]; sparse surplus exits and Type B routing where applicable	The homogeneous bottleneck realizes a sparse surplus exit, produces Type B fan data, or falls under fixed homogeneous caps.	Open consumer: each finite-classifier branch still requires its complete certificate.
Corollary 3.50	ENT/ SURP	downstream cap-to-spine consequence	inv 14, 15; Theorem 3.85	Fixed homogeneous caps also route the remaining geometric branch to the near-cubic spine.	This downstream route rejoins the already available node-[138] conclusion.
Proposition 3.87	ENT/ SURP	[144] (semantic bottleneck)	inv 14, 15; Corollary 3.50	Combines the sparse-pressure discharge with the cap-to-spine route.	Remains downstream of the open semantic consumer.
Lemma 4.7	ENT	[22] (many)	inv 1; Lemma 4.2	$\#\text{realizable states} \leq \#\text{labelled graphs}$.	The counting contradiction step.
Remarks 7.1 and 7.50	ENT/ EXC	[21] $\rightarrow$ [145] (weighted split)	exact packing order, current aggregate, the 91 safe/flat rows, graph-owned local interpretation, recoverable joint choices, and bounded injective spine code	Accepts one compatible extension or retains the identical window, prior aggregate, and extension failure; the sequential cold list executes the same-window surplus/corridor route.	Joint capacity, exact partition, and [151]–[153] provenance are verified; the normalized hot arithmetic and terminal cold closure remain open before [22]–[24].

Result	Stage	Diagram node(s)	Requires	What it does	Role
Remarks 7.34, 7.49 and 7.50	EXC	[153] (internal support)	exact cold-window residuals, literal local graph schedules, and predecessor provenance	Retains the finite same-window, short-return, long-prefix, component, and D4–D7 support stages as internal producers for original node [153].	No auxiliary diagram node or edge is introduced; open semantic consumers remain typed residuals inside [153].
Definitions 7.52, 7.54, 7.57, 7.58 and 7.63, Remarks 7.53 and 7.66, and Lemmas 7.55, 7.56, 7.60 to 7.62, 7.64 and 7.65	ENT/ EXC	[145]–[157] (cold)	inv 1, 8, 14, 15, 20, 35, 38; Definition 4.1; Remark 7.1; Theorem 15.48	The cold continuation at [153] is stage-major: literal attachment F1, then F2/F3/F4 at the same prefix, with only exhaustive F5 entering [154]. The F1 constructor is verified; the graph-owned F2–F5 producer and the germ bound remain open.	Node [153] is not green and supplies no terminal cold-window closure yet.
Remark 7.67	ENT	[24] (theta)	inv 20; Definition 4.1; Remark 7.1; Lemma 3.107	Independent induced- P_{13} density $\theta \leq \theta_{\text{win}}$ after the surplus-token branch supplies the spine estimate, with sharper high-entropy cap $\theta \leq 0.011985$.	Kills the high- P_{13} branch on the derived spine branch; forces a large remainder.

P_{13} packing and forced remainder.

Result	Stage	Diagram node(s)	Requires	What it does	Role
Corollary 5.2	REM	[15] (p13free)	Theorem 5.1	A counterexample contains an induced P_{13} .	Entry to the window/remainder split.
Lemma 6.1	REM	[18] (p13alg)	inv 7, 25	399 legal P_{13} labels; zero-defect quotient through lengths 1, 2, 3 is the identity.	The finite window-safety algebra.
Lemma 6.3	REM/ CURV	[21] (enum)	inv 25, 26	432672/543958 of safe wedges are curvature-positive.	Finite count feeding curvature supply.
Lemma 8.4	REM	[26] (R), [29] (stub)	Theorem 5.1, inv 21, 22, 23, 24	Remainder ($\geq 83.4\%$) gets ≤ 15 incidences/window; $\partial_+(R) \leq 15p_{13} + o(n)$ and $\partial_+(R) - \sigma_R \leq 15p_{13} + o(n)$.	Supply ceiling of the final collision.

Curvature supply, rank, two-budget.

Result	Stage	Diagram node(s)	Requires	What it does	Role
Lemma 9.1	CURV	[30] (wedge)	inv 21, 28; Corollary 8.7	Remainder has $\geq 2.543 \cdot R $ raw two-step curvature tests on the derived spine branch, with the sharper high-entropy value $\geq 2.574 \cdot R $.	Demand floor of the final collision.
Lemma 10.3	CURV	[31] (rank), [32] (drop)	Definitions 3.99 and 10.1	Extracts a finite functional target-dependence from a rank drop.	Makes the rank-drop interface explicit.
Lemma 10.4	CURV	[32] (drop), [33] (depopen)	inv 26, 28; Lemma 10.3	Localizes the three ways curvature dependence arises.	Organizes the forced-rank case analysis.
Lemma 10.5	CURV	[35] (dep)	inv 28, 38	Separated wedges tested by disjoint contexts.	Ensures curvature tests are independent.
Lemma 10.6	CURV	[40] (Z), [41] (properZ)	inv 34	Proper delocalization supports are forbidden.	Blocks the “spread it out” escape.

Result	Stage	Diagram node(s)	Requires	What it does	Role
Lemma 10.8	CURV	[44] (rankwall)	inv 35	Enlarged delocalization = repair networks.	Converts the other escape into killed structure.
Lemma 10.9	CURV	[43] (global)	inv 34, 35; Definitions 3.95 and 3.98	Whole-graph dependence is target-defective, a smaller closed counterexample, or exact on raw curvature labels.	Closes the delocalization loophole.
Lemma 10.10	CURV	[47] (fullrank)	inv 28, 29; Lemma 10.9	Curvature rank $\geq W_2 - o(R)$.	Forces near-maximal demand.
Corollary 10.1	CURV	[48] (cost)	inv 27, 28; Lemma 10.10	Curvature rank $\rightarrow$ entropy tax (c_Ω bits each).	Puts demand into budget currency.
Lemma 11.5	2BUD	[49] (eta), [50] (expensive)	inv 15, 26	A dominant local vertex-type appears in the low-entropy branch.	Handles the cheap-entropy regime.
Lemma 11.2	2BUD	[51] (expbranch)	node [49] positivity; node [50] high-power constructor	The realized remainder family contributes at least $(R /10) \log_2 n$ bits.	Supplies the exact arithmetic input for the node-[52] joint accounting problem.
Definition 11.3	2BUD	[52] (windowacct)	node [24] global completions; node [48] exact states; node [51] high-power bits	Same-context joint window-remainder accounting is realized or retained as an exact open residual.	Prevents an invalid product count; no entropy cap is inferred.
Lemma 11.4	2BUD	[53] (smallbudget)	node [48] realized ledger; node [49] exact state count; node [50] strict low constructor	The exact forced curvature power fits strictly inside the flat power times the low-entropy allowance, so the printed small-budget edge is impossible.	Routes the actual low constructor toward node [55], conditionally on its separate quarter-budget producer.
Remark 12.4	2BUD	[54] (entcap)	inv 8, 24, 29	Closure already holds at $K = 0$ (curvature machinery dispensable).	Over-determination / graceful degradation.

Large-budget residual and local exclusion (terminus).

Result	Stage	Diagram node(s)	Requires	What it does	Role
Definition 13.2	LB	[55] (thetabN)	inv 21, 22	Canonical decomposition of a support into atoms.	Sets up per-cell net-charge bookkeeping.
Definition 13.3	LB	[55] (thetabN)	inv 1, 6, 8, 22, 23, 24	Admissible support carrying all cumulative constraints.	The object the exclusion lemmas act on.
Lemma 13.5	LB	[56] (netcap), [58] (Ndef)	inv 13, 14, 24	Net charge localizes exactly over the connected components of R .	Cell-wise $\mathcal{N} \geq 0$ sums to a global deficiency bound.
Proposition 13.8	LB	[59] (globalnet)	inv 14, 24; Lemma 13.5	The large-budget residual has negative net charge.	Selects the supports to kill.
Definition 14.3	EXC	[88] (thresholds)	inv 23, 24	Defines receivers, missing-port count $q(w)$, routed load $L(w)$, and canonical load routing.	Supplies the bookkeeping variables for Type A.
Lemma 14.4	EXC	[93] (visible)	Definition 14.3	Defines anchored returns, receiver-entry returns, visible loads, and silent loads for one completion port.	Supplies the visible/silent split tested at node [93].
Lemma 14.7, Lemma 14.8	EXC	[93] (visible)	inv 23; Definition 14.5	Proves that the first entry of an anchored return is a receiver and records the degree-two entry budget.	Keeps receiver-entry returns inside the Type A receiver bookkeeping.
Lemma 14.9	EXC	[89] (saturated), [93] (visible), [94] (silent), [109] (route8)	inv 1, 30; Lemma 3.4	Every completion port has an anchored return.	Makes every saturated port test nonvacuous.

Result	Stage	Diagram node(s)	Requires	What it does	Role
Definition 14.10, Lemma 14.12, Definition 14.15	EXC	[109] (route8)	inv 1, 22, 23, 25, 30, 31	Defines route-8 residual data, target-complete-minimal trace basins, the declared u -supported response algebra, and connector bands forced by all internal channels.	Records the reduced silent trace-basin branch.
Definition 14.16	EXC	[95]–[109]	inv 1, 6, 8, 22, 25, 30, 31, 34, 35	Enumerates the eight alternatives tested after node [89] is saturated.	Provides the common exit language for nodes [95]–[109].
Definition 14.33	EXC	[101] (exit4)	inv 6, 8; Definition 14.16	Defines the peeling set for exit (4).	Names the loads that can be removed by the exit-(4) loop.
Lemma 14.34	EXC	[102] (peel)	inv 6, 8; Definition 14.33	Proves the exact 1/4 charge update for one peeled load.	Supplies the numerical update for the loop.
Lemma 14.38	EXC	[102] (peel), [89] (loop)	inv 6, 8; Lemma 14.34	Proves that one supported unpeeled load is removed at exit (4).	Makes the [102] return to [89] a finite strictly decreasing loop in $L_4(w)$.
Lemma 14.35	EXC	[93] (visible), [95]–[108]	inv 1, 8, 25, 30, 31; Definition 14.33	Four visible unpeeled receiver-entry returns realize one of exits (1)–(7), with exit (4) supporting an unpeeled load.	Routes the yes-branch of [93].
Lemma 14.36	EXC	[94] (silent), [101]–[109]	inv 6, 8, 34, 35; Definition 14.33, Definition 14.15	If no port has four visible unpeeled returns, the residual excess set realizes one of exits (4)–(8), with exit (4) supporting an unpeeled excess load.	Routes the no-branch of [93].
Lemma 14.37	EXC	[89] (saturated), [93] (visible), [94] (silent), [101] (exit4), [102] (peel), [109] (route8)	Lemma 14.35, Lemma 14.36	Combines the visible and silent residual routing lemmas after any previous peeling.	Iterates saturated routing until a closed exit, Type B handoff, route 8, or unsaturation occurs.
Definition 14.1	EXC	[107] (exit7)	inv 6, 8, 30, 31	Defines connector germs, first separators, and the finite continuation classes.	Supplies the vocabulary for the exit-(7) split.
Lemma 14.23	EXC	[107] (exit7); feeds [101], [103], [105] when absorbed	inv 6, 8, 30, 31; Definition 14.11	A nonabsorbed four-coordinate family either realizes exits (4)–(6) or has a surviving first separator.	Routes the absorbed subcases back to the already listed saturated exits.
Lemma 14.24	EXC	[107] (exit7)	inv 4, 6, 8; Definition 14.11	Cubic switches are absorbed unless a surviving first separator remains.	Removes the cubic part of exit (7).
Lemma 14.25	EXC	[108] (decorB)	inv 4, 6, 8; Lemma 14.24	A surviving separator has ambient degree at least four and produces a decorated Type B handoff envelope.	Converts exit (7) into the Type B input [108].
Lemma 14.39	EXC	[95]–[108]	inv 1, 6, 8, 25, 30, 31, 34, 35; Lemma 14.38, Lemma 14.25	Closes exits (1)–(3), (5), and (6), peels one load at exit (4), and transfers exit (7) to Type B.	Records the outcome of the explicit saturated-exit tests.
Lemma 14.27	EXC	[88] (thresholds)	inv 4, 24	Computes $H_0 \leq 4$, $H_1 \leq 8$, $H_2 \leq 12$.	Identifies the saturated branch-switch values.
Lemma 14.28	EXC	[93] (visible)	inv 1, 8, 25, 30, 31	Four visible receiver-entry returns force target, quotient, delocalization, or decorated handoff fan data.	Closes the visible saturated branch.
Lemma 14.32	EXC	[94] (silent)	inv 6, 8, 34, 35; Definition 14.19, Definition 14.15	Silent/mixed excess forces compression, delocalization, decorated handoff fan data, or indexed target-complete-minimal trace basins.	Closes the silent branch outside route 8.
Lemma 14.30, Lemma 14.31	EXC	[109] (route8)	inv 6, 8, 34, 35; Lemma 14.32	Visible-first excess gives $S_{\text{sil}}^{\text{exc}}(X) \geq 4D_A(X)$ and records the reduced route-8 residual.	Quantifies the silent route-8 residual.

Result	Stage	Diagram node(s)	Requires	What it does	Role
Lemma 14.40	EXC	[89] (saturated), [109] (route8)	Lemma 14.28, Lemma 14.32, Lemma 14.35, Lemma 14.36, Lemma 14.37, Lemma 14.39	Removes closed exits 1–3, 5, and 6 and exit-4 peeled loads from the 4/8/12 branch, hands exit 7 to Type B, and leaves reduced route 8 as a residual.	Leaves unsaturated Type A plus the target-defect exit, the Type B handoff, and the quantified route-8 branch.
Definition 14.46	EXC	[111] (deficit)	inv 14, 24; Definition 14.15, Lemma 14.31	Defines the Type A route-8 deficit D_A extracted from the reduced silent residual.	Names the route-8 part of the final pressure/route-8 closure.
Lemma 14.47	EXC	[112] (basins), [113] (burden)	inv 14, 24; Definition 14.46	Proves the indexed basin lower bound $N_{\text{basin}} \geq 4D_A$ and the large-budget lower bound for D_A .	Gives the lower side of the route-8 carrier count.
Definition 14.48	EXC	[110] (route8)	inv 8, 23, 30, 31; Definition 14.15	Defines indexed route-8 entries, active boundary carriers, essential carrier cores, and private carriers.	Supplies the carrier vocabulary.
Definition 14.50	EXC	[114] (core)	inv 8, 23, 30, 31; Definition 14.48	Defines true route-8 residual entries after exits (1)–(7) are absent.	States the residual condition used after [114].
Lemma 14.56	EXC	[114] (core)	inv 1, 8, 30, 31; Definition 14.48	A mixed target event in the carrier core crosses at least two carriers.	Gives the cut-parity constraint on essential carrier cores.
Lemma 14.58	EXC	[115] (smallcore), [116] (closedexits)	inv 1, 8, 30, 31; Lemma 14.56	Zero/one-essential-carrier entries realize exits (4)–(7).	Removes the small-core branch at [115].
Proposition 14.59	EXC	[117]–[123]	inv 14, 23, 24; Corollary 8.7, Lemma 14.47, Lemma 14.58	If no two-carrier entry exists, three private carriers per basin violate the boundary-incidence budget; hence a two-carrier entry exists.	Closes route 8 outside the terminal two-carrier entry.
Definition 14.51 Lemma 14.52, Lemma 14.53	EXC	[118] (two-carrier)	inv 6, 8, 23, 24; Definition 14.48	Inclusion-minimality of the essential carrier core gives a declared deletion witness for every essential carrier.	Supplies the witness data used at [124].
Lemma 14.54	EXC	[118] (two-carrier)	inv 1, 6, 8, 30, 31; Lemma 14.53	In a two-carrier entry, each declared deletion quotient lies in the canonical exit-(4) family.	Converts the terminal carrier data into canonical exit-(4) data.
Lemma 14.55	EXC	[124] (nogo)	inv 1, 6, 8, 30, 31; Lemma 14.54	Each canonical deletion quotient realizes exit (4).	Produces the forbidden absent exit in the terminal residual.
Definition 14.60 Theorem 14.61, Proposition 14.62	EXC	[124] (nogo)	Proposition 14.59, Lemma 14.54, Lemma 14.55	Packages the terminal two-carrier obstruction, proves it empty, and derives route-8 closure.	Closes the terminal route-8 interface.
Lemma 14.43	EXC	[90] (unsat)	Lemma 14.4, Lemma 14.27	Unsaturated receivers give the 3/7/11 charge inequality.	Supplies the Type A charge bound.
Lemma 14.45	EXC	[86], [87]–[109]	inv 1, 4, 8, 22, 23, 24, 25, 30, 31; Lemma 14.40, Lemma 14.43	Dense $\sigma = 0$ supports outside target-defect peeling, route 8, and the Type B handoff satisfy $\partial_+(X) \geq \frac{1}{4} V(X) $; a negative support has target-defect, route-8, or Type B handoff data.	Kills pure Type A after the split.
Remark 14.67	EXC	[70] (fsafe), [71] (certq)	inv 25	Certificate-marked fan-degree cap is a structural label-packing bound ($\leq 8/\text{window}$); fan-certificate residual centers are routed to Type B residual mass.	Supplies the Type B constant.
Definition 14.68 Definition 14.69, Definition 14.72	EXC	[72] (typeBfan)	inv 23, 24, 25	Defines fan-window incidences, fan-closed ports, and the hybrid window/non-window credit ledger used by one Type B fan.	Supplies the accounting objects used at the local B1 node.

Result	Stage	Diagram node(s)	Requires	What it does	Role
Definition 14.74, Lemma 14.75, Lemma 14.76	EXC	[72] (typeBfan)	inv 1, 23, 30, 31; Definition 14.68	Constructs the direct cycles forced by same-window and two-window closed incidences.	Removes direct fan-window and two-window cycle configurations before the incidence payment.
Lemma 14.77, Lemma 14.78, Lemma 14.79	EXC	[72] (typeBfan)	inv 23, 24, 25; Definition 14.72	Every certificate-marked positive-deficit residual uses $2c$ distinct non- h incidences, whose half-credit pays D_B with slack. Two fan-closed ports give $D_B \geq (k-3)/4 > 0$ and enter the local Type B fan ledger.	Quantifies and closes the local B1 Type B fan branch before the global B2 ledger argument. Closes a fan-closed local Type B center.
Proposition 14.80	EXC	[72] (typeBfan)	inv 4, 8, 25; Definition 14.69, Lemma 14.79	A compatible open pair produces the two fan-closed ports required by the local Type B fan ledger.	
Corollary 14.81	EXC	[69] (localgt), [72] (typeBfan)	inv 4, 8, 25; Lemma 14.70, Proposition 14.80	Degree-four triangular and heavy triangular routes produce the fan-closed ports required at node [72]; the heavy triangular route gives $D_B \geq (5k-19)/4 > 0$.	Routes compatible-pair data to node [72].
Corollary 14.82	EXC	[78]–[81], [72] (typeBfan)	inv 4, 8, 25; Lemma 3.18, Lemma 3.19, Proposition 14.80	Defines the exact receiver overload after deleting every actual high center and proves $\mathcal{N}_-(X) \leq \frac{21}{4}S_X + \frac{1}{4}\Omega_A(X)$ without local-entry or B2 hypotheses.	Routes triangular local data to the local Type B fan ledger.
Definition 14.90 and Proposition 14.96	EXC	[74]	inv 4, 14, 24; Hegde–Sandeep–Shashank internal-three-core exclusion		Gives the choice-free Type B-to-Type A quantitative boundary; unresolved and overlap branches cannot hide unbounded Type B charge.
Definition 14.83, Definition 14.84, Definition 14.85, Lemma 14.87, Lemma 14.89, Proposition 14.90, Definition 14.91, Proposition 14.97	EXC	[73]	inv 1, 4, 8, 22, 23, 24, 25, 34, 35; Lemma 14.78	Defines Type B carrier ledgers for supports with no fan-certificate residual center, proves B2 failures produce minimal overlap obstructions, and records the global-to-local constraints inherited by such obstructions.	Reduces local Type B closure to the B2 refined support ledger after local B1 is proved; disjoint-carrier failure is a globally admissible overlap obstruction.
Definition 14.19	EXC	[108] (decorB)	inv 1, 4, 8, 22, 25; Lemma 14.25	Defines the decorated fan envelope produced by Type A exit (7).	Records the incoming Type B support from node [108].
Lemma 14.26	EXC	[65] (typeB)	inv 1, 4, 8, 22, 25; Definition 14.19	Proves the decorated handoff is admissible Type B fan-envelope data.	Makes the node [108] handoff usable at Type B node [65].
Definition 14.18	EXC	[70] (fsafe)	inv 16, 25	Defines certificate-marked Type B fans.	Supplies the local certificate object for Type B fans.
Remark 14.67	EXC	[71] (certq)	inv 16, 25; Definition 14.18	Records the fan-degree label-packing cap.	Supplies the local certificate cap for Type B fans.
Definition 3.7	EXC	[78] (degree4)	inv 4, 8, 16, 25	Defines triangular ports at high-degree centers.	Names the triangular local configuration.
Definition 3.15	EXC	[79] (profile)	inv 4, 8, 16, 25; Definition 3.7	Defines triangular fan cores.	Supplies the profile object for triangular routing.
Lemma 3.9	EXC	[68] (degsplit), [69] (localgt)	inv 4, 8, 16, 25	Puts the high-degree neighbourhood into the normal form used by the local split.	Starts the Type B degree split.
Definition 3.10, Lemma 3.11	EXC	[69] (localgt)	inv 4, 8, 16, 25; Lemma 3.9	Defines fan-compatible open ports and proves compatibility for same-center open pairs.	Supplies the fan-compatible branch of the local split.

Result	Stage	Diagram node(s)	Requires	What it does	Role
Corollary 3.14	EXC	[78] (degree4)	inv 4, 8, 16, 25; Lemma 3.9	Activates the degree-four local branch.	Separates the degree-four case before triangular routing.
Lemma 3.12	EXC	[69] (localgt)	inv 4, 8, 16, 25; Definition 3.7, Definition 3.10	Splits heavy centers into a fan-compatible open pair or a triangular-port alternative.	Supplies the heavy-center local split.
Corollary 3.13	EXC	[79] (profile)	inv 4, 8, 16, 25; Lemma 3.12	Records the local dichotomy in the profile used downstream.	Supplies the local Type B case split before the incidence ledger.
Lemma 3.16	EXC	[78] (degree4)	inv 1, 4, 8, 25, 30, 31; Definition 3.15	Triangular shoulders have completion edges.	Sets up the triangular first-landing analysis.
Lemma 3.17	EXC	[79] (profile)	inv 1, 4, 8, 25, 30, 31; Lemma 3.16	Every relevant return enters through a shoulder.	Routes triangular returns to the landing analysis.
Lemma 3.18	EXC	[80] (landing)	inv 1, 4, 8, 25, 30, 31; Lemma 3.17	The first landing is central, cross-triangular, or outside.	Routes triangular ports toward fan-closed data.
Lemma 3.19	EXC	[81] (cross)	inv 1, 4, 8, 25, 30, 31; Lemma 3.18	Cross-shoulder multiplicity supplies the required second port.	Completes the triangular route to fan-closed data.
Lemma 14.70	EXC	[69] (localgt), [72] (typeBfan)	inv 4, 8, 25; Definition 3.10	A compatible open pair supplies two fan-closed ports at the same center.	Routes compatible-pair subcases into the Type B incidence payment.
Definition 14.7	EXC	[71]–[73], [82]–[85]	inv 4, 8, 16, 24, 25	Defines the certificate-marked multiclosed Type B residuals, including the B2-failure residual classes.	Names the Type B residuals that survive the local B1 ledger.
Lemma 14.98	EXC	[65] (typeB), [71]–[73], [82]–[85]	inv 4, 8, 16, 24, 25; Definition 14.71, Lemma 14.78, Lemma 14.79	Applies the refined Type B ledger: certificate-closed fans have nonnegative charge, positive-deficit fans are paid by B1, and B2 failures are isolated as residuals.	Kills Type B outside route 8 and outside the Type B bridge residual.
Definition 14.10	EXC	[75] (fanmass)	inv 4, 14, 16, 24, 35; Definition 4.1	Defines the Type B residual mass for fan-certificate residual centers, grouped decorated envelopes, and B2 failures.	Names the global bridge-mass quantity.
Lemma 14.104	EXC	[75] (fanmass)	inv 4, 14, 16, 24, 35; Definition 14.101	Extracts a route-8 non-window core from grouped decorated envelopes into D_A .	Keeps decorated-envelope mass separate from Type A route-8 deficit.
Lemma 14.100	EXC	[84] (local fanmass)	inv 4; nodes [75], [81]–[83]	Charges one exact selected center set in one fixed Type B scope to that scope's assigned high-center surplus.	Retains unresolved, nonnegative, route-8, and minimal-overlap outcomes without constructing a support family.
Lemma 14.106	EXC	[84] (global fanmass)	inv 4, 14, 16, 24, 35; Definition 14.101	Extracts a route-8 non-window core from Type B bridge residual data into D_A .	Keeps Type B bridge mass separate from Type A route-8 deficit.
Lemmas 14.102 and 14.105, Proposition 14.107	EXC	[84] (global fanmass)	inv 4, 14, 16, 24, 35; Definition 4.1; Lemma 14.79, Definition 14.91	Charges residual centers, processed-envelope boundary loss, and B2 failures to high-degree fan-center surplus and obtains $M_B \leq 416\sigma(G) = o(R)$.	Open producer contract: construct the canonical ordinary/grouped family and both coefficient-208 incidence bounds before this consequence is available.
Theorem 15.1	EXC	[76] (branch), [85] (typeB closed), [123] (pressure closed)	inv 20, 27, 28, 29, 35, 39	Closes the entropy-admissible branch outside explicit Type A target-defect/route-8 and Type B bridge residuals.	Combines the local exclusion theorems before the final pressure descent.

Result	Stage	Diagram node(s)	Requires	What it does	Role
Theorem 15.48 EXC		[123] (pressure closed)	inv 20, 27, 28, 29, 35, 39; Proposition 14.107 , Definition 15.45 , Lemma 15.46 , Lemma 15.44 , Lemma 15.47 , Proposition 14.59	Closes the unified Type A target-defect/route-8 ledger by finite exit-(4) descent and the terminal two-carrier route-8 no-go.	Sends only the terminal route-8 survivor to node [124], where it is impossible.
Theorem 1.1	(root goal)	[2] (bad) → terminal	Theorem 5.1 + surplus-token routing + all spine invariants	Every minimal counterexample either closes in the sparse-surplus branch, routes to Type B fan data, or reaches the near-cubic spine and then the terminal two-carrier route-8 no-go, which is empty.	The main closure theorem.

Invariant → lemmas reverse index (impact analysis). If invariant k needs revision, these results are affected; diagram-node numbers are shown in parentheses.

Inv.	Invariant	First node	Affected results	Impact note
1	Mersenne-return target algebra	Target [5]	Lemma 2.2 ([5]–[7]), Definition 2.5 ([5]), Lemma 3.101 ([11]), Lemma 3.102 ([12]), Lemma 3.103 ([13]), Lemma 3.107 ([18]), Lemma 4.7 ([22]), Definition 13.3 ([55]), Lemma 14.8 , Lemma 14.9 , Lemma 14.12 ([109]), Definition 14.16 , Lemma 14.39 , Lemma 14.28 ([93]), Lemma 14.75 ([72]), Lemma 14.44 , Lemma 14.45 ([86]), Definition 14.51 , Theorem 14.61 ([124])	Pervasive target predicate.
2	Minimal proper-subgraph condition	Core [8]	Lemma 3.1 ([8]); downstream use in Lemma 3.5 , Lemma 3.103 , Lemma 10.6	Excludes smaller $\delta \geq 3$ cores.
3	Edge-deletion criticality	Delete [9]	Lemma 3.2 ([9]); its high-degree corollary feeds inv 4	Edge-local minimality.
4	High-degree independence	High-degree [10]	Lemma 3.2 ([10]), Lemma 14.27 ([88]), Lemma 14.43 ([90]), Lemma 14.45 ([86]), Lemma 14.98 ([65]), inv 16	Controls surplus aggregation.
5	No passive degree-2 storage	Counterex. [2]	Lemma 8.4 ([29]), Lemma 10.8 ([44]), Lemma 10.9 ([43]), auxiliary repair-network closures	Prevents unrecorded subdivision storage.
6	Replacement dominance	Replace [13]	Lemma 3.101 ([11]), Lemma 3.103 ([13]), Definition 13.3 ([55]), Lemma 14.32 ([94]), Lemma 14.39 , Definition 14.51 , Definition 14.60 ([124])	Orders boundary profiles.
7	Lumpability defect	Univ. [12]	Lemma 3.102 ([12]), Lemma 3.103 ([13]), Lemma 6.1 ([18]); final spine step through inv 8	Detects invalid quotients.
8	Hereditary target-uncompressibility	Uncomp. [14]	Corollary 3.105 ([14]), Definition 13.3 ([55]), Remark 12.4 ([54]), Definition 14.15 ([109]), Definition 14.16 , Lemma 14.39 , Lemma 14.28 ([93]), Lemma 14.32 ([94]), Lemma 14.40 ([89]), Lemma 14.44 , Lemma 14.45 ([86]), Lemma 14.98 ([65]), Definition 14.51 , Definition 14.60 , Theorem 14.61 ([124]), Theorem 15.1 ([76], [123])	Main contradiction lever.
9	Minimum edge count	Counterex. [2]	Lemma 3.5 ([4]), Lemma 4.4 ([21])	Lower edge pressure.
10	Sparse upper envelope	Core [8]	Lemma 3.3 , Lemma 4.3 ([21]); feeds the Laman identity inv 12	Upper edge pressure.
11	Cycle-rank pressure	Minimality [4]	Lemma 3.5 ([4]); feeds inv 12 and inv 13	Cycle-rank lower bound.
12	Laman slack identity	Rank ids [4]	Lemma 4.4 ([21]), Definition 11.1	Links rank and sparsity.
13	Surplus-rank relation	Rank ids [4]	Lemma 13.5 ([56], [58]); net-charge notation throughout the large-budget branch	Translates surplus into charge.

Inv. Invariant	First node	Affected results	Impact note
14 High-degree surplus bound	Surplus branch [125]–[144] supplies spine estimate	Definition 4.1, Lemma 3.32, Definition 3.37, Lemma 3.38, Definition 3.39, Lemma 3.40, Definition 3.41, Lemma 3.44, Lemma 3.45, Definition 3.42, Lemma 3.43, Definition 3.46, Definition 3.47, Definition 3.74, Lemma 3.53, Lemma 3.84, Theorem 3.85, Corollary 3.50, Corollary 3.51, Lemma 3.64, Theorem 3.63, Definition 8.3, Lemma 8.5, Corollary 13.6, Proposition 3.87, Lemma 13.5 ([56], [58]), Proposition 13.8 ([59]), Proposition 14.59, Definition 14.60 ([124])	Keeps surplus sublinear on the derived spine branch, or derives the exact high-load alternative from the complete pair schedule: blocker-free pairs use their first selected port, blocked pairs use their canonical capacity token, and total role fibres count every pair exactly once. The P_{13} join identity quantifies window concentration, and any large homogeneous same-token fibre is a geometric bottleneck routed to sparse exits or Type B.
15 Near-cubic edge count	Surplus branch [125]–[144] supplies spine estimate	Definition 4.1, Theorem 3.85, Proposition 3.87, Lemma 4.4 ([21]), Lemma 4.2 ([21]), Lemma 11.5 ([49], [50]); entropy estimates using B_{skel}	Fixes the budget coefficient after the near-cubic bound is supplied by the exact surplus-token branch.
16 Certificate-marked fan degree	High-degree [67]	Lemma 14.98 ([65]), Remark 14.67 ([70], [71]), Lemma 14.66; related to inv 4	Auxiliary fan-label cap; fan-certificate residual centers are charged through surplus mass.
17 Spectral feasibility mask	Spectral [70]	Lemma 14.98 ([70]), Lemma 14.66, Remark B.5	Auxiliary spectral mask.
18 Spectral equality branch	Spectral [70]	Lemma 14.98 ([70]), Lemma 14.66, Remark B.5	Cubic equality subcase.
19 Spectral surplus branch	Spectral [70]	Lemma 14.98 ([70]), Lemma 14.66, Remark B.5	Surplus spectral subcase.
20 Low P_{13} density	Density [24]	Remark 7.1 ([21]–[22]), Remark 7.67 ([24]), Proposition 12.2 ([54]), Theorem 15.1 ([76], [123]), Lemma B.2	Window-density cap.
21 Forced remainder	Residual A [25]	Lemma 8.4 ([26], [29]), Lemma 9.1 ([30]), Definition 13.2 ([55]), Lemma B.2; all atom analysis	Supplies the large remainder.
22 P_{13} -free atom condition	Residual A [25]	Corollary 5.2 ([15]), Lemma 8.4 ([29]), Definition 13.2 ([55]), Definition 13.3 ([55]), Definition 14.19 ([108]), Definition 14.15 ([109]), Definition 14.16, Lemma 14.28 ([93]), Lemma 14.45 ([86]), Lemma 14.89	Local P_{13} -free structure inherited by overlap supports.
23 Window stub capacity	Stub accounting [29]	Lemma 8.4 ([29]), Definition 13.3 ([55]), Definition 14.3, Definition 14.5, Lemma 14.8, Definition 14.10, Definition 14.15 ([109]), Definition 14.48, Definition 14.60 ([124]), Definition 14.68, Definition 14.72, Lemma 14.77 ([72]), Lemma 14.78, Lemma 14.79, Definition 14.83, Definition 14.84, Lemma 14.89, Definition 14.91, Proposition 14.97, Lemma 14.45 ([86]); feeds inv 24	Outside-incidence supply; positive-deficit Type B residuals use hybrid half-credit on window and non-window incidences, and B2 is the required disjoint ledger.

Inv.	Invariant	First node	Affected results	Impact note
24	Remainder deficiency density	Stub accounting [29]	Lemma 8.4 ([29]), Definition 13.3 ([55]), Lemma 13.5 ([56], [58]), Proposition 13.8 ([59]), Remark 12.4 ([54]), Lemma 14.27 ([88]), Lemma 14.43 ([90]), Proposition 14.59, Definition 14.60 ([124]), Lemma 14.77 ([72]), Definition 14.83, Definition 14.84, Definition 14.85, Definition 14.91, Proposition 14.97, Lemma 14.44, Lemma 14.45 ([86])	Supply ceiling; Type B residual deficit is $D_B = c - \frac{11-k}{4}$ and bridge failure is represented by a minimal overlap obstruction.
25	Legal P_{13} labels	Label algebra [18]	Lemma 6.1 ([18]), Lemma 6.3 ([21]), Remark 14.67 ([70], [71]), Definition 14.15 ([109]), Definition 14.16, Lemma 14.39, Lemma 14.28 ([93]), Definition 14.73 ([72]), Definition 14.74, Lemma 14.75 ([72]), Lemma 14.76, Lemma 14.89, Lemma 14.44, Lemma 14.45 ([86]), Lemma 14.98 ([65]), Lemma 14.66	Finite label algebra and structural fan-window endpoint bookkeeping.
26	Two-step curvature density	Curvature enumeration [21]	Lemma 6.3 ([21]), Lemma 10.4 ([31], [32]), Lemma 11.5 ([49], [50]), Lemma 11.7 (separate repetitive low-entropy hypothesis), Section A	Curvature-positive wedge count.
27	Flatness entropy cost	Cost [48]	Corollary 10.11 ([48]), Remark 10.12, Proposition 12.2 ([54]), Theorem 15.1 ([76], [123])	Bits per independent test.
28	Raw atom curvature supply	Wedge lower bound [30]	Lemma 9.1 ([30]), Lemma 10.4 ([31], [32]), Lemma 10.5 ([35]), Lemma 10.10 ([47]), Corollary 10.11 ([48]), Theorem 15.1 ([76], [123])	Demand floor.
29	Curvature compression ratio	Full-rank squeeze [47]	Lemma 10.10 ([47]), Lemma 11.7 (separate repetitive low-entropy hypothesis), Remark 12.4 ([54]), Theorem 15.1 ([76], [123]); collides with inv 8	Independent-test cap.
30	Two-path criterion	Target algebra [5]	Lemma 2.2 ([5]–[7]), Lemma 14.12 ([109]), Definition 14.16, Lemma 14.39, Lemma 14.28 ([93]), Definition 14.51, Definition 14.60, Theorem 14.61 ([124]), Lemma 14.44, Lemma 14.45 ([86])	Pair-of-path target test.
31	Theta closure	Target algebra [5]	Lemma 3.102 ([12]), Lemma 3.103 ([13]), Definition 14.15 ([109]), Lemma 14.28 ([93]), Definition 14.51, Definition 14.60 ([124]); atom hazard profiles	Theta sums avoid targets.
32	Ear closure	Target algebra [5]	Lemma 3.102 ([12]), Lemma 3.103 ([13]); target-response profiles	Ear completions avoid targets.
33	Symmetric difference	Context univ. [12]	Lemma 3.102 ([12]), Lemma 3.103 ([13]), Lemma 10.9 ([43])	Overlap-cycle target test.
34	Rank-one dyadic support	Deloc. support [40]	Lemma 10.6 ([40], [41]), Lemma 10.9 ([43]), Lemma 14.28 ([93]), Lemma 14.32 ([94]), Lemma 14.39, Definition 14.60, Theorem 14.61 ([124]), Lemma 10.10 ([47]), Lemma B.3	Blocks rank-one smearing.
35	Composite cancellation	Repair id. [44]	Lemma 10.8 ([44]), Lemma 10.9 ([43]), Lemma 14.32 ([94]), Lemma 14.39, Definition 14.60, Theorem 14.61 ([124]), Lemma B.3	Controls cancellation repairs.
36	Odd carrier nonempty	Target algebra [5]	Lemma 2.2 ([5]–[7]), Lemma 3.102 ([12]), Lemma 10.9 ([43])	Power-of-two carrier test.
37	No common carrier	Carrier transport [35]	Lemma 10.5 ([35]), Lemma 10.6 ([40], [41]), Lemma 10.9 ([43])	Carrier-transport obstruction.
38	Polarized carrier transport	Separated testers [35]	Lemma 10.5 ([35]), Lemma 10.4 ([31], [32]), Lemma 10.6 ([40], [41]), Lemma 10.9 ([43])	Separates carrier testers.

Fast gap-check usage. To check the proof, confirm for each result that every invariant in its **Requires** cell is established at an earlier-or-equal node; a required invariant introduced at a strictly later node is a forward-reference gap. For the contradiction-producing results Lemma 14.45, Lemma 14.98, and Theorem 15.1, confirm the chain $\text{inv } 8 \leftarrow \text{Corollary 3.105} \leftarrow \text{Lemma 3.103}$ is intact and not circular with the result invoking it. The two-sided squeeze lives in exactly two cells: **demand** is Lemma 9.1 (invariant 28) and **supply** is Lemma 8.4 / Proposition 11.9 (invariants 24, 29); on the derived spine branch, any repair to a curvature estimate must keep demand $>$ supply, i.e. $2.543 > 0.611$ per $|R|$ (sharper high-entropy demand 2.574), or in the net-charge form $0.25 > 0.2282$.




Figure 12: Visual definition of the edge-rooted target algebra. The white circled vertices are the endpoints u and v , the grey arrow is the oriented root edge $e = uv$, and the black dots are internal vertices of the return path. The red path P lies in $G - e$, runs from v back to u , and has length $r = |P|$. The box labelled $r = 2^j - 1 \in \mathcal{M}$ records the target case: adjoining the root edge gives the simple cycle $C = P \cup e$ of length $|C| = r + 1 = 2^j \in \mathcal{D}$. The counterexample box records the equivalent local target condition $R_e(G) \cap \mathcal{M} = \emptyset$ for every oriented edge e .

2 Target cycles as Mersenne returns

Definition 2.1 (Mersenne return set). *Let G be a finite simple graph and let $e = uv$ be an oriented edge. Define*

$$R_e(G) = \{r : \text{there is a simple path of length } r \text{ from } v \text{ to } u \text{ in } G - e\}.$$

Let

$$\mathcal{M} = \{2^j - 1 : j \geq 2\} = \{3, 7, 15, 31, \dots\}.$$

Lemma 2.2 (Edge-rooted target equivalence). *A graph G contains a cycle of length 2^j if and only if there is an oriented edge e such that $2^j - 1 \in R_e(G)$. Hence G has no power-of-two cycle if and only if*

$$R_e(G) \cap \mathcal{M} = \emptyset \quad \text{for every oriented edge } e.$$

Proof. Suppose C is a cycle of length 2^j . Choose an oriented edge $e = uv$ on C . Removing e from C leaves a simple path of length $2^j - 1$ from v to u , and therefore $2^j - 1 \in R_e(G)$. Conversely, if $2^j - 1 \in R_e(G)$, then e together with the corresponding simple return path gives a simple cycle of length 2^j . $\square$

Lemma 2.3 (Two-path criterion). *Let P and Q be internally vertex-disjoint simple paths with the same endpoints. Then $P \cup Q$ is a simple cycle of length $|P| + |Q|$. In particular, if $|P| + |Q| \in \mathcal{D}$, then G contains a dyadic cycle.*

Proof. The two paths meet only in their common endpoints, so traversing P from one endpoint to the other and then traversing Q back gives a cycle with no repeated internal vertex. Its edge set is the disjoint union of the edge sets of P and Q , hence its length is $|P| + |Q|$. $\square$

Definition 2.4 (Counterexample). *A graph G is a counterexample if*

$$\delta(G) \geq 3 \quad \text{and} \quad R_e(G) \cap \mathcal{M} = \emptyset \quad \forall e \in \vec{E}(G).$$

Definition 2.5 (Dyadic-safe; target-safe). *A cycle is dyadic if its length is a power of two. A subgraph $X \subseteq G$ is dyadic-safe (equivalently target-safe) if G contains no dyadic cycle meeting $V(X)$. This is the contextual notion: such a cycle may run partly through $G - X$, so dyadic-safety of X is a property of how X is embedded in G , strictly stronger than the absence of a dyadic cycle internal to X . For $X = G$ it is exactly the no-power-of-two-cycle clause of the counterexample condition.*

For contradiction, throughout the paper we assume that a counterexample exists and choose one, denoted G , lexicographically minimal by

$$(|V(G)|, |E(G)|).$$

Set

$$n = |V(G)|, \quad m = |E(G)|.$$

Framework branch table. The target-cycle section supplies the first branch table.

Branch	Residual condition	Demand	Payer	Bounded multiplicity	Closure
Mersenne return	Some oriented edge e has $R_e(G) \cap \mathcal{M} \neq \emptyset$.	No residual demand remains.	The root edge and return path.	One oriented edge.	Closed by Lemma 2.2 .
Two-path closure	Two internally disjoint paths have dyadic total length.	No residual demand remains.	The two path lengths.	One unordered endpoint pair.	Closed by Lemma 2.3 .
Counterexample branch	All oriented edges avoid Mersenne returns.	The target-safe condition.	The counterexample definition.	Every oriented edge of G .	This is the standing branch of Definition 2.4 .

3 Minimality and replacement

Lemma 3.1 (No proper core). *Every proper subgraph $H \subsetneq G$ satisfies $\delta(H) \leq 2$.*

Proof. If $H \subsetneq G$ had $\delta(H) \geq 3$, then H would also have no power-of-two cycle, because every cycle of H is a cycle of G . Thus H would be a smaller counterexample, contradicting the minimality of G . $\square$

Lemma 3.2 (Deletion criticality). *Every edge of G has at least one endpoint of degree 3. In particular, the set $V_{\geq 4}(G)$ is independent.*

Proof. Suppose for contradiction that some edge $xy \in E(G)$ has $d_G(x), d_G(y) \geq 4$. Deleting xy lowers the degrees of x and y by one, to at least 3, and leaves every other degree unchanged, so $\delta(G - xy) \geq 3$. Moreover every cycle of $G - xy$ is a cycle of G , so deleting xy creates no new cycle and in particular no new power-of-two cycle; hence $R_e(G - xy) \subseteq R_e(G)$ for every oriented edge e , and $G - xy$ has no power-of-two cycle. Thus $G - xy$ is a counterexample with the same vertex set and one fewer edge, contradicting the edge-minimality of G . Therefore no edge joins two vertices of degree at least 4; equivalently, every edge has an endpoint of degree exactly 3, and the set $V_{\geq 4}(G)$ of vertices of degree at least 4 is independent. $\square$

Lemma 3.3 (Sparse upper envelope). *The minimal counterexample satisfies*

$$m \leq 2n - 2.$$

Proof. First G has at least one vertex of degree 3. Otherwise every vertex would have degree at least 4, and any edge of G would have two endpoints of degree at least 4, contradicting Lemma 3.2. Fix a degree-3 vertex v .

Every subgraph of $G - v$ is a proper subgraph of G . By Lemma 3.1, each such subgraph has a vertex of degree at most 2. Hence $G - v$ is 2-degenerate. A 2-degenerate graph on $N \geq 2$ vertices has at most $2N - 3$ edges: remove vertices in a degeneracy order, count at most 2 deleted edges at each step until two vertices remain, and then count at most one final edge.

Applying this with $N = n - 1$ gives

$$e(G - v) \leq 2(n - 1) - 3.$$

Since $d_G(v) = 3$,

$$m = e(G - v) + d_G(v) \leq 2(n - 1) - 3 + 3 = 2n - 2.$$

□

Lemma 3.4 (Bridgelessness of the minimal counterexample). *The graph G has no bridge. Consequently every edge of G lies on a cycle; equivalently, $R_e(G) \neq \emptyset$ for every oriented edge $e \in \vec{E}(G)$.*

Proof. Suppose for contradiction that $e = uv$ is a bridge. Let A and B be the components of $G - e$ containing u and v , respectively. Form G' from $G - e$ by identifying u and v to a single vertex z .

The graph G' is simple. Indeed, after deleting e , the neighbours of u lie in $A - \{u\}$ and the neighbours of v lie in $B - \{v\}$, which are disjoint. Thus identifying u and v creates no parallel edge, and deleting e before the identification creates no loop.

The minimum degree remains at least 3. Every vertex other than z has the same degree as in G . The identified vertex has

$$d_{G'}(z) = d_G(u) + d_G(v) - 2 \geq 4,$$

since $d_G(u), d_G(v) \geq 3$ and the bridge uv is counted once at each endpoint before deletion.

Next, G' has no power-of-two cycle. Any cycle of G' that uses z and has vertices on both sides would contain, after deleting z , a path from a vertex of $A - \{u\}$ to a vertex of $B - \{v\}$ in $G' - z$. This is impossible because $A - \{u\}$ and $B - \{v\}$ are disconnected in $G' - z$. Hence every cycle of G' lies wholly in the image of A or wholly in the image of B , and is therefore already a cycle of G . Since G has no power-of-two cycle, neither does G' .

Thus G' is a counterexample with $|V(G')| = |V(G)| - 1$, contradicting the vertex-minimality of G . Therefore G has no bridge. The final assertion follows because an edge is a bridge if and only if it lies on no cycle, and a cycle through an oriented edge $e = ab$ is exactly a return path from b to a in $G - e$ together with e .

For the finite formal audit, the contraction is represented without a graph search: the name (v) is removed and every former (v) -edge other than (uv) is reattached to (u) . The bridge condition makes the two endpoint neighbourhoods disjoint. They therefore inject into the new neighbourhood of the identified vertex, while the neighbour set of every other vertex injects under the same collapse map; this proves the minimum-degree assertion. To transport cycles, rotate a contracted cycle at the identified vertex. Its two incident cycle edges must both originate at (u) or both at (v) , since opposite origins would lift the cycle with its final edge removed to a (u) – (v) walk in $(G - uv)$. Replacing the identified vertex by the common origin then lifts the whole simple cycle injectively and preserves its length. Thus the reduction uses a single supplied bridge and no enumeration of components, paths, cycles, or graphs. □

Lemma 3.5 (Cycle-rank pressure). *The cycle rank*

$$\beta(G) = m - n + 1$$

satisfies

$$\beta(G) \geq \frac{n}{2} + 1.$$

Proof. Since $\delta(G) \geq 3$, we have $2m = \sum_v d(v) \geq 3n$. Therefore $m - n + 1 \geq n/2 + 1$. $\square$

Definition 3.6 (Surplus ports and excess selectors). *Let*

$$H = V_{\geq 4}(G).$$

By Lemma 3.2, H is independent, so every neighbour of a vertex in H has degree exactly 3.

A surplus port is an ordered edge

$$p = (h, x)$$

with $h \in H$ and $x \in N_G(h)$. We write

$$c(p) = h, \quad x(p) = x$$

for its high centre and cubic endpoint. Since $d_G(x) = 3$, there are two vertices, denoted a_p, b_p , such that

$$N_G(x(p)) = \{c(p), a_p, b_p\}.$$

The unordered pair

$$s(p) = \{a_p, b_p\}$$

is the shoulder pair of p , and the unordered pair

$$e_p^* = a_p b_p$$

is its shoulder chord.

An excess selector is a choice, for each $h \in H$, of a set

$$\mathcal{P}_{\text{exc}}(h) \subseteq \{(h, x) : x \in N_G(h)\}$$

of size $d_G(h) - 3$. We set

$$\mathcal{P}_{\text{exc}} = \bigcup_{h \in H} \mathcal{P}_{\text{exc}}(h).$$

Thus

$$|\mathcal{P}_{\text{exc}}| = \sum_{h \in H} (d_G(h) - 3) = \sigma(G).$$

The three unselected incident ports at a high vertex are its base ports; the selected ports are exactly the units of degree surplus that one would like to remove in order to make the graph cubic.

Definition 3.7 (High centers and port types). *A vertex $h \in V(G)$ is a high center if $d_G(h) \geq 4$, and a heavy center if $d_G(h) \geq 5$. Let $p = (h, x)$ be a surplus port at a high center. The port p is triangular if its shoulder chord lies in G :*

$$e_p^* \in E(G).$$

Equivalently, if $N_G(x) = \{h, a_p, b_p\}$, then $xa_p b_p x$ is a triangle not using h . The port is open if

$$e_p^* \notin E(G).$$

The word open records only the absence of the shoulder chord in G ; no suppressed graph is associated to the port.

Remark 3.8 (Finite selected-port audit). *Fix the declared vertex order and, at each center h , take the first $d_G(h) - 3$ neighbors as a canonical excess selector. The local information used below can then be exposed by five bounded scans followed by one proof-level interpretation. First, exclusion of a cycle of length 4 gives the neighborhood matching and common-neighbor conclusions of Lemma 3.9. Second, each selected port is classified as open or triangular from the adjacency of its two shoulder vertices. Third, the open selected ports are grouped by their center, producing either two at one center or the certificate that every center contains at most one. Fourth, only in the former branch, the two port endpoints are compared over the declared vertex set by their adjacency responses. Finally, one local witness per selected port certifies its exact two-shoulder contribution, so the resulting shoulder ledger has total $2|\mathcal{P}_{\text{exc}}|$. For the exact pair returned by the overloaded center fibre, the four-cycle-free common-neighbor conclusion then proves that nonadjacent endpoints have disjoint shoulder pairs and cannot occur in one another's shoulder pair. Thus this branch satisfies Definition 3.10; endpoint adjacency remains the explicit complementary branch. Finally, for any requested high center h , one scan of its declared neighbor schedule classifies all $d_G(h)$ incident ports. If the compatible-pair branch is absent, the matching-neighborhood conclusion forces the open-port count to be at most two, and the exact open/triangular cardinality partition gives at least $d_G(h) - 2$ triangular ports. This is the finite audit used in Lemma 3.12 and Corollary 3.13.*

These operations enumerate only the selected ports and the declared vertex set. In particular, they do not enumerate port pairs, vertex subsets, completion graphs, or response tables. In particular, the all-port step does not construct the port-pair universe: the three-port contradiction in the matching argument is proof-level. This audit supplies local structural data; it does not establish the Type B routing clauses, shoulder-completion routing, or the free/blocked sparse pair-response assertions.

Lemma 3.9 (High-neighbourhood normal form). *Let h be a high center. Then the following hold.*

- (a) *Every vertex in $N_G(h)$ has degree 3.*
- (b) *The graph induced by $N_G(h)$ is a matching.*
- (c) *If $x, y \in N_G(h)$ are nonadjacent, then x and y have no common neighbour outside $\{h\}$.*

Proof. For (a), apply Lemma 3.2 to the edge hx . Since $d_G(h) \geq 4$, the degree-3 endpoint of hx is x .

For (b), suppose that $G[N_G(h)]$ contained two adjacent edges xy and yz . Then

$$hxyzh$$

is a simple cycle of length 4, contradicting the assumption that G contains no power-of-two cycle. Thus every vertex of $G[N_G(h)]$ has degree at most 1, so $G[N_G(h)]$ is a matching.

For (c), suppose that nonadjacent $x, y \in N_G(h)$ had a common neighbour $z \neq h$. Then

$$hxzyh$$

is a simple cycle of length 4, again impossible. Hence no such z exists. $\square$

Definition 3.10 (Fan-compatible open ports). *Let h be a high center, and let $p = (h, x)$ and $q = (h, y)$ be distinct open ports at h . The pair $\{p, q\}$ is fan-compatible if*

$$x \notin s(q), \quad y \notin s(p), \quad s(p) \cap s(q) = \emptyset.$$

Equivalently, the two shoulder pairs are disjoint and neither port vertex is a shoulder endpoint for the other port; in this case the four non- h incidences can be assigned as local Type B fan incidences.

Lemma 3.11 (Same-center open ports are fan-compatible). *Let h be a high center, and let $p = (h, x)$ and $q = (h, y)$ be distinct open ports at h . If $xy \notin E(G)$, then $\{p, q\}$ is fan-compatible.*

Proof. Since $xy \notin E(G)$, the endpoint x is not in the shoulder pair $s(q)$, and y is not in the shoulder pair $s(p)$. Suppose that $s(p) \cap s(q)$ contained a vertex z . Then $z \neq h$, and the two nonadjacent neighbours $x, y \in N_G(h)$ would have the common neighbour z outside $\{h\}$. This contradicts Lemma 3.9. Hence the shoulder pairs are disjoint, and the pair is fan-compatible by definition. $\square$

Lemma 3.12 (Heavy-center triangular alternative). *Let h be a heavy center of degree k . If no two open ports at h form a fan-compatible open pair, then at least $k - 2$ ports at h are triangular. In particular, at least three ports at h are triangular.*

Proof. Let $U \subseteq N_G(h)$ be the set of vertices x such that the port (h, x) is open. If two vertices $x, y \in U$ were nonadjacent, then Lemma 3.11 would make the two ports (h, x) and (h, y) fan-compatible, contrary to the hypothesis. Therefore every two vertices in U are adjacent.

By Lemma 3.9, $G[N_G(h)]$ is a matching. A matching has no clique of size 3, so $|U| \leq 2$. All ports at h whose cubic endpoint is not in U are not open, hence triangular by Definition 3.7. Therefore the number of triangular ports at h is at least

$$k - |U| \geq k - 2.$$

Since $k \geq 5$, this is at least 3. $\square$

Corollary 3.13 (Heavy-center local dichotomy). *Let h be a heavy center. Then one of the following alternatives holds.*

- (i) *There are two open ports at h that are fan-compatible.*
- (ii) *At least $d_G(h) - 2$ ports at h are triangular; in particular, h has three triangular ports.*

Proof. If the first alternative fails, then Lemma 3.12 gives the second alternative. $\square$

Corollary 3.14 (Degree-four local activation). *Let h be a high center with $d_G(h) = 4$. Then one of the following alternatives holds.*

- (i) *There are two open ports at h that are fan-compatible.*
- (ii) *At least two ports at h are triangular.*

In case (i), any admissible Type B fan-window profile that records the two ports and assigns their four non- h incidences routes through Corollary 14.81. In case (ii), any admissible Type B fan-window profile that records two triangular ports and assigns their four triangle incidences routes through Proposition 14.80; in the certificate-marked B1 survivor it has

$$D_B(\mathfrak{F}_h) \geq 2 - \frac{11 - 4}{4} = \frac{1}{4} > 0.$$

Proof. Let $U \subseteq N_G(h)$ be the set of vertices x such that the port (h, x) is open. If two vertices $x, y \in U$ are nonadjacent, then Lemma 3.11 gives a fan-compatible open pair, which is alternative (i). Suppose no such pair exists. Then every two vertices of U are adjacent. By Lemma 3.9, the graph induced by $N_G(h)$ is a matching, so U has size at most 2. Since $d_G(h) = 4$, at least $4 - |U| \geq 2$ ports at h are not open. By Definition 3.7, those ports are triangular. This is alternative (ii).

The routing assertion in case (i) is exactly [Corollary 14.81](#), applied with $k = 4$. In case (ii), choose two triangular ports. For each of them the two triangle edges are precisely the two non- h incidences of the port vertex. Once those incidences are assigned to the fan envelope, both ports are fan-closed by [Definition 14.69](#). Applying [Proposition 14.80](#) with $r = 2$ and $k = 4$ gives the displayed deficit bound and the listed Type B alternatives. $\square$

Definition 3.15 (Triangular fan core and shoulder completions). *Let h be a heavy center and let*

$$\mathcal{T}_h = \{p_i = (h, x_i) : i \in I\}$$

be a nonempty set of triangular ports at h . For each $i \in I$, write

$$N_G(x_i) = \{h, a_i, b_i\}, \quad S_i = \{a_i, b_i\},$$

so that $x_i a_i b_i x_i$ is a triangle. The triangular fan core generated by $\mathcal{T}_h$ is

$$F_h(\mathcal{T}_h) = G\left[\{h\} \cup \{x_i : i \in I\} \cup \bigcup_{i \in I} S_i\right].$$

The vertices in S_i are the shoulders of p_i .

Let $s \in S_i$, and let s'_i denote the other vertex of S_i . A shoulder-completion incidence at s is an edge $sy \in E(G)$ with

$$y \notin \{x_i, s'_i\}.$$

It is central when $y = h$, cross-triangular when

$$y \in \bigcup_{j \in I, j \neq i} S_j,$$

and outside when

$$y \notin V(F_h(\mathcal{T}_h)) \cup N_G(h).$$

An outside incidence is the first edge by which the displayed triangular fan core attaches to the rest of the graph.

Lemma 3.16 (Shoulder-completion bookkeeping). *Let $p_i = (h, x_i)$ be a triangular port at a heavy center h , and let $S_i = \{a_i, b_i\}$. Then the following hold.*

- (a) *Each shoulder $s \in S_i$ has at least one shoulder-completion incidence.*
- (b) *At most one shoulder in S_i lies in $N_G(h)$.*
- (c) *If $s \in S_i \cap N_G(h)$, then $d_G(s) = 3$, its unique shoulder-completion incidence is the central edge sh , and s has no outside incidence.*
- (d) *No shoulder-completion incidence has endpoint in $N_G(h) \setminus \{h\}$.*

Proof. Fix $s \in S_i$, and let s'_i be the other shoulder. The two edges sx_i and ss'_i are present. Since $\delta(G) \geq 3$, s has at least one incident edge different from these two; this proves (a).

If both a_i and b_i lay in $N_G(h)$, then the vertex $x_i \in N_G(h)$ would be adjacent inside $G[N_G(h)]$ to both a_i and b_i . This contradicts the matching conclusion in [Lemma 3.9](#). Hence (b) holds.

Assume $s \in N_G(h)$. By [Lemma 3.9](#), $d_G(s) = 3$. Its three neighbours are already h , x_i , and s'_i . Thus the only shoulder-completion incidence at s is the central edge sh , proving (c).

Finally suppose that sy is a shoulder-completion incidence with $y \in N_G(h) \setminus \{h\}$. If y were adjacent to x_i , then $y \in S_i$, because the only neighbours of x_i are h, a_i, b_i ; this is excluded by the definition of shoulder-completion incidence. Hence x_i and y are nonadjacent vertices of $N_G(h)$. They have the common neighbour $s \neq h$, contradicting Lemma 3.9. This proves (d).

For the finite audit, the sites are the two shoulders of each actual triangular port at the requested center, in the declared port and shoulder order. At one site the witness scan ranges only over the declared vertex order and stops at the first endpoint satisfying the completion-incidence predicate. Thus there are at most $2|V(G)|$ sites and at most $2|V(G)|^2 + 2|V(G)| + 2$ primitive checks. The audit constructs no path, vertex subset, port pair, or auxiliary graph. $\square$

Lemma 3.17 (Port returns enter through shoulders). *Let $p_i = (h, x_i)$ be a triangular port at a heavy center h . Then there is a shoulder $s \in S_i$ and a simple s - h path Q in $G - x_i$. Moreover, the cycle obtained from Q by adding the two edges*

$$hx_i, \quad x_is$$

has length $|E(Q)| + 2$, and therefore

$$|E(Q)| \notin \{2^j - 2 : j \geq 2\}.$$

Exactly one of the following initial alternatives holds:

- (i) Q is the single central edge sh ;
- (ii) the first edge of Q is a noncentral shoulder-completion incidence at s ;
- (iii) Q begins with the shoulder edge ss'_i , followed by the central edge s'_ih , and has length 2;
- (iv) Q begins with ss'_i , followed by a noncentral shoulder-completion incidence at s'_i .

Proof. By Lemma 3.4, the edge hx_i lies on a cycle C . Deleting hx_i from C gives a simple x_i - h path in $G - hx_i$. The first edge of this path is x_is for some $s \in S_i$. Removing that first edge leaves a simple s - h path Q in $G - x_i$.

The edges of Q , together with x_is and hx_i , reconstruct the cycle C . Since G has no power-of-two cycle, $|E(Q)| + 2$ is not a power of two, which gives the displayed exclusion.

If the first edge of Q is sh , simplicity forces $Q = sh$, giving (i). Otherwise, if the first edge is not ss'_i , its endpoint is neither h , x_i , nor s'_i , so it is a noncentral shoulder-completion incidence at s , giving (ii). It remains to consider a first edge ss'_i . The next edge exists because $s'_i \neq h$. It cannot return to s , because Q is simple, and it cannot go to x_i , because $Q \subseteq G - x_i$. If it is s'_ih , simplicity again forces the path to stop, giving (iii). Every other possibility is a noncentral shoulder-completion incidence at s'_i , giving (iv).

The formal finite audit uses the non-bridge reachability theorem to obtain one simple path in $G - hx_i$ by proof-level choice. CT1 checks the resulting proof-carrying return certificate once. The construction enumerates no walks, paths, cycles, vertex subsets, subgraphs, or ambient graphs. $\square$

Lemma 3.18 (First landing exhaustion for triangular shoulders). *Let $\mathcal{T}_h$ be a set of triangular ports at a heavy center h . Every shoulder-completion incidence in the triangular fan core $F_h(\mathcal{T}_h)$ is exactly one of the following:*

- (i) a central incidence to h ;
- (ii) a cross-triangular incidence to a shoulder of another port in $\mathcal{T}_h$;

(iii) an outside incidence with endpoint in

$$V(G) \setminus (V(F_h(\mathcal{T}_h)) \cup N_G(h)).$$

No completion incidence lands at another neighbour of h , and no completion incidence lands at another port vertex x_j .

Proof. Let sy be a shoulder-completion incidence at a shoulder $s \in S_i$. If $y = h$, it is central. If $y \in N_G(h) \setminus \{h\}$, then [Lemma 3.16](#) forbids the incidence; in particular this excludes $y = x_j$ for every port vertex x_j , because each $x_j \in N_G(h)$.

It remains that $y \notin N_G(h)$. If $y \in V(F_h(\mathcal{T}_h))$, then y cannot be h , cannot be a port vertex x_j , and cannot be x_i or the other shoulder of p_i by the definition of completion incidence. Thus y is a shoulder of some other triangular port, and the incidence is cross-triangular. If $y \notin V(F_h(\mathcal{T}_h))$, then it is an outside incidence by definition. These cases are disjoint and exhaustive. $\square$

Lemma 3.19 (Cross-shoulder multiplicity). *Let p_i and p_j be distinct triangular ports at the same heavy center h . If there are two distinct cross-triangular incidences between S_i and S_j , then either G contains a 4-cycle or at least one shoulder in $S_i \cup S_j$ has ambient degree at least 4. Consequently, after routing high shoulders to the Type B surplus branch, the cross-triangular incidences between any two shoulder pairs form a matching of size at most one.*

Proof. Let the two cross-triangular edges be e_1 and e_2 . If they share a shoulder endpoint s , then s is incident with the two triangle edges inside its own port and with both e_1 and e_2 . Hence $d_G(s) \geq 4$.

Assume therefore that e_1 and e_2 are vertex-disjoint. Write $S_i = \{a_i, b_i\}$ and $S_j = \{a_j, b_j\}$. After relabelling the two shoulders inside each pair, the cross edges are $a_i a_j$ and $b_i b_j$. Then

$$a_i a_j b_j b_i a_i$$

is a simple cycle of length 4, using the cross edge $a_i a_j$, the triangular shoulder edge $a_j b_j$, the cross edge $b_j b_i$, and the triangular shoulder edge $b_i a_i$. This contradicts dyadic-safety unless the high-shoulder alternative already occurred in the shared-endpoint case.

If all shoulders have ambient degree 3, the shared-endpoint case is impossible and the vertex-disjoint case gives a forbidden 4-cycle. Thus at most one cross-triangular edge can join S_i to S_j , and the remaining incidences form a matching. $\square$

Remark 3.20 (Degree-greater-than-four residual leaf). [Corollary 3.13](#) isolates the degree-greater-than-four branch into two local configurations. In Leaf A, a fan-compatible open pair is recorded as two fan-closed ports by [Lemma 14.70](#); [Corollary 14.81](#) then gives a positive-deficit Type B fan whenever the fan is certificate-marked and survives the direct target, compression, and delocalization alternatives. If one of those direct alternatives occurs, the branch is already closed; if the fan is unmarked or the local entries cannot be globally disjoint, the branch is recorded as a fan-certificate residual center or a B2-overlap obstruction. In Leaf B, the triangular branch is routed by [Definition 3.15](#) and [Lemmas 3.16](#), [3.18](#) and [3.19](#) and by the Type B fan ledger in [Corollary 14.82](#): the guaranteed triangular ports are cubic-closed fan neighbours, their shoulder-completion incidences have only the central, cross-triangular, and outside first landings listed above, and the cross-triangular multiplicity left after high-shoulder routing is a matching survivor. The common primitive in both leaves is fan-closedness: two assigned non- h incidences at a cubic neighbour of h contribute one cubic-closed fan neighbour. Thus any heavy center of degree $k \geq 5$ supplies at least two fan-closed ports after the local dichotomy, and [Proposition 14.80](#) gives

$$D_B \geq 2 - \frac{11 - k}{4} = \frac{k - 3}{4} > 0.$$

Thus the active degree-greater-than-four closure is the Type B fan-closed ledger.

Remark 3.21 (Interface with the surplus-token branch). *Corollary 3.13* separates the degree-greater-than-four branch from the degree-4 branch. The degree-greater-than-four branch is routed through fan-closed Type B data by *Corollaries 14.81* and *14.82*. This local dichotomy is one input to the sparse surplus-token branch, not a replacement for it. Excess ports that do not close through the direct target, compression, delocalization, or fan-closed Type B alternatives are counted by the active surplus family $\mathcal{P}_{\text{exc}}$. The non-near-cubic branch then applies the free/blocked pair split and the capacity-token ledger of *Definition 3.41*. Consequently any attempted survivor above the near-cubic range first passes through the total pair route of *Definition 3.42*: blocked pairs retain their canonical blocker token and blocker role, while free pairs use their first selected port with role `freeAnchor`. The resulting three total-role-homogeneous same-token audit classes are window-incidence, remainder-surplus, and primitive-carrier load. The same-token bottleneck routing of *Lemma 3.84* then either realizes a sparse surplus exit, produces ordinary Type B fan data, or leaves fixed homogeneous caps and hence the near-cubic estimate.

Definition 3.22 (Open-port suppression). Let $p = (h, x)$ be an open surplus port, so

$$N_G(x) = \{h, a_p, b_p\}, \quad a_p b_p \notin E(G).$$

The suppression of p is the graph

$$G/p := G - x + a_p b_p.$$

A finite family $\mathcal{Q}$ of open surplus ports is suppressible if the following four conditions hold.

(a) The non-center supports

$$T(p) = \{x(p), a_p, b_p\} \quad (p \in \mathcal{Q})$$

are pairwise disjoint.

(b) No selected center is contained in any non-center support:

$$\{c(p) : p \in \mathcal{Q}\} \cap \bigcup_{p \in \mathcal{Q}} T(p) = \emptyset.$$

(c) The shoulder chords $e_p^* = a_p b_p$, $p \in \mathcal{Q}$, are pairwise distinct and none is an edge of G .

(d) For every high center h ,

$$|\{p \in \mathcal{Q} : c(p) = h\}| \leq d_G(h) - 3.$$

For a suppressible family $\mathcal{Q}$, define

$$G/\mathcal{Q} = G - \{x(p) : p \in \mathcal{Q}\} + \{a_p b_p : p \in \mathcal{Q}\}.$$

Thus $G/\{p\} = G/p$.

Lemma 3.23 (Suppression preserves simplicity and minimum degree). *If $\mathcal{Q}$ is a suppressible family of open surplus ports, then $G/\mathcal{Q}$ is finite and simple, and every vertex of $G/\mathcal{Q}$ has degree at least 3.*

Proof. Only the vertices $x(p)$, $p \in \mathcal{Q}$, are deleted, so the resulting graph is finite. No added edge is a loop, because $a_p \neq b_p$. No added edge duplicates an old edge, and no two added edges duplicate each other, by the definition of a suppressible family. Hence $G/\mathcal{Q}$ is simple.

Let v be a vertex that remains after the suppression. If $v = c(p)$ for some selected ports, then the only deleted neighbours of v are the port vertices $x(p)$ with center v : condition (b) prevents v from lying in any other selected non-center support. Thus

$$d_{G/\mathcal{Q}}(v) = d_G(v) - |\{p \in \mathcal{Q} : c(p) = v\}| \geq d_G(v) - (d_G(v) - 3) = 3$$

by condition (d).

If v is a shoulder endpoint for a selected port p , then condition (a) makes this selected port unique. The vertex v loses the edge $vx(p)$ and gains exactly the shoulder chord e_p^* . It is not deleted and is not a selected center by condition (b). Hence its degree is unchanged. Every other remaining vertex loses no incident edge and gains no incident edge, or only gains an added chord. Since $\delta(G) \geq 3$, all remaining vertices have degree at least 3. $\square$

Lemma 3.24 (Single-port suppression witness). *Let G be the lexicographically minimal counterexample, and let $p = (h, x)$ be an open surplus port. Then $G - x$ contains a simple a_p - b_p path of length $2^j - 1$ for some $j \geq 2$.*

Proof. The singleton $\{p\}$ is suppressible. By Lemma 3.23, G/p is a finite simple graph with minimum degree at least 3. It has one fewer vertex than G . By the minimality of G , the graph G/p contains a cycle C of length 2^j for some $j \geq 2$.

The cycle C must use the added shoulder chord $a_p b_p$. Otherwise all edges of C are old edges of G not incident with the deleted vertex x , and C is already a power-of-two cycle in G , impossible. Deleting the chord $a_p b_p$ from C leaves a simple a_p - b_p path in $G - x$, and its length is $2^j - 1$. $\square$

Lemma 3.25 (Suppressed-family critical cycle). *Let G be the lexicographically minimal counterexample, and let $\emptyset \neq \mathcal{Q}$ be a suppressible family of open surplus ports. Then $G/\mathcal{Q}$ contains a power-of-two cycle using at least one added shoulder chord. If a power-of-two cycle C in $G/\mathcal{Q}$ has length 2^j and uses exactly the added chords indexed by $\mathcal{S} \subseteq \mathcal{Q}$, then $\mathcal{S} \neq \emptyset$, and replacing each added chord $a_p b_p$ with the two-edge path*

$$a_p x(p) b_p$$

gives a simple cycle of G of length $2^j + |\mathcal{S}|$. Consequently

$$2^j + |\mathcal{S}|$$

is not a power of two.

Proof. By Lemma 3.23, $G/\mathcal{Q}$ is finite, simple, and has minimum degree at least 3. Since $\mathcal{Q} \neq \emptyset$, it has fewer vertices than G . Minimality gives a cycle C in $G/\mathcal{Q}$ whose length is a power of two. If C used no added shoulder chord, then C would be a power-of-two cycle already present in G , a contradiction. Hence the set $\mathcal{S}$ of added chords used by C is nonempty.

For each $p \in \mathcal{S}$, replace the edge $a_p b_p$ on C by the path $a_p x(p) b_p$. The supports $T(p)$ are pairwise disjoint, and none of the vertices $x(p)$ is present in $G/\mathcal{Q}$. Therefore these replacements are internally disjoint and introduce no repeated vertex. The resulting closed walk is a simple cycle in G . Each replacement adds exactly one edge, so the new length is $2^j + |\mathcal{S}|$. Since G has no power-of-two cycle, this integer is not a power of two. $\square$

Lemma 3.26 (Sparse slack identity). *For the standing parameters*

$$\lambda = 2n - 3 - m, \quad \sigma(G) = 2m - 3n,$$

one has

$$\sigma(G) = n - 6 - 2\lambda \quad \text{and} \quad m = \frac{3}{2}n + \frac{1}{2}\sigma(G).$$

Thus every additional unit of high-degree surplus consumes exactly one half-unit of the sparse slack λ .

Proof. Substituting $m = 2n - 3 - \lambda$ gives

$$2m - 3n = 2(2n - 3 - \lambda) - 3n = n - 6 - 2\lambda.$$

The second identity is the same equality solved for m . □

Definition 3.27 (Sparse surplus exits). *A sparse surplus exit is one of the following conclusions, each of which is defined without invoking the near-cubic Type A/Type B branch machinery:*

- (a) *a direct dyadic contradiction, equivalently a power-of-two cycle or a Mersenne return in the sense of Lemma 2.2;*
- (b) *a target-defective quotient, as defined by Lemma 3.102;*
- (c) *a nontrivial target-complete compression of a proper atom, forbidden by Lemma 3.103 and Corollary 3.105;*
- (d) *a proper or global delocalization coordinate governed by Lemmas 10.6 and 10.9;*
- (e) *an open-port suppression cycle whose chord set violates the arithmetic conclusion of Lemma 3.25.*

A graph survives the sparse surplus exits when none of these conclusions occurs for any selected surplus demand, any selected pair of surplus demands, or any baseline spine coordinate used in the entropy sandwich of Definition 3.69. The label name `def:named-surplus-exits` is retained for references in the dependency diagram; the exits in this branch are precisely the sparse exits listed above.

Lemma 3.28 (Excess-port extraction). *Let G be the lexicographically minimal counterexample, and let $\mathcal{P}_{\text{exc}}$ be the excess selector of Definition 3.6. Then*

$$|\mathcal{P}_{\text{exc}}| = \sigma(G).$$

Moreover, for every $p = (h, x) \in \mathcal{P}_{\text{exc}}$, the vertex h has degree at least 4, the vertex x has degree 3, and

$$N_G(x) = \{h, a_p, b_p\}$$

for two distinct vertices $a_p, b_p \neq h$.

Proof. By definition, each high-degree vertex h contributes exactly $d_G(h) - 3$ selected incident ports to $\mathcal{P}_{\text{exc}}$. The high-degree vertices are precisely the vertices contributing positive terms to $\sigma(G) = \sum_v (d_G(v) - 3)$, because every vertex of G has degree at least 3. Hence

$$|\mathcal{P}_{\text{exc}}| = \sum_{h \in V_{\geq 4}(G)} (d_G(h) - 3) = \sigma(G).$$

If $p = (h, x)$ is a selected port, then $h \in V_{\geq 4}(G)$. By Lemma 3.2, every edge has at least one degree-3 endpoint and $V_{\geq 4}(G)$ is independent. Since h is not degree 3, the other endpoint x must have degree 3. Thus x has exactly two neighbours different from h , denoted a_p and b_p . □

Remark 3.29 (Ordered CT6 audit). *The finite audit underlying the extraction is performed in the declared vertex order. At a vertex h , the monitored failure is*

$$d_G(h) \geq 4 \quad \text{and} \quad \exists x \in N_G(h) \text{ with } d_G(x) \neq 3.$$

The first-failure branch retains the first such centre. If no failure is found, the active ledger assigns contribution $d_G(h) - 3$ to h , so its total is exactly

$$\sum_{h \in V(G)} (d_G(h) - 3) = \sigma(G).$$

The dependent slot family $\{(h, i) : 0 \leq i < d_G(h) - 3\}$ has the same cardinality. Deletion criticality excludes the failure branch: on an edge hx with $d_G(h) \geq 4$, the degree-three endpoint must be x . Thus this audit uses one bounded neighbour scan per vertex, hence at most quadratic primitive adjacency/degree tests, and does not enumerate paths, subgraphs, or graphs. The return and suppression paths in [Lemma 3.30](#) are supplied separately by bridgelessness and minimality; the baseline demand of node [129] is fixed only after this activation stage.

Lemma 3.30 (Sparse port activation). *For every selected port $p = (h, x) \in \mathcal{P}_{\text{exc}}$, fix a proof-carrying finite response support $\Gamma(p)$ with the following properties.*

- (a) *The port support is $T(p) = \{x, a_p, b_p\}$.*
- (b) *There is a simple x - h path $R_p \subseteq G - hx$. Its first edge after x is either xa_p or xb_p .*
- (c) *If p is open, let $Q_p \subseteq G - x$ be one fixed path supplied by [Lemma 3.24](#); this path joins a_p to b_p and has length $2^{j(p)} - 1$ for some $j(p) \geq 2$.*
- (d) *The response support is the exact edge-labelled subgraph*

$$E(\Gamma(p)) = \begin{cases} E(R_p) \cup E(Q_p) \cup \{xa_p, xb_p\}, & p \text{ open,} \\ E(R_p) \cup \{xa_p, xb_p, a_pb_p\}, & p \text{ triangular,} \end{cases}$$

with vertex set consisting exactly of the endpoints of these edges. Thus in the triangular case the second summand is exactly the response triangle xa_pb_px , not an unspecified supergraph of it.

The witnesses R_p, Q_p are fixed once for each selected port. Every later local scan consumes these declared certificates and checks their edge lists; it does not search a path family.

Proof. By [Lemma 3.4](#), the edge hx is not a bridge. Therefore h and x remain connected in $G - hx$, and there is a simple x - h path in $G - hx$. Fix one such path and call it R_p . Since $d_G(x) = 3$ and $N_G(x) = \{h, a_p, b_p\}$, the first edge of R_p after x is one of xa_p, xb_p .

If p is open, [Lemma 3.24](#) supplies a simple a_p - b_p path $Q_p \subseteq G - x$ of length $2^{j(p)} - 1$; fix one such certificate. If p is triangular, then $a_pb_p \in E(G)$, so xa_pb_px is a triangle. In the open case take the displayed union of R_p, Q_p , and the two shoulder incidences; in the triangular case take the displayed union of R_p and the three triangle edges. Both unions are connected and finite, contain $T(p)$, and their edge and vertex sets are fixed by the display. Once the proof-carrying paths are fixed, the display determines $\Gamma(p)$ exactly. $\square$

Definition 3.31 (Active surplus demands). *Let G be a lexicographically minimal counterexample reaching node [18]. An active surplus demand is a selected surplus port*

$$p = (h, x) \in \mathcal{P}_{\text{exc}}$$

equipped with the fixed declared data $T(p)$, R_p , and $\Gamma(p)$ from Lemma 3.30, and not already removed by a sparse surplus exit of Definition 3.27. The response support $\Gamma(p)$ is used only as a carrier of target-response data; no disjointness between distinct response supports is assumed unless it is explicitly stated.

Lemma 3.32 (Surviving active family is the full excess-port family). *If G survives the sparse surplus exits, then*

$$\mathcal{A}_0 := \mathcal{P}_{\text{exc}}$$

is a finite family of active surplus demands and

$$|\mathcal{A}_0| = \sigma(G).$$

Proof. By Lemma 3.28, every selected port $p \in \mathcal{P}_{\text{exc}}$ has the canonical data $T(p)$, R_p , and $\Gamma(p)$ supplied by Lemma 3.30. Since G survives the sparse surplus exits, no selected surplus demand is removed by an exit of Definition 3.27. Thus every $p \in \mathcal{P}_{\text{exc}}$ is active in the sense of Definition 3.31. The cardinality is exactly $|\mathcal{P}_{\text{exc}}| = \sigma(G)$, again by Lemma 3.28. $\square$

Definition 3.33 (Canonical local connector support). *For two nonempty connected finite edge-labelled subgraphs $A, B \subseteq G$, choose a pair $(a, b) \in V(A) \times V(B)$ of minimum ambient distance and retain a proof-carrying shortest path $P_G(A, B)$ from a to b . Fix one such certificate for each ordered pair of supports. Define*

$$J_G(A, B) := A \cup P_G(A, B) \cup B.$$

If $A \cap B \neq \emptyset$, this path has length zero at the least common vertex. For an unordered active pair $\pi = \{p, q\}$, first order $p < q$ in the selected-port order and put

$$X_\pi := J_G(\Gamma(p), \Gamma(q)).$$

Thus X_π is a specified connected support, not a minimum Steiner subgraph. A producer may compute the certificate by one multi-source BFS and one predecessor trace. The verifier checks the retained path and the minimum-distance endpoint condition; it does not enumerate connected subgraphs or paths.

Definition 3.34 (Sparse surplus blockers). *Let $\mathcal{A}_0$ be a finite family of active surplus demands. A sparse surplus blocker for an unordered pair $\pi = \{p, q\} \subseteq \mathcal{A}_0$ is a tagged value $B = (\tau, b)$, where*

$$\tau \in \text{BlkTag} := \{\text{overlap}, \text{return}, \text{shoulder}, \text{profile}, \text{target}, \text{chord}\},$$

and b is one of the following concrete finite values of the indicated tag.

- (a) *For $\tau = \text{overlap}$, b is a vertex or edge-incidence contained in both exact demand supports*

$$\Gamma(p) \cap \Gamma(q).$$

- (b) *For $\tau = \text{return}$, b is a common vertex or common edge-incidence of the two canonical return paths R_p and R_q .*

(c) For $\tau = \text{shoulder}$, b is a vertex in

$$\{a_p, b_p, x(p)\} \cap \{a_q, b_q, x(q)\},$$

together with its two role labels in $\{\text{leftShoulder}, \text{rightShoulder}, \text{buffer}\}$.

(d) For $\tau = \text{profile}$, $b = (q, s, t, \xi, v_s, v_t)$, where q is a normalized quotient proposal on the finite declared signature of X_π , $q(s) = q(t)$, and the boundary-degree coordinate ξ has unequal values $v_s \neq v_t$. This is the finite audit certificate that the proposal does not preserve boundary fibres, in the sense of [Lemma 3.101](#).

(e) For $\tau = \text{target}$, $b = (q, s, t, K, \xi, v_s, v_t)$, where $q(s) = q(t)$, K is a supplied boundaried outside context, and the exact target-response coordinate ξ has unequal values $v_s \neq v_t$ after gluing K . This is the finite audit certificate that the proposal is not target-complete. A target-complete non-injective proposal is not represented by this tag: to be admissible it must instead carry the strictly smaller representative required by [Definition 3.98](#).

(f) For $\tau = \text{chord}$, $b = (\mathcal{Q}, C, \mathcal{S})$, where $\emptyset \neq \mathcal{Q} \subseteq \{p, q\}$ is suppressible, C is a proof-carrying simple power-of-two cycle in $G/\mathcal{Q}$, and $\emptyset \neq \mathcal{S} \subseteq \mathcal{Q}$ is exactly the set of added shoulder chords used by C , as in [Lemma 3.25](#).

Every audit or structural blocker record must be an object explicitly listed above. A pair is not declared blocked merely because the proof has not found a closure. Types (d) and (e) route the raw proposal through sparse exit (b) before admission; the admitted blocked ledger contains exactly the finite structural types (a)–(c) and (f).

Lemma 3.35 (Sparse blocker values are locally finite). *For every active pair π , the admitted structural carrier $\text{Blk}_{\text{adm}}(\pi)$, consisting of types (a)–(c) and the retained suppression certificates of type (f), is finite. Types (d) and (e) are audit outputs relative to one supplied raw quotient proposal; they are routed before the admitted pair ledger is formed and are empty for an admissible quotient. Membership is checked from X_π , the two exact response supports, and the supplied certificates; checking does not enumerate ambient subgraphs, paths, quotients, or outside contexts.*

Proof. The carriers in (a)–(c) are subsets of finite supports. For (f), activation retains at most one proof-carrying critical suppression cycle for each of p and q , so the pair scan adds at most two chord tags. The verifier checks simplicity, length, and chord use by one scan of each retained edge list. For a raw quotient proposal, a boundary mismatch or one supplied distinguishing context is a finite record of type (d) or (e). Boundary preservation and context-universal target completeness prove directly that no such record exists after the proposal has been admitted. $\square$

Definition 3.36 (Canonical sparse-blocker order). *Fix the vertex order used in the local schedules above. Admitted sparse surplus blockers are ordered as follows. First order by the type number (a)–(c), (f) in [Definition 3.34](#). Within a type, order vertices by the fixed vertex order, edge-incidences lexicographically by their ordered endpoint pair and incident end, and chord tuples by the lexicographic order of $(\mathcal{Q}, E(C), \mathcal{S})$. This gives a total order on every finite set of blockers arising from a fixed pair of active surplus demands.*

Definition 3.37 (Canonical blocker ledger). *Let $\mathcal{A}_0$ be a finite family of active surplus demands and put*

$$\Pi(\mathcal{A}_0) = \binom{\mathcal{A}_0}{2}.$$

For $\pi = \{p, q\} \in \Pi(\mathcal{A}_0)$, let $\text{Blk}_{\text{adm}}(\pi)$ be the finite admitted structural blocker set of [Lemma 3.35](#). Raw proposals that produce a type-(d) or type-(e) audit record have already left through the corresponding exit and do not enter this ledger. Define

$$\Pi_{\text{blk}} := \{\pi \in \Pi(\mathcal{A}_0) : \text{Blk}_{\text{adm}}(\pi) \neq \emptyset\}, \quad \Pi_{\text{free}} := \Pi(\mathcal{A}_0) \setminus \Pi_{\text{blk}}.$$

For every $\pi \in \Pi_{\text{blk}}$, the canonical blocker is

$$\Phi_{\text{can}}(\pi) := \min_{\prec} \text{Blk}_{\text{adm}}(\pi),$$

where $\prec$ is the order of [Definition 3.36](#). The canonical blocker set and incidence ledger are

$$\mathcal{B}_{\text{can}} := \Phi_{\text{can}}(\Pi_{\text{blk}}), \quad \mathcal{J}_{\text{can}} := \{(\pi, \Phi_{\text{can}}(\pi)) : \pi \in \Pi_{\text{blk}}\}.$$

For $B \in \mathcal{B}_{\text{can}}$, write

$$\mu(B) := |\Phi_{\text{can}}^{-1}(B)|.$$

Lemma 3.38 (Canonical blocker ledger has no overcounting). *For every finite family $\mathcal{A}_0$ of active surplus demands,*

$$|\Pi_{\text{blk}}| = |\mathcal{J}_{\text{can}}| = \sum_{B \in \mathcal{B}_{\text{can}}} \mu(B).$$

Consequently, if $\mu(B) \leq M$ for every $B \in \mathcal{B}_{\text{can}}$ and $|\mathcal{B}_{\text{can}}| \leq C_B n$, then

$$|\Pi_{\text{blk}}| \leq M C_B n.$$

Proof. The map

$$\pi \longmapsto (\pi, \Phi_{\text{can}}(\pi))$$

is a bijection from Π_{blk} to $\mathcal{J}_{\text{can}}$, so $|\Pi_{\text{blk}}| = |\mathcal{J}_{\text{can}}|$. The fibres $\Phi_{\text{can}}^{-1}(B)$, $B \in \mathcal{B}_{\text{can}}$, form a disjoint partition of Π_{blk} . Hence

$$|\Pi_{\text{blk}}| = \sum_{B \in \mathcal{B}_{\text{can}}} |\Phi_{\text{can}}^{-1}(B)| = \sum_{B \in \mathcal{B}_{\text{can}}} \mu(B).$$

The displayed upper bound follows by applying $\mu(B) \leq M$ to each summand and then using $|\mathcal{B}_{\text{can}}| \leq C_B n$. $\square$

Definition 3.39 (Primitive carrier of a sparse blocker). *An edge-incidence of G is an ordered pair (e, v) with $e \in E(G)$ and $v \in e$. The primitive sparse carrier universe is*

$$\mathfrak{U}_{\text{sp}}(G) := V(G) \sqcup I_E(G) \sqcup \mathcal{P}_{\text{exc}}, \quad I_E(G) := \{(e, v) : e \in E(G), v \in e\}.$$

Each sparse surplus blocker B has a primitive carrier

$$\kappa(B) \in \mathfrak{U}_{\text{sp}}(G)$$

defined as follows.

- (a) If $B = (\text{overlap}, b)$, where b is its vertex or edge-incidence value, then $\kappa(B) = b$.
- (b) If $B = (\text{return}, b)$, where b is its vertex or edge-incidence value, then $\kappa(B) = b$.
- (c) If $B = (\text{shoulder}, (v, \ell_p, \ell_q))$, then $\kappa(B) = v$; the finite role labels remain part of the blocker value.

- (d) If $B = (\text{profile}, (\xi, v))$ or $B = (\text{target}, (\xi, v))$, then $\kappa(B)$ is the least vertex, edge-incidence, or selected surplus port contained in the declared support of ξ . The finite value v is retained in the blocker value but does not alter its carrier. Paths and windows contribute their least contained vertex or edge-incidence. This carrier exists because every sparse pair-response coordinate is supported on $X_\pi \supseteq \Gamma(p) \cup \Gamma(q)$, and the declared signature of [Definition 3.94](#) generates quotient coordinates only from finite unions of declared supports.
- (e) If $B = (\text{chord}, (\mathcal{Q}, C, \mathcal{S}))$, then $\kappa(B)$ is the least selected surplus port whose shoulder chord belongs to $\mathcal{S}$. The set $\mathcal{S}$ is nonempty by [Lemma 3.25](#).

All minima use the fixed vertex order, the induced lexicographic order on edge-incidences, and the selected-port order used in [Definition 3.36](#).

Lemma 3.40 (Primitive carrier supply is linear). *For every lexicographically minimal counterexample reaching the sparse-surplus branch,*

$$|\mathfrak{U}_{\text{sp}}(G)| \leq 6n.$$

Proof.

$$|\mathfrak{U}_{\text{sp}}(G)| = |V(G)| + |I_E(G)| + |\mathcal{P}_{\text{exc}}| = n + 2m + \sigma(G).$$

By [Lemma 3.3](#), $m \leq 2n - 2$, so $2m \leq 4n - 4$. Also

$$\sigma(G) = 2m - 3n \leq n - 4 \leq n,$$

again using $m \leq 2n - 2$. Hence

$$|\mathfrak{U}_{\text{sp}}(G)| \leq n + (4n - 4) + n \leq 6n.$$

□

Definition 3.41 (Capacity tokens for blocked surplus pairs). *The capacity-token universe is the disjoint union*

$$\mathfrak{T}_{\text{cap}} = \mathfrak{T}_{\text{prim}} \sqcup \mathfrak{T}_R \sqcup \mathfrak{T}_W,$$

defined as follows.

$$\mathfrak{T}_{\text{prim}} := \mathfrak{U}_{\text{sp}}(G).$$

The remainder-surplus tokens are

$$\mathfrak{T}_R := \{(v, j) : v \in R, 1 \leq j \leq d_G(v) - 3\}.$$

Thus $|\mathfrak{T}_R| = \sigma_R$. The window-join tokens are

$$\mathfrak{T}_W := \{(e, RW) : e \in E(R, W)\} \sqcup \{(e, a), (e, b) : e = ab, a, b \text{ lie in distinct packed } P_{13} \text{ windows}\}.$$

Thus each window-remainder edge contributes one token and each cross-window edge contributes one token at each of its two window ends.

For a sparse surplus blocker $B = (\tau, b)$, let $\text{supp}(B)$ be its declared support: the singleton vertex or marked edge for an overlap or return value, the singleton carrier vertex for a shoulder value, the declared support of ξ for a profile or target value (ξ, v) , and the union of the shoulder chords indexed by $\mathcal{S}$ for a chord value $(\mathcal{Q}, C, \mathcal{S})$. Order the unordered pairs in Π_{blk} lexicographically by their two active surplus demands, using the selected-port order induced by the fixed vertex order. This is the canonical pair order. For $\pi \in \Pi_{\text{blk}}$, put $B_\pi = \Phi_{\text{can}}(\pi)$. The capacity token $\Theta_{\text{cap}}(\pi)$ is defined by the following canonical rule.

- (a) If $\text{supp}(B_\pi)$ contains a window-remainder edge, choose the least such edge and take the corresponding token $(e, RW) \in \mathfrak{T}_W$.
- (b) Otherwise, if $\text{supp}(B_\pi)$ contains a cross-window edge, choose the least such edge $e = ab$ and then the least of its two endpoint tokens in $\mathfrak{T}_W$.
- (c) Otherwise, if $\text{supp}(B_\pi)$ contains a vertex $v \in R$ with $d_G(v) > 3$, choose the least such vertex $v = v(\pi)$. Let $\mathcal{P}_v$ be the set of all $\pi' \in \Pi_{\text{blk}}$ for which the first two cases fail and the least high-degree remainder vertex in $\text{supp}(B_{\pi'})$ is v . Order $\mathcal{P}_v$ by the canonical pair order and let

$$\text{rk}_v(\pi) \in \{0, \dots, |\mathcal{P}_v| - 1\}$$

be the rank of π in this order. Set

$$j(\pi) = 1 + (\text{rk}_v(\pi) \bmod (d_G(v) - 3))$$

and take the token $(v, j(\pi)) \in \mathfrak{T}_R$.

- (d) If none of the preceding cases applies, set

$$\Theta_{\text{cap}}(\pi) := \kappa(B_\pi) \in \mathfrak{T}_{\text{prim}}.$$

In cases (a)–(c), $\Theta_{\text{cap}}(\pi)$ is the token selected there. For $t \in \mathfrak{T}_{\text{cap}}$, its token load is

$$\ell_{\text{cap}}(t) := |\Theta_{\text{cap}}^{-1}(t)|.$$

Definition 3.42 (Total surplus-pair token route). *Orient every pair in $\binom{\mathcal{A}_0}{2}$ by the canonical selected-port order, writing it as $\pi = (p_i, p_j)$ with $i < j$. Define the total token map*

$$\hat{\Theta}(\pi) = \begin{cases} \Theta_{\text{cap}}(\pi), & \pi \in \Pi_{\text{blk}}, \\ p_i, & \pi \in \Pi_{\text{free}}. \end{cases}$$

The second line is well typed because $p_i \in \mathcal{P}_{\text{exc}} \subseteq \mathfrak{T}_{\text{prim}}$. It is the free-anchor route. Its role is the new symbol **freeAnchor**. A blocked pair retains the role determined by its canonical blocker and the subtype of $\Theta_{\text{cap}}(\pi)$. Thus the total role alphabet is the disjoint union

$$\hat{\mathfrak{R}}_{\text{st}} = \{\text{blocked}(r) : r \in \mathfrak{R}_{\text{blk}}\} \sqcup \{\text{freeAnchor}\}, \quad |\mathfrak{R}_{\text{blk}}| = 4 \cdot 6 = 24, \quad |\hat{\mathfrak{R}}_{\text{st}}| = 25.$$

Here $\mathfrak{R}_{\text{blk}}$ is the full product of the four admitted structural blocker kinds (a)–(c), (f) and the six token subtypes; product entries that are not realized simply have empty fibres. Raw profile and target mismatches have already left as audit exits and do not consume blocked roles. Write $\hat{\rho}(\pi)$ for this role, and put

$$\hat{\Pi}_{t,r} = \{\pi : \hat{\Theta}(\pi) = t, \hat{\rho}(\pi) = r\}, \quad \hat{\ell}(t, r) = |\hat{\Pi}_{t,r}|, \quad \hat{\ell}_{\text{cap}}(t) = \sum_r \hat{\ell}(t, r).$$

Let $\hat{H}_{t,r}$ be the simple graph on $\mathcal{A}_0$ with edge set $\hat{\Pi}_{t,r}$. A matching or star in this graph is a total-role-homogeneous same-token pattern. When $r = \text{freeAnchor}$, every edge has first endpoint t , so the whole fibre is already a star centered at that selected port. At the node-[130] dispatch, before a blocked pair has been retokenized, the finite role is only its actual canonical blocker kind. Hence that executable dispatch has five roles: the four admitted structural blocker kinds and **freeAnchor**. The token subtype is computed only by the downstream capacity-token map.




Figure 13: Visual definition of the sparse surplus port, blocker, and capacity-token ledger. Panel I shows a high center h , a selected surplus port $p = (h, x)$, and the cubic endpoint x with its two shoulder vertices a_p, b_p . The dashed shoulder chord $e_p^* = a_p b_p$ illustrates an open port; the small triangle below illustrates the triangular case $e_p^* \in E(G)$. The selected ports in $\mathcal{P}_{\text{exc}}(h)$ are the $d_G(h) - 3$ surplus units at h , while the unselected incidences are base ports. Panel II shows two active demands p, q , their declared supports $T(p) \cup \Gamma(p)$ and $T(q) \cup \Gamma(q)$, and a blocker B for the pair $\pi = \{p, q\}$. A blocked pair is assigned its canonical blocker $\Phi_{\text{can}}(\pi)$; a free pair lies in Π_{free} and contributes an independent pair-response coordinate r_π . Panel III shows the capacity-token map Θ_{cap} . Each blocked pair is assigned exactly one token in the disjoint union $\mathfrak{T}_{\text{cap}} = \mathfrak{T}_W \sqcup \mathfrak{T}_R \sqcup \mathfrak{T}_{\text{prim}}$, where $\mathfrak{T}_W$ records window incidences, $\mathfrak{T}_R$ records surplus units in the remainder, and $\mathfrak{T}_{\text{prim}}$ records primitive carriers. Multiple blocked pairs may land on the same token t ; their number is the token load $\ell_{\text{cap}}(t) = |\Theta_{\text{cap}}^{-1}(t)|$.

Lemma 3.43 (Total pair route has no overcounting). *The total token–role fibres form a disjoint partition of the complete pair schedule:*

$$\binom{\mathcal{A}_0}{2} = \bigsqcup_{t \in \mathfrak{T}_{\text{cap}}} \bigsqcup_{r \in \widehat{\mathfrak{R}}_{\text{st}}} \widehat{\Pi}_{t,r}.$$

Consequently,

$$\binom{|\mathcal{A}_0|}{2} = \sum_{t \in \mathfrak{T}_{\text{cap}}} \sum_{r \in \widehat{\mathfrak{R}}_{\text{st}}} \widehat{\ell}(t, r).$$

If $\pi \in \widehat{\Pi}_{p, \text{freeAnchor}}$, then π is blocker-free, its first selected port is exactly p , and it retains the canonical connector and the two demand supports fixed at node [130].

Proof. The canonical blocker decision partitions the pair schedule into the disjoint sets Π_{blk} and Π_{free} . On the first set, Θ_{cap} and the blocked role are functions. On the second set, the canonical first endpoint and the distinct role **freeAnchor** are functions. The distinct role constructors make the two images disjoint, proving the partition and the cardinality identity. Membership in a free-anchor fibre unfolds to the negative side of the blocker decision and equality of the first endpoint with p . The connector and support data are the values already retained for that same free pair, so no new path or support is selected. $\square$

Lemma 3.44 (Capacity-token supply). *The capacity-token universe satisfies*

$$|\mathfrak{T}_{\text{cap}}| \leq 8n + \sigma(G).$$

More precisely,

$$|\mathfrak{T}_{\text{cap}}| = |\mathfrak{U}_{\text{sp}}(G)| + \sigma_R + e(R, W) + 2e_{\times}(W) = |\mathfrak{U}_{\text{sp}}(G)| + 15p_{13} + \sigma(G).$$

Proof. By construction,

$$|\mathfrak{T}_{\text{cap}}| = |\mathfrak{T}_{\text{prim}}| + |\mathfrak{T}_R| + |\mathfrak{T}_W| = |\mathfrak{U}_{\text{sp}}(G)| + \sigma_R + e(R, W) + 2e_{\times}(W).$$

The exact window-join identity [Lemma 8.5](#) gives

$$e(R, W) + 2e_{\times}(W) = 15p_{13} + \sigma_W.$$

Since $\sigma(G) = \sigma_R + \sigma_W$,

$$|\mathfrak{T}_{\text{cap}}| = |\mathfrak{U}_{\text{sp}}(G)| + 15p_{13} + \sigma(G).$$

By [Lemma 3.40](#), $|\mathfrak{U}_{\text{sp}}(G)| \leq 6n$. Also $13p_{13} \leq n$, because the packed P_{13} windows are vertex-disjoint. Hence $15p_{13} \leq (15/13)n < 2n$, and therefore

$$|\mathfrak{T}_{\text{cap}}| \leq 8n + \sigma(G).$$

□

Lemma 3.45 (Token ledger has no overcounting). *For every finite family $\mathcal{A}_0$ of active surplus demands,*

$$|\Pi_{\text{blk}}| = \sum_{t \in \mathfrak{T}_{\text{cap}}} \ell_{\text{cap}}(t).$$

Consequently, if $\ell_{\text{cap}}(t) \leq M_0$ for every $t \in \mathfrak{T}_{\text{cap}}$, then

$$|\Pi_{\text{blk}}| \leq M_0(8n + \sigma(G)).$$

Proof. The map $\Theta_{\text{cap}} : \Pi_{\text{blk}} \rightarrow \mathfrak{T}_{\text{cap}}$ assigns exactly one token to each blocked pair. The fibres $\Theta_{\text{cap}}^{-1}(t)$, $t \in \mathfrak{T}_{\text{cap}}$, form a disjoint partition of Π_{blk} . Therefore

$$|\Pi_{\text{blk}}| = \sum_{t \in \mathfrak{T}_{\text{cap}}} |\Theta_{\text{cap}}^{-1}(t)| = \sum_{t \in \mathfrak{T}_{\text{cap}}} \ell_{\text{cap}}(t).$$

If every token load is at most M_0 , then

$$|\Pi_{\text{blk}}| \leq M_0 |\mathfrak{T}_{\text{cap}}| \leq M_0(8n + \sigma(G))$$

by [Lemma 3.44](#).

□

Definition 3.46 (Token-fibre graphs and same-token patterns). *For $t \in \mathfrak{T}_{\text{cap}}$, the token-fibre graph H_t is the simple graph with vertex set $\mathcal{A}_0$ and edge set*

$$E(H_t) = \Pi_t = \Theta_{\text{cap}}^{-1}(t).$$

Thus

$$e(H_t) = \ell_{\text{cap}}(t).$$

A same-token K -matching at t is a matching of size K in H_t . A same-token K -star at t is a set of K edges of H_t incident with one common active demand. The token class of the pattern is the token class containing t : window-incidence, remainder-surplus, or primitive-carrier.

Definition 3.47 (Same-token blocker roles). *Let $t \in \mathfrak{T}_{\text{cap}}$, and let $\pi \in \Theta_{\text{cap}}^{-1}(t)$. Put $B_\pi = \Phi_{\text{can}}(\pi)$. The same-token role of π at t is the following finite label:*

$$\rho_t(\pi) = (\text{type}(B_\pi), \text{class}(t), \text{sub}(t)).$$

Here $\text{type}(B_\pi) \in \{a, b, c, d, e, f\}$ is the blocker type in [Definition 3.34](#). The token class $\text{class}(t)$ is one of W, R, P , according as

$$t \in \mathfrak{T}_W, \quad t \in \mathfrak{T}_R, \quad t \in \mathfrak{T}_{\text{prim}}.$$

The subtype is defined by

$$\text{sub}(t) = \begin{cases} RW, & t = (e, RW) \in \mathfrak{T}_W, \\ WW, & t = (e, a) \in \mathfrak{T}_W \text{ is a cross-window endpoint token,} \\ R, & t \in \mathfrak{T}_R, \\ V, & t \in V(G) \subseteq \mathfrak{T}_{\text{prim}}, \\ I, & t \in I_E(G) \subseteq \mathfrak{T}_{\text{prim}}, \\ P, & t \in \mathcal{P}_{\text{exc}} \subseteq \mathfrak{T}_{\text{prim}}. \end{cases}$$

Let $\mathfrak{R}_{\text{st}}$ be the set of all such labels. Then

$$|\mathfrak{R}_{\text{st}}| \leq 6(2 + 1 + 3) = 36.$$

We write $Q_{\text{st}} := 36$ for this uniform role bound. For $r \in \mathfrak{R}_{\text{st}}$, define the role fibre at t by

$$\Pi_{t,r} := \{\pi \in \Theta_{\text{cap}}^{-1}(t) : \rho_t(\pi) = r\}, \quad \ell(t, r) := |\Pi_{t,r}|.$$

The role-fibre graph $H_{t,r}$ is the spanning subgraph of H_t with edge set $\Pi_{t,r}$. The sets $\Pi_{t,r}$, $r \in \mathfrak{R}_{\text{st}}$, partition $\Theta_{\text{cap}}^{-1}(t)$, and therefore

$$\ell_{\text{cap}}(t) = \sum_{r \in \mathfrak{R}_{\text{st}}} \ell(t, r).$$

A same-token matching or same-token star is role-homogeneous if all of its edges π have the same label $\rho_t(\pi)$.

Definition 3.48 (Token-load threshold and homogeneous cap charge). *For $x \geq 0$, define*

$$\psi(x) := \begin{cases} 0, & x = 0, \\ \left\lceil \frac{1 + \sqrt{1 + 8x}}{4} \right\rceil, & x > 0. \end{cases}$$

Equivalently, $\psi(x)$ is the least integer $k \geq 0$ such that

$$x \leq k(2k - 1).$$

For an integer $L \geq 1$, define the role-homogeneous token cap charge

$$\text{Cap}_{\text{hom}}(L) := Q_{\text{st}}(L - 1)(2L - 3).$$

Thus $\text{Cap}_{\text{hom}}(L)$ is the uniform token load allowed by charging each of the at most Q_{st} role fibres separately when no role-homogeneous same-token L -matching or L -star occurs at that token.

Lemma 3.49 (Homogeneous extraction from a same-token overload). *If a token $t \in \mathfrak{T}_{\text{cap}}$ supports a same-token K -matching, then t supports a role-homogeneous same-token matching of size at least $\lceil K/Q_{\text{st}} \rceil$. If t supports a same-token K -star, then t supports a role-homogeneous same-token star of size at least $\lceil K/Q_{\text{st}} \rceil$.*

Proof. The edges of H_t are coloured by the role map $\rho_t : E(H_t) \rightarrow \mathfrak{R}_{\text{st}}$. By Definition 3.47, at most $Q_{\text{st}} = 36$ colours occur. If M is a same-token K -matching, then the sets

$$M_r := \{\pi \in M : \rho_t(\pi) = r\}, \quad r \in \mathfrak{R}_{\text{st}},$$

partition M . Hence one colour class has cardinality at least $\lceil K/Q_{\text{st}} \rceil$. That colour class is still a matching, because it is a subset of M , and it is role-homogeneous by construction.

The star case is identical. If S is a same-token K -star with centre p , then each colour class $S_r \subseteq S$ is still a star with the same centre p . One colour class has size at least $\lceil K/Q_{\text{st}} \rceil$. $\square$

Corollary 3.50 (Homogeneous same-token caps close the surplus branch). *Assume G has full active family $\mathcal{A}_0 = \mathcal{P}_{\text{exc}}$. Suppose there are fixed integers L_W, L_R, L_P such that:*

- (a) *no token in $\mathfrak{T}_W$ supports a total-role-homogeneous same-token L_W -matching or L_W -star;*
- (b) *no token in $\mathfrak{T}_R$ supports a total-role-homogeneous same-token L_R -matching or L_R -star;*
- (c) *no token in $\mathfrak{T}_{\text{prim}}$ supports a total-role-homogeneous same-token L_P -matching or L_P -star.*

Then

$$\sigma(G) = O(\sqrt{n}),$$

where the implicit constant depends only on L_W, L_R, L_P .

Proof. Let

$$M_0 := \max\{\text{Cap}_{\text{hom}}(L_W), \text{Cap}_{\text{hom}}(L_R), \text{Cap}_{\text{hom}}(L_P)\}.$$

We show that every token has total load at most M_0 . Fix $t \in \mathfrak{T}_W$. For each total role r , the graph $\hat{H}_{t,r}$ has no L_W -matching and no L_W -star; otherwise that matching or star would be a total-role-homogeneous same-token L_W -pattern at t , contrary to (a). By Lemma 3.53,

$$\hat{\ell}(t, r) \leq (L_W - 1)(2L_W - 3)$$

for every r . Summing the total role-fibre partition of Definition 3.42 gives

$$\hat{\ell}_{\text{cap}}(t) = \sum_{r \in \mathfrak{R}_{\text{st}}} \hat{\ell}(t, r) \leq Q_{\text{st}}(L_W - 1)(2L_W - 3) = \text{Cap}_{\text{hom}}(L_W) \leq M_0.$$

The argument for a token in $\mathfrak{T}_W$ used only the total role-fibre partition and the corresponding homogeneous cap. Replacing L_W by L_R or L_P gives the same bound for tokens in $\mathfrak{T}_R$ and $\mathfrak{T}_{\text{prim}}$. The hypotheses of Theorem 3.63 now hold with $c_1 = 1$, because $|\mathcal{A}_0| = \sigma(G)$ by Lemma 3.32. That theorem gives $\sigma(G) = O(\sqrt{n})$. $\square$

Corollary 3.51 (Forced homogeneous same-token scale). *Assume G survives the sparse surplus exits, admits a baseline spine demand with deficit $E_{\text{spine}}(n)$, and has full active family $\mathcal{A}_0 = \mathcal{P}_{\text{exc}}$. Let $\Lambda(G)$ be the quantity in Corollary 3.55. Then some capacity token supports a total-role-homogeneous same-token matching or total-role-homogeneous same-token star of size at least*

$$\psi\left(\frac{\Lambda(G)}{Q_{\text{st}}}\right).$$

In particular, if $E_{\text{spine}}(n) = O(n)$ and $\sigma(G)/\sqrt{n} \rightarrow \infty$, then the forced total-role-homogeneous same-token matching-or-star size tends to infinity.

Proof. Lemma 3.64 gives a token t with

$$\widehat{\ell}_{\text{cap}}(t) \geq \Lambda(G).$$

The total role fibres at t partition $\widehat{\Theta}^{-1}(t)$, so

$$\sum_{r \in \widehat{\mathfrak{R}}_{\text{st}}} \widehat{\ell}(t, r) = \widehat{\ell}_{\text{cap}}(t).$$

Since at most Q_{st} roles occur, some role r satisfies

$$\widehat{\ell}(t, r) \geq \frac{\Lambda(G)}{Q_{\text{st}}}.$$

Apply Lemma 3.53 to the total role-fibre graph $\widehat{H}_{t,r}$. If $s = \max\{\nu(\widehat{H}_{t,r}), \Delta(\widehat{H}_{t,r})\}$, then

$$\widehat{\ell}(t, r) = e(\widehat{H}_{t,r}) \leq s(2s - 1).$$

By the definition of ψ ,

$$s \geq \psi\left(\frac{\Lambda(G)}{Q_{\text{st}}}\right).$$

Thus $\widehat{H}_{t,r}$ contains a matching or star of that size. Because all edges of $\widehat{H}_{t,r}$ have the same total role r , this is a total-role-homogeneous same-token matching or star. The final assertion follows because Q_{st} is fixed. $\square$

Corollary 3.52 (Quantitative homogeneous overload lower bound). *Assume G survives the sparse surplus exits, has full active family $\mathcal{A}_0 = \mathcal{P}_{\text{exc}}$, and admits a baseline spine demand with deficit $E_{\text{spine}}(n) \leq C_E n$. Let $K_{\text{hom}}(G)$ be the largest integer K for which some capacity token supports a total-role-homogeneous same-token K -matching or a total-role-homogeneous same-token K -star. Then*

$$K_{\text{hom}}(G) \geq \psi\left(\frac{\max\left\{0, \binom{\sigma(G)}{2} - C_E n - \left(\frac{1}{2}\sigma(G) + 1\right) \log_2 n\right\}}{Q_{\text{st}}(8n + \sigma(G))}\right).$$

Consequently:

(a) If $\sigma(G) \geq A\sqrt{n}$, $A^2 \geq 6(C_E + 1)$, and n is sufficiently large in terms of A and C_E , then

$$K_{\text{hom}}(G) \geq \left\lceil \frac{A}{12\sqrt{Q_{\text{st}}}} \right\rceil.$$

(b) If $\sigma(G) \geq \alpha n$ for some fixed $\alpha > 0$, and n is sufficiently large in terms of α and C_E , then

$$K_{\text{hom}}(G) \geq \left\lceil \frac{\alpha}{12\sqrt{Q_{\text{st}}}} \sqrt{n} \right\rceil.$$

Proof. The displayed bound is Corollary 3.51 with $E_{\text{spine}}(n) \leq C_E n$ substituted into $\Lambda(G)$.

For (a), assume $\sigma(G) \geq A\sqrt{n}$. Since $\sigma(G) \leq n$ by Lemma 3.3, the denominator $8n + \sigma(G)$ is at most $9n$. For fixed A , the term $(\sigma(G)/2 + 1) \log_2 n$ is $o(n)$. For all sufficiently large n ,

$$\binom{\sigma(G)}{2} - C_E n - \left(\frac{1}{2}\sigma(G) + 1\right) \log_2 n \geq \left(\frac{A^2}{3} - C_E - 1\right) n \geq \frac{A^2}{6} n,$$

where the last inequality uses $A^2 \geq 6(C_E + 1)$. Hence

$$\frac{\binom{\sigma(G)}{2} - C_E n - \left(\frac{1}{2}\sigma(G) + 1\right) \log_2 n}{Q_{\text{st}}(8n + \sigma(G))} \geq \frac{A^2}{54Q_{\text{st}}}.$$

If $k < A/(12\sqrt{Q_{\text{st}}})$, then

$$k(2k - 1) < 2k^2 < \frac{A^2}{72Q_{\text{st}}} < \frac{A^2}{54Q_{\text{st}}}.$$

By the defining minimality of ψ , this gives (a).

For (b), assume $\sigma(G) \geq \alpha n$. For all sufficiently large n ,

$$\binom{\sigma(G)}{2} - C_E n - \left(\frac{1}{2}\sigma(G) + 1\right) \log_2 n \geq \frac{\alpha^2}{3} n^2.$$

Again $8n + \sigma(G) \leq 9n$. Therefore

$$\frac{\binom{\sigma(G)}{2} - C_E n - \left(\frac{1}{2}\sigma(G) + 1\right) \log_2 n}{Q_{\text{st}}(8n + \sigma(G))} \geq \frac{\alpha^2}{27Q_{\text{st}}} n.$$

If $k < \alpha\sqrt{n}/(12\sqrt{Q_{\text{st}}})$, then

$$k(2k - 1) < 2k^2 < \frac{\alpha^2}{72Q_{\text{st}}} n < \frac{\alpha^2}{27Q_{\text{st}}} n.$$

The defining minimality of ψ gives (b). □

Lemma 3.53 (Matching–star accounting). *Let H be a finite simple graph with $e(H)$ edges, matching number $\nu(H)$, and maximum degree $\Delta(H)$. Then*

$$e(H) \leq \nu(H)(2\Delta(H) - 1).$$

Consequently, if $K \geq 1$ and H contains neither a matching of size K nor a star of size K , then

$$e(H) \leq (K - 1)(2K - 3).$$

Equivalently, if $e(H) > (K - 1)(2K - 3)$, then H contains a matching of size K or a star of size K .

Proof. Choose a maximum matching M in H . Since M is maximal, every edge of H has at least one endpoint in the set U of vertices covered by M . Thus U is a vertex cover of H , and $|U| = 2\nu(H)$. Each vertex of U is incident with at most $\Delta(H)$ edges of H . Hence

$$\sum_{u \in U} d_H(u) \leq 2\nu(H)\Delta(H).$$

The $\nu(H)$ matching edges of M have both endpoints in U , and are therefore counted twice in the degree sum. Every other edge is counted at least once, because U is a vertex cover. Hence

$$e(H) \leq \sum_{u \in U} d_H(u) - \nu(H) \leq \nu(H)(2\Delta(H) - 1).$$

If H contains no matching of size K , then $\nu(H) \leq K - 1$. If it contains no star of size K , then $\Delta(H) \leq K - 1$. Substituting these two bounds gives $e(H) \leq (K - 1)(2K - 3)$. The final statement is the contrapositive. □

Corollary 3.54 (Same-token pattern caps close the surplus branch). *Assume G survives the sparse surplus exits, admits a baseline spine demand with $E_{\text{spine}}(n) \leq C_E n$, and has full active family $\mathcal{A}_0 = \mathcal{P}_{\text{exc}}$. Suppose there are fixed integers K_W, K_R, K_P such that:*

- (a) *no token in $\mathfrak{T}_W$ supports a total-role-homogeneous same-token K_W -matching or K_W -star;*
- (b) *no token in $\mathfrak{T}_R$ supports a total-role-homogeneous same-token K_R -matching or K_R -star;*
- (c) *no token in $\mathfrak{T}_{\text{prim}}$ supports a total-role-homogeneous same-token K_P -matching or K_P -star.*

Then

$$\sigma(G) = O(\sqrt{n}),$$

where the implicit constant depends only on C_E, K_W, K_R, K_P .

Proof. The three hypotheses are exactly those of [Corollary 3.50](#), with $L_W = K_W$, $L_R = K_R$, and $L_P = K_P$. That corollary gives $\sigma(G) = O(\sqrt{n})$. $\square$

Corollary 3.55 (Forced same-token overload scale). *Assume G survives the sparse surplus exits, admits a baseline spine demand with deficit $E_{\text{spine}}(n)$, and has full active family $\mathcal{A}_0 = \mathcal{P}_{\text{exc}}$. Put*

$$\Lambda(G) := \frac{\max \left\{ 0, \binom{\sigma(G)}{2} - E_{\text{spine}}(n) - \left(\frac{1}{2}\sigma(G) + 1 \right) \log_2 n \right\}}{8n + \sigma(G)}.$$

Then some capacity token supports a total-role-homogeneous same-token matching or total-role-homogeneous same-token star of size at least

$$\psi \left(\frac{\Lambda(G)}{Q_{\text{st}}} \right).$$

In particular, if $E_{\text{spine}}(n) = O(n)$ and $\sigma(G)/\sqrt{n} \rightarrow \infty$, then the forced same-token matching-or-star size tends to infinity.

Proof. [Lemma 3.64](#) gives a token t with

$$\widehat{\ell}_{\text{cap}}(t) \geq \Lambda(G).$$

The at most Q_{st} total role fibres partition this token fibre, so some total role r satisfies

$$\widehat{\ell}(t, r) \geq \frac{\Lambda(G)}{Q_{\text{st}}}.$$

Put $s = \max\{\nu(\widehat{H}_{t,r}), \Delta(\widehat{H}_{t,r})\}$. By [Lemma 3.53](#),

$$\widehat{\ell}(t, r) = e(\widehat{H}_{t,r}) \leq s(2s - 1).$$

The definition of ψ now gives $s \geq \psi(\Lambda(G)/Q_{\text{st}})$. Thus $\widehat{H}_{t,r}$ contains a matching or star of the claimed size, and all its edges have total role r . If $E_{\text{spine}}(n) = O(n)$ and $\sigma(G)/\sqrt{n} \rightarrow \infty$, then the numerator in $\Lambda(G)$ is $\frac{1}{2}\sigma(G)^2 - o(\sigma(G)^2)$, while $8n + \sigma(G) \leq 9n$ by $\sigma(G) \leq n$, which follows from [Lemma 3.3](#). Hence $\Lambda(G) \rightarrow \infty$, and so the displayed same-token pattern size tends to infinity. $\square$

Lemma 3.56 (Exact surplus-pair charge partition). *Let $\mathcal{A}_0$ be a finite family of active surplus demands. For $C \in \{W, R, P\}$, use the token classes*

$$\mathfrak{T}_W, \quad \mathfrak{T}_R, \quad \mathfrak{T}_P := \mathfrak{T}_{\text{prim}}.$$

Then the following disjoint union is exact:

$$\binom{\mathcal{A}_0}{2} = \bigsqcup_{C \in \{W, R, P\}} \bigsqcup_{t \in \mathfrak{T}_C} \bigsqcup_{r \in \hat{\mathfrak{R}}_{\text{st}}} \hat{\Pi}_{t,r}.$$

Consequently,

$$\binom{|\mathcal{A}_0|}{2} = \sum_{C \in \{W, R, P\}} \sum_{t \in \mathfrak{T}_C} \sum_{r \in \hat{\mathfrak{R}}_{\text{st}}} \hat{\ell}(t, r).$$

Every unordered pair of active demands appears in exactly one summand on the right-hand side.

Proof. By Lemma 3.43, every pair has one total token and one total role. The token is a single element of the disjoint universe

$$\mathfrak{T}_{\text{cap}} = \mathfrak{T}_W \sqcup \mathfrak{T}_R \sqcup \mathfrak{T}_{\text{prim}}.$$

Thus each pair lies in exactly one token class C , exactly one token $t \in \mathfrak{T}_C$. For that fixed t , the total role map $\hat{\rho}$ assigns exactly one role $r \in \hat{\mathfrak{R}}_{\text{st}}$, so $\pi \in \hat{\Pi}_{t,r}$. This proves that every pair appears in at least one displayed role-fibre summand.

The appearance is unique. The token classes are disjoint, $\hat{\Theta}$ is a function, and $\hat{\rho}$ is a function once t is fixed. Hence no pair can belong to two distinct triples (C, t, r) . The displayed union is disjoint and exhaustive. Taking cardinalities gives the second formula. $\square$

Theorem 3.57 (Sharp classwise homogeneous token budget). *Assume G survives the sparse surplus exits, admits a baseline spine demand with deficit $E_{\text{spine}}(n)$, and has full active family $\mathcal{A}_0 = \mathcal{P}_{\text{exc}}$. Put*

$$N_*(G) := \max \left\{ 0, \binom{\sigma(G)}{2} - E_{\text{spine}}(n) - \left(\frac{1}{2} \sigma(G) + 1 \right) \log_2 n \right\}.$$

For $C \in \{W, R, P\}$, set

$$\mathfrak{T}_C = \begin{cases} \mathfrak{T}_W, & C = W, \\ \mathfrak{T}_R, & C = R, \\ \mathfrak{T}_{\text{prim}}, & C = P, \end{cases} \quad \Pi_C := \left\{ \pi \in \binom{\mathcal{A}_0}{2} : \hat{\Theta}(\pi) \in \mathfrak{T}_C \right\},$$

and write $B_C := |\Pi_C|$, $S_C := |\mathfrak{T}_C|$. Then:

(a) *The total-pair demand satisfies the exact class split*

$$B_W + B_R + B_P = \binom{\sigma(G)}{2} \geq N_*(G).$$

(b) *The class supplies are*

$$S_W = 15p_{13} + \sigma_W, \quad S_R = \sigma_R, \quad S_P = |\mathfrak{U}_{\text{sp}}(G)| \leq 6n,$$

and hence

$$S_W + S_R + S_P \leq 8n + \sigma(G).$$

(c) If no token in $\mathfrak{T}_C$ supports a role-homogeneous same-token L_C -matching or role-homogeneous same-token L_C -star, then

$$B_C \leq \text{Cap}_{\text{hom}}(L_C)S_C = Q_{\text{st}}(L_C - 1)(2L_C - 3)S_C.$$

(d) If the three homogeneous caps L_W, L_R, L_P hold in the three token classes, then the necessary classwise budget inequality is

$$N_*(G) \leq \text{Cap}_{\text{hom}}(L_W)(15p_{13} + \sigma_W) + \text{Cap}_{\text{hom}}(L_R)\sigma_R + \text{Cap}_{\text{hom}}(L_P)|\mathfrak{U}_{\text{sp}}(G)|.$$

In particular,

$$N_*(G) \leq \text{Cap}_{\text{hom}}(L_W)(15p_{13} + \sigma_W) + \text{Cap}_{\text{hom}}(L_R)\sigma_R + 6 \text{Cap}_{\text{hom}}(L_P)n.$$

(e) If $N_*(G) > 0$, then for some nonempty token class C there is a token $t \in \mathfrak{T}_C$ supporting a role-homogeneous same-token matching or role-homogeneous same-token star of size at least

$$\psi\left(\frac{B_C}{Q_{\text{st}}S_C}\right),$$

and in particular at least

$$\psi\left(\frac{N_*(G)}{Q_{\text{st}}(8n + \sigma(G))}\right).$$

Proof. For (a), [Lemma 3.56](#) and $|\mathcal{A}_0| = \sigma(G)$ give

$$B_W + B_R + B_P = \binom{\sigma(G)}{2}.$$

The subtracted deficit and incremental-room terms in the definition of $N_*(G)$ are nonnegative, so $N_*(G) \leq \binom{\sigma(G)}{2}$. The three sets Π_W, Π_R, Π_P partition the complete pair family, because $\mathfrak{T}_{\text{cap}}$ is the disjoint union $\mathfrak{T}_W \sqcup \mathfrak{T}_R \sqcup \mathfrak{T}_{\text{prim}}$ and $\hat{\Theta}$ assigns exactly one token to each pair. Hence

$$B_W + B_R + B_P = \binom{\sigma(G)}{2} \geq N_*(G).$$

Part (b) is exactly [Lemma 3.44](#), with $\mathfrak{T}_{\text{prim}} = \mathfrak{U}_{\text{sp}}(G)$ and [Lemma 3.40](#) applied to S_P .

For (c), fix a class C and a token $t \in \mathfrak{T}_C$. For every role $r \in \hat{\mathfrak{R}}_{\text{st}}$, the total role-fibre graph $\hat{H}_{t,r}$ has no matching of size L_C and no star of size L_C ; otherwise that matching or star would be a role-homogeneous same-token L_C -pattern at t . By [Lemma 3.53](#),

$$\hat{\ell}(t, r) = e(\hat{H}_{t,r}) \leq (L_C - 1)(2L_C - 3).$$

The total role fibres partition the token fibre, and there are at most $25 \leq Q_{\text{st}} = 36$ of them, so

$$\hat{\ell}_{\text{cap}}(t) = \sum_{r \in \hat{\mathfrak{R}}_{\text{st}}} \hat{\ell}(t, r) \leq Q_{\text{st}}(L_C - 1)(2L_C - 3) = \text{Cap}_{\text{hom}}(L_C).$$

Summing the exact token fibres in class C gives

$$B_C = \sum_{t \in \mathfrak{T}_C} \hat{\ell}_{\text{cap}}(t) \leq \text{Cap}_{\text{hom}}(L_C)S_C.$$

Part (d) follows by summing the three estimates in (c) and using (a) and (b).

For (e), choose a class C for which B_C/S_C is maximal among nonempty classes with $S_C > 0$. Since the nonempty classes partition the complete pair family,

$$\frac{B_C}{S_C} \geq \frac{B_W + B_R + B_P}{S_W + S_R + S_P} \geq \frac{N_*(G)}{8n + \sigma(G)}.$$

Some token $t \in \mathfrak{T}_C$ then has load

$$\widehat{\ell}_{\text{cap}}(t) \geq B_C/S_C.$$

The total role fibres at t partition the token fibre, so some role $r \in \widehat{\mathfrak{R}}_{\text{st}}$ has

$$\widehat{\ell}(t, r) \geq \frac{B_C}{Q_{\text{st}}S_C}.$$

Apply [Lemma 3.53](#) to $\widehat{H}_{t,r}$. If $s = \max\{\nu(\widehat{H}_{t,r}), \Delta(\widehat{H}_{t,r})\}$, then

$$\widehat{\ell}(t, r) = e(\widehat{H}_{t,r}) \leq s(2s - 1).$$

Therefore

$$s \geq \psi\left(\frac{B_C}{Q_{\text{st}}S_C}\right).$$

The graph $\widehat{H}_{t,r}$ contains a matching or star of that size, and all its edges have the same role r . This gives the first displayed homogeneous lower bound. The second follows from

$$\frac{B_C}{S_C} \geq \frac{N_*(G)}{8n + \sigma(G)}$$

and the monotonicity of ψ . □

Corollary 3.58 (Quantified homogeneous class overload). *Assume the hypotheses of [Theorem 3.57](#), and fix proposed homogeneous caps L_W, L_R, L_P . Define the class capacities*

$$A_W := \text{Cap}_{\text{hom}}(L_W)(15p_{13} + \sigma_W), \quad A_R := \text{Cap}_{\text{hom}}(L_R)\sigma_R, \quad A_P := \text{Cap}_{\text{hom}}(L_P)|\mathfrak{U}_{\text{sp}}(G)|.$$

Put

$$D_{\text{exc}} := (N_*(G) - A_W - A_R - A_P)_+,$$

and, for $C \in \{W, R, P\}$, put

$$\Omega_C := (B_C - A_C)_+.$$

Then:

(a) *The total class overload satisfies*

$$\Omega_W + \Omega_R + \Omega_P \geq D_{\text{exc}}.$$

(b) *Hence at least one class C satisfies*

$$\Omega_C \geq \frac{1}{3}D_{\text{exc}}.$$

(c) *If $\Omega_C > 0$, then $S_C > 0$, and class C contains a role-homogeneous same-token matching or role-homogeneous same-token star of size at least*

$$\psi\left((L_C - 1)(2L_C - 3) + \frac{\Omega_C}{Q_{\text{st}}S_C}\right).$$

- (d) If $D_{\text{exc}} > 0$, and the two classes different from C do not overload, i.e. $B_{C'} \leq A_{C'}$ for $C' \neq C$, then class C overloads by at least the whole deficit:

$$\Omega_C \geq D_{\text{exc}},$$

and class C contains a role-homogeneous same-token matching or role-homogeneous same-token star of size at least

$$\psi\left((L_C - 1)(2L_C - 3) + \frac{D_{\text{exc}}}{Q_{\text{st}}S_C}\right).$$

- (e) The three pure-overload tests have the following sharp class-by-class lower bounds. Define

$$F_W := (N_*(G) - A_R - A_P)_+, \quad F_R := (N_*(G) - A_W - A_P)_+, \quad F_P := (N_*(G) - A_W - A_R)_+.$$

If $B_R \leq A_R$ and $B_P \leq A_P$, then

$$B_W \geq F_W, \quad \Omega_W \geq (F_W - A_W)_+ = D_{\text{exc}}.$$

If $S_W > 0$, the forced window-incidence homogeneous pattern has size at least

$$\psi\left(\frac{F_W}{Q_{\text{st}}S_W}\right).$$

If $B_W \leq A_W$ and $B_P \leq A_P$, then

$$B_R \geq F_R, \quad \Omega_R \geq (F_R - A_R)_+ = D_{\text{exc}}.$$

If $S_R > 0$, the forced remainder-surplus homogeneous pattern has size at least

$$\psi\left(\frac{F_R}{Q_{\text{st}}S_R}\right).$$

If $B_W \leq A_W$ and $B_R \leq A_R$, then

$$B_P \geq F_P, \quad \Omega_P \geq (F_P - A_P)_+ = D_{\text{exc}}.$$

If $S_P > 0$, the forced primitive-carrier homogeneous pattern has size at least

$$\psi\left(\frac{F_P}{Q_{\text{st}}S_P}\right).$$

The estimate is sharp for the scalar class-load ledger: the load vector

$$B_C = A_C + D_{\text{exc}}, \quad B_{C'} = A_{C'} \quad (C' \neq C)$$

has total load $A_W + A_R + A_P + D_{\text{exc}}$ and overloads only class C .

Proof. For (a), use the elementary inequality

$$\sum_C (B_C - A_C)_+ \geq \left(\sum_C B_C - \sum_C A_C \right)_+.$$

By [Theorem 3.57](#),

$$\sum_C B_C = B_W + B_R + B_P \geq N_*(G),$$

and therefore

$$\Omega_W + \Omega_R + \Omega_P \geq (N_*(G) - A_W - A_R - A_P)_+ = D_{\text{exc}}.$$

Part (b) is the pigeonhole principle applied to the three nonnegative numbers $\Omega_W, \Omega_R, \Omega_P$.

For (c), if $\Omega_C > 0$, then $B_C > A_C$, so $B_C > 0$, and hence $S_C > 0$. By [Theorem 3.57](#), class C contains a role-homogeneous same-token matching or star of size at least

$$\psi\left(\frac{B_C}{Q_{\text{st}}S_C}\right).$$

Since $B_C = A_C + \Omega_C$ and

$$A_C = \text{Cap}_{\text{hom}}(L_C)S_C = Q_{\text{st}}(L_C - 1)(2L_C - 3)S_C,$$

we have

$$\frac{B_C}{Q_{\text{st}}S_C} = (L_C - 1)(2L_C - 3) + \frac{\Omega_C}{Q_{\text{st}}S_C}.$$

This proves (c).

For (d), assume $B_{C'} \leq A_{C'}$ for the two classes $C' \neq C$. Then

$$B_C \geq N_*(G) - \sum_{C' \neq C} B_{C'} \geq N_*(G) - \sum_{C' \neq C} A_{C'} = A_C + (N_*(G) - A_W - A_R - A_P).$$

If $D_{\text{exc}} > 0$, this gives $B_C \geq A_C + D_{\text{exc}}$, hence $\Omega_C \geq D_{\text{exc}}$. Substituting this bound into (c) gives the displayed pattern-size lower bound. The final sentence is a direct substitution into the scalar identities defining the class-load ledger.

For (e), first assume $B_R \leq A_R$ and $B_P \leq A_P$. Then

$$B_W \geq N_*(G) - B_R - B_P \geq N_*(G) - A_R - A_P.$$

Since $B_W \geq 0$, this gives $B_W \geq F_W$. Subtracting A_W and taking positive parts gives

$$\Omega_W \geq (F_W - A_W)_+ = (N_*(G) - A_W - A_R - A_P)_+ = D_{\text{exc}}.$$

If $S_W > 0$, then [Theorem 3.57](#) gives a window-incidence role-homogeneous same-token matching or star of size at least

$$\psi\left(\frac{B_W}{Q_{\text{st}}S_W}\right) \geq \psi\left(\frac{F_W}{Q_{\text{st}}S_W}\right),$$

using monotonicity of ψ . The R - and P -estimates are the same argument with the class names permuted. $\square$

Proposition 3.59 (Exact 25-role coupled CT9 decision). *Assume that the same residual graph has reached node [136]. Fix integers $L_W, L_R, L_P \geq 2$, set*

$$b_C = (L_C - 1)(2L_C - 3) \quad (C \in \{W, R, P\}), \quad b_{\max} = \max\{b_W, b_R, b_P\},$$

and give every token-role label (t, r) the capacity $b_{C(t)}$, where $C(t)$ is the unique class of the token. Define the exact coupled capacity and its positive excess by

$$A_{25} := \sum_{C \in \{W, R, P\}} \sum_{t \in \mathfrak{T}_C} \sum_{r \in \widehat{\mathfrak{R}}_{\text{st}}} b_C, \quad D_{25} := \left(\binom{\sigma(G)}{2} - A_{25} \right)_+.$$

Then the following alternatives are exhaustive.

(a) If $D_{25} > 0$, CT9 returns a concrete label (t, r) satisfying

$$\widehat{\ell}(t, r) > b_{C(t)}.$$

The constructor of t routes this output uniquely through node [139] or node [141] to exactly one of the window-incidence, remainder-surplus, and primitive-carrier classes.

(b) If $D_{25} = 0$, then

$$\sigma(G)^2 \leq (450b_{\max} + 1)n.$$

In particular this side of the decision is the node-[138] near-cubic route, with an explicit constant depending only on the three fixed thresholds.

The decision uses only the selected graph's complete surplus-pair list and its finite token and role lists. Its number of label comparisons is at most $225n^3$.

Proof. Node [136] gives the exact partition

$$\binom{\sigma(G)}{2} = \sum_{C, t, r} \widehat{\ell}(t, r)$$

over the 25-element total-role alphabet. If $D_{25} > 0$, then $A_{25} < \binom{\sigma(G)}{2}$. Were every fibre bounded by its assigned capacity, summing those 25-role fibre inequalities over all tokens would give $\binom{\sigma(G)}{2} \leq A_{25}$, a contradiction. The first overloaded label in the fixed token–role order is therefore a concrete CT9 residual. The token universe is the disjoint union $\mathfrak{T}_W \sqcup \mathfrak{T}_R \sqcup \mathfrak{T}_{\text{prim}}$, so testing the token constructor gives the stated unique class route.

Suppose now that $D_{25} = 0$. Then $\binom{\sigma(G)}{2} \leq A_{25}$. Since each label capacity is at most $b_{\max}$, $|\widehat{\mathfrak{R}}_{\text{st}}| = 25$, and Lemmas 3.3 and 3.44 give

$$A_{25} \leq 25b_{\max}|\mathfrak{T}_{\text{cap}}| \leq 25b_{\max}(8n + \sigma(G)) \leq 225b_{\max}n.$$

Using the exact identity

$$\sigma(G)^2 = 2\binom{\sigma(G)}{2} + \sigma(G)$$

and again $\sigma(G) \leq n$, we obtain

$$\sigma(G)^2 \leq 450b_{\max}n + n = (450b_{\max} + 1)n.$$

Finally, the complete pair list has at most n^2 entries, the token list has at most $9n$ entries, and the role list has 25 entries. The literal product scan therefore uses at most $n^2 \cdot 9n \cdot 25 = 225n^3$ comparisons. $\square$

Corollary 3.60 (Coupled single-graph overload budget). *Assume the hypotheses of Theorem 3.57, and fix proposed homogeneous caps L_W, L_R, L_P . Put*

$$b_W := (L_W - 1)(2L_W - 3), \quad b_R := (L_R - 1)(2L_R - 3), \quad b_P := (L_P - 1)(2L_P - 3).$$

Using the uniform bound $Q_{\text{st}} = 36$, the conservative coupled capacity envelope for the three classes in the single graph G is

$$A_{\text{all}} = 36 \left[b_W(15p_{13} + \sigma_W) + b_R\sigma_R + b_P(4n + 2\sigma(G)) \right].$$

Equivalently, using $\sigma(G) = \sigma_W + \sigma_R$,

$$A_{\text{all}} = 36 \left[4b_P n + 15b_W p_{13} + (b_W + 2b_P)\sigma_W + (b_R + 2b_P)\sigma_R \right].$$

Define the coupled excess

$$D_{\text{all}} := (N_*(G) - A_{\text{all}})_+.$$

Then:

(a) The total role-fibre overload relative to this envelope is at least

$$D_{\text{all}}.$$

(b) The total number of role-fibre slots available to absorb this excess is at most

$$36(S_W + S_R + S_P) = 36(4n + 15p_{13} + 3\sigma(G)).$$

(c) If $D_{\text{all}} > 0$, then for some class $C \in \{W, R, P\}$ there is a role fibre in that class whose load is at least

$$b_C + \frac{D_{\text{all}}}{36(4n + 15p_{13} + 3\sigma(G))}.$$

Consequently G contains a role-homogeneous same-token matching or role-homogeneous same-token star of size at least

$$\psi \left(b_C + \frac{D_{\text{all}}}{36(4n + 15p_{13} + 3\sigma(G))} \right)$$

for that same class C .

Proof. The exact primitive-token supply is

$$S_P = |\mathfrak{U}_{\text{sp}}(G)| = n + 2m + \sigma(G).$$

Since $\sigma(G) = 2m - 3n$, we have $m = (3n + \sigma(G))/2$, and therefore

$$S_P = n + 2m + \sigma(G) = 4n + 2\sigma(G).$$

Together with $S_W = 15p_{13} + \sigma_W$ and $S_R = \sigma_R$, this gives the displayed formula for A_{all} ; the second displayed form is obtained by substituting $\sigma(G) = \sigma_W + \sigma_R$.

For each role fibre in class C , the cap threshold before a homogeneous L_C -matching or L_C -star is b_C . Padding the 25 actual roles to the uniform 36-role bound gives the capacity envelope

$$36b_W S_W + 36b_R S_R + 36b_P S_P = A_{\text{all}}.$$

Since the exact class theorem gives a total-pair demand at least $N_*(G)$, the excess above this coupled envelope is at least D_{all} . This proves (a).

The number of role-fibre slots is at most $36(S_W + S_R + S_P)$. Using the exact supplies,

$$S_W + S_R + S_P = (15p_{13} + \sigma_W) + \sigma_R + (4n + 2\sigma(G)) = 4n + 15p_{13} + 3\sigma(G),$$

which proves (b).

For a token $t \in \mathfrak{T}_C$ and a total role $r \in \hat{\mathfrak{R}}_{\text{st}}$, write

$$E_{t,r} := (\hat{\ell}(t, r) - b_C)_+.$$

The elementary inequality $\sum_i (x_i - y_i)_+ \geq (\sum_i x_i - \sum_i y_i)_+$, applied to the 25 actual role labels, and then weakened to the 36-role capacity envelope, gives

$$\sum_{C \in \{W, R, P\}} \sum_{t \in \mathfrak{T}_C} \sum_{r \in \hat{\mathfrak{R}}_{\text{st}}} E_{t,r} \geq \left(\sum_C B_C - A_{\text{all}} \right)_+ \geq D_{\text{all}}.$$

If $D_{\text{all}} > 0$, distribute these explicitly defined role-fibre excesses over the at most

$$36(4n + 15p_{13} + 3\sigma(G))$$

role-fibre slots. Some role fibre in some class C has excess at least the average displayed in (c), so its load is at least

$$b_C + \frac{D_{\text{all}}}{36(4n + 15p_{13} + 3\sigma(G))}.$$

Applying the matching-star bound of [Lemma 3.53](#) to that role-fibre graph gives the stated homogeneous matching-or-star lower bound by the definition of ψ . $\square$

Corollary 3.61 (Numerical normalized single-graph budget). *Assume the hypotheses of [Corollary 3.60](#). Set*

$$\theta := \frac{p_{13}}{n}, \quad w := \frac{\sigma_W}{n}, \quad r := \frac{\sigma_R}{n}, \quad s := \frac{\sigma(G)}{n} = w + r.$$

Then $0 \leq \theta \leq 1/13$ and $0 \leq s \leq 1$, and the coupled capacity envelope is

$$A_{\text{all}} = n \Gamma_L(\theta, w, r),$$

where

$$\Gamma_L(\theta, w, r) := 144b_P + 540b_W\theta + 36(b_W + 2b_P)w + 36(b_R + 2b_P)r.$$

Equivalently, the three token classes contribute the following single-graph budget:

class	S_C/n	role-fibre cap	A_C/n
W	$15\theta + w$	b_W	$36b_W(15\theta + w)$
R	r	b_R	$36b_Rr$
P	$4 + 2s$	b_P	$36b_P(4 + 2s)$

The corresponding envelope excess is

$$D_{\text{all}} = \left[\frac{\sigma(G)(\sigma(G) - 1)}{2} - \left(\frac{\sigma(G)}{2} + 1 \right) \log_2 n - E_{\text{spine}}(n) - n \Gamma_L(\theta, w, r) \right]_+.$$

If $E_{\text{spine}}(n) \leq C_E n$, then

$$D_{\text{all}} \geq \left[\frac{\sigma(G)(\sigma(G) - 1)}{2} - \left(\frac{\sigma(G)}{2} + 1 \right) \log_2 n - n(C_E + \Gamma_L(\theta, w, r)) \right]_+.$$

Moreover the total role-fibre slot count is at most

$$n \Delta(\theta, s), \quad \Delta(\theta, s) := 144 + 540\theta + 108s \leq \frac{3816}{13} = 293.5384615 \dots$$

Thus, if $D_{\text{all}} > 0$, some role fibre in one of the three classes has load at least

$$b_C + \frac{D_{\text{all}}}{n \Delta(\theta, s)} \geq b_C + \frac{13D_{\text{all}}}{3816n},$$

and hence supports a role-homogeneous same-token matching or star of size at least

$$\psi\left(b_C + \frac{13D_{\text{all}}}{3816n}\right).$$

For reference, the numerical cap charges are

L	$b(L) = (L-1)(2L-3)$	$36b(L)$
2	1	36
3	6	216
4	15	540
5	28	1008
6	45	1620
7	66	2376
8	91	3276

If $L_W = L_R = L_P = L$, the worst-case linear capacity coefficient is

$$\Gamma_L(\theta, w, r) = 36b(L)(4 + 15\theta + 3s) \leq \frac{3816}{13}b(L).$$

Proof. The identities for A_{all} and the three class contributions are the formula of [Corollary 3.60](#) divided by n , using $p_{13} = \theta n$, $\sigma_W = wn$, $\sigma_R = rn$, and $\sigma(G) = sn$. The inequality $\theta \leq 1/13$ follows from the vertex-disjointness of the packed P_{13} windows, and $s \leq 1$ follows from $\sigma(G) = 2m - 3n \leq n$ in [Lemma 3.3](#).

Since $D_{\text{all}} = (N_*(G) - A_{\text{all}})_+$ and $A_{\text{all}} \geq 0$, the identity $([X]_+ - A)_+ = [X - A]_+$ gives the displayed exact formula for D_{all} . Replacing $E_{\text{spine}}(n)$ by the upper bound $C_E n$ gives the next inequality.

The slot count in [Corollary 3.60](#) is

$$36(4n + 15p_{13} + 3\sigma(G)) = n(144 + 540\theta + 108s) = n\Delta(\theta, s).$$

Using $\theta \leq 1/13$ and $s \leq 1$ gives

$$\Delta(\theta, s) \leq 144 + \frac{540}{13} + 108 = \frac{3816}{13}.$$

Substitution into the role-fibre overload bound of [Corollary 3.60](#) proves the displayed load and matching-star estimates. The numerical table is obtained by substituting $L = 2, \dots, 8$ into $b(L) = (L-1)(2L-3)$. $\square$

Proposition 3.62 (Single-graph sparse-pressure routing). *Assume G reaches node [125], satisfies the structural constraints recorded at nodes [126]–[136], has $\mathcal{A}_0 = \mathcal{P}_{\text{exc}}$, and admits a baseline spine demand with $E_{\text{spine}}(n) \leq C_E n$. Fix proposed homogeneous caps L_W, L_R, L_P , and let b_W, b_R, b_P be as in [Corollary 3.60](#). Put*

$$\Gamma_{\max} := 144b_P + \frac{540}{13}b_W + 36 \max\{b_W + 2b_P, b_R + 2b_P\}.$$

Define

$$R_L(n) := \frac{1 + \log_2 n + \sqrt{(1 + \log_2 n)^2 + 8(C_E + \Gamma_{\max})n + 8 \log_2 n}}{2}.$$

Then the sparse-pressure survivor has the following exhaustive alternatives.

- (a) If no token in $\mathfrak{T}_W, \mathfrak{T}_R, \mathfrak{T}_{\text{prim}}$ supports the corresponding role-homogeneous L_W, L_R, L_P matching or star, then

$$\sigma(G) \leq R_L(n).$$

In particular, for all sufficiently large n , this alternative routes to node [138], the near-cubic spine.

- (b) If $\sigma(G) > R_L(n)$, then $D_{\text{all}} > 0$. Hence, by [Corollary 3.61](#), G contains a role-homogeneous same-token matching or star in one of the three token classes. According to the class of the token, G is routed to node [140], [142], or [143].
- (c) Therefore no graph can remain at node [137]. An alleged survivor which is neither routed to [138] nor to one of [140], [142], [143] must violate at least one of the structural inputs used in the derivation: the sparse envelope $m \leq 2n - 2$ at [126], the full active family and baseline demand at [129], the exact window-join identity at [135], or the exact role-fibre partition at [136].

Proof. Assume first that the three homogeneous caps hold. Then each role fibre in class C has load at most b_C , by [Lemma 3.53](#); hence the total pair load in the extended token-role fibres is at most A_{all} . Since [Theorem 3.57](#) gives total-pair demand at least $N_*(G)$, we have $N_*(G) \leq A_{\text{all}}$. Thus

$$\binom{\sigma(G)}{2} \leq E_{\text{spine}}(n) + \left(\frac{\sigma(G)}{2} + 1 \right) \log_2 n + A_{\text{all}}.$$

By [Corollary 3.61](#),

$$A_{\text{all}} = n\Gamma_L(\theta, w, r), \quad \Gamma_L(\theta, w, r) = 144b_P + 540b_W\theta + 36(b_W + 2b_P)w + 36(b_R + 2b_P)r.$$

The structural constraints at [126] and the disjointness of the packed P_{13} windows give

$$0 \leq \theta \leq \frac{1}{13}, \quad w + r = \frac{\sigma(G)}{n} \leq 1.$$

Therefore

$$\Gamma_L(\theta, w, r) \leq 144b_P + \frac{540}{13}b_W + 36 \max\{b_W + 2b_P, b_R + 2b_P\} = \Gamma_{\text{max}}.$$

Using $E_{\text{spine}}(n) \leq C_E n$ gives

$$\frac{\sigma(G)(\sigma(G) - 1)}{2} \leq (C_E + \Gamma_{\text{max}})n + \left(\frac{\sigma(G)}{2} + 1 \right) \log_2 n.$$

After multiplying by 2 and moving the $\sigma(G)$ -terms to the left,

$$\sigma(G)^2 - (1 + \log_2 n)\sigma(G) - 2(C_E + \Gamma_{\text{max}})n - 2\log_2 n \leq 0.$$

The positive root of the corresponding quadratic is precisely $R_L(n)$, so $\sigma(G) \leq R_L(n)$. This proves (a). Since C_E and Γ_{max} are fixed once the caps are fixed, $R_L(n) = O(\sqrt{n} + \log n) = O(\sqrt{n})$; for large n , this is the near-cubic route [138].

For (b), take the contrapositive of (a). If $\sigma(G) > R_L(n)$, the three homogeneous caps cannot all hold. Equivalently the coupled cap inequality is violated. Quantitatively, the displayed quadratic inequality is false, so

$$\binom{\sigma(G)}{2} - E_{\text{spine}}(n) - \left(\frac{\sigma(G)}{2} + 1 \right) \log_2 n > A_{\text{all}},$$

and hence $D_{\text{all}} > 0$. [Corollary 3.61](#) then produces a role fibre in class W , R , or P with load above its cap; [Lemma 3.53](#) turns that load into the corresponding role-homogeneous matching or star. The class labels are exactly the diagram branches [140], [142], and [143].

Part (c) is the logical union of (a) and (b). Every inequality used above is one of the stated structural inputs: $m \leq 2n - 2$ gives $\sigma(G)/n \leq 1$, the full active family and baseline demand give the lower demand $N_*(G)$, the window-join identity gives the W -token supply, and the total role-fibre partition gives the exact complete-pair sum. If a graph avoids all four routed outcomes, at least one of those inputs must fail. $\square$

Theorem 3.63 (Tokenized surplus accounting closure). *Let $\mathcal{A}_0$ be a finite family of active surplus demands such that*

$$|\mathcal{A}_0| \geq c_1 \sigma(G)$$

for a fixed $c_1 > 0$. Suppose the total token ledger satisfies

$$\widehat{\ell}_{\text{cap}}(t) \leq M_0 \quad (t \in \mathfrak{T}_{\text{cap}})$$

for a fixed constant M_0 . Then

$$\sigma(G) = O(\sqrt{n}),$$

where the implicit constant depends only on c_1 and M_0 .

Proof. By [Lemma 3.43](#) and the assumed load cap,

$$\binom{|\mathcal{A}_0|}{2} = \sum_{t \in \mathfrak{T}_{\text{cap}}} \widehat{\ell}_{\text{cap}}(t) \leq M_0 |\mathfrak{T}_{\text{cap}}| \leq M_0 (8n + \sigma(G)).$$

The last inequality is [Lemma 3.44](#). Moreover $\sigma(G) \leq n$ by [Lemma 3.3](#), so the right-hand side is at most $9M_0n$. For all sufficiently large $\sigma(G)$,

$$\binom{|\mathcal{A}_0|}{2} \geq \frac{c_1^2}{4} \sigma(G)^2.$$

Combining the displays gives $\sigma(G)^2 \leq 36M_0c_1^{-2}n$ outside a fixed finite range. Enlarging the constant to cover that range proves $\sigma(G) = O(\sqrt{n})$. $\square$

Lemma 3.64 (Exact capacity-token load alternative). *Let $\mathcal{A}_0$ be a finite family of active surplus demands, and put*

$$s := |\mathcal{A}_0|, \quad L_{\text{max}} := \max_{t \in \mathfrak{T}_{\text{cap}}} \widehat{\ell}_{\text{cap}}(t).$$

Then

$$\binom{s}{2} \leq L_{\text{max}} |\mathfrak{T}_{\text{cap}}|.$$

Equivalently,

$$L_{\text{max}} \geq \frac{\binom{s}{2}}{|\mathfrak{T}_{\text{cap}}|}.$$

Using [Lemma 3.44](#), this implies the coarser but denominator-free bound

$$L_{\text{max}} \geq \frac{\binom{s}{2}}{8n + \sigma(G)}.$$

In particular, for the full active family $\mathcal{A}_0 = \mathcal{P}_{\text{exc}}$, if $\sigma(G) \geq 2$, then

$$L_{\text{max}} \geq \frac{\sigma(G)^2}{4(8n + \sigma(G))}.$$

Proof. The total ledger gives the exact identity

$$\binom{s}{2} = \sum_{t \in \mathfrak{T}_{\text{cap}}} \hat{\ell}_{\text{cap}}(t) \leq L_{\text{max}} |\mathfrak{T}_{\text{cap}}|.$$

This is [Lemma 3.43](#). Rearranging gives the exact lower bound for L_{max} . The denominator-free version follows from $|\mathfrak{T}_{\text{cap}}| \leq 8n + \sigma(G)$. For the full active family, [Lemma 3.32](#) gives $s = \sigma(G)$, and $\binom{\sigma(G)}{2} \geq \sigma(G)^2/4$ for $\sigma(G) \geq 2$. $\square$

Definition 3.65 (Sparse surplus-pair response coordinates). *Let $\mathcal{A}_0$ be a finite family of active surplus demands. For an unordered pair $\pi = \{p, q\} \subseteq \mathcal{A}_0$, let X_π be the proof-selected minimum-distance connector support of [Definition 3.33](#). Since $T(p) \subseteq \Gamma(p)$ and $T(q) \subseteq \Gamma(q)$, it contains all declared data of both demands. Let*

$$\partial X_\pi = \{v \in V(X_\pi) : v \text{ is incident with an edge of } G - X_\pi\}.$$

The sparse pair-response coordinate r_π is the declared coordinate of $\rho_{\partial X_\pi}^{\text{ex}}(X_\pi)$ whose label is π , whose support is X_π , and whose value is the exact target-response data of the two labelled demands p and q inside X_π . This coordinate is part of the declared signature by clauses (D2), (D7), and (D8) of [Definition 3.94](#).

For a nonempty set $\Pi \subseteq \binom{\mathcal{A}_0}{2}$, order its pairs lexicographically as $\pi_1 < \dots < \pi_k$, put $Y_1 = X_{\pi_1}$, and set

$$Y_{i+1} = J_G(Y_i, X_{\pi_{i+1}}), \quad X_\Pi := Y_k.$$

Thus the joint support uses $k-1$ deterministic multi-source BFS joins and no minimum-connected-subgraph search. For $\Pi = \emptyset$, take $X_\Pi = \emptyset$ and the coordinate family empty. Each r_π is then viewed as a labelled coordinate of $\rho_{\partial X_\Pi}^{\text{ex}}(X_\Pi)$ by restriction to its declared support.

Remark 3.66 (Finite surplus-slot stratification). *Before constructing the response coordinates of [Definition 3.65](#), one may separate the degenerate state space from the genuine pair branch using only the labelled surplus-slot family $\mathcal{P}_{\text{slot}}$. Since*

$$|\mathcal{P}_{\text{slot}}| = \sigma(G),$$

either $\sigma(G) \leq 1$, or there are two distinct labelled slots $p, q \in \mathcal{P}_{\text{slot}}$. This preliminary split requires one scan of the slot list; it does not enumerate $\binom{\mathcal{P}_{\text{slot}}}{2}$, build the supports $X_{\{p,q\}}$, or decide whether the resulting pair is free or blocked. Those semantic constructions belong to the subsequent branches of node [130].

Lemma 3.67 (Sparse pair quotient audit and admissible independence). *Let $\mathcal{A}_0$ be a finite family of active surplus demands and let $\Pi \subseteq \binom{\mathcal{A}_0}{2}$. Suppose the coordinate family*

$$\mathcal{R}_\Pi = \{r_\pi : \pi \in \Pi\}$$

is evaluated in $\rho_{\partial X_\Pi}^{\text{ex}}(X_\Pi)$. A raw quotient proposal that identifies two distinct members of $\mathcal{R}_\Pi$ either fails the boundary or context-universal audit, producing a record of type (d) or (e), or is admissible and carries the certified strictly smaller representative required by [Definition 3.98](#). In the latter case minimality gives a contradiction. Consequently every admitted quotient is label-injective on $\mathcal{R}_\Pi$.

Proof. Let q identify distinct coordinates r_π and $r_{\pi'}$. If their boundary-degree functions differ, choose one unequal coordinate; together with the two compared states and their values this is exactly the type-(d) audit record. Otherwise the boundary profile is preserved. If some supplied outside context distinguishes their target responses, that context and the unequal response values form the type-(e) audit record. This is precisely the negation of context-universal target completeness.

It remains to consider a proposal that passes both audits. By [Definition 3.98](#), a non-injective proposal is admitted as a rank reduction only together with a strictly smaller proper or closed representative. The representative preserves the minimum-degree baseline and transports every target certificate back to G . Since G itself avoids the target, the representative also avoids it. It is therefore a strictly smaller counterexample, contradicting the minimality of G . Thus an admitted proposal cannot identify two distinct coordinates. Notice that this argument uses only the supplied representative certificate and minimality; it does not enumerate proposals, contexts, graphs, or paths. $\square$

Proposition 3.68 (Sparse pair quotient-rank dichotomy). *Let $\mathcal{A}_0$ be a finite family of active surplus demands. If G survives the sparse surplus exits and no pair in $\binom{\mathcal{A}_0}{2}$ has a sparse surplus blocker, then the full family*

$$\{r_{\{p,q\}} : \{p,q\} \in \binom{\mathcal{A}_0}{2}\}$$

is label-injective under every admitted rank quotient. No entropy or simultaneous Boolean-realization conclusion is asserted or used.

Proof. Apply [Lemma 3.67](#) with $\Pi = \binom{\mathcal{A}_0}{2}$. The no-blocker hypothesis places every pair on the free side of the canonical ledger. The conclusion is label injectivity for every admitted quotient. The downstream proof uses the literal pair partition and capacity-token ledger, not an entropy interpretation of this quotient-rank statement. $\square$

Definition 3.69 (Baseline spine demand). *Put*

$$N = \binom{n}{2}, \quad m_0 = \left\lceil \frac{3}{2}n \right\rceil, \quad B_0(n) = \log_2 \binom{N}{m_0}.$$

A family $\mathcal{I}_{\text{spine}}$ of declared target coordinates is a baseline spine demand with deficit $E_{\text{spine}}(n)$ if $\mathcal{I}_{\text{spine}}$ is independently target-testable and

$$|\mathcal{I}_{\text{spine}}| \geq B_0(n) - E_{\text{spine}}(n).$$

In applications, $\mathcal{I}_{\text{spine}}$ is the union of the independent target coordinates already forced by the near-cubic spine before sparse surplus-pair coordinates are added. The deficit $E_{\text{spine}}(n)$ measures the unused room in the near-cubic skeleton budget after the window, remainder, curvature, and local-residual demands have been imposed.

Lemma 3.70 (Exact cubic baseline budget). *With*

$$N = \binom{n}{2}, \quad m_0 = \left\lceil \frac{3}{2}n \right\rceil, \quad B_0(n) = \log_2 \binom{N}{m_0},$$

one has

$$B_0(n) = \frac{3}{2}n \log_2 n + O(n).$$

Proof. For the upper bound, use

$$\binom{N}{m_0} \leq \left(\frac{eN}{m_0}\right)^{m_0}.$$

Since $N = \frac{1}{2}n^2 + O(n)$ and $m_0 = \frac{3}{2}n + O(1)$, the ratio eN/m_0 is $\Theta(n)$, and therefore

$$\log_2 \binom{N}{m_0} \leq \left(\frac{3}{2}n + O(1)\right) (\log_2 n + O(1)) = \frac{3}{2}n \log_2 n + O(n).$$

For the lower bound, the standard product estimate

$$\binom{N}{k} = \prod_{i=0}^{k-1} \frac{N-i}{k-i} \geq \left(\frac{N}{k}\right)^k \quad (1 \leq k \leq N)$$

applied with $k = m_0$ gives

$$\log_2 \binom{N}{m_0} \geq m_0 \log_2 \left(\frac{N}{m_0}\right) = \left(\frac{3}{2}n + O(1)\right) (\log_2 n + O(1)) = \frac{3}{2}n \log_2 n - O(n).$$

□

Lemma 3.71 (Incremental edge-count room above cubic). *Let $m_0 = \lceil 3n/2 \rceil$, $N = \binom{n}{2}$, and let*

$$m = m_0 + s$$

with $s \geq 0$ and $m \leq 2n - 2$. Then

$$\log_2 \binom{N}{m} - \log_2 \binom{N}{m_0} \leq s \log_2 n.$$

For the minimal counterexample,

$$s \leq \frac{1}{2}\sigma(G) + 1.$$

Proof. For $s = 0$ there is nothing to prove. If $s > 0$, then

$$\frac{\binom{N}{m_0+s}}{\binom{N}{m_0}} = \prod_{i=1}^s \frac{N - m_0 - i + 1}{m_0 + i} \leq \left(\frac{N}{m_0}\right)^s \leq n^s,$$

because $N \leq n^2/2$ and $m_0 \geq 1$. Taking logarithms gives the first claim. Finally,

$$\sigma(G) = 2m - 3n = 2(m_0 + s) - 3n = 2s + (2m_0 - 3n),$$

where $2m_0 - 3n \in \{0, 1\}$. Hence $s \leq \sigma(G)/2 + 1$. □

Lemma 3.72 (Mixed sparse-spine quotient audit). *Let $\mathcal{A}_0$ be a finite family of active surplus demands, let $\mathcal{I}_{\text{spine}}$ be any family of declared baseline spine coordinates, and put*

$$\mathcal{R}_{\mathcal{A}_0} = \{r_{\{p,q\}} : \{p,q\} \in \binom{\mathcal{A}_0}{2}\}.$$

Let q be a raw quotient proposal on the exact response profile carrying

$$\mathcal{I}_{\text{spine}} \cup \mathcal{R}_{\mathcal{A}_0}.$$

If q is rank-reducing on this family, then a boundary-profile or context-universality failure supplies sparse surplus exit (b), a certified proper target-complete reduction supplies sparse surplus exit (c), and a certified closed target-complete reduction supplies a strictly smaller admissible closed representative. Consequently, in a lexicographically minimal counterexample surviving the sparse surplus exits, every functional admissible rank quotient is label-injective on the displayed family.

Proof. If the proposal identifies states with different boundary degree profiles, then [Lemma 3.101](#) forbids a target-complete identification. If the determined coordinate is a pair coordinate $r_{\{p,q\}}$, the offending boundary-degree entry is the raw type-(d) audit record. In every case the non-fibrewise quotient is target-defective, which is sparse surplus exit (b).

Assume from now on that the boundary degree profile is preserved. If the proposal is not target-complete against all contexts, then [Lemma 3.102](#) makes it target-defective. This is sparse surplus exit (b). For a pair coordinate, the supplied distinguishing context is exactly the raw type-(e) audit record routed by that exit.

It remains that the proposal is target-complete. If its support is proper in G , then admissibility requires a strictly smaller proper representative in the sense of [Definition 3.96](#). By [Lemma 3.103](#) and [Corollary 3.105](#), this is impossible in the minimal counterexample and is sparse surplus exit (c).

If the support is all of G , the closed exact-profile clause of [Definition 3.98](#) applies. A non-injective closed quotient is admissible only when represented by a strictly smaller admissible closed graph H . Then H is finite and simple, satisfies $\delta(H) \geq 3$, and has closed profile contained in $H_\emptyset(G)$. A dyadic cycle in H would add the corresponding empty-context target event to $H_\emptyset(H)$, impossible because $H_\emptyset(G)$ contains no such event. Hence H is a strictly smaller counterexample, contradicting minimality. If no such representative exists, the proposal is not an admissible rank reduction. Thus every functional admissible rank quotient in a surviving minimal counterexample is label-injective on the declared family. This conclusion concerns quotient rank only; it does not assert that all Boolean response vectors are realized. $\square$

Theorem 3.73 (Sharp surplus-overload audit). *Let G reach node [125], survive the sparse surplus exits, have full active family $\mathcal{A}_0 = \mathcal{P}_{\text{exc}}$, and admit a baseline spine demand with deficit $E_{\text{spine}}(n)$. Define $N_*(G)$ as in [Theorem 3.57](#). Let*

$$x := e_\times(W)/n, \quad \theta := p_{13}/n, \quad w := \sigma_W/n, \quad r := \sigma_R/n, \quad s := \sigma(G)/n = w + r.$$

For each subtype $u \in \{RW, WW, R, V, I, P\}$, let S_u be the number of capacity tokens of subtype u , and let B_u be the number of all active pairs whose total token $\hat{\Theta}$ has subtype u . Thus free-anchor pairs contribute to subtype P , while blocked pairs contribute to the subtype of their canonical capacity token. Then:

(a) *The subtype supplies are exact:*

subtype	S_u	S_u/n
RW	$e(R, W)$	$15\theta + w - 2x$
WW	$2e_\times(W)$	$2x$
R	σ_R	r
V	n	1
I	$2m$	$3 + s$
P	$\sigma(G)$	s

where $0 \leq x \leq (15\theta + w)/2$, $0 \leq \theta \leq 1/13$, and $0 \leq s \leq 1$.

(b) The complete pair demand has the exact subtype split

$$B_{RW} + B_{WW} + B_R + B_V + B_I + B_P = \binom{\sigma(G)}{2} \geq N_*(G).$$

(c) The total token supply and total role-fibre slot supply are

$$\sum_u S_u = 4n + 15p_{13} + 3\sigma(G) = n(4 + 15\theta + 3s),$$

and

$$Q_{\text{st}} \sum_u S_u = 36n(4 + 15\theta + 3s) \leq \frac{3816}{13} n.$$

(d) If $S_u > 0$, then some token of subtype u supports a role-homogeneous same-token matching or star of size at least

$$\psi\left(\frac{B_u}{Q_{\text{st}}S_u}\right).$$

Consequently, if $N_*(G) > 0$, some subtype u supports such a pattern of size at least

$$\psi\left(\frac{N_*(G)}{36(4n + 15p_{13} + 3\sigma(G))}\right) \geq \psi\left(\frac{13N_*(G)}{3816n}\right).$$

Proof. The exact window split follows from [Definition 3.41](#): RW tokens are the window-remainder edges and WW tokens are the two endpoint tokens on each cross-window edge. Hence

$$S_{RW} = e(R, W), \quad S_{WW} = 2e_{\times}(W).$$

By [Lemma 8.5](#),

$$e(R, W) = 15p_{13} + \sigma_W - 2e_{\times}(W),$$

which gives the first two normalized rows. The bound on x is the nonnegativity of $e(R, W)$. The R -row is the definition of $\mathfrak{T}_R$:

$$S_R = \sum_{v \in R} (d_G(v) - 3) = \sigma_R.$$

For primitive tokens,

$$\mathfrak{T}_{\text{prim}} = V(G) \sqcup I_E(G) \sqcup \mathcal{P}_{\text{exc}},$$

so the subtype supplies are

$$S_V = n, \quad S_I = 2m, \quad S_P = |\mathcal{P}_{\text{exc}}| = \sigma(G).$$

Since $\sigma(G) = 2m - 3n$, one has $2m = 3n + \sigma(G)$, giving $S_I/n = 3 + s$. The inequalities $0 \leq \theta \leq 1/13$ and $0 \leq s \leq 1$ follow respectively from the vertex-disjointness of the packed P_{13} windows and from [Lemma 3.3](#).

The total map $\hat{\Theta}$ assigns every active pair to exactly one capacity token by [Definition 3.42](#). Its blocked side is Θ_{cap} , and its free side is the first selected surplus port, which has subtype P . The six subtypes listed in [Definition 3.47](#) partition the capacity-token universe. Hence their fibres partition the complete active-pair family. Since $|\mathcal{A}_0| = \sigma(G)$, [Lemma 3.43](#) gives

$$\sum_u B_u = \binom{\sigma(G)}{2} \geq N_*(G).$$

Summing the six supply rows gives

$$\sum_u S_u = e(R, W) + 2e_x(W) + \sigma_R + n + 2m + \sigma(G).$$

Using Lemma 8.5, $\sigma(G) = \sigma_W + \sigma_R$, and $2m = 3n + \sigma(G)$, this is

$$15p_{13} + \sigma_W + \sigma_R + n + (3n + \sigma(G)) + \sigma(G) = 4n + 15p_{13} + 3\sigma(G).$$

Multiplication by $Q_{\text{st}} = 36$ gives the role-fibre slot count. Since $\theta \leq 1/13$ and $s \leq 1$,

$$36(4 + 15\theta + 3s) \leq 36\left(4 + \frac{15}{13} + 3\right) = \frac{3816}{13}.$$

Fix a subtype u with $S_u > 0$. Among the S_u tokens of this subtype, some token t has load at least B_u/S_u . Its role fibres partition its load into at most Q_{st} parts, so some role fibre has size at least $B_u/(Q_{\text{st}}S_u)$. Applying Lemma 3.53 to that role-fibre graph gives a role-homogeneous same-token matching or star of size at least $\psi(B_u/(Q_{\text{st}}S_u))$.

If $N_*(G) > 0$, choose a subtype with maximal ratio B_u/S_u among nonempty subtypes. The preceding partition gives

$$\frac{B_u}{S_u} \geq \frac{\sum_v B_v}{\sum_v S_v} \geq \frac{N_*(G)}{4n + 15p_{13} + 3\sigma(G)}.$$

The displayed global lower bound follows by monotonicity of ψ . $\square$

Definition 3.74 (Same-token routing germs). *Let $t \in \mathfrak{T}_{\text{cap}}$, let $r \in \hat{\mathfrak{R}}_{\text{st}}$, and let $\pi = \{p, q\} \in \hat{\Pi}_{t,r}$. The same-token routing support*

$$Z(\pi; t, r)$$

is constructed from the token support and $\Gamma(p), \Gamma(q)$, together with the declared support of the tagged canonical blocker $\Phi_{\text{can}}(\pi)$ on the blocked-role branch. On the freeAnchor branch, use instead the canonical connector support X_π already retained at node [130]. Order these finite supports and all their connected components by their declared labels, start with the first component, and join each subsequent component by J_G from Definition 3.33. The resulting connected BFS closure is $Z(\pi; t, r)$; hence it is exact and requires no subgraph search.

The token root $c(t)$ is the window endpoint of an RW-token, the marked endpoint of a cross-window token, the vertex v of a remainder token (v, j) , the vertex of a primitive vertex token, the marked end of a primitive edge-incidence token, or the buffer $x(p)$ of a primitive selected-port token. For $u \in \{p, q\}$, the routing germ $g(\pi, t, u)$ is the lexicographically first shortest path in $Z(\pi; t, r)$ from $c(t)$ to $T(u)$, stopped at its first vertex in $T(u)$. It is obtained by one BFS inside the supplied support.

For two routing germs with the same root, delete their maximal common initial edge word. If both residual words are nonempty and their first incidences are distinct, the last common vertex is their first separator. Otherwise one whole germ is an initial segment of the other (equality included), and the germs are parallel. These alternatives are disjoint and exhaustive for finite rooted paths; in particular the case in which one germ reaches its labelled port before the other is not omitted.

The routing label of a pair in $\Pi_{t,r}$ records the same-token role r , the subtype of t , the ordered endpoint of the pair under discussion, the local open/triangular status of the corresponding selected ports, the terminal role of each canonical germ in $\{\text{leftShoulder}, \text{rightShoulder}, \text{buffer}\}$, the boundary-degree profile of the bounded port supports $T(p), T(q)$, the P_{13} -label entries appearing in the bounded part of the support, and the suppressed-chord flag when the blocked role has type (f), and a distinct free-anchor flag otherwise. These labels form a finite set; denote its cardinality by Q_{geom} .

Lemma 3.75 (Fixed coarse bottleneck classification). *Let $\mathcal{M} = (\pi_0, \dots, \pi_{48})$ be the literal role-homogeneous matching or star supplied by one of nodes [140], [142], or [143]. Give each pair its coarse structural code consisting of its token subtype, endpoint role, open/triangular bit, and terminal germ role. This code has exactly 48 values. Then the deterministic scan of $\mathcal{M}$ returns two distinct pairs π_i, π_j with the same coarse code. It retains their literal selected-window attachment maps and the two canonical shortest rooted germs.*

Proof. Scan the fixed list in its inherited order, comparing each code with the remaining suffix. If no equality occurred, the 49 observed codes would be distinct elements of a 48-element set. Thus the scan returns its first collision. The homogeneous-pattern list is duplicate free, so the two source pairs are distinct. Attachment maps are projections of those two literal pairs, and each germ is obtained by the canonical BFS in its supplied local support. No family of paths, contexts, quotients, subgraphs, graphs, or colorings is inspected. $\square$

This classification does not prove equality of the attachment maps, equality of external responses, a sparse-exit certificate, or a decorated Type B entry. Its exact residual is passed to the finite classifier at node [144].

Lemma 3.76 (Finite semantic-bottleneck classification). *For the exact collision and two rooted germs retained at node [144], scan the declared coordinates*

$$\mathfrak{A} = \mathfrak{W} \times [13] \times \{0, 1\} \times \{\text{freeAnchor}, \text{blockedPrimitive}, \text{blockedCapacity}\}.$$

Thus $|\mathfrak{A}| = 78p_{13}$. The scan returns either its first coordinate at which the two literal attachment predicates differ, or a proof that they agree on every coordinate of $\mathfrak{A}$. In the latter case, the stored tree-path comparison of the two canonical germs is exactly one of

$$\text{leftPrefix}, \quad \text{rightPrefix}, \quad \text{divergeAtRoot}, \quad \text{divergeAfterEdge}.$$

These five outcomes form an exhaustive CT10 table with exactly one accepted tag. The complete execution uses

$$234p_{13} + 7 \leq 234|V(G)| + 7$$

primitive checks.

Proof. Compare the two retained adjacency predicates in the inherited order on $\mathfrak{A}$. A first failed Boolean equality supplies the mismatch coordinate. If no comparison fails, the scan's exhaustive alternative gives coordinatewise equality and hence equality of the retained attachment maps on their whole declared domain. The canonical rooted-germ comparison is an inductive tree-path comparison with precisely the four constructors displayed above. CT10 scans the fixed five-tag table and accepts only the tag returned by this classifier. Each attachment coordinate costs the two predicate projections and their equality check, while the fixed table contributes seven checks. Finally $p_{13} \leq |V(G)|$ because the selected induced paths are vertex-disjoint. No ambient family of paths, contexts, quotients, subgraphs, graphs, or colorings is enumerated. $\square$

Remark 3.77 (Scope of the finite semantic classifier). *Node [144] proves only the preceding five-way local classification. It does not turn a mismatch or a rooted-germ shape into a sparse exit, CT3 response equivalence, a decorated Type B handoff, a fixed homogeneous cap, or the near-cubic spine. Node [144] performs only the verified local separator split below; the stronger branchwise implications remain open at node [144].*

Lemma 3.78 (Local semantic-bottleneck separator split). *Consume node [144]’s exact five-way result. A mismatch or either prefix constructor is retained verbatim. In a root-divergence constructor, the two literal root incidences and the first actual third neighbour are retained; in an after-edge constructor, three pairwise-distinct literal incidences are retained. In either divergent case the separator is classified as cubic or of degree at least four. This exhaustive local decision uses at most $|V(G)| + 1$ primitive checks.*

Proof. Inspect the constructor returned at node [144]. The first three constructors require no new search. Root divergence scans the actual ordered neighbour row only until the first neighbour distinct from the two retained incidences; minimum degree three guarantees its existence. After-edge distinctness follows from the two stored canonical tree paths and their no-repetition proofs. The final comparison is the literal separator degree against three. No ambient family is enumerated. $\square$

Lemma 3.79 (Cubic switch normalization). *Consume node [144]’s exact local frontier. Mismatch and aligned-prefix leaves are retained unchanged. Each cubic root or after-edge leaf determines the literal four-vertex cubic-star datum carried by its already proved three pairwise-distinct incidences, while each high leaf retains only degree at least four. The normalization is exhaustive and performs no new primitive check.*

Lemma 3.80 (Local separator projection). *Consume node [144] exactly. Each cubic leaf projects its literal cubic-star datum to the corresponding switch-boundary shape with the same four-vertex support. Each high leaf exposes only the identical center’s declared port row, whose cardinality is its degree and is at least four. Other leaves pass through unchanged. Work is at most $|V(G)|$, with no ambient scan.*

Lemma 3.81 (Strong-semantic obligation frontier). *Consume node [144] exactly and retain its literal payload. Tag mismatch, aligned-prefix, cubic, and high-center leaves respectively with the pending sparse-exit, fixed-cap, CT3, and Type B obligations. These tags are requests, not certificates. The construction inspects one local constructor and proves none of the four requested semantic conclusions.*

Lemma 3.82 (First local separator clause). *Consume node [144] exactly. Mismatch and prefix leaves retain their evidence. Each cubic leaf exposes its three injective adjacent boundary incidences, and each high leaf exposes the first four distinct declared ports and their adjacent endpoints. At most four fixed local positions are inspected. No pending semantic certificate is discharged.*

Lemma 3.83 (Pairwise local separator clause). *Consume node [144] exactly and retain its pending obligation unchanged. On a cubic leaf, the three boundary vertices are pairwise distinct and the internal vertex differs from each boundary vertex. On a high leaf, the four exposed endpoints are pairwise distinct and the centre differs from each endpoint. Mismatch and aligned-prefix leaves pass through unchanged. These conclusions use only inherited injectivity and adjacency, require no visible search, and do not discharge the retained semantic obligation.*

Lemma 3.84 (Same-token bottleneck routing). *Assume G survives the sparse surplus exits of Definition 3.27. Let $t \in \mathfrak{T}_{\text{cap}}$ and $r \in \hat{\mathfrak{R}}_{\text{st}}$. If a total-role-homogeneous same-token matching or star in $\hat{H}_{t,r}$ has more than Q_{geom} edges, then one of the following occurs:*

- (a) *a sparse surplus exit occurs; or*
- (b) *the support produces a decorated Type B handoff fan envelope, hence is routed to the Type B fan ledger.*

Proof. Let $\mathcal{M}$ be the matching or star. Since there are only Q_{geom} routing labels, two distinct edges $\pi_1, \pi_2 \in \mathcal{M}$ have the same routing label. If $\mathcal{M}$ is a star, choose the two noncentral endpoints of π_1 and π_2 ; if $\mathcal{M}$ is a matching, choose an endpoint from each edge. In both cases we obtain two distinct selected surplus demands whose connector data have the same token, the same total role, and the same boundary-degree fibre. On a blocked-role fibre this includes the same blocker type and token subtype. On the free-anchor fibre the common token is literally the first selected port of every pair, so the selected edges already form a star at that port.

Consider the two canonical routing germs from the token root $c(t)$ toward those two selected demands. If the germs are parallel, then the two declared response coordinates have the same bounded port data and terminal role and lie in the same boundary-degree fibre. By the exhaustive definition of parallel, either the rooted paths agree or one ends first. In the latter case restrict the longer path to the shorter terminal vertex; this is the finite endpoint-labelled connector coordinate recorded in the declared signature. Identifying the two endpoint-labelled coordinates is a nontrivial quotient of declared coordinates. If some outside context distinguishes the two responses, the quotient is target-defective by [Lemma 3.102](#), which is a sparse surplus exit. If the identification is target-complete on a proper support, then [Lemma 3.103](#) and [Corollary 3.105](#) give the target-complete compression exit. If the identification becomes valid only after adjoining a larger connected support, then [Lemmas 10.6](#) and [10.9](#) give the proper or global delocalization exit. In the closed exact-profile case, [Definition 3.98](#) permits a non-injective quotient only when there is a smaller admissible closed representative; that representative would be a smaller counterexample. Thus the parallel case is impossible in a survivor.

It remains that the two germs have a first separator z . If $d_G(z) = 3$, the root incidence of the common prefix together with the two distinct next incidences exhausts all incidences at z . The finite switch support at z therefore has no unused ambient incidence through which the two coordinates can be distinguished without using the declared quotient data. There are then exactly three target-completeness possibilities for the switch quotient. If some compatible outside context distinguishes the two switched coordinates, the quotient is target-defective by [Lemma 3.102](#); this is sparse surplus exit (b) in [Definition 3.27](#). If every compatible context preserves the identification on the finite switch support, then the two same-interface coordinates give a nontrivial target-complete compression of a proper support, contradicting [Lemma 3.103](#) and [Corollary 3.105](#) and entering sparse surplus exit (c). If context-universality holds only after adjoining a larger connected support, the support is a proper or global delocalization; these are excluded by [Lemmas 10.6](#) and [10.9](#) and enter sparse surplus exit (d). Thus a cubic first separator cannot survive, and any surviving first separator must satisfy

$$d_G(z) \geq 4.$$

The high-degree separator z , together with the separated connector tails from z to the two selected surplus demands, is exactly the decorated handoff data of [Definition 14.19](#). The fan-safety condition for this envelope says that every pair of separated tails has no dyadic two-tail closure, has compatible P_{13} -window labels, and lies in one boundary-degree fibre of the declared response algebra. If a two-tail closure has dyadic length, then sparse surplus exit (a) of [Definition 3.27](#) occurs. If the closure is not dyadic but the two P_{13} -labels are incompatible, the P_{13} -label collision is the finite label exit already included in sparse surplus exit (a). If the labels are compatible but the boundary-degree fibre or target response is separated by an outside context, the quotient is target-defective, giving exit (b). If the same response is target-complete on a proper finite support, [Lemma 3.103](#) and [Corollary 3.105](#) give exit (c). If it becomes target-complete only after adjoining a larger connected support, [Lemmas 10.6](#) and [10.9](#) give exit (d). These alternatives exhaust every failure of fan-safety. For a blocker of type (f), the only additional possibility is that

the common support is the suppressed open-port chord set; if the chord arithmetic produces a power-of-two length, exit (e) of [Definition 3.27](#) occurs, and otherwise the two open-port shoulders are assigned as fan incidences by [Lemma 3.11](#) and [Corollary 14.81](#) or, in the triangular alternative, by [Corollary 14.82](#). Thus every surviving separated case enters the Type B fan ledger. $\square$

Theorem 3.85 (Geometric closure of homogeneous same-token overloads). *Assume G reaches node [125], survives the sparse surplus exits, has full active family $\mathcal{A}_0 = \mathcal{P}_{\text{exc}}$, and admits a baseline spine demand with $E_{\text{spine}}(n) \leq C_E n$. Put*

$$L_{\text{geom}} := Q_{\text{geom}} + 1.$$

Then every total-role-homogeneous same-token matching or star of size L_{geom} in the window-incidence, remainder-surplus, or primitive-carrier token classes realizes either a sparse surplus exit or a decorated Type B handoff fan envelope.

Consequently, on the subbranch in which sparse surplus exits are absent and all decorated Type B handoff data have been routed into the Type B fan ledger, the three fixed homogeneous caps

$$L_W = L_R = L_P = L_{\text{geom}}$$

hold. On that subbranch,

$$\sigma(G) = O(\sqrt{n}) \quad \text{and} \quad m = \frac{3}{2}n + O(\sqrt{n}).$$

Proof. The first assertion is exactly [Lemma 3.84](#), applied to the role fibre containing the alleged matching or star.

If sparse exits are absent and decorated Type B handoff data have already been routed to the Type B fan ledger, the first assertion says that no capacity token in any of the three classes can support a total-role-homogeneous same-token L_{geom} -matching or L_{geom} -star. Therefore the hypotheses of the bounded alternative in [Proposition 3.62](#) hold with $L_W = L_R = L_P = L_{\text{geom}}$. That alternative gives $\sigma(G) = O(\sqrt{n})$. Since

$$m = \frac{3}{2}n + \frac{1}{2}\sigma(G)$$

by [Lemma 3.26](#), the edge-count estimate follows. $\square$

Remark 3.86 (Exact surplus frontier). *The sparse surplus bookkeeping has no disjointness loss. The active family is $\mathcal{A}_0 = \mathcal{P}_{\text{exc}}$ by [Lemma 3.32](#). Free pairs use the free-anchor route of [Definition 3.42](#). Blocked pairs are assigned a canonical blocker and then one capacity token by [Definition 3.41](#). The total token-role fibres partition the complete pair set by [Lemma 3.43](#). Therefore any attempted non-near-cubic surplus survivor must first pass through one of the six subtype total loads in [Theorem 3.73](#), grouped into the three diagram classes $W = \{RW, WW\}$, $R = \{R\}$, and $P = \{V, I, P\}$. A large load in any one of those fibres is not terminal: [Theorem 3.85](#) routes it to a sparse surplus exit, to Type B fan data, or to the fixed homogeneous caps that imply the near-cubic spine.*

Proposition 3.87 (Non-near-cubic sparse-pressure routing). *Let G be a lexicographically minimal counterexample that reaches node [18]. If G survives the sparse surplus exits, then the exact sparse-pressure audit has only the following outcomes:*

- (a) G satisfies the near-cubic spine estimate $\sigma(G) = O(\sqrt{n})$;




Figure 14: Visual definition of the homogeneous same-token bottleneck. Panel I shows active surplus pairs assigned by the total token map to one overloaded token $t \in \mathfrak{T}_{\text{cap}}$. The total-role partition of the fibre over t produces a role r for which $\hat{H}_{t,r}$ contains a large matching or a large star. Panel II displays the two possible homogeneous patterns: every shown edge has the same token t and the same total role r . Panel III shows the geometric routing of two selected demands from the same homogeneous pattern. Parallel routing germs give one of the quotient, compression, or delocalization exits already included among the sparse surplus exits. If the germs first separate at a high-degree vertex z , the separated tails form decorated Type B handoff fan data. When neither a sparse surplus exit nor Type B fan data remains, all three token classes have fixed homogeneous caps, and [Corollary 3.50](#) gives $\sigma(G) = O(\sqrt{n})$.

(b) a sparse surplus exit occurs; or

(c) a total-role-homogeneous same-token bottleneck produces decorated Type B fan data and is routed to the Type B fan ledger.

In particular, the window-incidence, remainder-surplus, and primitive-carrier same-token classes are not terminal non-near-cubic leaves.

Proof. If a sparse surplus exit occurs, there is nothing to route. Otherwise [Lemma 3.32](#) gives $\mathcal{A}_0 = \mathcal{P}_{\text{exc}}$, and [Lemma 3.43](#) supplies the exact complete pair demand used in the classwise overload audit. The blocked-pair ledger and its capacity-token subroute are exact by [Lemmas 3.38](#) and [3.45](#), and the free pairs enter the primitive selected-port class.

Apply [Theorem 3.85](#). If no homogeneous same-token pattern reaches size L_{geom} , the fixed caps $L_W = L_R = L_P = L_{\text{geom}}$ hold and [Corollary 3.50](#) gives $\sigma(G) = O(\sqrt{n})$. If such a pattern exists, then [Lemma 3.84](#) gives either a sparse surplus exit or decorated Type B handoff data. The total subtype identity in [Theorem 3.73](#) records where the load came from: RW and WW are window-incidence loads, R is a remainder-surplus load, and V, I, P are primitive-carrier loads. In all cases the load is either closed, routed to Type B, or bounded by the fixed cap that yields the near-cubic estimate. $\square$

3.1 Boundaried atoms, target-complete quotients, and replacement

Definition 3.88 (Boundaried graph and gluing). *Let T be a finite label set. A T -boundaried graph is a graph X together with an injective labelling of a subset $\partial X \subseteq V(X)$ by T . The vertices in ∂X are boundary vertices; all others are internal. If X and Y are T -boundaried graphs with disjoint*

internal vertex sets, their gluing is denoted

$$X \oplus_T Y,$$

and is obtained by identifying equally labelled boundary vertices and taking the union of edges.

Definition 3.89 (Atom). *An atom of G is a connected T -boundaried subgraph X occurring as one side of a decomposition $G = X \oplus_T Y$ along a boundary $T \subseteq V(G)$; the other side Y is its context. The atom is proper if $X \neq G$. The replacement and uncompressibility arguments of this section quantify over proper atoms.*

Definition 3.90 (Boundary degree profile). *Let X be a T -boundaried graph. For $t \in T$, let v_t be the boundary vertex of X labelled by t . The boundary degree profile of X is the vector*

$$\mathbf{d}_\partial(X) = (d_X(v_t))_{t \in T}.$$

All target-response states in this paper are taken fibrewise over $\mathbf{d}_\partial$. In particular, two boundaried states with different boundary degree profiles are never eligible to be identified by a target-complete quotient.

Definition 3.91 (Obstruction profile). *For a T -boundaried graph X , define its target obstruction profile*

$$\mathbf{H}_T(X) = \{Y \in \text{Ctx}_T : X \oplus_T Y \text{ contains a power-of-two cycle}\},$$

where Ctx_T is the class of finite T -boundaried contexts that can be glued to X . If X' and X are T -boundaried graphs, write

$$X' \preceq_T X$$

if

$$\mathbf{H}_T(X') \subseteq \mathbf{H}_T(X).$$

Thus X' is at most as obstructing as X .

Definition 3.92 (Target-complete quotient). *Let X be a T -boundaried atom, and let $\mathcal{R}(X)$ be a collection of local target-response coordinates of X , such as terminal trace lengths, edge-rooted Mersenne-return tests, curvature-flatness tests, boundary degree data, and other marked response data. A quotient of $\mathcal{R}(X)$ is target-complete if every identification made by the quotient preserves both:*

- (a) *the boundary degree profile $\mathbf{d}_\partial(X)$; and*
- (b) *the target predicate after gluing X to every T -boundaried context.*

Thus two coordinates may be identified only inside a fixed boundary-degree fibre and only when no context $Y \in \text{Ctx}_T$ creates a power-of-two cycle from one coordinate and not from the other. An identification failing this context-universal test is target-defective.

Lemma 3.93 (Composition of target-complete quotients). *Let*

$$\rho \longrightarrow Q \longrightarrow Q'$$

be boundary-degree-preserving response quotients tested against the same class of compatible boundaried contexts. If both displayed quotients are target-complete, then the composite quotient $\rho \rightarrow Q'$ is target-complete. Consequently, if $\rho \rightarrow Q$ is target-complete and the composite $\rho \rightarrow Q'$ is not target-complete, then $Q \rightarrow Q'$ is not target-complete. Moreover, any compatible context that distinguishes a realization of Q' from the state Q also distinguishes that realization from ρ after composing with the target-complete quotient $\rho \rightarrow Q$.

Proof. Fix a compatible outside context Y . Target-completeness of $Q \rightarrow Q'$ says that every realization of Q' , glued to Y , has the same target predicate as the corresponding Q -state. Target-completeness of $\rho \rightarrow Q$ says that the Q -state, glued to Y , has the same target predicate as ρ . Hence every realization of Q' has the same target predicate as ρ after gluing to Y . Since Y was arbitrary, the composite is target-complete. The contrapositive gives the second assertion. The last assertion is the same transitivity statement applied to a fixed distinguishing context. $\square$

Definition 3.94 (Declared response-coordinate signature). *The response-coordinate signature used in the proof is fixed once and for all. For a T -boundaried support X , a declared coordinate is a tuple*

$$(k, \iota, \text{supp}_X(r), \text{val}_X(r)),$$

where k is its kind, ι is its label, $\text{supp}_X(r)$ is a finite set of vertices, edges, boundary incidences, paths, or windows of X , and $\text{val}_X(r)$ is the corresponding value in the embedded support. The declared coordinates are generated by the following clauses.

- (D1) *boundary-degree entries at labelled boundary vertices;*
- (D2) *edge-rooted return data, completion-port data, first-entry receivers, connector lengths, receiver-entry channels, and the corresponding Mersenne return tests;*
- (D3) *induced- P_{13} window labels, legal-label relations, packed-window incidences, and cross-window incidence data;*
- (D4) *raw curvature coordinates indexed by internal length-two wedges;*
- (D5) *Type A trace data: canonical traces, trace-incidence coordinates, connector-band constraints, cross-port theta constraints, silent-basin response coordinates, carrier restrictions, and their labelled subcoordinates;*
- (D6) *Type B fan and bridge data: fan centers, fan-safe pairs, certificate labels, closed fan-window pairs, hybrid incidence entries, candidate ledger entries, overlap demands, decorated handoff fan response coordinates, and decorated handoff-arm coordinates;*
- (D7) *sparse surplus data: selected surplus ports, canonical port returns, open-port suppression paths, triangular-port response triangles, and sparse surplus-pair response coordinates;*
- (D8) *finite products, labelled copies, restrictions, and quotient images of coordinates already generated in clauses (D1)–(D7), with support equal to the union of the supports of the entries used.*

No other local response coordinate is available to a quotient. Whenever a later definition names a coordinate family, it names the subfamily generated by this signature with the displayed labels. In particular, a phrase such as “needed to evaluate a return” means that the relevant port, first-entry incidence, connector endpoint, channel, window, or boundary incidence is a member of the declared support specified by the corresponding clause above.

Definition 3.95 (Exact response profile). *Let X be a T -boundaried support. Let $\mathfrak{S}_T(X)$ be the finite labelled family of coordinates generated on X by [Definition 3.94](#). The exact response profile of X is*

$$\rho_T^{\text{ex}}(X) = (\mathbf{d}_\partial(X), \mathbf{H}_T(X), \mathcal{R}_T^{\text{ex}}(X)),$$

where $\mathcal{R}_T^{\text{ex}}(X) = \mathfrak{S}_T(X)$, including the raw curvature coordinates carried by X when they are present. The profile is exact in the following sense: two distinct coordinate labels remain distinct entries of $\mathcal{R}_T^{\text{ex}}(X)$ even if their numerical values in the embedded graph coincide.

A quotient

$$q : \rho_T^{\text{ex}}(X) \longrightarrow Q$$

is label-injective on a subfamily $\mathcal{A} \subseteq \mathcal{R}_T^{\text{ex}}(X)$ if every coordinate of $\mathcal{A}$ is retained by q and

$$q(a) = q(b) \quad \text{with } a, b \in \mathcal{A} \quad \implies \quad a = b.$$

It is rank-reducing on $\mathcal{A}$ if it is not label-injective on $\mathcal{A}$.

Definition 3.96 (Proper representative of an exact-profile quotient). *Let $G = X \oplus_T Y$, where $X \subsetneq G$ is a proper T -boundaried support, and let*

$$q : \rho_T^{\text{ex}}(X) \rightarrow Q$$

be a quotient of its exact profile. A T -boundaried graph X' properly represents q if there is a quotient

$$q_{X'} : \rho_T^{\text{ex}}(X') \rightarrow Q$$

with the following properties.

- (a) $\mathbf{H}_T(X') \subseteq \mathbf{H}_T(X)$.
- (b) $\mathbf{d}_{\partial}(X') = \mathbf{d}_{\partial}(X)$.
- (c) X' contains no internal power-of-two cycle.
- (d) Every internal vertex of X' has degree at least 3.
- (e) Every labelled coordinate of Q is the image under $q_{X'}$ of a coordinate generated by [Definition 3.94](#), and the quotient relations among the coordinates under discussion are the same as those imposed by q .

The representative is strictly smaller if X' is smaller than X in the lexicographic order used for the minimal counterexample.

Definition 3.97 (Closed representative of an exact-profile quotient). *Let*

$$q : \rho_{\emptyset}^{\text{ex}}(G) \rightarrow Q$$

be a quotient of the closed exact profile. A finite simple closed graph H represents q if there is a quotient

$$q_H : \rho_{\emptyset}^{\text{ex}}(H) \rightarrow Q$$

with the following properties.

- (a) The obstruction-profile component of q_H is no larger than that of G :

$$\mathbf{H}_{\emptyset}(H) \subseteq \mathbf{H}_{\emptyset}(G).$$

- (b) The boundary-degree component is the empty boundary profile.

- (c) Every labelled coordinate of Q is the image under q_H of a coordinate generated by the declared signature of [Definition 3.94](#), and the quotient relations among raw curvature coordinates recorded in Q are the same quotient relations as those imposed by q .

The representative is strictly smaller when H is smaller than G in the lexicographic minimality order. It is admissible closed when, in addition, $\delta(H) \geq 3$.

Definition 3.98 (Admissible rank quotient). *Let $\mathcal{A}$ be a family of declared response coordinates carried by a connected support $X \subseteq G$, and put*

$$T = \{v \in V(X) : v \text{ is incident with an edge of } G - X\}.$$

An admissible rank quotient for $\mathcal{A}$ is a quotient

$$q : \rho_T^{\text{ex}}(X) \longrightarrow Q$$

which preserves the boundary degree profile and is target-complete against all T -boundaried contexts in the sense above.

A map on the finite coordinate labels is first a raw quotient proposal. It becomes an admissible rank quotient only after the two preservation properties above and the representative clause below have been certified. Thus a boundary-profile mismatch and a distinguishing outside context are audit failures of a raw proposal, not branches of an already admissible quotient.

If $X \subsetneq G$, then a rank-reducing admissible quotient must be represented by a strictly smaller proper representative in the sense of [Definition 3.96](#). Equivalently, a target-complete proper-support correlation that has no smaller graph representative is not an admissible rank reduction; it is only an abstract identification of labels and does not reduce the graph-realizable target rank used in this proof.

If $T = \emptyset$ and $X = G$, the quotient is evaluated in the closed exact profile $\rho_G^{\text{ex}}(G)$. In that closed case, absence of an outside context is not an identification rule: an admissible rank quotient must be label-injective on the declared coordinates under discussion unless it is represented by a strictly smaller admissible closed representative in the sense of [Definition 3.97](#).

Thus the empty-boundary profile may record exact global information, but exact global information does not collapse two labelled response coordinates. A closed non-injective quotient can affect target rank only when it supplies a strictly smaller closed representative; such a representative would be a smaller counterexample if the target predicate is still false.

Definition 3.99 (Functional rank quotient). *Let $q : \rho_T^{\text{ex}}(X) \rightarrow Q$ be an admissible rank quotient for a finite coordinate family $\mathcal{A}$. The quotient q is functional on $\mathcal{A}$ if the following rank axiom holds. Whenever $\mathcal{I} \subseteq \mathcal{A}$ and $a \in \mathcal{A} \setminus \mathcal{I}$ are such that q is label-injective on $\mathcal{I}$, but is not label-injective on $\mathcal{I} \cup \{a\}$, there is a finite subfamily $\mathcal{B} \subseteq \mathcal{I}$ and a function ϕ on the realized q -image vectors of $\mathcal{B}$ such that*

$$q(a) = \phi((q(b))_{b \in \mathcal{B}})$$

for every realization of the exact response profile under q . If the quotient sends a to a single quotient value, then $\mathcal{B} = \emptyset$. If the quotient identifies a with a coordinate $b \in \mathcal{I}$, then $\mathcal{B} = \{b\}$ and $\phi(q(b)) = q(b)$.

The admissible quotient system used to compute target-rank consists only of admissible rank quotients that are functional on the coordinate family under discussion.

Remark 3.100 (Rank coordinates and entropy charge). *The curvature target-rank r_Ω is always computed from the declared coordinates of the exact response profile and from admissible rank quotients in the sense of [Definitions 3.95](#), [3.98](#) and [3.99](#). Thus a rank coordinate contributes to the*

entropy count only when it remains a labelled, independently target-testable coordinate after all admissible quotients have been applied. This is the interface used by [Lemma 3.107](#) and [Proposition 11.9](#): raw curvature tests are not counted twice, and an abstract identification that is not an admissible rank quotient does not reduce r_Ω .

Lemma 3.101 (Boundary degree profiles cannot be quotient-merged). *Let X_1 and X_2 be T -boundaried supports. If*

$$\mathbf{d}_\partial(X_1) \neq \mathbf{d}_\partial(X_2),$$

then no target-complete quotient identifies X_1 and X_2 .

Proof. Condition (a) in the definition of a target-complete quotient requires the quotient to preserve the boundary degree profile. If $\mathbf{d}_\partial(X_1) \neq \mathbf{d}_\partial(X_2)$, an identification of X_1 with X_2 would identify two different boundary-degree profiles, so it violates condition (a). Thus two supports lying in different boundary-degree fibres are not eligible for identification, independently of the context-universal target test in condition (b). $\square$

Lemma 3.102 (Context-universality of target-complete identifications). *Let X be a T -boundaried atom, and let $\mathcal{R}(X)$ be any collection of local target-response coordinates of X . Suppose that two coordinates $r_1, r_2 \in \mathcal{R}(X)$ are identified in a target-complete quotient of X . Then r_1 and r_2 have the same target response against every T -boundaried context Y .*

Equivalently, for every $Y \in \text{Ctx}_T$, gluing Y to X cannot distinguish r_1 from r_2 by creating a power-of-two cycle in one case and not in the other. Consequently any identification valid only for the actual outside context $G - X$, but not for all T -boundaried contexts, is target-defective.

Proof. This is precisely the meaning of target-completeness. If some context Y_0 distinguished r_1 and r_2 , then one of the two gluings would contain a power-of-two cycle and the other would not. The quotient would fail to preserve the target predicate for Y_0 and therefore would not be target-complete. $\square$

Lemma 3.103 (Replacement lemma for boundaried atoms). *Let $G = X \oplus_T Y$, where X is a proper atom of G . Suppose there exists a T -boundaried graph X' such that:*

- (i) $X' \preceq_T X$;
- (ii) $\mathbf{d}_\partial(X') = \mathbf{d}_\partial(X)$;
- (iii) X' contains no internal power-of-two cycle;
- (iv) every internal vertex of X' has degree at least 3;
- (v) X' is strictly smaller than X in the same lexicographic order used for G .

Then G was not minimal. Consequently no such replacement exists in G .

Proof. Form

$$G' = X' \oplus_T Y.$$

Because $\mathbf{d}_\partial(X') = \mathbf{d}_\partial(X)$, every boundary vertex has the same final degree in G' as it had in G . Vertices of Y are unchanged, and internal vertices of X' have degree at least 3 by hypothesis. Hence $\delta(G') \geq 3$.

We claim that G' has no power-of-two cycle. Let C be a cycle of G' , and suppose for contradiction that its length is a power of two. If C is wholly contained in Y , then C is also a cycle of

G , contradicting that G is a counterexample. If C is wholly contained in X' , then C is an internal power-of-two cycle of X' , forbidden by hypothesis (iii). Otherwise C uses internal vertices of both X' and Y ; since X' and Y meet only in the boundary ∂ , the cycle C lives in $G' = X' \oplus_T Y$ and contains a power-of-two cycle, so by definition of the obstruction profile $Y \in H_T(X')$. Hypothesis (i) gives $X' \preceq_T X$, that is $H_T(X') \subseteq H_T(X)$, whence $Y \in H_T(X)$. By definition this means $X \oplus_T Y = G$ contains a power-of-two cycle, contradicting that G is a counterexample. Hence G' has no power-of-two cycle. Together with $\delta(G') \geq 3$ and hypothesis (v), G' is a counterexample strictly smaller than G in the lexicographic order, contradicting the minimality of G . $\square$

Definition 3.104 (Target-complete compression). *Let X be a proper atom. A nontrivial target-complete compression of X is a smaller T -boundaried representative X' satisfying the hypotheses of Lemma 3.103. A quotient or identification of local response coordinates that is not target-complete is called a target-defective quotient.*

Corollary 3.105 (Hereditary target-uncompressibility). *No proper atom of G admits a nontrivial target-complete compression. Moreover, a non-target-complete identification cannot be used as a replacement that preserves the counterexample predicate: by definition of target completeness and Lemma 3.102, some outside context distinguishes the identified states by producing a power-of-two cycle from one and not from the other.*

Proof. The first assertion is exactly Lemma 3.103. The second assertion is Lemma 3.102. $\square$

Definition 3.106 (Target tester and independent target coordinates). *Let X be a T -boundaried graph. A target tester for a local coordinate r of X is a T -boundaried context Y_r such that the gluing $X \oplus_T Y_r$ contains a power-of-two cycle if and only if the coordinate r takes the obstructing value, while all other coordinates under discussion are held fixed.*

A finite family of local coordinates $\mathcal{I} = \{r_1, \dots, r_k\}$ is independently target-testable if its target responses realize 2^k distinct target-complete states; equivalently, for every subset $J \subseteq \{1, \dots, k\}$ there is a target-complete state of the atom in which precisely the coordinates indexed by J are obstructing. The target rank of a coordinate family is the maximum size of an independently target-testable subfamily.

Lemma 3.107 (Independent target coordinates consume skeleton entropy). *Let $\mathcal{I}$ be an independently target-testable family of size k arising canonically from graphs in a labelled graph class $\mathcal{G}$. Then*

$$2^k \leq |\mathcal{G}|.$$

In particular,

$$k \leq \log_2 |\mathcal{G}|.$$

Proof. By independent target-testability, the family $\mathcal{I}$ realizes 2^k distinct target-complete states. If two distinct such states came from the same labelled graph, the canonical target-response map of that graph would be identical in both cases, contradiction. Hence the number of states is at most the number of labelled graphs in the ambient class. $\square$

Framework branch table. The target-algebra and replacement section instantiates the framework branch table as follows.

Branch	Residual condition	Demand	Payer	Bounded multiplicity	Closure
Target hit	An edge-rooted Mersenne return is present.	No residual demand remains.	The target predicate.	Not applicable.	It is a dyadic cycle by Lemma 2.2 .
Deletion	Deleting an edge would leave a smaller counterexample.	The deleted edge.	Lexicographic minimality.	One edge at a time.	Contradicts Lemma 3.2 .
Boundary fibre	A quotient identifies supports with different boundary-degree profiles.	The boundary-degree discrepancy.	The fixed boundary fibre.	One canonical boundary profile per fibre.	Forbidden by Lemma 3.101 .
Target defect	An outside context distinguishes two identified coordinates.	A target-defective quotient.	The defect account attached to its declared support.	The finite declared coordinate signature.	Routed by Lemma 3.102 .
Compression	A proper atom has a smaller target-complete representative.	The proper support size.	Minimality.	One replacement per proper atom.	Contradicts Lemma 3.103 and Corollary 3.105 .
Entropy coordinate	Independent target coordinates are realized.	The number of target states.	The labelled skeleton count.	The canonical state map.	Bounded by Lemma 3.107 .

4 Sparse rank parameters

Define

$$\beta = m - n + 1, \quad \lambda = 2n - 3 - m.$$

Then

$$\boxed{\beta + \lambda = n - 2.}$$

Define the surplus above cubic degree

$$\sigma = \sum_{v \in V(G)} (d(v) - 3) = 2m - 3n.$$

Then

$$\boxed{\beta = \frac{n}{2} + 1 + \frac{\sigma}{2}}, \quad \boxed{\lambda = \frac{n}{2} - 3 - \frac{\sigma}{2}}.$$

The sparse skeleton entropy in the sparse (λ -parametrized) branch is

$$B_{\text{skel}}(n, \lambda) = \log_2 \binom{\binom{n}{2}}{2n - 3 - \lambda}.$$




Figure 15: Visual definition of bounded replacement and target-completeness. The left panel shows a proper atom X glued to its context Y along the labelled boundary $T = \{t_1, t_2, t_3\}$, whose orange vertices are identified in $G = X \oplus_T Y$. The right panel shows a smaller T -bordered representative X' glued to the same context. The hypotheses $\mathbf{d}_\partial(X') = \mathbf{d}_\partial(X)$ and $X' \preceq_T X$ mean that the boundary degrees are in the same fibre and every context obstructed by X' is also obstructed by X . If X' has no internal dyadic cycle and all internal vertices keep degree at least 3, then $X' \oplus_T Y$ would be a smaller counterexample, contradicting minimality. The lower inset records the complementary failure mode: identifying two response coordinates r_1, r_2 is target-defective when some context Y_0 creates a dyadic cycle for one state and not for the other, so the quotient is not target-complete.

Definition 4.1 (Near-cubic spine hypothesis). *The near-cubic spine hypothesis is the assertion that every lexicographically minimal counterexample surviving to the entropy/net-charge spine satisfies*

$$m = \frac{3}{2}n + O(\sqrt{n}), \quad \sigma(G) = O(\sqrt{n}).$$

Equivalently, the global surplus is $o(n)$, and the labelled skeleton budget available to the spine is

$$\frac{3}{2}n \log_2 n + o(n \log n).$$

In the final proof, this estimate is supplied by Proposition 3.87 before either the normalized skeleton budget $1.5n \log_2 n + o(n \log n)$ or the fan-mass estimate $\sigma(G) = o(|R|)$ is invoked. When a later local lemma states Definition 4.1 among its hypotheses, that line records the branch state needed for the normalized estimate, not an additional global assumption. The local surplus bookkeeping relevant to deriving this hypothesis is recorded in Definitions 3.6 and 3.10. The degree-greater-than-four part is separated by Corollary 3.13 and Remark 3.20; the degree-four local activation is Corollary 3.14; and the non-near-cubic surplus branch is routed by the exact token ledger and the same-token bottleneck closure Lemma 3.84, Theorem 3.85, and Proposition 3.87.

The following lemmas make explicit why this is a safe *upper* bound on every degree of freedom

available to a near-cubic graph. The entropy budget is not an encoding assumption: it is merely the number of labelled adjacency matrices in the ambient class.

Lemma 4.2 (Labelled skeleton budget dominates all auxiliary data). *Fix n and m . Let $\mathcal{G}_{n,m}$ be the set of all finite simple labelled graphs on vertex set $[n]$ with exactly m edges. Then*

$$|\mathcal{G}_{n,m}| = \binom{\binom{n}{2}}{m}.$$

Consequently, if a collection of target-complete states is assigned canonically to graphs in $\mathcal{G}_{n,m}$, then the number of realized states is at most

$$\binom{\binom{n}{2}}{m}.$$

In particular, any target-complete response algebra that realizes more than

$$2^{B_{\text{skel}}(n,\lambda)}$$

states among graphs with $m = 2n - 3 - \lambda$ edges cannot be realized by such graphs.

Proof. A labelled simple graph on $[n]$ is uniquely determined by its edge set, a subset of the $\binom{n}{2}$ possible unordered pairs. Thus the number with m edges is exactly $\binom{\binom{n}{2}}{m}$.

All auxiliary objects used in the proof—a maximal induced- P_{13} packing, atom decomposition, boundary profiles, spanning-tree choices, and curvature-response data—are functions of the labelled adjacency matrix once a deterministic tie-breaking rule is fixed. We fix throughout the lexicographically first object among all objects with the required extremal property. Thus no auxiliary structure supplies independent degrees of freedom beyond the adjacency matrix. Hence the number of realized target-complete states is bounded by the number of labelled skeletons. $\square$

Lemma 4.3 (Variable edge-count budget). *Let $I \subseteq \{0, 1, \dots, \binom{n}{2}\}$ be any interval of possible edge counts. Then the number of labelled graphs on $[n]$ whose edge count lies in I is at most*

$$|I| \max_{m \in I} \binom{\binom{n}{2}}{m}.$$

Consequently the logarithmic budget differs from the largest fixed- m budget by at most $\log_2 |I| \leq \log_2 n^2 = 2 \log_2 n$ whenever I has length at most n^2 .

Proof. The graphs with different edge counts are disjoint, so the total number is the sum of $\binom{\binom{n}{2}}{m}$ over $m \in I$. This sum is at most $|I|$ times its largest summand. Taking logarithms gives the claim. $\square$

Lemma 4.4 (Near-cubic specialization of the skeleton budget). *Suppose*

$$m = \frac{3}{2}n + O(\sqrt{n}).$$

Then

$$\log_2 \binom{\binom{n}{2}}{m} = \frac{3}{2}n \log_2 n + O(n) + O(\sqrt{n} \log n).$$

In particular,

$$B_{\text{skel}}(n, \lambda) = \frac{3}{2}n \log_2 n + o(n \log n)$$

in the near-cubic branch.

Proof. Let $N = \binom{n}{2}$. The standard binomial estimate

$$\binom{N}{m} \leq \left(\frac{eN}{m}\right)^m$$

gives

$$\log_2 \binom{N}{m} \leq m \log_2 \left(\frac{eN}{m}\right).$$

If $m = \frac{3}{2}n + O(\sqrt{n})$, then $N = \frac{1}{2}n^2 + O(n)$ and therefore

$$\frac{eN}{m} = \Theta(n).$$

Hence

$$\log_2 \binom{N}{m} \leq \left(\frac{3}{2}n + O(\sqrt{n})\right) (\log_2 n + O(1)) = \frac{3}{2}n \log_2 n + O(n) + O(\sqrt{n} \log n).$$

The matching lower asymptotic is not needed in the proof; the displayed upper bound is what is used. Since $O(n) + O(\sqrt{n} \log n) = o(n \log n)$, the final statement follows. $\square$

Remark 4.5 (Scope of the near-cubic budget). *Lemma 4.4 is the conversion step from $m = \frac{3}{2}n + O(\sqrt{n})$ to the entropy budget used later. The surplus-token routing branch supplies this edge-count estimate before the lemma is invoked in the final proof.*

Remark 4.6 (Robustness of the skeleton budget). *The budget counts all labelled simple graphs with the given edge count. The near-cubic, minimality, atom, curvature, and target constraints only restrict this set. Automorphisms only reduce the number of unlabelled graphs, so labelled counting is an overcount. Degree sequences are not additional data: they are determined by the edge set. Likewise, packings, atoms, trees, and profiles are treated canonically and are deterministic functions of the graph. Thus one cannot obtain extra graph-theoretic degrees of freedom by varying these auxiliary choices; doing so would only count the same labelled graph multiple times.*

Lemma 4.7 (State-count comparison). *Let $\mathcal{S}$ be a family of target-complete states realized by labelled graphs in a class $\mathcal{G}$. If the state map is canonical, then*

$$|\mathcal{S}| \leq |\mathcal{G}|.$$

Equivalently, if the proof establishes K independently testable Boolean target coordinates, so that at least 2^K target-complete states are realized, then necessarily

$$K \leq \log_2 |\mathcal{G}|.$$

Proof. The canonical state map sends each graph to one state. Therefore the number of realized states cannot exceed the number of graphs. If K independent Boolean target coordinates are realized, they give at least 2^K distinct target-complete states. Hence $2^K \leq |\mathcal{G}|$ and the logarithmic inequality follows. $\square$

Framework branch table. The sparse-rank section supplies the global budget branch table.

Branch	Residual condition	Demand	Payer	Bounded multiplicity	Closure
Fixed edge count	States are assigned to a fixed- m labelled graph class.	Target-complete states.	Labelled edge sets.	One canonical state per labelled graph.	Bounded by Lemma 4.2 .
Variable edge count	The edge count varies over an interval.	Extra edge-count choices.	The largest fixed- m class.	One edge-count value per graph.	Absorbed by Lemma 4.3 .
Near-cubic budget	$m = \frac{3}{2}n + O(\sqrt{n})$.	Skeleton entropy.	$\frac{3}{2}n \log_2 n + o(n \log n)$.	The labelled graph count.	Computed in Lemma 4.4 .
State overflow	Independent target states exceed the skeleton count.	Independent Boolean coordinates.	The labelled skeleton budget.	One canonical graph-to-state map.	Contradicts Lemma 4.7 .
Non-near-cubic route	The surplus-token branch survives the sparse exits.	Same-token load.	The capacity-token ledger.	The token-fibre partition.	Closed or routed by Proposition 3.87 .

5 Induced P_{13} windows

We use the following external theorem.

Theorem 5.1 (Hegde–Sandeep–Shashank). *Every P_{13} -free graph of minimum degree at least 3 contains a cycle whose length is a power of two.*

Reference. This is Theorem 3 of Hegde, Sandeep, and Shashank [1]. □

Corollary 5.2. *The minimal counterexample G contains an induced P_{13} .*

Proof. If G were P_{13} -free, then since $\delta(G) \geq 3$, [Theorem 5.1](#) would give a power-of-two cycle in G , contradicting that G is a counterexample. □

Definition 5.3 (Maximum induced- P_{13} packing). *Let p_{13} be the maximum size of a vertex-disjoint family of induced copies of P_{13} in G , and put*

$$\theta = \frac{p_{13}}{n}.$$

Framework branch table. The induced-window section contributes the external P_{13} -free input and the maximal packing notation consumed by the window entropy branch.

Branch	Residual condition	Demand	Payer	Bounded multiplicity	Closure
--------	--------------------	--------	-------	----------------------	---------

External P_{13} -free input	G is P_{13} -free.	Minimum degree at least 3.	The HSS theorem.	The theorem is consumed only through its stated interface.	Closed by Theorem 5.1 .
Induced window exists	The minimal counterexample is tested against the external input.	$\delta(G) \geq 3$ and no dyadic cycle.	The counterexample condition.	If G were P_{13} -free, HSS gives a forbidden target cycle.	Corollary 5.2 .
Maximal window packing	A vertex-disjoint induced- P_{13} family is chosen maximal.	The packing size p_{13} .	The complement R .	Every unchosen induced window meets the packing.	Supplies $\theta = p_{13}/n$ for Remark 7.67 and the remainder analysis.

6 The P_{13} curvature algebra

Let

$$P = v_0 v_1 \cdots v_{12}$$

be an induced P_{13} . An outside vertex x has attachment label

$$S(x) = \{i : xv_i \in E(G)\} \subseteq \{0, 1, \dots, 12\}.$$

A nonempty label S is *legal* if a single outside vertex x with attachment label S forms no power-of-two cycle of length 4 or 8 together with P . Indeed, if x is adjacent to v_i and v_j with $i < j$, then the subpath $v_i v_{i+1} \cdots v_j$ of P , which has length $j - i$, together with the two edges xv_i and xv_j forms a cycle of length $(j - i) + 2$. As $1 \leq j - i \leq 12$, this length lies in $\{3, \dots, 14\}$ and is a power of two precisely when $(j - i) + 2 \in \{4, 8\}$, that is $j - i \in \{2, 6\}$. Hence S is legal if and only if

$$|i - j| \notin \{2, 6\} \quad \forall i, j \in S.$$

Let

$$\mathcal{L} = \{S \subseteq \{0, \dots, 12\} : S \neq \emptyset \text{ and } |i - j| \notin \{2, 6\} \forall i, j \in S\}$$

be the set of legal nonempty labels.

Lemma 6.1 (Legal label count). *The number of legal nonempty labels is*

$$|\mathcal{L}| = 399.$$

The distribution by size is

$$13, 60, 122, 122, 63, 17, 2$$

for sizes 1, 2, 3, 4, 5, 6, 7, respectively.

Proof. This is a direct enumeration of nonempty subsets of $\{0, \dots, 12\}$ avoiding pairwise differences 2 and 6. The verification code in [Section A](#) computes the stated values. $\square$




Figure 16: Visual definition of the P_{13} label algebra and the two-step curvature tensor. The black horizontal path is the induced P_{13} , with indices $0, \dots, 12$ displayed below its vertices. An outside vertex has label equal to the set of path indices to which it is adjacent. The lower red vertex z illustrates an illegal label: attachments at indices 2 and 4 close a 4-cycle through the subpath $v_2v_3v_4$, so legal labels must avoid pairwise gaps 2 and 6. The upper vertices carry the singleton labels $S = \{1\}$, $A = \{8\}$, and $T = \{5\}$. The blue outside edges show the two individually safe length-1 relations $C_1(S, A) = 1$ and $C_1(A, T) = 1$. The dashed red two-step outside path from S to T , together with the red path segment $v_1 \cdots v_5$ and the two attachment edges, has length $2 + 2 + |1 - 5| = 8$; hence $C_2(S, T) = 0$. Thus $\Omega_2(S, A, T) = 1$: two locally safe outside steps compose into an unsafe dyadic test.

Definition 6.2 (P_{13} safety and two-step curvature). *For $s \geq 0$, define a safety relation C_s on legal labels by*

$$C_s(S, T) = 1$$

if and only if

$$s + 2 + |i - j| \notin \mathcal{D} \quad \forall i \in S, j \in T.$$

Thus an outside path of length s between vertices carrying labels S, T is safe through the induced path P precisely when $C_s(S, T) = 1$.

Define the two-step curvature tensor

$$\Omega_2(S, A, T) = C_1(S, A)C_1(A, T)(1 - C_2(S, T)).$$

So $\Omega_2 = 1$ means two individually safe outside edges compose into an unsafe outside path of length 2.

Lemma 6.3 (Curvature enumeration). *For the 399 legal labels,*

$$|\{(S, A, T) : C_1(S, A) = C_1(A, T) = 1\}| = 543958,$$

and

$$|\{(S, A, T) : \Omega_2(S, A, T) = 1\}| = 432672.$$

Hence the number of curvature-flat locally safe wedges is

$$111286.$$

Therefore the flatness entropy cost of one independent two-step curvature test is

$$c_\Omega = \log_2 \frac{543958}{111286} = 2.28922315244\dots$$

Proof. The proof is the same direct finite enumeration as in [Lemma 6.1](#), now applied to ordered triples of legal labels. A locally safe wedge is a triple (S, A, T) with $C_1(S, A) = C_1(A, T) = 1$. It is curvature-positive when $C_2(S, T) = 0$. The code in [Section A](#) performs precisely these tests and returns the displayed integer counts. The entropy value is then the elementary calculation

$$\log_2 \frac{543958}{111286} = 2.28922315244071\dots$$

□

A finite-scale multi-relation curvature package through length 14 yields the constant

$$c_{13} = 118.108581006\dots$$

bits per independent P_{13} -window package per dyadic scale. This value is the sum of the 91 finite curvature barriers indexed by ordered pairs of outside path lengths $a, b \geq 1$ with $a + b \leq 14$. The number of such pairs is

$$|\{(a, b) : a, b \geq 1, a + b \leq 14\}| = \sum_{s=2}^{14} (s - 1) = \sum_{k=1}^{13} k = 91,$$

since for each total $s = a + b \in \{2, \dots, 14\}$ there are exactly $s - 1$ admissible pairs $(a, s - a)$ with $1 \leq a \leq s - 1$; the individual barriers are computed in [Section A](#).

Framework branch table. The curvature-algebra section is a finite computation interface: every quantity it exports is either an integer count or a named entropy rate verified in [Section A](#).

Branch	Residual condition	Demand	Payer	Bounded multiplicity	Closure
Illegal label	An outside vertex attaches at path-index distance 2 or 6.	The attachment label S .	The target predicate.	The induced path segment plus two attachment edges is a 4- or 8-cycle.	Excluded by the counterexample condition.
Legal label enumeration	All nonempty labels avoiding gaps 2 and 6.	The finite set $\mathcal{L}$.	Direct enumeration.	The code tests every subset of $\{0, \dots, 12\}$.	Lemma 6.1 .
Two-step curvature test	$C_1(S, A) = C_1(A, T) = 1$ but $C_2(S, T) = 0$.	The triple (S, A, T) .	The finite curvature tensor.	The same predicates define both the proof and the code check.	Lemma 6.3 .

Multi-scale package constant	Ordered path-length pairs $a, b \geq 1$, $a + b \leq 14$.	The 91 barrier rates.	The compu- tational certificate.	The exported value is the displayed sum $c_{13} =$ 118.108581006 . . .	Consumed by Remark 7.1 .
------------------------------------	--	--------------------------	--	--	---

7 Entropy bound on independent P_{13} windows

The finite curvature computation gives the local safe and flat counts. The paper uses them as conditional fibre costs, not as coordinates of a Boolean cube. The exact live/cold interface is recorded next.

Remark 7.1 (Weighted multi-scale P_{13} -window package). *Fix one selected window. At the separated dyadic scales, order the declared barrier tests and write*

$$\mathcal{S}_0 \supseteq \mathcal{S}_1 \supseteq \cdots \supseteq \mathcal{S}_r$$

for the successive fibres of graph-owned connector states. If the j -th test has barrier index $i(j)$, a live package carries the local inequality

$$W_{i(j)}|\mathcal{S}_{j+1}| \leq F_{i(j)}|\mathcal{S}_j|,$$

where W_i and F_i are the verified safe and flat counts. It also carries a nonempty final fibre, an injective graph-owned connector code, and the occurrence of every one of the 91 barriers at $\log_2 n - O(1)$ separated scales. Multiplying the displayed local inequalities gives the full variable-factor product cost and hence the constant

$$c_{13} = \sum_i \log_2(W_i/F_i) = 118.10858100609198 \dots$$

A live package is not tested independently of the packages already retained. Process the fixed packing order while carrying one jointly compatible hot aggregate. At a window P , accept only an extension which adjoins the complete local package to that aggregate, gives recoverable simultaneous choices for every retained owner, preserves their graph responses, and has an injective code in the fixed spine family. If no such extension exists, keep the previous aggregate and retain that identical P as a cold residual. Thus dependence is the cold outcome prescribed by nodes [145]–[157], not a third branch, and no coordinatewise activity is promoted to simultaneous Boolean realization.

The Lean interface verifies the variable-factor telescope, positive scale multiplicity with bounded scale loss, graph-semantic connector tests, and the packing-order compatible-extension ledger. Its final aggregate carries a finite dependent joint-state family, a left inverse recovering every retained local choice, response commutation for each recovered owner, and an injective code bounded by the fixed spine capacity. Rejection records the exact prior aggregate and the failure of its extension type. The accepted and rejected lists partition the original packing exactly. On the negative side, this sequential rejection list is filtered by the graph-owned ambient-cubic predicate without replacing it by the ambient packing: the cubic and non-cubic sublists partition that same list. Every retained cubic cold window has exactly fifteen literal external stubs, its first two stubs are the transit stubs, and the remaining thirteen form an exact proof-carrying branch-excess schedule whose entries retain their cold window of origin. Each entry is passed unchanged to the outside-component corridor runner. This formalizes the local handoffs through nodes [151] and [152]. Moreover, the cold list

has no repeated selected window. Its non-cubic sublist therefore injects into the packing's non-cubic windows, and the local surplus payment gives

$$|\mathcal{P}_{\text{cold}}^{\text{noncub}}| \leq \sigma(G), \quad 13|\mathcal{P}_{\text{cold}}| \leq |\mathcal{E}_{\text{br}}| + 13\sigma(G).$$

These are exact finite inequalities; no asymptotic estimate or ambient-window substitution is used.

Since node [21] retains the exact inequality $\sigma(G)^2 \leq Kn$, the discarded loss also satisfies $|\mathcal{P}_{\text{cold}}^{\text{noncub}}|^2 \leq Kn$. This is the finite square-root form of the paper's $o(n)$ assertion and uses only the verified node-[21] predecessor.

The remaining hot-side obligation is the exact arithmetic passage from the per-package weighted rate bounds to the dependent product cardinality, and from the fixed spine-state capacity to the printed normalized hot bound. On the cold side, the literal stage-major F1 constructor is verified; node [153] still requires the graph-owned F2–F5 semantic producer described in [Remark 7.59](#). Until those obligations are discharged, the density conclusion is not used as a proved input.

Lemma 7.2 (Pointwise actual-attachment cold fork). *Fix one selected induced P_{13} window. The finite classifier whose states are the actual outside vertices and whose thirteen response coordinates are their adjacency bits to the path returns a canonical omitted assignment. Its cold residual retains the exact same selected window.*

Proof. The classifier scans the 2^{13} Boolean assignments against the actual outside vertices. Its hot outcome would realize the all-true assignment. The realizing outside vertex would then be adjacent to path positions (0) and (2); together with the intervening two path edges these adjacencies form a four-cycle, contrary to target avoidance. Hence the exhaustive classifier returns its canonical first missing assignment. The selected window is passed unchanged into the cold residual. $\square$

Remark 7.3 (Scope of the actual-attachment fork). *Node [153] is a pointwise thirteen-coordinate adjacency scan. It is distinct from the still-open simultaneous realization problem for the 91 multi-scale coordinates on the edge [21]→[22], and it yields no density estimate. It retains only the exact selected window as the structural trigger for node [153].*

Lemma 7.4 (Same-window structural frontier). *Fix the node-[153] cold residual on one selected induced P_{13} window. The graph-owned continuation returns exactly one of the following four proof-carrying outputs:*

- (i) *a path position whose ambient degree is greater than three;*
- (ii) *a canonical external stub on the same window and a power-of-two cycle obtained by restoring that stub to its deleted-edge return;*
- (iii) *a canonical external stub and the first high-degree vertex on its return corridor, together with the clean earlier prefix; or*
- (iv) *a canonical external stub and a quiet structural germ whose support has size at most n , whose root-cycle length is rejected, and whose corridor vertices have degree at most three.*

The visible certificate checks for one supplied window are at most

$$26 + 15p_{13} + \sigma_W + n.$$

Proof. First inspect the thirteen path positions. If one has degree different from three, the minimum-degree hypothesis makes the first such position a strict surplus output. Otherwise the exact fifteen-stub identity makes the window ambient-cubic and nonempty in the external-incidence schedule; choose its first scheduled stub. Bridgelessness supplies a proof-carrying simple return after deleting this stub. Scan the vertices of that return in order. The first accepted event is either the restored root cycle with power-of-two length or the first vertex of degree greater than three. If no event occurs, the negative scan constructs precisely the structural germ in (iv). The scan checks at most two passes over the thirteen positions, one pass over the $15p_{13} + \sigma_W$ external-incidence schedule, and at most n vertices of the supplied simple return. $\square$

Lemma 7.5 (Dyadic constructor closes). *If node [153] returns its dyadic constructor, then the constructor's exact stub position and restored root cycle form a cold dyadic hit on the same selected window. The certificate-driven CT1 run reaches its C1 terminal in one check, and this branch contradicts target avoidance.*

Proof. The first-event payload already contains the canonical external stub, the clean prefix, and the proof that restoring the deleted edge gives a simple cycle of accepted power-of-two length. Package that same cycle with the stub's Fin(13) position. The existing CT1 target-certificate runner validates this single supplied cycle and reaches C1. Hence the ambient selected graph contains the forbidden target, contrary to the inherited avoiding field. $\square$

Remark 7.6 (Scope of the same-window frontier). *Node [153] is pointwise and structural. Its window-surplus and corridor-high constructors are typed surplus handoffs, not contradictions. Its quiet output is only an ambient-finite structural germ and is not promoted to a uniformly bounded germ. Thus node [155] is green only for the computed dyadic constructor; no cold-mass, density, bounded-overlap, or nodes [153]–[157] aggregate conclusion follows.*

Lemma 7.7 (D1–D3 base-scale split). *Suppose the computed node-[153] output is its quiet constructor, with its exact fork, selected window, canonical stub, no-event proof, and structural germ. Set*

$$Q_{\text{base}} := 4^2 13^2 2^{13}.$$

One support-length comparison returns exactly one of

$$|\text{supp}(\Gamma)| \leq Q_{\text{base}} \quad \text{or} \quad Q_{\text{base}} < |\text{supp}(\Gamma)|,$$

while retaining all of that quiet-constructor data.

Proof. The factors in Q_{base} count two capped boundary-degree roles, two window offsets, and thirteen target-response bits. Equivalently, this is the exact normalized fixed two-boundary state cardinal when the additional local coordinate type is empty. Compare the literal support length with this fixed natural number. The non-strict outcome is the short residual and the negated outcome is exactly the strict long residual. $\square$

Remark 7.8 (Scope of the D1–D3 split). *Node [153] does not assert that an empty local-coordinate alphabet captures D4–D7. It proves no repeated state, bounded-germ semantics, CT3 compression, or density estimate. Its exact short and long residuals are consumed locally in two internal stages of node [153], without adding those missing semantics. This auxiliary split is independent of the weighted live/cold handoff in Remark 7.50 and supplies no window entropy.*

Lemma 7.9 (Short-return third incidence). *Suppose the computed node-[153] result is its short constructor. At the root of its supplied deleted-edge return, the declared-order third incidence is well defined. Its endpoint lies either on that return support or outside it, and these two outcomes are exhaustive. The visible work is at most*

$$2n + 3 + Q_{\text{base}}.$$

Proof. The minimum-degree hypothesis gives degree at least three at the return root, while the quiet germ says that this root is not high; hence its degree is exactly three. The first return step and the restored deleted edge are two distinct actual incidences. The graph-owned root classifier therefore selects the unique third incidence in the declared neighbour order. Test its endpoint once against the supplied short support. Membership gives the on-support constructor and nonmembership gives the outside constructor. The declared-neighbour scan, return-certificate allowance, three root checks, and support-membership scan give the displayed bound. $\square$

Remark 7.10 (Scope of the short-root split). *Node [153] inspects only this one root incidence. Its on-support and outside outcomes feed the separately verified internal node-[153] branch consumers, but not their normalized one-return rejoin [153]. It supplies no cold-skeleton boundary-response semantics, which remain further downstream at [153]. No support-wide scan, D4–D7 coordinate, CT3 execution, or density claim occurs.*

Lemma 7.11 (On-return chord resolution). *Suppose the computed node-[153] result is its on-support constructor, and let i be the first position of the selected third endpoint on the supplied return. The literal chord cut out by positions 1 and i has length $i + 1$. If this length is a power of two, the existing CT1 target runner reaches C1 and contradicts target avoidance. Otherwise the surviving output retains the exact deleted-edge return obtained by replacing the initial return prefix with the selected third incidence, and this return is strictly shorter than the supplied return. The visible work is at most*

$$Q_{\text{base}} + 1.$$

Proof. The input stores equality with node [153]’s actual on-support result. Scan the same short support once to find the canonical index i . The selected third incidence is distinct from both the first return step and restored endpoint, so $1 < i < |\Gamma|$. The standard chord construction therefore gives a simple cycle of length $i + 1$. A power-of-two length supplies a concrete target certificate to the existing CT1 runner and is impossible in the inherited target-avoiding context. On rejection, retain the suffix beginning at i and prepend the selected third edge. Its length is $|\Gamma| - i + 1 < |\Gamma|$, and the runner records equality with this exact return. One bounded support scan and one length decision give the displayed work bound. $\square$

Remark 7.12 (Scope of the chord resolution). *Node [153] closes only the accepted target-length constructor and otherwise returns one exact strictly shorter deleted-edge return. It does not normalize the two node-[153] branches into a common residual, iterate the shortening, prove termination, reconstruct D4–D7, execute CT3, or derive density. Its exact shorter residual is consumed by the verified proof-level rejoin [153] and packed-support transition [153]; its two computed outcomes feed the separately verified internal node-[153] branch consumers.*

Lemma 7.13 (Outside-return cubic boundary star). *Suppose the computed node-[153] result is its outside constructor. Orient the selected third incidence from the return root r to its endpoint x_3 . Then*

$$r \in \text{supp}(\Gamma), \quad x_3 \notin \text{supp}(\Gamma), \quad rx_3 \in E(G).$$

Together with the first return step x_1 and the restored deleted-edge endpoint x_2 , this gives three distinct leaves at the cubic root r , and every neighbour of r is one of x_1, x_2, x_3 . Thus these leaves form a finite three-boundary, one-internal-owner star shape. This projection performs zero additional primitive checks beyond node [153].

Proof. The input includes equality with the actual node-[153] runner, not merely an independently supplied nonmembership proof. Its outside constructor therefore retains the same return support, selected third incidence, and proof that its endpoint lies outside. The return root lies on the supplied path, and the selected incidence supplies adjacency and distinction from both the first return step and restored endpoint. Node [153] has already proved that the root has degree three. The generic cubic-star construction packages these three incidences, and equality of their three-element boundary with the full neighbour set proves the ownership assertion. All operations project this proof-carrying data, so the additional-check ledger is zero. $\square$

Remark 7.14 (Scope of the cubic boundary star). *Node [153] certifies only this one oriented support crossing and the exact three-incidence ownership at its cubic root. It does not construct an all-inside consumer, successor or second stub, component path, any D4–D7 coordinate, a CT3 input, or a density estimate. The exact outside residual is consumed by the verified proof-level rejoin [153] and packed-support transition [153]; its two computed outcomes feed the separately verified downstream internal node-[153] branch consumers.*

Lemma 7.15 (Normalized one-return boundary). *Fix the common computed node-[153] short residual. Suppose the input is either the exact rejected-chord computation from node [153] or the exact outside computation from node [153]. There is a selected deleted-edge return Γ' , an outside endpoint x_{out} , and the retained cubic star at the return root r such that*

$$|\text{supp}(\Gamma')| \leq Q_{\text{base}}, \quad |\Gamma'| \leq |\Gamma|, \quad x_{\text{out}} \notin \text{supp}(\Gamma'), \quad rx_{\text{out}} \in E(G).$$

The star's three leaves own every incidence at r . On the rejected-chord branch $|\Gamma'| < |\Gamma|$; on the outside branch $|\Gamma'| = |\Gamma|$. This normalization performs zero additional primitive checks.

Proof. The input is indexed by its complete predecessor computation. On the rejected-chord branch select node [153]'s exact shorter return. Its old first return step is absent from the new support, remains adjacent to the common root, and becomes x_{out} ; strict decrease is inherited from the resolver. On the outside branch keep node [153]'s original return and selected third endpoint; its nonmembership and equality of lengths are inherited unchanged. In both cases the same cubic root divergence supplies the three-leaf star and exhausts every root incidence. Walk support has one more vertex than walk length, so the original Q_{base} support bound and the branchwise length comparison give the selected support bound. The construction only projects proof-carrying predecessor fields. $\square$

Remark 7.16 (Scope of the normalized return). *Node [153] is a proof-level normalization of one return only. It does not itself scan packed support. Its exact computed output is consumed by node [153], whose exact all-inside branch subsequently feeds verified owner-change node [153]. That internal support stage supplies no owner table or cross-window token, successor or second stub, component path [153], iteration or termination theorem, D4–D7 coordinate, CT3 input, or density estimate.*

Lemma 7.17 (Normalized-return packed-support transition). *Let Γ' be the exact computed normalized return from node [153], and let*

$$U_{\text{ac}} := \bigcup_{W: W \text{ is an ambient-cubic selected window}} \text{supp}(W).$$

Exactly one of the following proof-carrying outcomes is returned:

- (i) every vertex of $\text{supp}(\Gamma')$ belongs to U_{ac} ;
- (ii) the first edge whose endpoints have different U_{ac} -membership is oriented from its inside endpoint to its outside endpoint and supplies the literal **BoundaryStub**, its exact outside endpoint, and the exact connected component of the induced outside remainder containing that endpoint.

If p is the selected-window packing number, $n = |V(G)|$, and $L = |\Gamma'|$, the visible work is exactly

$$13pn + 13p + 26pL$$

and is at most

$$n^2 + (2Q_{\text{base}} + 1)n.$$

Proof. View the selected deleted-edge return in the ambient graph. Its final vertex is the retained selected-window position, and that window is ambient-cubic, so the final vertex belongs to U_{ac} . Scan the return edge indices in their declared order and test membership of both consecutive endpoints in U_{ac} . If no membership change occurs, final membership propagates backward across every edge, proving full support containment. Otherwise the first hit has consecutive endpoints with opposite membership; orient it from inside to outside. Membership in the explicit union chooses an ambient-cubic selected window and its exact $\text{Fin}(13)$ position. Adjacency and outside nonmembership then give the literal boundary stub, whose endpoint determines the stated induced-remainder component.

Computing thirteen degree tests for each of the p selected windows costs $13pn$, recording their cubic flags costs $13p$, and two scans of the $13p$ position list at each of the L path edges cost $26pL$. The disjoint selected-window packing bound gives $13p \leq n$, while node [153] gives $L \leq Q_{\text{base}}$, yielding the polynomial bound. $\square$

Remark 7.18 (Scope of the packed-support transition). *The union U_{ac} contains all ambient-cubic selected windows. It is not the manuscript's selected cold subfamily, so node [153] proves no cold aggregate or cold-mass statement. Its exact all-inside output is consumed by verified owner-change node [153]. The transition output constructs one exact stub, endpoint, and component, but no successor, second stub, or path through that component; verified node [153] computes those objects on the exact first-transition branch. There is no iteration, termination, $D4$ – $D7$ reconstruction, $CT3$ execution, or density conclusion.*

Lemma 7.19 (Component boundary schedule and BFS path). *Suppose node [153] returns its exact first-transition constructor, with boundary stub b and outside endpoint in the induced remainder component C . There is a computed finite list $\mathcal{B}_C$ of precisely the boundary stubs whose outside endpoint lies in C , and:*

- (i) $b \in \mathcal{B}_C$, the list has no duplicates, and $|\mathcal{B}_C| \geq 2$;
- (ii) the stored successor b^+ is the genuine cyclic `List.next` value of b in $\mathcal{B}_C$, so $b^+ \neq b$ and both stubs have outside endpoint in C ;
- (iii) a declared-order BFS in C returns a path from the outside endpoint of b to that of b^+ , and this path has length equal to their graph distance in C .

The second stub is obtained by an explicit first-hit scan of $\text{WindowIndex} \times \text{Fin}(13)$. The scan stores the exact hit and proves that every preceding slot is absent. The component predicate used by this scan, the complete incident-stub filter, and the path object is the same computed outside BFS finset. If S denotes the sum of one plus the cardinalities of all explicitly scanned local collections, the complete visible work is at most $50S^3$.

Proof. Retain equality with node [153]’s actual first-transition run, so the anchor is its literal boundary stub and the returned component is unchanged. Delete the anchor edge and use the inherited nonbridge certificate to compute its return. Along that return, scan the single stored outside-component predicate until its first exit. The proof-carrying prefix remains in the same computed BFS finset, while the crossing edge reaches a deleted-window vertex outside it.

Scan the finite window-slot list for the first external incidence entering that same finset. The first-hit certificate supplies the selected slot and absence at all earlier slots. Its stub differs from the anchor because its crossing edge avoids the deleted anchor edge. Filtering the actual finite stub enumeration by membership in the same BFS finset is complete and duplicate-free, and contains both distinct stubs. Hence its length is at least two. The generic fixed-point-free theorem for cyclic `List.next` on a duplicate-free list of length at least two gives $b^+ \neq b$. Restricting to the computed component object and running the declared-order BFS gives the stated path and shortest-length equality.

The visible ledger includes the window ledger, $13p$ slot checks, the return-length times outside-component lookup, the outside BFS budget, the quadratic component restriction, the actual token filter, and the final component BFS budget. The two generic BFS bounds and domination by the local scale give $50S^3$. $\square$

Remark 7.20 (Scope of the component boundary schedule). *This internal support stage lies only on the short first-transition branch of node [153]. Its exact schedule and shortest path feed the one-state component observation internal to that node. The long-support branch and its open full semantic consumer are incompatible alternatives; no diagram edge is added between these internal branches. This stage proves no $D4$ – $D7$ label, repetition, Boolean or cold-family semantics, return iteration, target closure, $CT3$ execution, or density estimate.*

Lemma 7.21 (One component $D1$ – $D3$ observation). *Let node [153] return its exact component boundary schedule with anchor stub b , cyclic successor b^+ , and returned component C . Independently run the declared-order BFS tree in the computed finite component object and let P_C be its shortest path from the outside endpoint of b to that of b^+ . Equality with this one computed path defines a rank taking value zero on P_C and one otherwise. It packages P_C for the observation interface without generating, ordering, or scanning a family of paths.*

The two literal capped boundary degrees, the two literal $\text{Fin}(13)$ window offsets, and $|P_C|$ determine one genuine state

$$s_C \in \text{State}(\text{Fin}(0)).$$

Its thirteen target-response bits are the power-of-two tests at $|P_C| + j$, and its local response is the unique function out of $\text{Fin}(0)$. The output retains exactly $\text{MissingD4D7}(C, P_C)$. If $n = |V(G)|$, materializing this one state uses exactly $2n + 13$ visible checks and at most $15(n + 1)$.

Proof. Use node [153]’s actual graph input, not a reconstructed component or supplied path. The graph observation package runs the existing declared-order BFS-tree constructor on that input. Its path theorem and shortest-distance equality certify the resulting P_C . The equality-to- P_C rank is zero on the computed path and one on any unequal supplied candidate, so the required declared-first inequality follows by simplification; this compares one candidate with one stored value and performs no path-family scan.

The existing two-boundary observation map reads the anchor and successor degrees and offsets literally and reads the connector length from P_C . Project it with the unique $\text{Fin}(0)$ local response and the power-of-two predicate. This gives the displayed state definitionally. The residual constructor has no D4–D7 field and therefore returns the genuine missing-reconstruction witness. Two degree rows cost $2n$ checks and the fixed target-offset table costs thirteen; all remaining fields are stored projections, proving the exact and linear bounds. $\square$

Remark 7.22 (Scope of the component D1–D3 observation). *This internal node-[153] stage produces one state, not a sequence or repetition theorem. Its missing-reconstruction residual feeds the cyclic component D1–D3 ledger split internal to the same node. Only that split may feed the open D4–D7 reconstruction or coarse-repeat consumer. These internal routes do not enter the mutually exclusive long-branch classifier or its semantic consumer. This stage proves no repetition, D4–D7 reconstruction, CT3 compression, Boolean or cold-family semantics, target closure, or density estimate.*

Lemma 7.23 (Cyclic component D1–D3 ledger split). *Let $\mathcal{S}_C$ be node [153]’s complete duplicate-free incident-stub schedule, with its stored cyclic successor, and let s_C be the exact node-[153] residual. There is one locally computed D1–D3 row s_x for each $x \in \mathcal{S}_C$ such that the anchor row is exactly s_C , every row retains its computed connector path and MissingD4D7, and*

$$(\exists x \neq y \in \mathcal{S}_C : s_x = s_y) \quad \text{or} \quad |\mathcal{S}_C| \leq Q_{\text{base}} = 4^2 13^2 2^{13}.$$

If Λ_C is the maximum of the local finite input sizes used by the schedule and its connector computations, the visible work is at most $100\Lambda_C^4$.

Proof. The ledger input retains the typed node-[153] residual together with equality to its actual execution. At the anchor stub, the row constructor returns this retained state definitionally. At every other stub it re-anchors node [153]’s exact component-boundary input, uses the stored `List.next` successor, and runs the same declared-order local BFS constructor. Thus the row list is indexed exactly by $\mathcal{S}_C$; it does not range over an ambient state space.

Apply the finite code-collision classifier to this observed row list. A collision supplies two distinct scheduled rows with equal coarse $\text{State}(\text{Fin}(0))$ values. In the collision-free case, injectivity into that finite code gives $|\mathcal{S}_C| \leq |\text{State}(\text{Fin}(0))|$, and the proved cardinality identity rewrites the right-hand side to $4^2 13^2 2^{13} = Q_{\text{base}}$. The explicit ledger charges only the incident schedule, component restrictions, and local BFS clauses; its proved bound is $100\Lambda_C^4$. $\square$

Remark 7.24 (Scope of the cyclic D1–D3 ledger). *This internal node-[153] stage compares only the coarse states of rows actually produced from the local incident schedule. It does not enumerate the state universe. Equality of two coarse rows is not yet full response equivalence or a CT8 removal certificate, and the bounded branch supplies no D4–D7 reconstruction. Both residuals feed the verified finite classifier internal to node [153]. This stage proves no CT3 compression, second connector, Boolean or cold-family realization, target closure, return iteration, termination, or density estimate, and it has no diagram edge to the mutually exclusive long-branch support stages.*

Lemma 7.25 (D4–D7 availability and coarse-repeat classifier). *Consume node [153]’s exact generic result together with its equality to the specialized P_{13} execution on the same branch context. Exactly one of the following occurs:*

- (a) *two distinct retained rows have the same coarse D1–D3 state, and CT10 promotes the first of this exact pair as a missing refinement obligation;*

- (b) the schedule has length at most Q_{base} and every actual row has graph-derived D4–D7 clauses;
or
- (c) the schedule has length at most Q_{base} , the local scan returns its first actual row without graph-derived D4–D7 clauses, and CT10 consumes that exact row on the unchanged branch context.

The graph scan and either CT10 route use at most $3\Lambda_C$ primitive checks, where Λ_C is the declared local scale.

Proof. Run the graph-owned classifier on node [153]’s retained result. Its repeated constructor already contains the two distinct scheduled rows and their equal coarse states. The associated CT10 input has exactly these two classes and no invented D4–D7 datum, so the public runner reaches its promoted terminal. In the bounded constructor, scan the duplicate-free list of actual incident stubs in its inherited order with the decidable graph-derived-clause predicate. The finite reconstruction runner returns either a family indexed by every scanned row or its first missing row. The latter is passed literally to the second CT10 route. Both route inputs store the original branch context definitionally. The work ledger charges only the actual schedule and the two retained CT10 inputs; it never materializes a state, response, context, graph, or universe family. $\square$

Remark 7.26 (Scope of the D4–D7 classifier). *Node [153] proves availability or exposes one exact missing refinement. It does not turn coarse equality into compatible-response equality, produce a CT8 removal or certified smaller object, construct a second connector or cycle, or prove target closure. Node [153] performs only the exact readiness split below; the semantic implications remain white at node [153].*

Lemma 7.27 (Component D4–D7 semantic readiness). *Consume node [153]’s exact execution. The bounded complete-reconstruction constructor is impossible because the actual anchor stub belongs to the stored schedule while its D4–D7 clause type is empty. Hence the result is either the routed coarse pair, retaining both missing-clause witnesses, or the bounded ledger with its routed first actual missing row. One predecessor constructor inspection suffices, and its work is at most $\Lambda_C + 1$.*

Proof. The source stores equality with the specialized node-[153] run. Membership of the head anchor in the actual stub list contradicts a reconstructed family at that row. The other two constructors already retain, respectively, both coarse missing witnesses or the exact first-missing marker and their unchanged CT10 routes. No ambient family is enumerated. $\square$

Lemma 7.28 (Fixed D4–D7 clause schedule). *Consume node [153]’s exact result. Retain each dependent missing-clause marker and attach the fixed noduplicated obligation list [D4, D5, D6, D7]. The coarse constructor emits two such ledgers and hence eight slots; the bounded constructor emits one and hence four. No clause truth, response equivalence, removal operation, smaller object, or CT8 result is asserted.*

Lemma 7.29 (D4 obligation cursor). *Consume node [153]’s exact dependent ledgers. Focus (D4) as the next obligation while retaining the exact unevaluated tail ([D5, D6, D7]) and every marker. The actual returned tails contain six slots in the coarse case and three in the bounded case. No clause truth or response semantics is asserted.*

Lemma 7.30 (Local D4 evaluation request). *Consume node [153] exactly. For each actual dependent cursor, emit the singleton request [D4] while retaining its marker and the literal tail [D5, D6, D7]. Thus the coarse branch emits two requests and the bounded branch one. No predicate or truth value is supplied, and no response or CT8 conclusion follows.*

Lemma 7.31 (D4 evaluator residual). *Consume node [153] exactly. Retain every dependent marker, singleton D4 request, and literal D5–D7 tail, and expose the two missing inputs: a graph-local predicate and a proof of its provenance. These are obligation tags only; no predicate, Boolean, evaluator, response, or CT8 result is supplied. The actual output has at most four tags.*

Lemma 7.32 (D4 evaluator construction residual). *Consume node [153] exactly. Retain each singleton D4 request, its dependent marker, and the literal D5–D7 tail. Refine the missing evaluator obligation into the ordered graph-owned requirements: component-local data, a graph-owned predicate definition, and a derivation of that predicate. The coarse branch carries two such residuals and the bounded branch one, for at most six tags. No predicate, Boolean, evaluator, compatible response, or CT8 conclusion is supplied.*

Lemma 7.33 (Graph-owned D4 evaluation after node [153]). *For every construction residual returned inside node [153], the retained marker determines the same component input and canonical component path as in the preceding internal support stages. On the literal support of that path, let a D4 coordinate consist of one of the two boundary-window roles and an ordered pair of distinct support neighbours of a support vertex. Its value is the Boolean test that the curvature relation Ω_2 holds for the three actual attachment labels of this wedge.*

The coarse constructor returns this graph-owned D4 evaluation for each of its two retained markers, and the bounded constructor returns it for its one retained marker. In every case the output retains equality with the exact node-[153] residual, the requested slot is D4, and the unevaluated tail is exactly [D5, D6, D7]. If $n = |V(G)|$, the visible local work is at most $8n^3$ in the coarse case and at most $4n^3$ in the bounded case.

Proof. The dependent missing-clause marker stores the component boundary input and the canonical shortest path, so no component, path, or predicate is chosen by the evaluator. Enumerate the vertices in the support of that one path and the canonically ordered pairs of distinct support vertices. Retain a triple precisely when its center is adjacent to both endpoints, and evaluate the two boundary-window copies using the graph-derived attachment labels and the already fixed power-of-two length predicate. This is the declared D4 response by definition, and the cursor equalities give the requested slot and the unchanged tail.

The support has at most n vertices. For each of at most the cube of the support size in raw triples, classification uses at most two adjacency decisions. A retained wedge has at most two further response evaluations, one for each boundary role. Thus the conservative full ledger has at most $4n^3$ checks per marker. The coarse branch has two markers and the bounded branch one. No ambient response, context, state, graph, or universe family is enumerated. $\square$

Remark 7.34 (Open full component semantics). *This is an internal residual interface of original node [153]. Lemma 7.33 discharges its first local clause without adding a diagram node. It does not construct D5 or D6, a Boolean semantics for D7, compatible-response equality, a removal operation, a smaller object, or a CT8 execution. Any continuation must obtain those inputs on the same component support and may not enumerate ambient responses, contexts, states, graphs, or universes.*

Lemma 7.35 (First ambient-cubic owner change). *Suppose node [153] returns its exact all-inside constructor on the computed normalized return Γ' . Assign to every vertex of Γ' its unique ambient-cubic selected-window owner and store this owner sequence. Then there is a first edge of Γ' whose endpoints have distinct owners. The returned package retains the exact two endpoint windows and positions, the literal adjacent vertices, and two opposite tokens, one owned by each endpoint window and incident with the other endpoint. Both token subtypes are exactly `crossWindow`.*

Writing p for the selected-window packing number, $n = |V(G)|$, and $L = |\Gamma'|$, the local work ledger is exactly

$$13pn + |\text{supp}(\Gamma')| 13p + L$$

and is at most

$$n^2 + Q_{\text{base}}(n + 1).$$

Proof. The all-inside equality is the literal output of node [153], so it supplies containment of this same return support in the prepared union of ambient-cubic selected-window supports. Prepare the finite list of ambient-cubic (W, j) slots once. For each of the $|\text{supp}(\Gamma')|$ path vertices, perform one finite slot lookup and store the resulting window and position. Disjointness of the selected packing makes the owner window unique. Scan consecutive entries of this stored table in path order.

If the scan found no change, every return vertex would have the final vertex's selected-window owner. That owner is the original stub window, whereas the first vertex is the original stub neighbour and lies outside the support of that window, a contradiction. Hence the scan returns its first owner change. Each endpoint belongs to its stored selected window and is adjacent to the other endpoint. Since the two owner windows are distinct and the selected supports are disjoint, the other endpoint is external to the first endpoint's window in each orientation. These two oriented incidences are the asserted tokens; because the opposite endpoints remain in the packed selected-window union, both have exact subtype **crossWindow**.

Preparing the ambient-cubic inventory costs $13pn$. The stored owner table uses one scan through at most $13p$ prepared slots per support vertex, and the first-change scan makes one stored-owner comparison per edge. Thus the exact ledger is $13pn + |\text{supp}(\Gamma')| 13p + L$. The packing bound $13p \leq n$ and the inherited bounds $|\text{supp}(\Gamma')| \leq Q_{\text{base}}$ and $L \leq Q_{\text{base}}$ give $n^2 + Q_{\text{base}}(n + 1)$. $\square$

Remark 7.36 (Scope of the owner change). *The owner table ranges over all ambient-cubic selected windows, not the manuscript's selected cold subfamily. Node [153] retains exactly one first cross-window edge and its two opposite tokens. It constructs no boundary successor, target-cycle closure, cold aggregate, D4–D7 coordinate, CT3 input, or density estimate. Its exact token pair feeds the verified zero-check residual at [153].*

Lemma 7.37 (Exact cross-window token-pair residual). *Let node [153] return its complete computed first-cross-window package. Node [153] returns two tokens t_L, t_R , their exact endpoint windows and positions, and adjacent endpoint vertices x_L, x_R , such that*

- (i) t_L and t_R are exactly the two tokens already stored by node [153], and $t_L \neq t_R$;
- (ii) both token subtypes are exactly **crossWindow**;
- (iii) t_L is owned at x_L and points to x_R , whereas t_R is owned at x_R and points to x_L ;
- (iv) the two oriented contributions have the same literal undirected edge, $\{x_L, x_R\} = \{x_R, x_L\}$.

The residual performs exactly zero additional primitive checks after node [153].

Proof. Apply the reusable cross-window token-pair route to node [153]'s retained first-cross-window structure. All token, window, position, and endpoint fields are projections of that same proof-carrying source. Their exact cross-window subtypes and opposite neighbour equalities are inherited verbatim. If the two tokens were equal, equality of their owner projections would identify the two endpoint windows, contradicting the certified owner change. Symmetry of the unordered endpoint pair gives equality of the two literal edges. The route performs no finite search or predicate evaluation, so its additional-check ledger is definitionally zero. $\square$

Remark 7.38 (Scope of the token-pair residual). *Node [153] is the exact terminal residual of the all-inside branch. It constructs no repeated owner, second connector, cycle, cold-family membership, demand or capacity bound, boundary successor, target closure, CT execution, or density estimate.*

Lemma 7.39 (Forced long-support prefix). *Suppose the computed node-[153] result is its long constructor. The identical corridor support contains a canonical initial segment of exactly $Q_{\text{base}} + 1$ positions. The first Q_{base} positions form its base and the unique overflow position has literal index Q_{base} . Every supplied support position is classified exhaustively as lying in this prefix or after it.*

Proof. The long constructor retains the strict inequality $Q_{\text{base}} < |\text{supp}(\Gamma)|$ on the same branch context. Thus the order-preserving inclusion of indices $0, \dots, Q_{\text{base}}$ into the literal support is well defined. Elementary finite-index arithmetic splits a prefix position into an index below Q_{base} or the unique equality case, and splits any supplied support position according to whether its index is below $Q_{\text{base}} + 1$. The handoff itself performs no scan; either local classification uses one natural number comparison. $\square$

Remark 7.40 (Scope of the long prefix). *Node [153] supplies only literal finite positions and their exact inclusion. It assigns no normalized state label, proves no repetition, and constructs no D4–D7 response, CT17 target, offset, compatibility relation, or density estimate. Its exact prefix feeds the verified first-nine classifier [153].*

Lemma 7.41 (First-nine observed-label refinement). *Suppose node [153] is the actual forced-prefix result on the computed long branch. Map its first nine positions, in order, to the identical literal corridor support. Label each resulting vertex by*

$$(d_G(v) \bmod 4, \mathbf{1}_{v \in V(\mathcal{P}_{13})}) \in \text{Fin}(4) \times \{0, 1\}.$$

Then two distinct first-nine positions have the same label. The two actual occurrences form a CT10 datum table whose coarse-label row is populated and whose first missing semantic class is the compatible-response layer. The complete local work is at most $144(|V(G)| + 1) + 9$ primitive checks.

Proof. The exact node-[153] equality identifies every prefix index with the corridor entry at the same literal list position. The displayed label is computed from that graph entry and has exactly eight possible values. The declared order on the first nine positions therefore has an actual repeated label; the executable scan retains its first collision and never enumerates the eight-label universe. CT10 receives exactly the two collided positions. Both lie in the coarse-label class, so the next class, compatible response contexts, is the first missing row and is promoted. There are at most thirty-six pair comparisons. Charging both uncached label evaluations per comparison gives at most $144(|V(G)| + 1)$ checks, and the three-class, two-datum CT10 scan adds at most nine. $\square$

Remark 7.42 (Scope of first-nine refinement). *Node [153] proves only equality of the displayed graph-derived coarse label at two distinct literal prefix occurrences. It proves no equality of D4–D7 or target responses, no common compatible-context completeness, and no CT8 removal or certified smaller object. Node [153] performs only the exact degree refinement below; the stronger obligations remain white at [153].*

Lemma 7.43 (Two-occurrence exact degree refinement). *Consume node [153]’s exact collided pair and promoted CT10 **responseContexts** obligation. Read the full degrees of precisely the two retained literal corridor vertices. The result records either equality of these degrees or their inequality together with the already proved common residue modulo four. Both branches retain distinct occurrences and equal selected-packing membership bits. The visible work is $2|V(G)| + 1 \leq 2(|V(G)| + 1)$.*

Proof. The node-[153] equality fixes both occurrences and their literal vertices. Compute their two degree rows and compare the resulting natural numbers once. The coarse-label equality supplies the residue and packing-bit equalities in both constructors. Only these two local rows are inspected. $\square$

Lemma 7.44 (First-nine local-clause alignment). *Consume node [153]’s exact degree result and retained CT10 `responseContexts` obligation. Compare adjacency of the two literal vertices on precisely the same first nine prefix coordinates. The result is the first mismatching coordinate, together with equality on every preceding coordinate, or equality on all nine coordinates. Both degree constructors are retained without interpretation. Exactly eighteen adjacency evaluations are charged.*

Lemma 7.45 (First-eighteen local-clause alignment). *Consume node [153] exactly. An inherited first-nine mismatch passes through unchanged. Otherwise compare only literal positions 9, ..., 17, returning the first second-block mismatch with its clean prefix or alignment on both fixed nine-coordinate blocks. At most eighteen new adjacency decisions are made; the retained CT10 obligation is not interpreted as response equivalence.*

Lemma 7.46 (First-twenty-seven local-clause alignment). *Consume node [153] exactly. Either inherited mismatch passes through with the same dependent payload. Otherwise compare only literal positions 18, ..., 26, returning the first third-block mismatch with its clean prefix or equality on all first twenty-seven coordinates. At most eighteen new adjacency decisions are made, and the retained degree and CT10 provenance is unchanged.*

Lemma 7.47 (First-thirty-six local-clause alignment). *Consume node [153] exactly. Each inherited mismatch passes through unchanged. Otherwise compare only literal positions 27, ..., 35, returning the first fourth-block mismatch with its clean prefix or equality on all first thirty-six coordinates. At most eighteen new adjacency decisions are made, with all degree and CT10 provenance retained.*

Lemma 7.48 (Long-prefix compatible-response frontier). *Consume node [153] exactly and retain its four possible mismatch leaves and its first-thirty-six-aligned leaf, together with the inherited degree and CT10 provenance. For a mismatch in block $j \in \{0, 1, 2, 3\}$, retain its original first-hit certificate and embed its local coordinate with offset $9j$ into the same 36-position corridor schedule. The two literal adjacency responses differ at that exact coordinate, so every mismatch leaf now gives a proved local distinguishing response. On the first-thirty-six-aligned leaf, retain all four alignment proofs and record exactly two missing producers: completeness of these clauses for the manuscript’s D4–D7 response semantics, and a certified reduction indexed by the exact collided pair. The current consumer inspects one retained constructor; it executes neither the reserved CT8 response scan nor any removal.*

Remark 7.49 (Open full long-prefix semantics). *This is a verified internal residual interface of original node [153]. Its four mismatch leaves now have typed distinguishing contexts on the fixed schedule. The aligned leaf still requires the complete compatible D4–D7 response layer and a certified exact-pair reduction before CT8 may run. It may inspect only explicit local clauses and may not enumerate an ambient response, context, state, graph, or universe family.*

Remark 7.50 (Formal weighted live/cold handoff). *The formal handoff processes the exact node-[21] packing in its declared order. Starting from the original graph as the empty aggregate, it either accepts a graph-semantic compatible extension of the current aggregate or rejects that identical window while leaving the aggregate unchanged. A rejection stores the current aggregate and $\neg \text{Nonempty}(\text{Extension}(A, P))$; it does not assert that P lacks a standalone local package. The*

accepted aggregate stores recoverable simultaneous local choices and a bounded injective spine code. The split is proof-level and enumerates no ambient graph, context, state subset, or Boolean universe. The pointwise thirteen-bit classifier at internal stages of node [153] remains auxiliary same-window geometry.

Definition 7.51 (Hot/cold window ledger). *Set*

$$c_{\text{hot}} := 118.10858100609198\dots, \quad \theta_{\text{win}} := \frac{3}{2c_{\text{hot}}} = 0.012700177982179218\dots$$

For a packing density $\theta = p_{13}/n$, set

$$\tau(\theta) := \frac{15\theta}{1 - 13\theta}.$$

Then

$$\tau(\theta) < \frac{1}{4} \iff \theta < \frac{1}{73} = 0.0136986301369863\dots,$$

and

$$\tau(\theta) < \frac{3}{13} \iff \theta < \frac{1}{78} = 0.0128205128205128\dots$$

The value $1/78$ is the route-8 private-carrier threshold, and $1/73$ is the negative-net-charge threshold.

Fix the declared order on a maximal vertex-disjoint induced- P_{13} packing $\mathcal{P}$. A packed window is hot when its compatible extension is accepted into the aggregate accumulated from the preceding windows. It is cold when that extension type is empty at its exact turn. Rejection retains the same aggregate and the same window. The one sequential run gives the exact partition

$$\mathcal{P} = \mathcal{P}_{\text{hot}} \sqcup \mathcal{P}_{\text{cold}}, \quad C := |\mathcal{P}_{\text{cold}}|.$$

Only a hot certificate contributes the full conditional-fibre product cost. A cold certificate retains its exact window and is routed to the local first-failure branch.

Definition 7.52 (Surviving cold branch). *The surviving cold branch is the branch on which:*

- (i) no target cycle has been found;
- (ii) no sparse surplus exit, target-defective quotient, target-complete compression, or delocalization exit is active;
- (iii) no unpaid exit-(4) peel remains;
- (iv) no Type B bridge or handoff residual remains outside its ledger;
- (v) no Type A route-8 residual remains outside [Theorem 15.48](#); and
- (vi) the sparse surplus branch has already supplied the spine estimate in the form $m = \frac{3}{2}n + o(n)$ and $\sigma(G) = 2m - 3n = o(n)$.

On this branch, $|V_{\geq 4}(G)| \leq \sigma(G) = o(n)$. Hence all but $o(n)$ packed P_{13} -windows are ambient-cubic.

Remark 7.53 (Hot-failure counting consequence). *After the weighted package and hot-family aggregation of Remark 7.1 have been constructed, the following counting implication applies on Definition 7.52: if the live-hot entropy comparison does not close, then*

$$C \geq (\theta - \theta_{\text{win}})n - o(n).$$

In particular, at the route-8 threshold $\theta = 1/78$,

$$C \geq 0.000120334838333602 n - o(n),$$

and at the negative-net-charge threshold $\theta = 1/73$,

$$C \geq 0.000998452154807083 n - o(n).$$

Definition 7.54 (Cold skeleton and branch excess). *Let*

$$\mathcal{P}_{\text{cold}}^{\text{cub}} := \{P \in \mathcal{P}_{\text{cold}} : V(P) \cap V_{\geq 4}(G) = \emptyset\}.$$

By Definition 7.52,

$$|\mathcal{P}_{\text{cold}} \setminus \mathcal{P}_{\text{cold}}^{\text{cub}}| = o(n).$$

For each $P \in \mathcal{P}_{\text{cold}}^{\text{cub}}$, contract the path P to one skeleton vertex p and keep one incident half-edge for every edge of G leaving P . The resulting incidence object is the ambient-cubic cold skeleton $\mathfrak{S}_{\text{cold}}$. If $s(P)$ is the number of external stubs of P , define the branch-excess contribution of P by

$$b(P) := \max\{0, s(P) - 2\},$$

and set

$$b(\mathfrak{S}_{\text{cold}}) := \sum_{P \in \mathcal{P}_{\text{cold}}^{\text{cub}}} b(P).$$

The branch-excess units are selected concretely as follows. The global lexicographic order on vertices and edges induces an order on the external stubs of every ambient-cubic cold window. The first two stubs of P are called the transit stubs; the remaining $s(P) - 2$ stubs, when $s(P) > 2$, are the selected branch-excess half-edges of P . Let

$$\mathcal{E}_{\text{br}}(P)$$

be this selected set and let

$$\mathcal{E}_{\text{br}} := \bigcup_P \mathcal{E}_{\text{br}}(P).$$

Thus

$$|\mathcal{E}_{\text{br}}| = b(\mathfrak{S}_{\text{cold}}).$$

All later cold-branch counts are counts of these selected half-edges. Hence a stub can be charged at most once at its cold-window endpoint.

Lemma 7.55 (Cold window stub excess). *On Definition 7.52,*

$$b(\mathfrak{S}_{\text{cold}}) \geq 13C - o(n).$$

Proof. Let P be an ambient-cubic induced P_{13} with vertices $v_0, \dots, v_{12}$. The internal path edges contribute

$$2 \cdot 12 = 24$$

to the sum of degrees inside P , while the ambient cubic degrees contribute

$$13 \cdot 3 = 39.$$

Thus P has exactly $39 - 24 = 15$ external stubs. In the skeleton, a degree-two corridor through p can absorb at most two of these stubs without branching. Hence $b(P) = 15 - 2 = 13$ for every ambient-cubic cold window. By [Definition 7.52](#), only $o(n)$ cold windows are not ambient-cubic. Summing over the remaining windows gives the claim. $\square$

Lemma 7.56 (Literal short self-return filter). *Let J be an initial segment of a cold return corridor, starting at a boundary stub attached to offset $a \in \{0, \dots, 12\}$, and let x be its current endpoint. For an offset $r \in \{0, \dots, 12\}$, suppose that $xv_r \in E(G)$. Then J , the two attachment edges, and the segment of the induced P_{13} between v_a and v_r form a simple cycle of length*

$$|E(J)| + 2 + |a - r|.$$

Consequently, if this length is a power of two for some offset with the stated attachment edge, the corridor closes at case (F1). If no such offset exists, case (F1) is absent at this stage and the same stage is tested for (F2), (F3), and (F4).

Proof. The corridor prefix is a simple path in the graph obtained by deleting the cold-window interiors, so its support is disjoint from the anchor window. Its initial boundary stub supplies the edge to v_a , and the hypothesis supplies the edge from x to v_r . Appending these edges and the unique segment of the induced path between v_a and v_r gives a simple cycle. The three pieces have lengths $|E(J)|$, 2, and $|a - r|$, respectively. The final statement is the stage-major priority rule in [Definition 7.58](#). $\square$

Definition 7.57 (Cold bounded germ). *A cold bounded germ is a finite boundaried support with two boundary interfaces x, y and two same-interface x, y representatives $Q[x, y]$ and E . In a terminal corridor these representatives are the two bounded completion strands between the same interfaces; in a repeated-state corridor one representative is the actual corridor segment and the other is the canonical representative determined by the repeated cold corridor state. The germ also carries the inherited boundary degree profile, P_{13} -window labels, and target-response profile. Its increment is*

$$\delta := |E| - |Q[x, y]|.$$

It is length-changing if $\delta \neq 0$, and equal-length if $\delta = 0$. The word bounded means that the number of internal vertices and all boundary data are bounded by the fixed finite cut-state constant used by the local target algebra; consequently there are only finitely many germ types.

Definition 7.58 (Cold corridor states and first failures). *Let X_{cold} be the union of the ambient-cubic cold windows, with their internal path edges retained. Delete the interiors of these windows and look at a connected component K of the remaining outside graph. The boundary stubs of K are the edges from K to ambient-cubic cold windows. By [Lemma 3.4](#), no such component has only one boundary stub, since that unique boundary edge would be a bridge of G .*

Order the boundary stubs of K lexicographically as

$$h_1, \dots, h_m \quad (m \geq 2)$$

and give them the cyclic successor relation $h_i \mapsto h_{i+1}$, with $h_{m+1} = h_1$. If $h_i = \epsilon \in \mathcal{E}_{\text{br}}(P)$ is a selected branch-excess half-edge, choose inside K the lexicographically first simple path joining the outside endpoint of h_i to the outside endpoint of h_{i+1} . Together with the two boundary stubs h_i, h_{i+1} , this path is the cold return corridor of ϵ . Thus each selected branch-excess half-edge has exactly one corridor, and each boundary stub is the successor of at most one selected half-edge.

For an initial segment J of this corridor, let $T(J)$ be its two active boundary interfaces. Its cold corridor state is the finite two-boundary cut-state obtained from $\rho_{T(J)}^{\text{ex}}(J)$ by retaining exactly the boundary-degree profile, the two active boundary half-edges, the cold-window offsets met at the two interfaces, and the declared local coordinates of [Definition 3.94](#) whose support is contained in the bounded active interface. It is not the full labelled prefix. If two prefixes have the same finite cut-state but differ in exact target response against some compatible context, that discrepancy is recorded as a first failure of type (F2) below. Thus, after excluding (F2), equality of cold corridor states is equality for every target-response coordinate used by the local replacement. Since the boundary has size two and the retained window and local-coordinate labels are finite, the set of possible cold corridor states is bounded by a constant Q_{cold} depending only on the fixed declared signature. Set

$$M_{\text{cold}} := Q_{\text{cold}} + 30, \quad B_{\text{cold}} := 15(1 + 3(2^{M_{\text{cold}}+2} - 1)), \quad D_{\text{cold}} := M_{\text{cold}}B_{\text{cold}} + 1.$$

The additive 30 only covers the two P_{13} -window interfaces and the two boundary stubs. Thus M_{cold} is a uniform upper bound for the number of vertices in a first-failure cold exchange, and B_{cold} is a uniform upper bound for the number of selected cold half-edges whose first-failure support can contain a fixed subcubic vertex.

The first failure of ϵ is computed in stage-major order. Initial segments are read in their order along the corridor. At a fixed initial segment, the thirteen window offsets are inspected in the order $0, 1, \dots, 12$ for (F1), after which (F2), (F3), and (F4) are tested at that same segment:

- (F1) the current endpoint has an actual attachment edge to a displayed cold-window offset and the literal completion of [Lemma 7.56](#) has power-of-two length;
- (F2) two prefixes with the same displayed boundary data have different target response against some compatible outside context;
- (F3) two prefixes have the same exact target response against every outside context and one gives a strictly smaller proper representative;
- (F4) the corridor first enters a declared Type B handoff envelope or the route-8 carrier support already recorded in the branch state;
- (F5) every initial segment is negative for (F1)–(F4), and either the corridor reaches its successor boundary stub before $Q_{\text{cold}} + 1$ states have been read, or a cold corridor state repeats.

In case (F5), the bounded support is the whole terminal corridor in the first subcase, and the support between the two equal states in the second subcase. With its two boundary interfaces, this bounded support is the first-failure cold exchange of ϵ .

Remark 7.59 (Pending cold semantic producer). The literal F1 constructor and the stage-major finite runner have been formalized. On the locally surviving branch, the structural corridor runner also proves the exhaustive F5 support split: after exclusion of an actual F2 distinction, it returns either the whole bounded terminal corridor or the first repeated structural-state span; an F3 compression is already excluded by minimality. The runner is not yet an unconditional implementation

of node [153]: it still requires the graph-owned composition of these pieces into one stage-major F2–F5 producer, including the exact Type B or route-8 key for F4, and it must promote the terminal or repeated support to the full two-boundary cold bounded germ with its finite target-response table and replacement data. These certificates are local to the stored corridor prefixes; their construction does not enumerate the ambient context space.

Lemma 7.60 (Cold corridor first-failure routing). *Assume Definition 7.52. For every $\epsilon \in \mathcal{E}_{\text{br}}$, the cold return corridor of ϵ has a first failure. The first failure is routed as follows:*

- (i) case (F1) is a dyadic cycle in G ;
- (ii) case (F2) is a target-defective quotient, hence belongs to the sparse exit or to the exit-(4) ledger;
- (iii) case (F3) is a target-complete compression of a proper support;
- (iv) case (F4) is an already named Type B or route-8 handoff;
- (v) case (F5) is a cold bounded germ in the sense of Definition 7.57.

Consequently, on the surviving cold branch, every selected branch-excess half-edge that is not in the already discarded $o(n)$ handoff or surplus ledgers supplies a canonical cold bounded-germ incidence.

Proof. The return corridor exists by Lemma 3.4. Follow its initial segments in stage-major order as specified in Definition 7.58. At each segment, absence of every literal (F1) attachment advances to (F2), then (F3), then (F4), before the next segment is read. If the successor boundary stub is reached before $Q_{\text{cold}} + 1$ states have been read, then the whole corridor, together with the two boundary stubs and the two P_{13} -window interfaces, has at most $M_{\text{cold}} = Q_{\text{cold}} + 30$ vertices and gives the terminal subcase of (F5). Otherwise, before the successor is reached, $Q_{\text{cold}} + 1$ states are read; two states are equal by the definition of Q_{cold} , and the first such equality gives the repeat subcase of (F5). Hence a first failure always exists.

If the first failure is (F1), its recorded attachment edge and corridor prefix give the literal cycle of Lemma 7.56; its length is a power of two, which is impossible in a counterexample. If it is (F2), the actual quotient is valid only for the current outside context and fails for another compatible context. By Lemma 3.102, this is not target-complete, so it is a target-defective quotient. In the global branch ledger such defects are exactly the sparse exits; in the Type A location they are the exit-(4) peels, and both are excluded in Definition 7.52.

In (F3), context-universality holds and the displayed representative is strictly smaller while preserving the boundary-degree profile and the exact target response. Thus it satisfies Definition 3.96 and is forbidden by Lemma 3.103 and Corollary 3.105. In (F4), the corridor has reached precisely one of the declared interfaces of Definition 14.19 or the route-8 carrier support, so the charge is transferred to the already existing Type B or route-8 ledger.

It remains to consider (F5). In the terminal subcase, the whole corridor has bounded size and two boundary interfaces. In the repeat subcase, the two repeated states have the same two boundary interfaces, the same inherited boundary-degree profile, the same window labels, and the same exact target-response coordinates; the support between them is bounded by Q_{cold} . In both subcases the support carries two same-interface representatives: the actual corridor representative and the terminal or repeated-state exchange representative with the same retained cut-state. Thus it is a cold bounded germ in the sense of Definition 7.57. All cases except (F5) are excluded on the surviving branch, up to the $o(n)$ windows and corridors already charged to the non-ambient-cubic, handoff, or surplus ledgers. Hence every remaining selected branch-excess half-edge supplies a canonical bounded-germ incidence. $\square$

Lemma 7.61 (Cold branch-excess extracts bounded germs). *Assume Definition 7.52. If the cold skeleton has branch excess $b = b(\mathfrak{S}_{\text{cold}})$, then it contains a family of pairwise vertex-disjoint candidate cold bounded germs of size*

$$N_{\text{germ}} \geq \frac{b}{D_{\text{cold}}} - o(n).$$

In particular, after hot failure,

$$N_{\text{germ}} \geq \frac{13C}{D_{\text{cold}}} - o(n) \geq \frac{13(\theta - \theta_{\text{win}})}{D_{\text{cold}}} n - o(n).$$

Proof. By Definition 7.54, there are exactly $b = |\mathcal{E}_{\text{br}}|$ selected branch-excess half-edges, and each is selected once. By Lemma 7.60, after deleting the already routed $o(n)$ surplus, non-ambient-cubic, Type B, and route-8 handoff incidences, each remaining selected half-edge supplies one canonical first-failure cold bounded germ. Let $\mathcal{G}_{\text{cand}}$ be this candidate family. Then

$$|\mathcal{G}_{\text{cand}}| \geq b - o(n).$$

It remains only to extract a vertex-disjoint subfamily without double-counting overlapping local supports. Each member of $\mathcal{G}_{\text{cand}}$ has at most M_{cold} vertices by Definition 7.58. If a candidate support contains a vertex of degree at least 4, then the corresponding corridor first enters the high-degree handoff ledger and was already removed; hence every remaining candidate lies in the subcubic part of the surviving branch. In a subcubic graph the ball of radius $M_{\text{cold}} + 2$ around a vertex has at most

$$1 + 3(2^{M_{\text{cold}}+2} - 1)$$

vertices, and every ambient-cubic cold window contributes at most 15 external stubs. Therefore a fixed vertex belongs to the support of at most

$$B_{\text{cold}} = 15(1 + 3(2^{M_{\text{cold}}+2} - 1))$$

candidate germs. A fixed candidate support has at most M_{cold} vertices, so it meets at most $M_{\text{cold}}B_{\text{cold}}$ other candidate supports. The intersection graph of $\mathcal{G}_{\text{cand}}$ thus has maximum degree at most $M_{\text{cold}}B_{\text{cold}}$.

A greedy maximal independent set in a graph of maximum degree Δ has size at least the number of vertices divided by $\Delta + 1$. With $\Delta = M_{\text{cold}}B_{\text{cold}}$, this gives a vertex-disjoint candidate germ family of size

$$N_{\text{germ}} \geq \frac{|\mathcal{G}_{\text{cand}}|}{M_{\text{cold}}B_{\text{cold}} + 1} \geq \frac{b}{D_{\text{cold}}} - o(n).$$

The remaining inequalities use Lemma 7.55 and Remark 7.53. □

Lemma 7.62 (Bounded germ trichotomy). *No length-changing cold bounded germ survives on Definition 7.52. More explicitly, every such germ falls into one of the following three cases:*

- G1. Hit-realized. *Some compatible live completion and window offset close a dyadic cycle. This contradicts the counterexample condition.*
- G2. Hit-distinguished. *Some compatible outside context distinguishes the two representatives by dyadic truth value without already realizing the cycle in the current graph. The induced quotient is target-defective, so it is routed to the sparse exit or exit-(4) ledger.*

G3. Silent. *No compatible outside context and no relevant scale distinguishes the two representatives. Then replacing the longer representative by the shorter one preserves the boundary degree profile and the target response against every context, creates no dyadic cycle, and strictly decreases the support. This is a nontrivial target-complete compression of a proper support.*

Each case is impossible on the surviving cold branch.

Proof. The cases are exhaustive by whether a compatible completion realizes a dyadic hit, distinguishes dyadic truth without realization in G , or never distinguishes the two representatives.

In G1, the completed representative plus the relevant window segment is a cycle whose length is a power of two. This contradicts the definition of a counterexample.

In G2, the two local responses agree in the actual quotient but disagree in a compatible context. By Lemma 3.102, such an identification is not target-complete; equivalently it is a target-defective quotient. This is one of the named surplus exits or, inside Type A, an exit-(4) peel. Both are excluded by Definition 7.52.

In G3, all compatible contexts give the same target response for the two representatives. Because $\delta \neq 0$, one representative is strictly longer. Replacing the longer representative by the shorter one keeps the same boundary vertices, the same boundary degree profile, and the same target response against every context. It cannot create a dyadic cycle, since any such cycle would distinguish the two responses. Thus the replacement is a smaller target-complete representative of a proper support, forbidden by Lemma 3.103 and Corollary 3.105. Hence no length-changing bounded germ survives. $\square$

Definition 7.63 (Finite same-interface cold table). *The same-interface cold table is the finite table whose rows are equal-length cold bounded germs and the short self-return rows that reach exhaustive case (F5) after the literal attachment test of Lemma 7.56, recorded with:*

(T1) *the two boundary vertices and their boundary-degree profile;*

(T2) *the two terminal cold-window stubs and their offsets in $\{0, \dots, 12\}$;*

(T3) *the exact response profile generated by Definition 3.94;*

(T4) *the target truth value of every compatible completion represented by that exact profile.*

The table is finite because the support size, boundary size, window labels, and declared coordinate labels are bounded by Definition 7.58. A row is realizing if some compatible completion gives a power-of-two cycle, distinguishing if two same-interface representatives have different target truth value in some compatible context, and neutral otherwise.

Lemma 7.64 (Same-interface cold table closure). *Every row of the same-interface cold table is routed to one of the already closed outcomes: a dyadic cycle, a target-defective quotient, an existing Type B or route-8 handoff, or a target-complete proper-support compression. In particular, an equal-length cold bounded germ and an F5 short self-return cannot be a terminal cold residual.*

Proof. Let R be a row of the table. If R is realizing, the corresponding completion is a power-of-two cycle in G , which contradicts the counterexample condition. If R first meets a declared Type B or route-8 interface, then by the definition of the surviving cold branch the charge is transferred to that already closed ledger.

Assume R is not realizing and does not enter a handoff interface. If there is a compatible context distinguishing the two same-interface representatives, then Lemma 3.102 says that their

identification is not target-complete. This is a target-defective quotient, hence a sparse exit or an exit-(4) peel, both excluded by [Definition 7.52](#).

It remains to handle a neutral row. Neutrality means that the two representatives have the same boundary-degree profile and the same target response against every compatible context. The row is a first-failure row: otherwise it would not have entered the cold table. Therefore the quotient identifying the two same-interface representatives removes at least one declared branch-excess coordinate while preserving every coordinate needed by the target response. By [Definition 3.98](#), such a proper-support target-complete quotient is admissible only when it has a strictly smaller proper representative in the sense of [Definition 3.96](#). That representative satisfies the hypotheses of [Lemma 3.103](#), and is forbidden by [Corollary 3.105](#). Hence no neutral row survives either. $\square$

Lemma 7.65 (Increment arithmetic for cold germs). *Let a homogeneous cold germ have increment δ , oriented so $\delta \geq 0$. In this statement an achievable length is a literal completion represented in the germ's exact response table; membership in an arithmetic interval or residue class alone does not supply an attachment edge or a graph cycle. The following cases exhaust the finite arithmetic:*

(a) *If $1 \leq \delta \leq 12$, then the achievable length blocks*

$$[L + j\delta, L + j\delta + 12] \quad (0 \leq j \leq N)$$

overlap. If their union spans a dyadic scale, the branch is G1; otherwise the whole family lies in one bounded dyadic gap and is a bounded germ type handled by [Lemma 7.62](#).

(b) *If $\delta \geq 13$ and the doubling orbit modulo δ hits one of the thirteen smear residues $L, L + 1, \dots, L + 12$, the branch is G1.*

(c) *If $\delta \geq 13$ and the doubling orbit avoids all thirteen smear residues, the odd part of δ gives a periodic carrier. A target-visible carrier is G2, and a target-invisible carrier is G3.*

(d) *If $\delta = 0$, the equal-length switch belongs to the finite same-interface cold table; each table row is closed by [Lemma 7.64](#).*

For odd δ , the sufficient hit criterion

$$\text{ord}_\delta(2) > \delta - 13$$

forces case (b). For even $\delta = 2^a u$, the same criterion applies to the odd part u after the bounded transient $k < a$.

Proof. For (a), consecutive blocks overlap because the next block begins at $L + (j+1)\delta \leq L + j\delta + 12$. Hence their union is an interval. An interval containing a power of two gives G1; if it contains none, all corresponding completions live inside one bounded gap between consecutive powers of two, so only finitely many target-response types occur and [Lemma 7.62](#) applies.

For (b), a congruence $2^k \equiv L + r \pmod{\delta}$ with $0 \leq r \leq 12$ means that, after adding the appropriate number of homogeneous copies, one attainable length is 2^k . This is exactly a hit-realized germ.

For (c), avoidance of all smear residues means the orbit of powers of two is confined to the complement of the thirteen target residues modulo the odd part of δ . The residue class is then a periodic carrier. If some compatible context sees the carrier, the two representative responses are target-distinguished and G2 applies. If no compatible context sees it, the carrier is target-equivalent and G3 applies.

For (d), equal-length switches cannot be charged as length-changing increments. They are finite because the germ support and boundary data are bounded, and [Lemma 7.64](#) closes every row of the finite same-interface table.

If δ is odd and $\text{ord}_\delta(2) > \delta - 13$, the doubling orbit has more residues than the complement of the thirteen smear residues. Therefore it must hit a smear residue, giving (b). If $\delta = 2^a u$, the initial powers with $k < a$ form a bounded transient, and for $k \geq a$ the congruence reduces to the odd modulus u . $\square$

Table 22: Cold branch quantitative residual ledger.

Case	Trigger	Quantitative constraint	Charged resource	Outcome
0	Sparse non-spine branch	$\sigma(G)$ routes to sparse exit, Type B, or $o(n)$	sparse capacity-token ledger	closed or spine supplied
1	$\theta < 1/78$	$\tau(\theta) < 3/13$	route-8 carrier collision	closed
2	$1/78 \leq \theta < 1/73$	delicate density interval	hot/cold branch	routed through rows 4–16
3	$\theta \geq 1/73$	$\tau(\theta) \geq 1/4$	hot/cold branch	routed through rows 4–16
4	Hot package live	$(c_{\text{hot}} - o(1)) \mathcal{P}_{\text{hot}} \log_2 n$ bits	entropy budget	density cap
5	Hot branch fails	$C \geq (\theta - \theta_{\text{win}})n - o(n)$	cold skeleton	linear cold family
6	Non-cubic cold windows	at most $\sigma(G) = o(n)$	surplus ledger	sublinear loss
7	Ambient-cubic cold window	exactly 15 external stubs	stub count	13 branch-excess units
8	Cold skeleton extraction	$N_{\text{germ}} \geq 13C/D_{\text{cold}} - o(n)$	first-failure and bounded-overlap constant D_{cold}	bounded germs
9	Short self-return	F1 requires an actual endpoint–offset attachment and dyadic literal completion	stage-major local scan	F1, then same-stage F2/F3/F4; only exhaustive F5 enters the finite table
10	G1	live completion hits a shifted power	target algebra	dyadic cycle
11	G2	context distinguishes the representatives	target-defect ledger	sparse exit or exit-(4)
12	G3	no context distinguishes the representatives	replacement	compression contradiction
13	$1 \leq \delta \leq 12$	overlapping length blocks	target algebra or finite table	G1, G2, or G3

Case Trigger	Quantitative constraint	Charged resource	Outcome
14 $\delta \geq 13$ orbit hits a represented offset residue	$2^k \equiv L + r \pmod{\delta}$, $0 \leq r \leq 12$, with a literal compatible completion	target algebra	G1
15 $\delta \geq 13$ orbit avoids smear	periodic carrier on odd part	carrier quotient	G2 or G3
16 $\delta = 0$	finite same-interface cold table	target response table	hit, defect, handoff, or compression

Remark 7.66 (Pending cold-branch closure). *The intended cold route has no terminal survivor after the near-cubic spine estimate has been supplied. On Definition 7.52, either the route-8 carrier inequality closes $\theta < 1/78$, or the live-hot entropy comparison closes, or the hot failure produces a cold family whose branch excess yields a bounded germ; that germ is routed by Lemmas 7.60, 7.62, 7.64 and 7.65 to a dyadic cycle, a target-defective quotient or exit-(4) route, a Type B handoff already in its ledger, a route-8 closure, or a target-complete compression. Completing this route requires the graph-owned corridor and bounded-germ producers, the overlap bound D_{cold} , and the finite context-reflection theorem described above. Nodes [151]–[152] establish the ambient-cubic filtering, inject every discarded non-cubic cold window into the inherited surplus ledger, and prove the exact 15-stub/13-excess inequality, including the exact square-root loss bound inherited from node [21]. They do not supply the remaining semantic producers.*

Remark 7.67 (Pending near-cubic P_{13} window density bounds). *The intended window-only conclusion is*

$$\theta \leq \theta_{\text{win}} + o(1), \quad \theta_{\text{win}} := \frac{1.5}{118.108581006} = 0.01270017798 \dots$$

In the high-entropy branch of Proposition 11.9, where the remainder also contributes $\frac{1}{10} \log_2 n$ bits per remainder vertex, the sharper bound is

$$\theta \leq 0.01198542083 \dots + o(1).$$

Together with the high-entropy remainder contribution, the intended sharper conclusion and its remainder consequences are

$$|R| \geq (1 - 13\theta_{\text{win}})n - o(n) = 0.83489768623 \dots n - o(n),$$

and

$$\frac{15\theta}{1 - 13\theta} \leq \tau_{\text{win}} + o(1), \quad \tau_{\text{win}} := \frac{15\theta_{\text{win}}}{1 - 13\theta_{\text{win}}} = 0.22817486846 \dots < \frac{1}{4}.$$

These inequalities follow algebraically once the weighted live-package aggregation and the terminal cold routes of Remarks 7.1 and 7.66 are proved. The normalized high-entropy comparison would then be

$$\frac{1 - 13\theta}{10} + 118.108581006 \theta = 0.1 + 116.808581006 \theta \leq 1.5 + o(1). \quad (1)$$

Node [24] is the complementary branch handoff: it retains the exact finite ceiling and the cross-multiplied window-only cap selected at node [22]. It also names the normalized high-entropy refinement $116.808581006 p_{13} \leq 1.4n$ as a downstream requirement. Proving that refinement and closing the strict-quarter hot/cold route remain later obligations; they are not inserted into the node-[24] handoff.

Framework branch table. The P_{13} -window entropy and hot/cold closure uses the following branch table.

Branch	Residual condition	Demand	Payer	Bounded multiplicity	Closure
Window package	Packed windows are active after the sparse surplus exits.	Independent window coordinates.	The near-cubic skeleton budget.	The vertex-disjoint window packing.	Remarks 7.1 and 7.67 .
Route-8 threshold	$\theta < 1/78$.	Private-carrier load.	The route-8 carrier inequality.	The carrier ledger.	Closed by Theorem 15.48 .
Hot branch	The graph-owned realization interface is available.	Hot window bits.	Skeleton entropy.	Packed windows.	Pending the realization and commuting-gluing producers of Remark 7.1 .
Cold branch	The classifier returns explicit omitted assignments.	Branch-excess stubs.	First-failure extraction.	The bounded constant D_{cold} .	Pending the corridor, overlap, germ-coverage, and target-defect continuation producers of Remark 7.66 .
High entropy	The remainder pays $ R \log_2 n/10$ bits.	Window plus remainder bits.	The same skeleton entropy.	One target-coordinate family.	The sharper cap in Remark 7.67 .

8 The P_{13} -free remainder and positive deficiency

Let $\mathcal{P}$ be a maximal vertex-disjoint family of induced copies of P_{13} in G . Let

$$W = \bigcup_{P \in \mathcal{P}} V(P), \quad R = G - W.$$

Definition 8.1 (Exact P_{13} -remainder residual). *The finite output carried from the window-density branch into Residual A is an integer window ceiling U_{13} together with the packing certificate*

$$p_{13} \leq U_{13}.$$

Put $L_R := n - 13U_{13}$, with natural-number truncated subtraction. Residual A retains this certificate, the selected maximum packing, and its literal complement $R = G - W$. Since the packed supports are pairwise disjoint,

$$|R| + 13p_{13} = n,$$




Figure 17: Visual definition of the packed-window and remainder ledger. The blue region is $W = \bigcup_{P \in \mathcal{P}} V(P)$, where $\mathcal{P}$ is a maximal vertex-disjoint family of induced P_{13} windows. Black vertices and black edges are window vertices and their internal path edges; the displayed indices $0, \dots, 12$ name the vertices of $P^{(1)} = p_0 \dots p_{12}$. The green region is the remainder $R = G - W$; white vertices are remainder vertices, green edges are edges internal to R , and R_1, R_2 mark components of R . The packing size is $p_{13} = |\mathcal{P}|$, its density is $\theta = p_{13}/n$, $|W| = 13p_{13}$, and $|R| = (1 - 13\theta)n - o(n)$. Maximality of $\mathcal{P}$ makes every component of R induced- P_{13} -free; the HSS theorem then forbids any internal subgraph of minimum degree at least 3 inside a component of R . Orange arrows are window–remainder incidences, counted by $e(R, W)$ and by the oriented boundary ledger $\partial_E R$. Since each packed window is induced, $G[W]$ has exactly $12p_{13}$ path edges plus the cross-window edges counted by $e_{\times}(W)$. The exact degree-sum identity is $e(R, W) + 2e_{\times}(W) = 15p_{13} + \sigma_W$, so every surplus unit on a window must be spent either as an additional window–remainder incidence or as two half-incidences in a cross-window join.

and therefore the local meaning of “ R is large” is the exact conclusion

$$L_R = n - 13U_{13} \leq |R|.$$

The same tuple crosses unchanged from node [25] to node [26]. No packing, remainder, or size parameter is selected again at that panel boundary.

By the near-cubic window-only part of [Remark 7.67](#),

$$|R| \geq 0.83489768623 \dots n - o(n).$$

Equivalently, the density branch supplies a choice of U_{13} with $n - 13U_{13} \geq 0.83489768623 \dots n - o(n)$; the exact local handoff itself uses only the certified natural-number packing inequality. By the maximality of $\mathcal{P}$, the graph R contains no induced P_{13} , since any such copy would extend $\mathcal{P}$; hence every component of R is P_{13} -free. We claim that no component of R contains a subgraph of minimum degree at least 3. Indeed, suppose $H \subseteq R$ had $\delta(H) \geq 3$. Passing to the induced subgraph $G[V(H)] \supseteq H$ only adds edges, so $\delta(G[V(H)]) \geq \delta(H) \geq 3$; and $G[V(H)]$, being an induced subgraph of a P_{13} -free component of R , is itself P_{13} -free. By [Theorem 5.1](#), $G[V(H)]$ then contains a power-of-two cycle, which is also a cycle of G , contradicting that G is a counterexample.

Definition 8.2 (Positive deficiency and surplus). *For a subgraph $X \subseteq G$, define its positive deficiency by*

$$\partial_+(X) = \sum_{v \in V(X)} \max\{0, 3 - d_X(v)\}.$$

Define the ambient surplus carried by vertices of X by

$$\sigma(X) = \sum_{v \in V(X)} \max\{0, d_G(v) - 3\}.$$

For the whole graph, write

$$\sigma(G) = \sum_{v \in V(G)} (d_G(v) - 3).$$

Positive deficiency, not signed cubic charge, measures external degree demand: high-degree vertices inside X cannot cancel the fact that vertices of internal degree below 3 need outside incidences.

Definition 8.3 (Window-remainder surplus split). *Let*

$$\sigma_W := \sum_{w \in W} (d_G(w) - 3), \quad \sigma_R := \sum_{v \in R} (d_G(v) - 3).$$

Since $\delta(G) \geq 3$, these sums are nonnegative,

$$\sigma_R = \sigma(R), \quad \sigma(G) = \sigma_W + \sigma_R.$$

Let $e_\times(W)$ be the number of edges of G whose endpoints lie in two distinct packed P_{13} -windows. The packed windows are induced, so the edges of $G[W]$ are exactly the $12p_{13}$ path edges inside the windows together with these $e_\times(W)$ cross-window edges.

Lemma 8.4 (External-incidence supply for the P_{13} -free remainder). *Assume [Definition 4.1](#). The P_{13} -free remainder satisfies*

$$\partial_+(R) \leq e(R, W) \leq 15p_{13} + o(n)$$

and hence

$$\partial_+(R) - \sigma_R \leq 15p_{13} + o(n).$$

Equivalently,

$$\frac{\partial_+(R) - \sigma_R}{|R|} \leq \frac{15\theta}{1 - 13\theta} + o(1).$$

Proof. Since $\delta(G) \geq 3$, every vertex $v \in R$ has $d_G(v) \geq 3$, so v is deficient in R only because some of its incidences leave R . Writing $d_G(v) = d_R(v) + e_v$, where e_v is the number of edges from v to $W = G - R$, and using $\max\{0, 3 - d_G(v)\} = 0$,

$$\max\{0, 3 - d_R(v)\} \leq \max\{0, 3 - d_G(v)\} + e_v = e_v.$$

Summing over $v \in R$,

$$\partial_+(R) = \sum_{v \in R} \max\{0, 3 - d_R(v)\} \leq \sum_{v \in R} e_v = e(R, W),$$

the number of edges between R and W .

It remains to bound $e(R, W)$ by the outside-incidence capacity of the windows. Each $P \in \mathcal{P}$ is an *induced* P_{13} , so the subgraph of G on $V(P)$ has exactly its 12 path edges. The number of edges

leaving $V(P)$ (to R or to another window) equals the degree sum on $V(P)$ minus twice the internal edges:

$$\sum_{w \in V(P)} d_G(w) - 2 \cdot 12 = \left(39 + \sum_{w \in V(P)} (d_G(w) - 3) \right) - 24 = 15 + \sigma(P),$$

where $\sigma(P) = \sum_{w \in V(P)} (d_G(w) - 3)$ is the surplus carried by P and $39 = 13 \cdot 3$ is the cubic degree sum. Every R – W edge leaves some window, so summing over the p_{13} windows,

$$e(R, W) \leq \sum_{P \in \mathcal{P}} (15 + \sigma(P)) = 15p_{13} + \sigma_W \leq 15p_{13} + \sigma(G),$$

where $\sigma_W = \sum_{P \in \mathcal{P}} \sigma(P) \leq \sigma(G)$ since the windows are vertex-disjoint and surplus is nonnegative. Combining the two displays gives

$$\partial_+(R) \leq e(R, W) \leq 15p_{13} + \sigma(G).$$

By [Definition 4.1](#), $\sigma(G) = O(\sqrt{n}) = o(n)$, so

$$\partial_+(R) \leq e(R, W) \leq 15p_{13} + o(n).$$

Since $\sigma_R \geq 0$, this implies

$$\partial_+(R) - \sigma_R \leq 15p_{13} + o(n).$$

Dividing by $|R| = (1 - 13\theta)n + o(n)$ and using $p_{13} = \theta n$ gives the normalized form. $\square$

Lemma 8.5 (Exact window-join identity). *With notation as in [Definition 8.3](#),*

$$e(R, W) + 2e_{\times}(W) = 15p_{13} + \sigma_W.$$

Proof. Sum degrees over the window vertices. Since $|W| = 13p_{13}$,

$$\sum_{w \in W} d_G(w) = 3|W| + \sigma_W = 39p_{13} + \sigma_W.$$

The same degree sum counts twice the edges internal to W and once the edges from W to R . By [Definition 8.3](#), the internal edges of $G[W]$ are the $12p_{13}$ path edges inside the packed windows and the $e_{\times}(W)$ cross-window edges. Hence

$$39p_{13} + \sigma_W = 2(12p_{13} + e_{\times}(W)) + e(R, W).$$

Rearranging gives the identity. $\square$

Lemma 8.6 (Surplus-aware window-stub inequality). *No near-cubic hypothesis is required for the following estimates:*

$$\partial_+(R) \leq e(R, W) \leq 15p_{13} + \sigma_W$$

and

$$\partial_+(R) - \sigma_R \leq 15p_{13} + \sigma_W - \sigma_R.$$

Proof. The proof of [Lemma 8.4](#) gives the first inequality

$$\partial_+(R) \leq e(R, W)$$

using only $\delta(G) \geq 3$. By [Lemma 8.5](#),

$$e(R, W) = 15p_{13} + \sigma_W - 2e_{\times}(W) \leq 15p_{13} + \sigma_W,$$

because $e_{\times}(W) \geq 0$. Subtracting σ_R from the first displayed upper bound gives

$$\partial_+(R) - \sigma_R \leq 15p_{13} + \sigma_W - \sigma_R.$$

$\square$

Corollary 8.7 (Boundary-incidence supply). Assume [Definition 4.1](#). The total positive deficiency of the P_{13} -free remainder satisfies

$$\partial_+(R) \leq 15p_{13} + o(n).$$

Equivalently,

$$\frac{\partial_+(R)}{|R|} \leq \tau_{\text{win}} + o(1), \quad \tau_{\text{win}} = 0.22817486846 \dots$$

Proof. The proof of [Lemma 8.4](#) gives

$$\partial_+(R) \leq e(R, W) \leq 15p_{13} + \sigma(G).$$

By [Definition 4.1](#), $\sigma(G) = O(\sqrt{n}) = o(n)$. Hence $\partial_+(R) \leq 15p_{13} + o(n)$. Dividing by $|R| = (1 - 13\theta)n + o(n)$ and using $\theta \leq \theta_{\text{win}} + o(1)$ from [Remark 7.67](#) gives the normalized form. $\square$

Framework branch table. The packed-window remainder and deficiency supply close or route residuals as follows.

Branch	Residual condition	Demand	Payer	Bounded multiplicity	Closure
Forced remainder	The maximal P_{13} -packing leaves R .	Unpacked vertices.	Maximality of the packing.	Componentwise in R .	R is P_{13} -free.
Internal 3-core	A component of R has a subgraph of minimum degree at least 3.	No budget is available.	Theorem 5.1 .	One component at a time.	Gives a dyadic cycle.
Window incidence	Positive deficiency occurs in R .	$\partial_+(R)$.	Window-remainder stubs.	At most 15 cubic exits per packed window, plus surplus.	$\partial_+(R) \leq 15p_{13} + o(n)$ by Lemma 8.4 and Corollary 8.7 .
Surplus-adjusted supply	Remainder surplus is tracked separately.	$\partial_+(R) - \sigma_R$.	The exact window-join identity.	The stub and cross-window partition.	Lemmas 8.5 and 8.6 .

9 Curvature supply in the atom remainder

For a component C of R , define the number of internal length-two wedges by

$$W_2(C) = \sum_{v \in V(C)} \binom{d_C(v)}{2}.$$

Let

$$W_2(R) = \sum_{C \subseteq R} W_2(C).$$

Lemma 9.1 (Wedge lower bound). *For every component C ,*

$$W_2(C) \geq 3|V(C)| - 2\partial_+(C).$$

In particular, under [Definition 4.1](#), in the window-only low-deficiency branch, where $\partial_+(R) \leq (\tau_{\text{win}} + o(1))|R|$ with $\tau_{\text{win}} = 0.22817486846\dots$,

$$W_2(R) \geq \omega_{\text{win}}|R| - o(|R|), \quad \omega_{\text{win}} = 2.54365026308\dots$$

In the sharper high-entropy branch, where $\partial_+(R) \leq (0.21296321056\dots + o(1))|R|$, this improves to

$$W_2(R) \geq \omega|R| - o(|R|), \quad \omega = 2.57407357888\dots$$

Proof. For a component C , let

$$n_i = |\{v \in V(C) : d_C(v) = i\}| \quad (i \geq 0).$$

Then

$$|V(C)| = \sum_{i \geq 0} n_i,$$

$$\partial_+(C) = 3n_0 + 2n_1 + n_2,$$

and

$$W_2(C) = \sum_{i \geq 0} \binom{i}{2} n_i.$$

A direct calculation gives

$$W_2(C) - (3|V(C)| - 2\partial_+(C)) = 3n_0 + n_1 + \sum_{i \geq 3} \left(\binom{i}{2} - 3 \right) n_i \geq 0.$$

Thus

$$W_2(C) \geq 3|V(C)| - 2\partial_+(C).$$

Summing over the connected components C of R , and using $\sum_C |V(C)| = |R|$ and $\sum_C \partial_+(C) = \partial_+(R)$ (within a component $d_C = d_R$, so the per-component deficiencies add up to $\partial_+(R)$),

$$W_2(R) = \sum_{C \subseteq R} W_2(C) \geq 3|R| - 2\partial_+(R).$$

In the near-cubic window-only low-deficiency branch, [Corollary 8.7](#) gives

$$\partial_+(R) \leq (\tau_{\text{win}} + o(1))|R|,$$

with $\tau_{\text{win}} = 0.22817486846\dots$. Therefore

$$W_2(R) \geq (3 - 2 \cdot 0.22817486846\dots - o(1))|R| = (\omega_{\text{win}} - o(1))|R|,$$

where

$$\omega_{\text{win}} = 3 - 2 \cdot 0.22817486846\dots = 2.54365026308\dots$$

If the sharper high-entropy cap is available, then $\partial_+(R)/|R| \leq 0.21296321056\dots + o(1)$. Substituting this value into $W_2(R) \geq 3|R| - 2\partial_+(R)$ gives

$$W_2(R) \geq (3 - 2 \cdot 0.21296321056\dots - o(1))|R| = (\omega - o(1))|R|,$$

where the constant is

$$\omega = 3 - 2 \cdot 0.21296321056\dots = 2.57407357888\dots$$

□

Framework branch table. The curvature-supply section instantiates the framework table as follows.

Branch	Residual condition	Demand	Payer	Bounded multiplicity	Closure
Component wedge count	C is a component of R .	$W_2(C)$.	Positive deficiency.	The degree-count identity by internal degree.	$W_2(C) \geq 3 V(C) - 2\partial_+(C)$ by Lemma 9.1 .
Window-only demand	$\partial_+(R) \leq (\tau_{\text{win}} + o(1)) R $.	Linear wedge supply.	The window-stub deficiency cap.	Summation over components.	$W_2(R) \geq \omega_{\text{win}} R - o(R)$.
High-entropy demand	The sharper high-entropy deficiency cap is available.	Larger wedge supply.	The high-entropy density cap.	The same degree-count identity.	$W_2(R) \geq \omega R - o(R)$.

10 Curvature rank and rank forcing

Definition 10.1 (Curvature target-rank). *For an atom C , let $\mathcal{W}_2(C)$ be the set of raw internal length-two curvature tests in C . A subfamily $\mathcal{A} \subseteq \mathcal{W}_2(C)$ survives an admissible rank quotient q when q is label-injective on $\mathcal{A}$ in the sense of [Definition 3.95](#). It survives the admissible quotient system when it survives every functional admissible rank quotient of [Definitions 3.98](#) and [3.99](#). The curvature target-rank $r_\Omega(C)$ is the maximum size of such a surviving subfamily. Equivalently, it is the maximum size of an independently target-testable subfamily after all target-complete identifications with proper support have been imposed. For the atom remainder R , let $r_\Omega(R)$ be the target rank of the union of the raw atom curvature tests. Closed whole-graph data are evaluated in the exact profile $\rho_\mathcal{O}^{\text{ex}}(G)$, so the empty context by itself creates no curvature-rank identifications.*

Definition 10.2 (Curvature target-dependence). *Let $\mathcal{A}$ be a finite family of curvature tests. A target-dependence in $\mathcal{A}$ is a pair*

$$(a, \mathcal{B}), \quad a \in \mathcal{A}, \quad \mathcal{B} \subseteq \mathcal{A} \setminus \{a\},$$

such that a functional admissible rank quotient of the exact response profile determines the target-response value of a from the target-response vector on $\mathcal{B}$. In this definition, the phrase “adjoining a to $\mathcal{B}$ does not increase the number of target-complete states” means precisely that the projection to the $\mathcal{B}$ -coordinates is the graph of the function determining the a -coordinate under the admissible quotient. The dependence is proper when this determination is not a tautological equality of the local coordinate a with one of the coordinates in $\mathcal{B}$.

A determination certificate for $(a, \mathcal{B})$ is a tuple

$$\mathbf{c} = (X, T, q, \mathcal{P})$$

with the following data.

- (a) $X \subseteq G$ is a connected T -boundaried support carrying all coordinates in $\{a\} \cup \mathcal{B}$.

- (b) $q : \rho_T^{\text{ex}}(X) \rightarrow Q$ is a functional admissible rank quotient.
- (c) For any two realizations of $\rho_T^{\text{ex}}(X)$ with the same q -image on the coordinates of $\mathcal{B}$, the q -image of a is the same. Equivalently, the q -value of a is a function of the q -value vector on $\mathcal{B}$.
- (d) $\mathcal{P}$ is the finite set of declared paths, boundary incidences, window segments, or connector incidences recorded by the certificate as the support data for the context-universal determination.

Here a realization means a T -boundaried support whose exact response profile maps to the same quotient data Q under q . The support of $\mathbf{c}$ is the lexicographically first smallest connected subgraph of G containing the declared supports of a , of all coordinates in $\mathcal{B}$, and of every member of $\mathcal{P}$. The support of $(a, \mathcal{B})$ is the support of a certificate with minimal support. A target-dependence is inclusion-minimal if no proper connected subsupport carries a proper target-dependence with the same determined coordinate.

Lemma 10.3 (Target-rank circuit extraction). *Let $\mathcal{A}$ be a finite family of declared coordinates evaluated under the functional admissible quotients of a fixed exact response profile. If a subfamily $\mathcal{I} \subseteq \mathcal{A}$ is maximal among subfamilies that survive every functional admissible rank quotient and $a \in \mathcal{A} \setminus \mathcal{I}$, then there is a subfamily $\mathcal{B} \subseteq \mathcal{I}$ such that $(a, \mathcal{B})$ is a target-dependence in the sense of Definition 10.2. In particular, if no proper target-dependence exists among the coordinates of $\mathcal{A}$, then the whole family $\mathcal{A}$ survives every functional admissible rank quotient and is independently target-testable in the rank sense of Definition 10.1.*

Proof. Since $\mathcal{I}$ is maximal for survival, $\mathcal{I} \cup \{a\}$ fails to survive some functional admissible rank quotient

$$q : \rho_T^{\text{ex}}(X) \rightarrow Q.$$

The subfamily $\mathcal{I}$ survives q . Hence the failure of label-injectivity for $\mathcal{I} \cup \{a\}$ involves a . By Definition 3.99, the quotient q is functional on $\mathcal{A}$. Therefore there is a finite subfamily $\mathcal{B} \subseteq \mathcal{I}$ whose q -image vector determines the q -image of a . Choose such a subfamily inclusion-minimal. If a is sent to a single quotient value after the quotient forgets its separate label, then $\mathcal{B} = \emptyset$. If a is identified with a coordinate of $\mathcal{I}$, then the functional certificate may be taken with $\mathcal{B} = \{b\}$ for such a coordinate b . The quotient q , together with the declared supports of a , of $\mathcal{B}$, and of the finitely many coordinates used by q , is a determination certificate.

If the dependence is tautological, then a is the same local coordinate as an element of $\mathcal{B}$. Exact coordinate labels keep distinct declared raw coordinates distinct, so a rank loss among distinct raw curvature coordinates yields a proper dependence. The contrapositive gives the final assertion. $\square$

If $r_\Omega(C) < |\mathcal{W}_2(C)|$, then some coordinate outside a maximal independently target-testable subfamily is target-dependent on that subfamily in the sense of Definition 10.2 and Lemma 10.3. This is the only meaning of “dependence” used below.

Lemma 10.4 (Localization of curvature dependence). *Let C be a proper atom and let $\mathcal{W}_2(C)$ be its raw curvature tests. If*

$$r_\Omega(C) < |\mathcal{W}_2(C)|,$$

then some target-dependence among raw curvature tests falls into one of the following cases:

- (i) a target-defective quotient;
- (ii) an admissible nontrivial target-complete compression of a proper atom;
- (iii) a target-dependence that becomes valid only after adjoining a connected support $Z \supsetneq C$.

Proof. A strict rank drop means that some raw curvature test is target-dependent on the others in the sense of [Definition 10.2](#) and [Lemma 10.3](#). Choose a determination certificate with inclusion-minimal connected support. The certificate has an admissible quotient q and a finite support set $\mathcal{P}$. Either the determination certified by q fails for some T -boundaried outside context, or it is context-universal. In the context-universal case, the connected support of the certificate either equals C or strictly contains C . These alternatives are mutually exclusive and exhaust all certificates:

- (i) if the determination fails for some outside context, it is not preserved under context-universality and hence is a target-defective quotient;
- (ii) if it holds for every outside context already with support C , then q is a target-complete rank-reducing quotient of the proper atom C . Since q is admissible, [Definition 3.98](#) supplies a strictly smaller proper representative. Thus this case is a nontrivial target-complete compression of C ;
- (iii) otherwise it holds for every outside context but only once additional structure outside C is adjoined; taking an inclusion-minimal connected such support yields a connected $Z \supsetneq C$.

□

Lemma 10.5 (Separated wedges are tested by disjoint contexts). *Let $u, v \in V(R)$ have identical rooted radius- r neighbourhood types, $B_r(u) \cong B_r(v)$, and suppose the closed balls $\overline{B}_r(u)$ and $\overline{B}_r(v)$ are vertex-disjoint. Let $w(u), w(v)$ be the corresponding internal length-two wedge tests at the two roots. Then any target tester separating $w(u)$ from $w(v)$ has its support in the complement of $\overline{B}_r(u) \cup \overline{B}_r(v)$. Consequently, any quotient identifying $w(u)$ with $w(v)$ is either context-universal or target-defective.*

Proof. A target tester for a wedge coordinate is a boundaried context glued to the support of that wedge. The wedge itself is an internal length-two path at the root of its radius- r ball. Any closure witnessing a dyadic target that is not already internal to the ball must leave the ball through its boundary; if the witness stayed wholly inside the ball, the ball would already contain a target cycle, contrary to internal target-safety of the remainder. Hence every tester that separates the two wedge coordinates is supported outside the corresponding closed balls.

Since the rooted balls are isomorphic and vertex-disjoint, the two local wedge states have identical internal data and identical boundary degree profiles. If an outside context distinguishes them, then the identification fails the context-universality condition of target-completeness and is target-defective. If no outside context distinguishes them, the identification is context-universal. These are the only two possibilities. □

Lemma 10.6 (Proper delocalization supports are forbidden). *Let $C \subsetneq G$ be a proper atom. Suppose a target-dependence among curvature tests of C becomes target-complete only after adjoining an inclusion-minimal connected support $Z \supseteq C$. If*

$$Z \subsetneq G,$$

then the dependence is impossible.

Proof. Regard Z as a boundaried graph with boundary

$$\partial Z = \{v \in V(Z) : v \text{ is incident with an edge of } G - Z\}.$$

The response state of Z includes its boundary degree profile and all target-response data inherited from C . Since $Z \subsetneq G$, it is a proper boundaried support. If the dependence fails against some outside ∂Z -context, it is target-defective. If it succeeds against every outside context, it is a nontrivial target-complete compression of the proper support Z , forbidden by [Corollary 3.105](#). $\square$

Definition 10.7 (Repair-network terminology). *Let Z be the support of a determination certificate. A vertex of Z is a passive degree-two subdivision vertex if it has degree 2 in Z , is not a boundary vertex of Z , and is not contained in the declared support of any coordinate, path endpoint, packed-window incidence, trace incidence, fan decoration, or carrier used by the certificate. Such a vertex records only an unlabelled subdivision of a response path.*

After deleting the original coordinate supports of the dependence from Z , each remaining connected component that meets the deleted part only in boundary leaves is a delayed compensation component. If such a component has p boundary leaves and all non-leaf vertices have degree 3 except for the recorded surplus contribution σ_Z , then it is a 1–3 repair network up to surplus. Its internal vertex count is denoted s , and its cycle rank is denoted β_Z .

Lemma 10.8 (Enlarged delocalization supports are repair networks). *Let Z be a connected support on which a formerly nonlocal dependence becomes internal. Suppose Z is target-safe, contains no proper subgraph of minimum degree at least 3, and has no passive degree-two subdivision vertices. Then each delayed compensation component of Z is a 1–3 repair network up to surplus. More precisely, if such a connected component has p boundary leaves, s internal vertices, cycle rank β_Z , and surplus σ_Z , then*

$$s = p - 2 + 2\beta_Z - \sigma_Z.$$

Proof. The component has s internal vertices and p boundary leaves, hence $s + p$ vertices in all. Let e be the number of edges joining two internal vertices; each boundary leaf is attached by a single edge, so the component has $e + p$ edges in total. Each internal vertex has degree 3 up to its surplus, so the internal vertices contribute $3s + \sigma_Z$ to the degree sum, and each of the p leaves has degree 1. The handshake identity $\sum_v d(v) = 2 \cdot (\text{number of edges})$ therefore reads

$$(3s + \sigma_Z) + p = 2(e + p),$$

which simplifies to

$$3s + \sigma_Z = 2e + p. \tag{\dagger}$$

The component is connected, so its cycle rank is

$$\beta_Z = (e + p) - (s + p) + 1 = e - s + 1, \quad \text{equivalently} \quad e = s - 1 + \beta_Z.$$

Substituting $e = s - 1 + \beta_Z$ into $(\dagger)$ gives $3s + \sigma_Z = 2(s - 1 + \beta_Z) + p = 2s - 2 + 2\beta_Z + p$, and solving for s yields

$$s = p - 2 + 2\beta_Z - \sigma_Z.$$

$\square$

Lemma 10.9 (No silent whole-support delocalization). *Let a curvature target-dependence have inclusion-minimal connected support equal to the whole graph G , and let $q : \rho_{\mathcal{O}}^{\text{ex}}(G) \rightarrow Q$ be the functional admissible rank quotient in its determination certificate. Then q is injective on the declared raw curvature coordinates. In particular, no proper curvature rank-drop dependence has whole support.*

Proof. The quotient q is not a raw proposal: admissibility is part of the determination certificate in [Definitions 3.98](#) and [10.2](#). Hence its boundary-profile and target-response obligations have already been checked on the fixed finite exact response system. Since $T = \emptyset$, there is no additional outside-support family to enumerate at this node.

Suppose that q were not injective on the declared raw coordinates. The closed clause of [Definition 3.98](#) then supplies a strictly smaller admissible closed representative G_q . It is finite and simple, has minimum degree at least 3, and satisfies $H_\emptyset(G_q) \subseteq H_\emptyset(G)$. If G_q contained a dyadic cycle, its empty-context target response would also occur in G , contrary to the choice of G . Thus G_q is a strictly smaller counterexample, contradicting lexicographic minimality. Therefore q is injective. A proper rank drop would identify distinct declared raw coordinates by [Lemma 10.3](#), so no such whole-support dependence remains. $\square$

Lemma 10.10 (Forcing of full admissible quotient-rank). *In a minimal counterexample, the manuscript's admissible quotient-rank satisfies*

$$r_\Omega(R) \geq W_2(R) - o(W_2(R)).$$

Proof. If not, then

$$r_\Omega(R) < W_2(R) - o(W_2(R)).$$

By [Lemma 10.3](#), a rank loss among the raw curvature coordinates gives a proper target-dependence. Choose such a dependence with a determination certificate of inclusion-minimal connected support, and call that support Z . The minimality of Z gives two possibilities.

First suppose $Z \subsetneq G$. Regard Z as a proper boundaried atom with boundary ∂Z . If the dependence is already visible inside a smaller proper atom $C \subseteq Z$, then [Lemma 10.4](#) applies to C . Its three alternatives are respectively: a target-defective quotient, a target-complete compression of a proper atom, or a dependence that requires an enlarged connected support. The first is excluded by the definition of target-completeness; the second is excluded by [Corollary 3.105](#); and the third, because the enlarged support is still proper in G , is excluded by [Lemma 10.6](#). If no smaller $C \subsetneq Z$ carries the dependence, then Z itself is the first proper support on which it becomes target-complete, and [Lemma 10.6](#) again excludes it.

It remains that the inclusion-minimal support is $Z = G$. Then the dependence becomes context-universal only globally. By [Lemma 10.9](#), its already-admitted quotient is injective on the declared raw curvature coordinates. It therefore cannot identify two distinct local curvature tests, so no proper admissible target-dependence survives. By [Lemma 10.3](#), the raw curvature tests have no admissible rank loss. Hence

$$r_\Omega(R) = W_2(R),$$

which implies the stated weaker bound $r_\Omega(R) \geq W_2(R) - o(W_2(R))$. $\square$

Corollary 10.11 (Conditional finite curvature magnitude and product cost). *Assume the node-[24] packing cap and retain the node-[47] full-rank ledger on the same remainder. Writing $s(G)$ for the total degree surplus, one has the exact finite inequality*

$$250825743018 |R| \leq 98608581006 r_\Omega(R) + 197217162012 s(G).$$

On the high-entropy density branch the left coefficient may be replaced by 253825743018.

If, in addition, a finite ordered family of pairwise distinct graph-completion states, an injective code of those states into the labelled baseline-skeleton family, and a conditional-fibre ledger for every curvature coordinate are supplied on this same residual, then that ledger gives

$$543958^{r_\Omega(R)} \leq 111286^{r_\Omega(R)} N_{\text{skel}},$$




Figure 18: Visual proof of the curvature supply–budget squeeze. Panel I shows the P_{13} -free remainder R producing raw internal length-two wedge tests; the wedge lower bound gives $W_2(R) \geq \omega_{\text{win}}|R| - o(|R|)$, with $\omega_{\text{win}} = 2.543\dots$ on the window-only branch. Panel II shows the rank-forcing step: if the raw tests did not survive as target-rank, a dependence would route to target defect, proper compression, or delocalization, all of which are closed in the minimal counterexample. Thus $r_\Omega(R)$ is linear in $W_2(R)$. Panel III records the resulting collision with the independent-test budget: the full-rank branch demands at least $2.543|R| - o(|R|)$ tests, while the supply side has capacity at most $0.611|R| + o(|R|)$.

where N_{skel} is the exact number of labelled baseline skeletons. Constructing this simultaneous graph-owned realization from node [47] is a separate unproved obligation of node [48]; no complementary outcome or edge to node [55] is introduced.

Proof. The first inequality is the integer form of the wedge lower bound after the node-[24] density cap is substituted; its error is twice the retained total surplus, rather than an unproved asymptotic term. The sharper coefficient is obtained by the same calculation from the high-entropy cap. Under the stated additional realization hypothesis, each local conditional-fibre inequality

$$543958 |S_{i+1}| \leq 111286 |S_i|$$

is multiplied along the declared coordinate order. The final fibre is nonempty and the initial fibre injects into the labelled skeleton family, which gives the displayed product bound. This conditional calculation does not construct the realization. $\square$

Remark 10.12 (Provenance of the curvature product). *The ingredients of Corollary 10.11 play distinct roles. The independence of the remainder curvature tests—that they realize full target-rank (Lemma 10.10)—is forced by the routing of Lemma 10.4: any dependence among the tests would be a target-defective quotient, a target-complete compression of a proper atom (excluded by Corollary 3.105), or a delocalization. Proper delocalization is excluded by Lemma 10.6; whole-graph delocalization is excluded by the closed exact-profile argument in Lemma 10.9. Thus each rank-reduction route is unavailable in the main spine. Full target rank alone, however, does not imply simultaneous product realization of the finite responses. The graph-owned response semantics together with the conditional-fibre ledger is the additional ingredient that supplies the multiplicative factor, and its absence is the explicit open successor of node [48]. As Remark 12.4 records, the large-budget continuation does not spend this factor.*

Framework branch table. The curvature-rank forcing section closes the rank-drop alternatives by the following table.

Branch	Residual condition	Demand	Payer	Bounded multiplicity	Closure
Rank drop	$r_\Omega(R) < W_2(R) - o(W_2(R))$.	A dependent curvature coordinate.	The target-dependence route.	A finite functional circuit.	Lemmas 10.3 and 10.4 .
Proper support	The dependence localizes in $Z \subsetneq G$.	The support profile.	Replacement or target defect.	The proper boundary profile.	Lemma 10.6 .
Whole graph	The dependence delocalizes to all of G .	The closed exact profile.	Closed representative minimality.	The empty-boundary exactness rule.	Lemma 10.9 .
Full rank	All rank-drop alternatives are excluded.	Raw wedge tests.	Curvature target rank.	All but $o(W_2)$ tests.	Lemma 10.10 and Corollary 10.11 .

11 Per-vertex remainder entropy and two-budget routing

Throughout this section, R is the P_{13} -free remainder of the minimal counterexample, with $|R| = (1 - 13\theta)n - o(n)$. The aim is to make the entropy branch explicit on the finite state family actually supplied by the realized node-[48] ledger.

Definition 11.1 (Finite realized-state entropy of the remainder). *On the realized successor of node [48], let $\mathcal{S}_R$ be the exact finite ordered family of graph-completion states carried by its conditional-fibre certificate. Put*

$$N_R = |\mathcal{S}_R|, \quad b_R = \lfloor \log_2 N_R \rfloor, \quad \eta(R) = \frac{\log_2 N_R}{|R|}.$$

The certificate proves $N_R > 0$ and injects $\mathcal{S}_R$ into the literal family of labelled baseline edge sets. Thus this definition reuses the supplied local state schedule and does not enumerate labelled graphs, subsets, contexts, or Boolean assignments. It is not identified with the class of all graphs on $V(R)$ satisfying the remainder constraints. Node [50] uses the denominator-free natural-number comparison

$$n^{|R|} \leq N_R^{10} \quad \text{or} \quad N_R^{10} < n^{|R|}.$$

This is the exponentiated form of the entropy threshold and is a single local arithmetic comparison on the node-[49] count; it does not evaluate real logarithms or generate any state family.

Lemma 11.2 (High-power remainder bit contribution). *On the node-[50] high branch,*

$$n^{|R|} \leq N_R^{10},$$

the realized node-[49] state family satisfies

$$\frac{|R|}{10} \log_2 n \leq \log_2 N_R.$$

Proof. Node [49] gives $N_R > 0$, and the minimal counterexample has $n > 0$ vertices. Apply the increasing function $\log_2$ to the node-[50] inequality and use the power law to obtain

$$|R| \log_2 n \leq 10 \log_2 N_R.$$

Division by 10 gives the claim. This argument consumes only the two natural powers already compared at node [50]; it performs no finite-family enumeration. $\square$

Definition 11.3 (Joint-accounting frontier). *At node [52], a joint accounting realization consists of a finite multiscale window-state schedule built from the node-[24] commuting local completions, the exact node-[48] remainder-state schedule, and one map sending each window-remainder pair to an actual node-[24] global completion and then to a labelled baseline edge set. Both coordinates must be recoverable from the global completion, and the resulting baseline-skeleton code must be injective.*

For any such supplied realization, the verified finite-capacity theorem gives the corresponding product bound without materializing the Cartesian product. The construction of the multiscale window schedule, joint commutation, injective baseline encoding, and the authored demand bound remains the open obligation at node [52]. No absence constructor or conditional edge to node [55] is part of the original diagram. The conditional capacity theorem does not by itself assert the sharper density cap or the contradiction at node [54]; those require the separate joint-capacity arithmetic displayed between nodes [52] and [54]. Its computation enumerates no state product, graph family, subset family, context universe, or Boolean cube.

Lemma 11.4 (Exact low-entropy forced-cost fit). *On the strict low constructor of node [50], let*

$$r = |\Omega|, \quad A = 543958^r, \quad B = 111286^r, \quad U = n^{|R|}.$$

Then

$$A^{10} < B^{10}U.$$

Equivalently, the reverse small-budget comparison $B^{10}U \leq A^{10}$ at node [53] is impossible.

Proof. The exact conditional-fibre ledger at node [48], together with its nonempty terminal fibre, gives

$$A \leq BN_R.$$

This uses the actual node-[49] state count, not the larger family of all remainder graphs. Raising to the tenth power and using the strict node-[50] low inequality $N_R^{10} < U$ gives

$$A^{10} \leq (BN_R)^{10} = B^{10}N_R^{10} < B^{10}U,$$

because $B > 0$. The proof transports these supplied natural-number inequalities symbolically; it evaluates no power and scans no finite family. The surviving large-budget output reaches node [55] only after the separate same-context strict-quarter budget required there has been proved. $\square$

Lemma 11.5 (Dominant local type in the low-entropy branch). *Fix $r \geq 1$. Suppose the radius- r local-type coordinate of the remainder is structurally repetitive, in the sense that*

$$\log_2 |\{ (B_r(v))_{v \in V(R)} \}| = o(|R|),$$

i.e. the number of realizable radius- r type vectors is subexponential in $|R|$. Then for every $\varepsilon > 0$ a single rooted radius- r type covers all but at most $\varepsilon|R|$ vertices of R , for all large n . To derive this hypothesis from the low-state-count regime of [Definition 11.1](#), a later consumer must still interpret the states as remainder graphs, construct their canonical radius- r type vectors, and prove the required relabelling closure. None of those semantic properties follows from the cardinality bookkeeping at node [49].

Proof. There are only finitely many rooted subcubic radius- r types; call the number $t = t(r)$, a constant independent of n . Suppose, for contradiction, that no single type covered a $(1 - \varepsilon)$ -fraction of $V(R)$. Then the empirical type distribution $\pi = (\pi_1, \dots, \pi_t)$ over the t types satisfies $\max_a \pi_a \leq 1 - \varepsilon$, so its Shannon entropy obeys $H(\pi) \geq h(\varepsilon) > 0$, where $h(\varepsilon) = -\varepsilon \log_2 \varepsilon - (1 - \varepsilon) \log_2 (1 - \varepsilon)$ is the binary entropy function. The labelled remainder class is closed under relabeling of the fixed vertex set $V(R)$, and the radius- r type vector is a canonical labelled coordinate. Therefore every permutation of one realized type vector with empirical distribution π is again realized by a relabelled graph in the same class. Distinct assignments of types to the labelled vertices give distinct type vectors, so the number of type vectors with empirical distribution π is at least the multinomial coefficient, which by Stirling's approximation satisfies

$$\binom{|R|}{\pi_1 |R|, \dots, \pi_t |R|} = 2^{H(\pi)|R| - o(|R|)} \geq 2^{h(\varepsilon)|R| - o(|R|)}.$$

Thus the number of realizable type vectors is $2^{\Theta(|R|)}$, exponential in $|R|$, contradicting the hypothesis that it is subexponential. Hence some rooted radius- r type covers a $(1 - \varepsilon)$ -fraction of $V(R)$. $\square$

Remark 11.6. *The structural-repetitiveness hypothesis of Lemma 11.5 is stronger than the single-scalar low-entropy regime $N_R^{10} < n^{|R|}$. The radius- r type coordinate is one canonical coordinate of the remainder skeleton, so its bit-content is bounded by $\eta(R)|R|$. When that coordinate is $o(|R|)$ the lemma applies. When it is not $o(|R|)$ but the total remainder entropy is still below $\frac{1}{10}|R| \log_2 n$, no dominance conclusion is asserted here; that nonrepetitive low-entropy case is passed unchanged to the large-budget net-charge analysis, which is independent of the precise curvature-rank fraction by Remark 12.4.*

Lemma 11.7 (Dominant type forces independent curvature translates). *Suppose a rooted radius- r type with $r \geq 2$ occurs on all but $o(|R|)$ vertices of R and contains an internal length-two wedge at its root. Then*

$$r_\Omega(R) \geq c_r |R| - o(|R|)$$

for a constant $c_r > 0$ depending only on r .

Proof. Let $D \subseteq V(R)$ be the set of vertices carrying the dominant rooted radius- r type, so $|D| \geq |R| - o(|R|)$. Subcubicity bounds each closed radius- r ball by

$$|\overline{B}_r(v)| \leq 1 + 3(2^r - 1),$$

a constant depending only on r . Greedily choose a maximal subset $S \subseteq D$ whose members are pairwise at distance greater than $2r$; then the closed balls $\{\overline{B}_r(u) : u \in S\}$ are pairwise vertex-disjoint. By maximality every vertex of D lies within distance $2r$ of some $u \in S$, and each ball of radius $2r$ has at most

$$b' = 1 + 3(2^{2r} - 1)$$

vertices. Hence

$$|S| \geq |D|/b' \geq |R|/b' - o(|R|).$$

Each ball carries one internal length-two wedge at its root, so these give at least $|S|$ raw curvature tests. Hence $W_2(R) \geq |S|$. Also subcubicity gives $W_2(R) \leq 3|R|$, because a vertex of degree at most 3 is the center of at most three unordered length-two wedges. By Lemma 10.10,

$$r_\Omega(R) \geq W_2(R) - o(W_2(R)).$$

Since $W_2(R) \geq |S| \geq |R|/b' - o(|R|)$, the error term is $o(|R|)$, and

$$r_\Omega(R) \geq |S| \geq |R|/b' - o(|R|).$$

Take $c_r = 1/b'$. □

Remark 11.8 (A positive fraction of independent tests suffices). *The separated-translate construction already yields independent wedge tests on a positive fraction of the remainder, $|S| \geq |R|/b' - o(|R|)$ with b' a constant. For the only downstream use—the forced cost of [Corollary 10.11](#) feeding the threshold $\Theta(n)$ —any positive linear rank $r_\Omega(R) \geq c|R|$ supplies a positive forced cost $cc_\Omega|R|$, and by [Remark 12.4](#) the closure outside the explicit residuals holds for every nonnegative value of that cost. The conclusion is therefore insensitive to the precise rank fraction: full rank is the sharpest form, but a positive fraction is all that the remainder of the argument consumes.*

Proposition 11.9 (Two-budget routing). *Assume [Definition 4.1](#). Let R be the P_{13} -free remainder of a minimal counterexample and let $\eta(R)$ be its per-vertex skeleton entropy ([Definition 11.1](#)). Exactly one of the following holds.*

- (a) **High-entropy branch:** $n^{|R|} \leq N_R^{10}$. The remainder then contributes at least

$$\frac{|R|}{10} \log_2 n$$

bits. Node [52] then returns either a same-context joint accounting realization or the open joint-accounting residual of [Definition 11.3](#). Only a separately proved joint-capacity overflow may enter node [54]; every nonclosing leaf is routed to the large-budget analysis.

- (b) **Structurally repetitive low-entropy branch:** $N_R^{10} < n^{|R|}$ and the radius-2 local-type coordinate has $o(|R|)$ bit-content. Then [Lemma 11.5](#) yields a dominant radius-2 type. If that type contains an internal length-two wedge, then [Lemma 11.7](#) supplies a positive linear curvature-rank lower bound.
- (c) **Nonrepetitive low-entropy branch:** $N_R^{10} < n^{|R|}$ and the radius-2 local-type coordinate has nonvanishing linear bit-content. This proposition makes no curvature-rank claim in this case; the branch is passed to the large-budget net-charge analysis.

In every case the surviving residual is subsequently passed to the large-budget net-charge analysis of [Proposition 13.8](#), with any additional curvature cost used only when it has been established by the preceding rank-forcing lemmas.

Proof. The exact natural-number inequalities $n^{|R|} \leq N_R^{10}$ and $N_R^{10} < n^{|R|}$ partition all remainders. In the low-entropy case, the radius-2 local-type coordinate either has $o(|R|)$ bit-content or it does not. This gives the three listed cases, which are exhaustive and mutually exclusive.

In branch (a), [Lemma 11.2](#) shows that the realized state family contributes at least

$$\frac{|R|}{10} \log_2 n = \frac{n - 13p_{13}}{10} \log_2 n - o(n \log n)$$

bits. The node-[52] dichotomy of [Definition 11.3](#) records whether the missing same-context window-remainder injection has actually been supplied. This proposition does not infer the feasibility inequality from the two separate ledgers.

In branch (b), the local-type coordinate is subexponential, so [Lemma 11.5](#) produces a dominant radius-2 type. In the large-budget range the degree-3 bulk has positive density by [Lemmas 8.4](#)




Figure 19: Visual proof of the two-budget routing funnel. Panel I splits the P_{13} -free remainder by the exact denominator-free comparison $n^{|R|} \leq N_R^{10}$. Panel II shows the high-entropy branch: taking logarithms spends at least $(|R|/10) \log_2 n$ bits, after which node [52] asks for the missing same-context joint window-remainder accounting. Only a proved joint-capacity overflow closes; the open or fitting outcomes route forward. Panel III shows the structurally repetitive low-entropy branch: subexponential radius-2 type data gives a dominant local type, and an internal length-two wedge in that type supplies positive linear curvature rank. Panel IV records the routing convention of [Proposition 11.9](#): nonrepetitive low-entropy data is passed to the large-budget net-charge branch, and curvature cost is carried forward only on branches where the preceding rank-forcing lemmas have established it.

and [9.1](#); if the dominant type contains an internal length-two wedge at its root, then [Lemma 11.7](#) gives a positive linear curvature-rank lower bound. Branch (c) asserts no such dominance and is simply routed forward. $\square$

Remark 11.10 (Role of the forced curvature cost). *When [Lemma 10.10](#) is invoked, the window-only forced curvature cost is $K_{\text{win}}|R| - o(|R|)$. In the sharper high-entropy low-deficiency branch this improves to $K|R| - o(|R|)$, and that sharper value is the numerical constraint used in the entropy cap of [Proposition 12.2](#). The weaker separated-translate input of [Lemma 11.7](#) supplies only a positive linear rank in the structurally repetitive low-entropy branch. The final large-budget closure outside the explicit residuals does not require the exact curvature constant: [Remark 12.4](#) records that the window-density and net-charge accounting already give the needed normalized deficiency gap.*

Framework branch table. The two-budget routing funnel is summarized by the following framework table.

Branch	Residual condition	Demand	Payer	Bounded multiplicity	Closure
High entropy	$n^{ R } \leq N_R^{10}$.	Remainder bits.	Joint skeleton capacity, if realized.	One exact window-remainder pair family.	Lemma 11.2 and Definition 11.3 .

Repetitive low entropy	A radius-2 coordinate has $o(R)$ bits.	A dominant local type.	Separated curvature translates.	Bounded-radius overlap.	Lemmas 11.5 and 11.7.
Nonrepetitive low entropy	The low-entropy data do not give dominance.	No local entropy closure here.	The large-budget route.	Not applicable.	Passed to Section 13.
Cost separation	Curvature cost is available only after rank forcing.	Independent tests.	The entropy budget.	The rank-forced branch.	Remark 11.10.

12 The admissible entropy cap

In this section assume [Definition 4.1](#). The entropy cap is admissible only when the remaining non-curvature budget is smaller than the forced curvature cost. We make this precise. In the high-entropy accounting of [Section 7](#), the window and remainder contributions consume, per factor $n \log_2 n$, the normalized amount

$$0.1 + 116.808581006 \theta$$

of the near-cubic labelled skeleton budget

$$1.5 n \log_2 n + o(n \log n)$$

of [Lemma 4.4](#). Hence the *remaining non-curvature budget* is

$$(1.5 - (0.1 + 116.808581006 \theta)) n \log_2 n = (1.4 - 116.808581006 \theta) n \log_2 n.$$

On top of this, [Corollary 10.11](#) imposes the linear forced curvature cost

$$K|R| - o(|R|) = K(1 - 13\theta)n - o(n), \quad K = 5.89262883286 \dots$$

For a residual graph to exist, the window, remainder, and forced-curvature demands together must fit inside the skeleton budget; equivalently the forced curvature cost must fit inside the remaining non-curvature budget,

$$K(1 - 13\theta)n \leq (1.4 - 116.808581006 \theta) n \log_2 n + o(n \log n).$$

Dividing by n , no residual graph exists once

$$(1.4 - 116.808581006 \theta) \log_2 n < K(1 - 13\theta). \tag{2}$$

Definition 12.1 (The finite-size packing threshold). *Define*

$$\Theta(n) = \frac{1.4 - \frac{K}{\log_2 n}}{116.808581006 - \frac{13K}{\log_2 n}},$$

where

$$K = 5.89262883286 \dots$$

Proposition 12.2 (Entropy cap closes the high- θ branch). *Assume Definition 4.1. In the high-entropy branch of Proposition 11.9, if*

$$\theta > \Theta(n) + o(1),$$

then no minimal residual graph exists.

Proof. Suppose $\theta > \Theta(n) + o(1)$. By the definition of $\Theta(n)$ this is precisely inequality (2), that is, the remaining non-curvature budget is strictly smaller than the forced full-rank curvature cost of Corollary 10.11. Then the window package of Remark 7.1, the remainder bits, and the forced-curvature bits together strictly exceed the near-cubic skeleton budget $1.5 n \log_2 n + o(n \log n)$. These bits form one independently target-testable coordinate family, so the number of realized target-complete states would exceed the number of labelled skeletons, contradicting Lemmas 3.107 and 4.2. Hence no residual graph exists in this branch. $\square$

Remark 12.3 (Order of the entropy argument). *The entropy cap is not used to assert an a priori linear upper bound on $r_\Omega(R)$. Instead, the rank-forcing argument first forces full curvature rank. Entropy is then applied only in the region where the remaining budget is smaller than the forced curvature cost.*

Remark 12.4 (Robustness of the closure to the forced curvature cost). *The forced curvature cost sharpens the high-entropy threshold $\Theta(n)$ of Definition 12.1, but the large-budget closure outside the explicit residuals does not need that sharpening. The near-cubic window-only part of Remark 7.67 gives*

$$\theta \leq \theta_{\text{win}} + o(1), \quad \theta_{\text{win}} = 0.01270017798 \dots,$$

and therefore

$$\frac{15\theta}{1 - 13\theta} \leq \tau_{\text{win}} + o(1), \quad \tau_{\text{win}} = 0.22817486846 \dots < \frac{1}{4}.$$

Thus the curvature-rank and forced-cost machinery is not required for the net-charge closure outside the explicit residuals, which already follows from the window-only density bound together with the local analysis of Sections 13 and 14. Consequently the final closure uses neither the exact value of c_Ω (Remark 10.12) nor the precise rank fraction (Remark 11.8) as an additional hypothesis.

n	$\log_2 n$	$\Theta(n)$	$\tau_\Theta = 15\Theta(n)/(1 - 13\Theta(n))$
10^3	9.9658	0.007411	0.123019
10^4	13.2877	0.008614	0.145505
10^6	19.9316	0.009776	0.167991
10^9	29.8974	0.010529	0.182982
10^{12}	39.8631	0.010899	0.190477
∞	∞	0.011985	0.212963

Framework branch table. The admissible entropy cap closes only the branch where its numerical inequality is active.

Branch	Residual condition	Demand	Payer	Bounded multiplicity	Closure
--------	--------------------	--------	-------	----------------------	---------

Admissible cap	The remaining non-curvature budget is smaller than the forced curvature cost.	Window, remainder, and curvature bits.	The near-cubic skeleton budget.	One target-complete coordinate family.	Closed by Proposition 12.2 .
Below threshold	The entropy inequality does not close.	The large-budget residual.	Net-charge accounting.	Not applicable.	Routed to Section 13 .
Robust route	The exact curvature magnitude is not used.	Window-only deficiency gap.	Net-charge and local discharging.	The same residual ledger.	Remark 12.4 .

13 The large-budget residual

Throughout this section assume [Definition 4.1](#). Regardless of whether the high-entropy cap of [Proposition 12.2](#) applies, the near-cubic window-only bound in [Remark 7.67](#) gives

$$\theta \leq \theta_{\text{win}} + o(1), \quad |R| = n - 13p = (1 - 13\theta)n \geq (1 - 13\theta_{\text{win}})n - o(n).$$

By maximality of the induced- P_{13} packing, every component of R is P_{13} -free. By [Theorem 5.1](#), no component of R has an internal subgraph of minimum degree at least 3.

Definition 13.1 (Internal 3-core). *For a graph or support X , the internal 3-core of X is the maximal induced subgraph of X with minimum degree at least 3, obtained equivalently by repeatedly deleting vertices of internal degree at most 2. The internal 3-core is empty if no nonempty subgraph of X has minimum internal degree at least 3.*

The supply inequality [Lemma 8.4](#) gives

$$\partial_+(R) - \sigma_R \leq 15p_{13} + o(n), \quad \sigma_R = \sigma(R).$$

Dividing by $|R| = (1 - 13\theta)n + o(n)$ yields

$$\frac{\partial_+(R) - \sigma_R}{|R|} \leq \frac{15\theta}{1 - 13\theta} + o(1).$$

Since $\theta \leq \theta_{\text{win}} + o(1)$, every large-budget residual satisfies

$$\boxed{\frac{\partial_+(R) - \sigma_R}{|R|} \leq \tau_{\text{win}} + o(1),}$$

where

$$\tau_{\text{win}} = \frac{15\theta_{\text{win}}}{1 - 13\theta_{\text{win}}} = 0.22817486846 \dots < \frac{1}{4}.$$

Definition 13.2 (Canonical support decomposition). *The canonical support decomposition of the residual consists of:*

- (i) *the connected components $Y_1, \dots, Y_k$ of R , ordered by the deterministic tie-breaking rule fixed in [Lemma 4.2](#); and*

(ii) an assignment of every ambient surplus unit $d_G(h) - 3$ with $h \in V_{\geq 4}(G) \cap V(R)$ to exactly one piece.

All surplus units of $h \in V_{\geq 4}(G) \cap V(R)$ are assigned to the unique piece containing h . Surplus carried by vertices in the packed windows is not assigned to local supports; it has already been absorbed into the $o(n)$ term in [Lemma 8.4](#). For the support associated to Y_i , write

$$|V(X_i)| := |Y_i|, \quad \sigma(X_i) := \text{the number of surplus units assigned to } Y_i.$$

Thus

$$\sum_i |V(X_i)| = |R|, \quad \sum_i \sigma(X_i) = \sigma(R).$$

When an assigned surplus unit is represented by an actual high-degree fan vertex in the local Type B analysis, the corresponding fan envelope records that vertex as decoration data; it is not counted a second time in $|V(X_i)|$. Because the pieces are the connected components of R , there are no edges of R between distinct pieces.

Definition 13.3 (Admissible support). An admissible support X is a connected remainder piece Y_i of the canonical support decomposition, together with its assigned surplus units and any local decoration data needed to realize those units in the Type B fan analysis. Its vertex count and positive deficiency are those of the remainder piece:

$$|V(X)| = |Y_i|, \quad \partial_+(X) = \sum_{v \in Y_i} \max\{0, 3 - d_{Y_i}(v)\}.$$

Its surplus $\sigma(X)$ is the assigned surplus credit from [Definition 13.2](#). As a boundaried support of the minimal counterexample G , X carries the inherited constraints used in the local analysis: it is P_{13} -free, has empty internal 3-core, is dyadic-safe in the contextual sense ([Definition 2.5](#)), has its deficient vertices supplied from the packing W ([Lemma 8.4](#)), and is hereditarily target-uncompressible ([Corollary 3.105](#)). The quantities $\partial_+(X)$, $\sigma(X)$, and $\mathcal{N}(X)$ below are therefore accounting data for the embedded support, not invariants of an isolated graph.

Definition 13.4 (Net charge). For an admissible support X , define

$$\mathcal{N}(X) = \partial_+(X) - \sigma(X) - \frac{1}{4}|V(X)|.$$

Lemma 13.5 (Net-charge localization under support splitting). Let R be decomposed canonically into connected admissible supports $X_1, \dots, X_k$. Then

$$\sum_{i=1}^k \mathcal{N}(X_i) = \partial_+(R) - \sigma(R) - \frac{1}{4}|R|.$$

Consequently, if

$$\partial_+(R) - \sigma(R) - \frac{1}{4}|R| < -\varepsilon|R|$$

for some fixed $\varepsilon > 0$, then some connected support X_i has $\mathcal{N}(X_i) < 0$.

Proof. The remainder vertex sets Y_i form a partition of $V(R)$, and the ambient surplus units form a disjoint assigned credit partition. By [Definition 13.2](#),

$$\sum_i |V(X_i)| = |R|, \quad \sum_i \sigma(X_i) = \sigma(R).$$




Figure 20: Visual definition of net-charge localization. The green rectangle is the P_{13} -free remainder R , decomposed into connected admissible supports X_i by the canonical support decomposition. Green ellipses show supports with nonnegative net charge; the red ellipse is the localized support $X = X_i$ with $\mathcal{N}(X) < 0$. Each support carries positive deficiency $\partial_+(X_i)$, assigned ambient surplus $\sigma(X_i)$, and net charge $\mathcal{N}(X_i) = \partial_+(X_i) - \sigma(X_i) - \frac{1}{4}|V(X_i)|$. The top budget box records additivity over connected components, $\sum_i \mathcal{N}(X_i) = \partial_+(R) - \sigma(R) - \frac{1}{4}|R| < 0$. The solid arrow labelled “localize” selects a negative support, the dashed arrow records inherited admissibility data, and the two outgoing solid arrows route that support to the Type A low-deficiency atom branch or the Type B high-degree fan-safe support branch. The inherited data are contextual dyadic-safety, empty internal 3-core, componentwise P_{13} -freeness, and window-supplied deficient boundary.

For each vertex $u \in Y_i$, the component property gives $d_{X_i}(u) = d_R(u)$. Hence the positive deficiencies add exactly:

$$\sum_i \partial_+(X_i) = \partial_+(R).$$

Therefore

$$\sum_i \mathcal{N}(X_i) = \sum_i \partial_+(X_i) - \sigma(R) - \frac{1}{4}|R| = \partial_+(R) - \sigma(R) - \frac{1}{4}|R|,$$

the displayed identity. If the right-hand side is negative by a linear amount and all summands were nonnegative, their sum could not be negative; hence some connected support has negative net charge. $\square$

Corollary 13.6 (Global window-join pressure). *Let R be decomposed canonically into connected admissible supports $X_1, \dots, X_k$. Then*

$$\sum_{i=1}^k \mathcal{N}(X_i) \leq 15p_{13} + \sigma_W - \sigma_R - \frac{1}{4}|R|.$$

Consequently, if every connected admissible support in this decomposition has nonnegative net charge, then

$$\sigma_W - \sigma_R \geq \frac{1}{4}|R| - 15p_{13} = \frac{n - 73p_{13}}{4}.$$

Equivalently, if $p_{13} \leq \theta n + o(n)$, then the no-negative-support alternative forces

$$\sigma_W - \sigma_R \geq \frac{1 - 73\theta}{4} n - o(n).$$

In particular, under the near-cubic window-only cap $\theta \leq \theta_{\text{win}} + o(1)$,

$$\sigma_W - \sigma_R \geq \kappa_{\text{win}} n - o(n), \quad \kappa_{\text{win}} := \frac{1 - 73\theta_{\text{win}}}{4} = 0.0182217518254 \dots$$

Proof. By Lemma 13.5,

$$\sum_i \mathcal{N}(X_i) = \partial_+(R) - \sigma_R - \frac{1}{4}|R|.$$

The surplus-aware window-stub inequality Lemma 8.6 gives

$$\partial_+(R) - \sigma_R \leq 15p_{13} + \sigma_W - \sigma_R.$$

Combining the two displays proves the first inequality.

If every X_i has $\mathcal{N}(X_i) \geq 0$, then the left-hand side is nonnegative. Therefore

$$0 \leq 15p_{13} + \sigma_W - \sigma_R - \frac{1}{4}|R|,$$

which is the displayed lower bound on $\sigma_W - \sigma_R$. Since $|R| = n - 13p_{13}$, this lower bound equals

$$\frac{n - 13p_{13}}{4} - 15p_{13} = \frac{n - 73p_{13}}{4}.$$

The density forms follow by substituting $p_{13} \leq \theta n + o(n)$ and then $\theta = \theta_{\text{win}} + o(1)$. $\square$

Remark 13.7 (Meaning of the join-pressure inequality). *Corollary 13.6 is the exact accounting form of the non-near-cubic pressure. Surplus carried by R appears with a negative sign in $\mathcal{N}(R)$, so it helps produce a negative admissible support. Avoiding such a support requires a linear excess of window surplus over remainder surplus. By Lemma 8.5, that excess is not abstract: it is realized by additional window-remainder incidences and cross-window joins.*

Proposition 13.8 (Large-budget residual is negative-net-charge). *If the large-budget branch survives and the local lower bound fails by a fixed linear margin, i.e. if for some $\varepsilon > 0$,*

$$\partial_+(R) - \sigma(R) \leq \frac{1}{4}|R| - \varepsilon|R|$$

for all sufficiently large n , then there is a connected admissible support X with

$$\mathcal{N}(X) < 0.$$

Proof. By Lemma 13.5,

$$\sum_i \mathcal{N}(X_i) = \partial_+(R) - \sigma(R) - \frac{1}{4}|R|.$$

By the displayed hypothesis, the right-hand side is negative for all sufficiently large n . Hence at least one connected admissible support X_i has negative net charge. Choose such a support and call it X . $\square$

Framework branch table. The large-budget residual localizes negative mass by the following branch table.

Branch	Residual condition	Demand	Payer	Bounded multiplicity	Closure
Net-charge cap	$\partial_+(R) - \sigma(R) < \frac{1}{4} R $.	Negative total charge.	The support decomposition.	The superadditive component sum.	Lemma 13.5.
Localization	The total net charge is negative.	A connected deficit support.	The admissible support class.	The component partition.	Proposition 13.8.
Window pressure	No negative support appears without surplus pressure.	$\sigma_W - \sigma_R$.	The window-join identity.	The stub partition.	Corollary 13.6.

14 Local analysis of a negative-net-charge support

Let X be a connected admissible support with

$$\mathcal{N}(X) < 0.$$

We show that X is forced into one of two local obstruction classes.

Remark 14.1 (How to read the local lemmas). *The support X is admissible ([Definition 13.3](#)): a piece of the minimal counterexample G , so it carries every global constraint, not merely its own induced structure. In particular it is dyadic-safe in the strong, contextual sense ([Definition 2.5](#))— G has no power-of-two cycle meeting $V(X)$, including cycles that leave X through its boundary and return through $G - X$ —its deficient vertices are supplied from the packing W ([Lemma 8.4](#)), and it is hereditarily target-uncompressible ([Corollary 3.105](#)). Consequently the inequalities $\mathcal{N}(X) \geq 0$ proved below are not intrinsic properties of subcubic, empty-3-core, P_{13} -free graphs taken in isolation—such graphs may violate them—but properties of supports that admit a completion to a graph of minimum degree three with no power-of-two cycle. Every step uses this admissibility; a lemma read with only its bare graph hypotheses has been read with strictly fewer hypotheses than it states.*

14.1 Type A: low-deficiency P_{13} -free atom

If X carries no high-degree surplus, then $\sigma(X) = 0$ and

$$\mathcal{N}(X) < 0$$

means

$$\partial_+(X) < \frac{1}{4}|V(X)|.$$

If X is subcubic, then $\partial_+(X) = \sum_{v \in V(X)} (3 - d_X(v)) = 3|V(X)| - 2|E(X)|$, so $\partial_+(X) < \frac{1}{4}|V(X)|$ is equivalent to $3|V(X)| - 2|E(X)| < \frac{1}{4}|V(X)|$, that is

$$\frac{2|E(X)|}{|V(X)|} > 2.75.$$

Since X lies in the P_{13} -free remainder, it is P_{13} -free. We first bound its diameter. A shortest path between two vertices is always induced, so if X had two vertices at distance 12, the shortest path between them would be an induced path on 13 vertices, i.e. an induced P_{13} , contradicting P_{13} -freeness. Hence $\text{diam}(X) \leq 11$. Consequently breadth-first layering from any root v_0 has at most $\text{diam}(X) + 1 = 12$ layers; since X is subcubic, the first layer (the neighbours of v_0) has at most 3 vertices, and each vertex in a later layer contributes at most 2 vertices to the next layer, one of its at most three incident edges leading back toward the root. Therefore

$$|V(X)| \leq 1 + 3 \sum_{i=0}^{10} 2^i = 1 + 3(2^{11} - 1) = 6142.$$

If the root has degree at most 2, this improves to $1 + 2(2^{11} - 1) = 4095$; in either case the Type A class is finite:

$$|V(X)| \leq 6142, \quad \partial_+(X) < |X|/4, \quad X \text{ is } P_{13}\text{-free and dyadic-safe.}$$

The exclusion of this class is established by a contextual density lemma. The lemma is stated separately because the bound is a statement about completions in G , not about isolated P_{13} -free graphs.

Definition 14.2 (Type A support). *A Type A support is a connected admissible support X with $\sigma(X) = 0$. Since G has minimum degree at least 3, this condition gives $d_G(v) = 3$ for every $v \in V(X)$. Hence X is subcubic, the deficient vertices of X are exactly the vertices of internal degree at most 2, and the P_{13} -freeness, empty internal 3-core, contextual dyadic-safety, supplied deficiency, and hereditary target-uncompressibility of X are inherited from [Definition 13.3](#). In the negative-net-charge branch a Type A support also satisfies $\partial_+(X) < |V(X)|/4$.*

Definition 14.3 (Receivers, ports, and routed load). *Let X be a Type A support. A receiver is a vertex $w \in V(X)$ with $d_X(w) \leq 2$, and*

$$q(w) = 3 - d_X(w).$$

Since $\sigma(X) = 0$, every vertex of X has ambient degree exactly 3, so $q(w)$ is also the number of ambient edges from w to $G - X$. An oriented edge $\vec{e} = (w, h)$ with $wh \in E(G) \setminus E(X)$ is a completion port of w .

Let X_3 be the subgraph induced by the vertices of X with internal degree 3. The empty internal 3-core condition implies that every component of X_3 has an edge to at least one receiver. For a cubic vertex u , let T_u be the lexicographically first path in X that starts at u , stays in X_3 until its last edge, and ends at a receiver; write this terminal receiver as $r(u)$. The path T_u is the canonical trace of u . The routed load at w is

$$L(w) = |\{u \in V(X) : d_X(u) = 3, r(u) = w\}|.$$

A receiver is saturated when $L(w) \geq 4q(w)$ and unsaturated when $L(w) \leq 4q(w) - 1$. For $j = 0, 1, 2$, let

$$\mathcal{R}_j(X) = \{w \in V(X) : d_X(w) = 2 - j\}.$$

Lemma 14.4 (Canonical receiver loads). *For every Type A support X , the map $u \mapsto r(u)$ is defined for each degree-3 vertex of X , and every degree-3 vertex is assigned to exactly one receiver.*

Proof. The empty internal 3-core hypothesis implies that no component of the induced subgraph X_3 is closed under all three incident edges of its vertices. Therefore each component of X_3 has an edge to a vertex of internal degree at most 2, i.e. to a receiver. The canonical trace T_u is chosen by a fixed lexicographic tie-breaking rule among these receiver-reaching paths, so its terminal receiver $r(u)$ is defined and unique. Thus every degree-3 vertex is assigned to exactly one receiver. $\square$

Definition 14.5 (Receiver-entry returns and visible load). *Let $\vec{e} = (w, h)$ be a completion port. An anchored return through $\vec{e}$ is a simple path $P : h \rightsquigarrow w$ in $G - wh$. Its first vertex in X , traversed from h to w , is denoted $\text{ent}_X(P)$.*

If r, w are receivers, a receiver-entry channel from r to w is a simple path $Q : r \rightsquigarrow w$ in X . An anchored return is a receiver-entry return if it decomposes as

$$P = \Gamma \circ Q,$$

where Γ runs from h to a receiver r with internal vertices outside X , and Q is a receiver-entry channel from r to w .

A receiver-entry return $P = \Gamma \circ Q$ through $\vec{e}$ is visible for a routed cubic vertex u with $r(u) = w$ if the canonical trace T_u is a subpath of the channel Q , or if Q is the lexicographically first channel assigned to the terminal receiver edge of T_u . The port $\vec{e}$ carries t visible receiver-entry returns when it carries such returns for t distinct routed cubic vertices. Let $L_{\text{vis}}(w, \vec{e})$ be this number. A load routed to w is visible if it is visible at some completion port of w ; if more than one port sees it, it is assigned to the first such port in the canonical order. The total number of visible loads routed to w is $L_{\text{vis}}(w)$; the remaining loads are silent and have count

$$L_{\text{sil}}(w) = L(w) - L_{\text{vis}}(w).$$

Lemma 14.6 (Common-port return cycle). *Let $\vec{e} = (w, h)$ be a completion port, and let P_1, P_2 be two anchored returns through $\vec{e}$. If P_1 and P_2 are internally vertex-disjoint as paths in $G - wh$, then $P_1 \cup P_2$ is a simple cycle of length $|P_1| + |P_2|$. In particular, if $|P_1| + |P_2| \in \mathcal{D}$, then G contains a dyadic cycle.*

Proof. By definition, each anchored return through $\vec{e}$ is a simple path from h to w in $G - wh$. Thus the two paths have the same endpoints. The internal-disjointness hypothesis is exactly the hypothesis of [Lemma 2.3](#), so their union is a simple cycle of the stated length. $\square$

Lemma 14.7 (First entry is a receiver). *Let P be an anchored return through a completion port of a Type A support X . Then $\text{ent}_X(P)$ is a receiver.*

Proof. Let $r = \text{ent}_X(P)$. The predecessor of r on P lies outside X , so r has an ambient incidence not counted in $d_X(r)$. Since $\sigma(X) = 0$, every vertex of X has ambient degree 3. Hence $d_X(r) \leq 2$, and r is a receiver. $\square$

Lemma 14.8 (Degree-two entry exclusion and entry budget). *Let w be a receiver with $d_X(w) = 2$, and let $\vec{e} = (w, h)$ be its unique completion port. Every anchored return through $\vec{e}$ first enters X at a receiver $r \neq w$. Moreover the number of distinct external edges through which such returns first enter X is at most*

$$\sum_{\substack{r \text{ receiver} \\ r \neq w}} q(r).$$




Figure 21: Visual definition of a Type A receiver. The shaded region is the Type A support X , and $\sigma(X) = 0$ records that the support has no ambient surplus. White circled vertices are receivers; the terminal receiver is w , and $r = \text{ent}_X(P)$ is the first vertex of X met by the anchored return P . Black filled vertices are internal vertices of the support or of a displayed trace. The dashed circled vertex h lies outside X , and the grey arrow $\vec{e} = (w, h)$ is the completion port being tested. The red path is the anchored return $P = \Gamma \circ Q$: the outside connector Γ runs from h to r , while the internal receiver-entry channel $Q : r \rightsquigarrow w$ runs inside X . The solid blue path T_u is the canonical trace of the routed cubic vertex u and makes u visible at $\vec{e}$ because $T_u \subseteq Q$; the dashed blue trace from u' illustrates a routed load not seen by this port. The three legend boxes define the receiver data, visibility criterion, and receiver capacities. Here $d_X(w)$ is the degree of w inside X , $q(w) = 3 - d_X(w)$ is the number of missing ambient incidences at w , and $L(w) = |\{u : r(u) = w\}|$ is the number of routed cubic vertices assigned to w . The receiver is saturated when $L(w) \geq 4q(w)$, and the displayed capacities 3, 7, 11 are the values of $4q(w) - 1$ for $d_X(w) = 2, 1, 0$.

Proof. An anchored return through $\vec{e}$ is a path from h to w in $G - wh$. By Lemma 14.7, its first vertex r in X is a receiver. If $r = w$, then the return enters w from outside X through an ambient edge different from the deleted edge wh . This is impossible because $d_X(w) = 2$ and $\sigma(X) = 0$ give exactly one ambient edge from w to $G - X$, namely wh . Hence $r \neq w$. For each first-entry receiver r , the entering edge is one of the $q(r)$ completion ports of r , giving the displayed bound. $\square$

Lemma 14.9 (Ports have anchored returns). *Every completion port of a Type A support has at least one anchored return.*

Proof. A completion port is an oriented edge of G . By Lemma 3.4, the underlying edge lies on a cycle in G . Removing the port edge from that cycle leaves a simple return path in the required orientation. $\square$

Definition 14.10 (Connector length and channel spectrum). *Let $\vec{e} = (w, h)$ be a completion port and let*

$$P = \Gamma \circ Q$$

be a receiver-entry return through $\vec{e}$, where Γ runs from h to the first-entry receiver r with internal vertices outside X , and $Q : r \rightsquigarrow w$ lies in X . The path Γ is the connector, and its length is denoted $g(\Gamma)$.

For receivers r, w define the internal channel spectrum

$$\Lambda_X(r, w) = \{|Q| : Q \text{ is a simple path in } X \text{ from } r \text{ to } w\}.$$

The interval length $I_X(r, w)$ is the largest cardinality of an interval of consecutive integers contained in $\Lambda_X(r, w)$.

Definition 14.11 (Connector germs and continuation classes). *Let X be a Type A support, let $\vec{e} = (w, h)$ be a completion port, and let*

$$P = \Gamma \circ Q$$

be a receiver-entry return through $\vec{e}$. Write

$$\Gamma = (x_0, x_1, \dots, x_g), \quad x_0 = h, \quad x_g = \text{ent}_X(P),$$

where $x_1, \dots, x_{g-1} \notin X$. The outside connector germ of P is the ordered path Γ together with the first-entry receiver x_g , the channel Q , and the declared response coordinate of [Definition 3.94](#) determined by $(\vec{e}, x_g, g(\Gamma), Q)$.

Let κ_i, κ_j be two declared response coordinates through the same completion port, with connector germs Γ_i, Γ_j . They have the same continuation class up to an outside vertex z when z occurs in both germs, the two germs have the same ordered prefix from h to z , and the two coordinates have the same image in the relevant boundary-degree fibre. They separate at z when they have the same continuation class up to z and their next incidences after z are distinct. Here the next incidence may be an edge from z to an outside vertex or the first-entry edge from z into X .

For a finite family $\mathcal{K}$ of declared response coordinates through $\vec{e}$, a vertex $z \notin X$ is a first separator of $\mathcal{K}$ if some pair of coordinates in $\mathcal{K}$ separates at z and no pair separates at an earlier vertex on their common prefix from h .

If κ_i, κ_j separate at z , the associated switch support $S_z(\kappa_i, \kappa_j)$ is the finite connected support consisting of the common prefix from h to z , the two connector tails from z to their first-entry receivers, the two receiver-entry channels in X , the completion port boundary datum, and the declared supports of the two response coordinates. The separator z is absorbed when the response identification on this switch support is target-defective, target-complete on a nontrivial response quotient, or target-complete only after adjoining a larger connected support. It is surviving when it is not absorbed.

Lemma 14.12 (Spectral interval pressure). *Let $\vec{e} = (w, h)$ be a completion port and let Γ be a connector from h to a first-entry receiver r . Then*

$$g(\Gamma) + \lambda \notin \mathcal{M} \quad \text{for every } \lambda \in \Lambda_X(r, w).$$

Consequently, if $\Lambda_X(r, w)$ contains the interval $[a, b] \cap \mathbb{Z}$, then

$$g(\Gamma) \notin \bigcup_{\mu \in \mathcal{M}} ([\mu - b, \mu - a] \cap \mathbb{Z}_{\geq 1}).$$

Proof. For any simple channel $Q : r \rightsquigarrow w$ in X , the concatenation $\Gamma \circ Q$ is a simple return through $\vec{e}$: the internal vertices of Γ avoid X , while Q lies in X . If $g(\Gamma) + |Q| \in \mathcal{M}$, then $\Gamma \circ Q$ together with the port edge gives a power-of-two cycle by [Lemma 2.2](#), contradicting the counterexample condition. Varying Q over all channels gives the first assertion. The displayed band exclusion is the same statement applied to every $\lambda \in [a, b] \cap \mathbb{Z}$. $\square$

Remark 14.13 (The receiver-5 witness and the band constraint). *For the explicit 15-vertex near-miss atom described in the referee material, the receiver-channel spectra from the two other degree-2 receivers to the saturated receiver are both the full interval $[3, 14]$. Thus a connector length g into either entry receiver must avoid*

$$\bigcup_{\mu \in \mathcal{M}} [\mu - 14, \mu - 3].$$

In particular, below 112 the allowed positive lengths are

$$[13, 16] \cup [29, 48] \cup [61, 112],$$

and $[1, 12]$ is forbidden. This is not an additional external theorem; it is the universally quantified target condition $R_e(G) \cap \mathcal{M} = \emptyset$ applied to all channel lengths at once.

Definition 14.14 (Silent trace basins). *Let u be a routed cubic vertex of a Type A support, and let T_u be its canonical trace ending at $r(u)$. For a connected subgraph $S \subseteq X$ with $T_u \subseteq S$, let*

$$\partial S = \{v \in V(S) : v \text{ is incident with an edge of } G - S\},$$

and regard S as a ∂S -boundaried graph with the inherited labels. Each local target-response coordinate used in this paper has a declared finite support: for example, a return coordinate is supported on its port and channel, a P_{13} label coordinate on its window, a curvature coordinate on its length-two wedge, and a boundary-degree coordinate on the corresponding boundary vertex. A coordinate is u -supported if its declared support meets T_u , or if it is one of the completion-port labels, receiver-entry channel lengths, connector-band constraints, boundary-degree entries, or P_{13} labels whose declared support in [Definition 3.94](#) belongs to a receiver-entry return whose internal channel contains T_u or is canonically assigned to the terminal receiver edge of T_u . Let $\mathcal{R}_u(S)$ be the finite family of u -supported coordinates carried by S . The trace-incidence coordinate recording the labelled path T_u is included whenever $T_u \subseteq S$; hence a trace basin always has at least one internal u -supported coordinate.

For $S \subseteq X$, let

$$\text{res}_{X,S} : \mathcal{R}_u(X) \dashrightarrow \mathcal{R}_u(S) \cup \{\mathbf{d}_{\partial}(S)\text{-entries}\}$$

be the partial restriction map which is defined on a coordinate precisely when the coordinate's declared support is carried by S , or when its value is fixed by the boundary degree profile $\mathbf{d}_{\partial}(S)$. A connected $S \subseteq X$ containing T_u is trace-complete for u if $\text{res}_{X,S}$ is defined on every coordinate of $\mathcal{R}_u(X)$ and the restricted value equals the value of that coordinate in X . Since X itself is trace-complete for u , the finite set of trace-complete connected subgraphs has lexicographically first inclusion-minimal elements. The trace basin B_u is the lexicographically first such element.

The trace-response state of B_u is

$$\rho_u(B_u) = (\mathbf{d}_{\partial}(B_u), \mathcal{R}_u(B_u), \mathbf{H}_{\partial B_u}(B_u)).$$

The family $\mathcal{R}_u(B_u)$ is the complete declared coordinate family for the u -supported target events used in the route-8 branch: once the boundary degree profile and all entries of $\mathcal{R}_u(B_u)$ are fixed, every compatible outside context has the same truth value for each such declared event. Carrier quotients below are tested only against this declared u -supported target algebra. A response quotient of $\rho_u(B_u)$ is obtained by identifying or forgetting entries of the finite coordinate family $\mathcal{R}_u(B_u)$ while preserving the boundary degree profile. A realization of such a quotient is a boundaried response state with the same boundary degree profile whose image under the quotient map is the given quotient. The quotient is target-complete for $\rho_u(B_u)$ if, for every outside ∂B_u -context compatible with the

boundary profile, all realizations of the quotient give the same target predicate as $\rho_u(B_u)$ after gluing. Thus a target-complete quotient identifies only target-indistinguishable response states. The quotient is nontrivial if it forgets or identifies a coordinate whose declared support meets $B_u - \partial B_u$, or whose declared support contains an internal edge of B_u , that is, an edge of B_u with both endpoints in $V(B_u)$.

A response quotient of $\rho_u(B_u)$ is trace-local and non-carrier if it identifies or forgets entries of the fixed coordinate family $\mathcal{R}_u(B_u)$, preserves the full boundary degree profile, and does not delete an ambient boundary carrier from the carrier set used later in [Definition 14.48](#). Thus carrier-restriction and carrier-deletion quotients are excluded from trace-complete-minimality; they are tested only in the two-carrier clause of [Definition 14.49](#).

The trace basin B_u is target-complete-minimal when none of the following four alternatives occurs:

- (a) a trace-local non-carrier quotient or identification inside $\rho_u(B_u)$ is distinguished by some outside ∂B_u -context, hence is target-defective;
- (b) $\rho_u(B_u)$ has a nontrivial target-complete response quotient with the same boundary degree profile; when this quotient is realized by a smaller connected representative, it is a target-complete compression of a proper atom;
- (c) an equality among coordinates of $\rho_u(B_u)$ becomes target-complete only after adjoining a larger connected support $Z \supsetneq B_u$, either with $Z \subsetneq G$ or with $Z = G$;
- (d) two declared outside connector germs of $\rho_u(B_u)$, with the same boundary-degree image, have a surviving first separator in the sense of [Definition 14.11](#); by [Lemma 14.24](#) this separator has ambient degree at least 4.

Thus a trace basin that is not target-complete-minimal realizes, respectively, one of the quotient, response-compression, delocalization, or decorated handoff fan alternatives used below in exits (4)–(7).

Definition 14.15 (Silent-core residual profile). A silent-core residual profile is the data

$$(X, w, A(w), E(w), \mathcal{U}(w), (B_u)_{u \in \mathcal{U}(w)}, \{\Lambda_X(r, w)\}_r, \mathcal{B})$$

obtained from a saturated Type A receiver w when no port carries four visible receiver-entry returns and none of the target, quotient, compression, delocalization, or decorated handoff fan alternatives applies. Here $A(w)$ is the canonical payable set of [Definition 14.29](#), $E(w)$ is the excess basin, $\mathcal{U}(w)$ is the set of unpaid silent routed vertices in that basin, $(B_u)_{u \in \mathcal{U}(w)}$ is the indexed family of trace basins from [Definition 14.14](#), and $\mathcal{B}$ records the completion-port labels and connector-length bands forced by [Lemmas 14.9](#) and [14.12](#). Such a profile is admissible only if every B_u with $u \in \mathcal{U}(w)$ is target-complete-minimal and all connector bands, P_{13} labels, boundary degree profiles, and cross-port theta constraints are compatible with the standing counterexample invariants.

Definition 14.16 (Saturated receiver exits). Let w be a saturated receiver in a Type A support. The saturated receiver branch has the following eight exits:

- (1) an anchored return through a completion port of w has length in $\mathcal{M}$;
- (2) two anchored receiver-entry returns through one completion port are internally vertex-disjoint as anchored paths and their lengths sum to a power of two;

- (3) a shared P_{13} window violates the corresponding legal-label relation C_s ;
- (4) a quotient in the canonical exit-(4) family $\mathcal{Q}_4(w)$ of [Definition 14.49](#) is target-defective;
- (5) the response data give a nontrivial target-complete response compression; when this compression is realized by a smaller proper atom, it is a target-complete compression in the sense of [Lemma 3.103](#);
- (6) the equality of responses delocalizes to a proper support or to the whole graph;
- (7) a high-degree decorated handoff fan envelope is produced;
- (8) none of the previous alternatives occurs, and the data form the silent-core residual profile of [Definition 14.15](#).

Exits (1)–(3), (5), and (6) are closed exits. Exit (4) is the target-defect peeling exit: it removes exactly the routed load whose declared coordinate is used by the canonical target-defective quotient, as in [Definition 14.33](#) and [Lemma 14.38](#). Exit (7) is the Type B handoff exit. Exit (8) is the route-8 residual exit.

Handoff interface. The next definitions are the Type B objects used by the Type A exit-(7) handoff. They are stated here so that the saturated Type A analysis can invoke the interface without using the later Type B exclusion theorem.

Definition 14.17 (Fan-safe relation). *For a high-degree vertex h , the fan-safe graph $\mathcal{F}_h$ has vertex set $N(h)$. Distinct vertices $u, v \in N(h)$ are adjacent in $\mathcal{F}_h$ when the pair satisfies all of the following five conditions:*

- (i) there is no simple u – v return in $G - h$ of length $2^j - 2$;
- (ii) the P_{13} labels carried by the pair satisfy the relevant legal-label curvature relation, in particular $C_2(S(u), S(v)) = 1$ in the length-two fan case;
- (iii) the quotient identifying the two boundary-response coordinates is not target-defective;
- (iv) no target-complete compression of the fan pair preserves the boundary degree profile in a smaller proper representative;
- (v) the equality of the two response coordinates does not delocalize to a proper larger support or to the whole graph.

Definition 14.18 (Marked Type B fan). *A fan-certificate labelling of a high-degree center h is a map*

$$S_h : N(h) \longrightarrow \mathcal{L}$$

from the fan neighbours to the legal nonempty P_{13} labels of [Lemma 6.1](#), such that

$$C_2(S_h(u), S_h(v)) = 1 \quad (u \neq v \in N(h)).$$

A Type B fan is certificate-marked when this labelling is part of the assigned support data. A high-degree center assigned to a Type B support but not certificate-marked is a fan-certificate residual center; it is included in the Type B bridge-residual mass of [Definition 14.101](#) and is not used in the certificate-closed local discharging step.

A marked Type B fan centered at a high-degree vertex h is the data of h , the neighbour set $N(h)$, and a fan-certificate labelling S_h . Since $d_G(h) \geq 4$, [Lemma 3.2](#) forces every neighbour of h to have degree three in G . A neighbour $u \in N(h)$ is cubic-closed in the marked fan if, after assigning the fan support, the two non- h incidences of u are also assigned to that support. Equivalently, u has internal degree 3 in the assigned fan envelope and contributes charge $-\frac{1}{4}$. Thus

$$N_G(u) = \{h, x_u, y_u\}.$$

When x_u and y_u lie in the packed P_{13} windows, their incidences are recorded in the window-incidence profile below; when they lie in the same window, the corresponding fan-certificate label is non-singleton. A neighbour that is not cubic-closed has internal degree at most 2 in the assigned fan support and contributes charge at least $\frac{3}{4}$.

Definition 14.19 (Decorated handoff fan envelope). A decorated handoff fan envelope is a pair

$$\mathfrak{X} = (Y, H),$$

together with the handoff-arm data described below. Here $Y \subseteq R$ is a connected P_{13} -free remainder core and $H \subseteq V_{\geq 4}(G)$ is a set of high-degree decorations assigned to Y .

For each $h \in H$ there is a nonempty assigned first-neighbour set

$$K_h \subseteq N_G(h).$$

For every $a \in K_h$ the envelope records a simple handoff arm

$$A_{h,a} = a = x_0, x_1, \dots, x_{\ell(a)} = y_{h,a}, \quad y_{h,a} \in V(Y),$$

whose internal vertices avoid $Y \cup H \cup \{h\}$. The full handoff path from h to Y is the edge ha followed by $A_{h,a}$. The first neighbours in K_h are distinct. The arm union for a fixed h is finite and is recorded as part of the decorated response support; different arms may share a later outside suffix, in which case the shared vertices and edges are counted once in that declared support. The ordinary adjacent fan is the special case $\ell(a) = 0$ for all a , so $K_h \subseteq V(Y)$.

The set K_h is required to be a clique in the fan-safe graph $\mathcal{F}_h$ of [Definition 14.17](#). This condition is imposed on the actual neighbours of h , not on the terminal vertices $y_{h,a}$. Thus any return from a to b in $G - h$ of length $2^j - 2$ would close with the two edges ha, hb to form a cycle of length 2^j ; the return is allowed to run through the recorded arms and through Y .

The branch data producing the envelope include the boundary response coordinates incident with h , the handoff arms, their terminal coordinates in Y , and the boundary-degree profile of the decorated support. Its net charge is

$$\mathcal{N}(\mathfrak{X}) = \partial_+(Y) - \sum_{h \in H} (d_G(h) - 3) - \frac{1}{4}|V(Y)|.$$

The P_{13} -free and empty-3-core hypotheses are imposed on the counted core Y ; the vertices of H supply surplus and fan-return constraints, while the handoff arms are declared response data used to route the Type A separator into the Type B fan branch.

Definition 14.20 (Decorated Type B envelope supports). Let $\mathcal{Y}$ be a finite family of pairwise vertex-disjoint Type A cores which produce exit-(7) handoffs. A decorated Type B envelope support for $\mathcal{Y}$ is obtained from the bipartite core-center incidence graph $\mathcal{I}(\mathcal{Y})$. The left vertices of $\mathcal{I}(\mathcal{Y})$ are the cores $Y \in \mathcal{Y}$; the right vertices are the ambient high-degree handoff centers; and Y is adjacent

to h when one of the exit-(7) handoffs from Y has center h . For a connected component $\mathfrak{C}$ of this incidence graph, write

$$\mathcal{Y}_{\mathfrak{C}} = \{Y \in \mathcal{Y} : Y \in V(\mathfrak{C})\}, \quad H_{\mathfrak{C}} = \{h : h \in V(\mathfrak{C})\}.$$

The grouped envelope for $\mathfrak{C}$ has counted core

$$Y_{\mathfrak{C}}^* = \bigcup_{Y \in \mathcal{Y}_{\mathfrak{C}}} Y$$

and center-token sum

$$\omega(\mathfrak{C}) = \sum_{h \in H_{\mathfrak{C}}} \omega(h), \quad \omega(h) = d_G(h) - 3.$$

This uses the ambient surplus of h whether h lies in the remainder or in a packed P_{13} -window. The grouped envelope uses $\omega(h)$ once inside its incidence component. If the same surplus unit is also canonically assigned to the ordinary support containing h , that ordinary use is counted separately in the at-most-twice global fan-mass bound of [Definition 14.101](#).

The net charge of the grouped envelope is

$$\mathcal{N}(\mathfrak{X}_{\mathfrak{C}}^*) = \sum_{Y \in \mathcal{Y}_{\mathfrak{C}}} \partial_+(Y) - \omega(\mathfrak{C}) - \frac{1}{4} \sum_{Y \in \mathcal{Y}_{\mathfrak{C}}} |V(Y)|.$$

The handoff arms and fan-safe neighbour sets are the union of the corresponding decorated handoff fan envelopes in $\mathfrak{C}$, with shared outside vertices and edges counted once.

Lemma 14.21 (Decorated envelope transfer has no double counting). *Let $\mathcal{Y}$ be a family of pairwise disjoint Type A cores which produce exit-(7) handoffs, and form the grouped decorated Type B envelope supports of [Definition 14.20](#). Then the Type A core deficiencies are counted exactly once in the grouped envelope family, each ambient handoff-center token is counted exactly once in that family, and*

$$\sum_{\mathfrak{C}} \mathcal{N}(\mathfrak{X}_{\mathfrak{C}}^*) = \sum_{Y \in \mathcal{Y}} \left(\partial_+(Y) - \frac{1}{4} |V(Y)| \right) - \sum_{\mathfrak{C}} \omega(\mathfrak{C}).$$

Consequently a negative Type A handoff contribution is not discarded when the branch leaves the Type A calculation; it is transferred to the grouped Type B envelope ledger.

Proof. The cores in $\mathcal{Y}$ are components of the canonical decomposition, so their vertex sets and positive deficiencies are disjoint summands. The connected components of the core-center incidence graph partition the set of cores and the set of handoff centers. Thus a core with handoffs to several centers is counted once, in the single component containing all those centers, and a center serving several cores contributes its token once to that same component. Summing the displayed definition of $\mathcal{N}(\mathfrak{X}_{\mathfrak{C}}^*)$ over the incidence components gives the identity. The identity is an equality of the transferred ledger entries, so no core deficiency or grouped-envelope center token is counted twice inside the handoff envelope family. $\square$

Lemma 14.22 (Window handoff center accounting). *Let h be an exit-(7) handoff center lying in a packed P_{13} -window. Then either exit (3) occurs at the saturated receiver that produced the handoff, or the grouped envelope containing h is charged to the surplus token $d_G(h) - 3$ of the window vertex h .*

Proof. The handoff arms through a packed window carry the legal P_{13} -labels of Lemma 6.1. If two such arms violate the relevant relation C_s , then the saturated receiver realizes the label-collision exit (3). Otherwise the labels are legal and the handoff survives as a decorated Type B envelope. By Lemma 14.24, the handoff center has $d_G(h) \geq 4$, so it supplies the positive surplus token $d_G(h) - 3$. That token is one of the center tokens of the incidence component containing h in Definition 14.20. $\square$

Lemma 14.23 (Continuation routing at a completion port). *Let X be a Type A support, let $\vec{e}$ be a completion port, and let $\mathcal{K}$ be a finite family of declared response coordinates through $\vec{e}$ lying in one boundary-degree fibre. Suppose the coordinates in $\mathcal{K}$ are pairwise distinct in every target-complete quotient and are distinguished by the actual outside context. Then either one of exits (4)–(6) of Definition 14.16 occurs, or $\mathcal{K}$ has a surviving first separator in the sense of Definition 14.11.*

Proof. Each coordinate in $\mathcal{K}$ has a finite outside connector germ by Definition 14.11. If two coordinates have the same germ through their first entry into X and the same image in the boundary-degree fibre, then identifying them is an identification of declared coordinates on a finite response state. By Lemma 3.102, either some outside context distinguishes that identification, which is exit (4), or the identification is target-complete. In the latter case it is nontrivial because the two declared coordinates are pairwise distinct in every surviving target-complete quotient; hence it is exit (5) unless its validity requires adjoining a larger connected support, which is exit (6).

Thus, if exits (4)–(6) do not occur, some pair of germs must separate before their first-entry response data are identified. Since $\mathcal{K}$ is finite and each germ is finite, there is a first such separator in the prefix order. If that separator were absorbed, the definition of absorbed gives exactly one of exits (4)–(6). Therefore, outside exits (4)–(6), the first separator is surviving. $\square$

Lemma 14.24 (Cubic switch absorption). *Let z be a surviving first separator for two declared response coordinates through one completion port of a Type A support. Then*

$$d_G(z) \geq 4.$$

Proof. Let the two connector germs have common prefix ending at z . If z is the initial outside vertex h of the completion port, the port edge wh is the root incidence at z ; otherwise the last edge of the common prefix is the root incidence. Since the two germs separate at z , they use two distinct next incidences after z . Hence $d_G(z) \geq 3$, and $d_G(z) \leq 2$ is impossible.

Assume $d_G(z) = 3$. Then the root incidence and the two next incidences used by the separated germs are all incidences of z . Consequently the switch support S_z of Definition 14.11 has no unused ambient incidence at z . The two separated responses therefore form a finite declared bounded response state with the same boundary-degree profile: the boundary records the root incidence, the two connector tails to their first entries in X , and the two receiver-entry channels. Consider the quotient that identifies the two separated response coordinates on this finite state. If some compatible outside context distinguishes the two responses, the quotient is target-defective by Lemma 3.102, which is exit (4). Otherwise the identification is target-complete. If all declared support needed for this target-complete identification is already contained in S_z , then the quotient is nontrivial, because the two response coordinates are distinct, and this is exit (5). If some declared support needed for the identification lies outside S_z , choose the lexicographically first smallest connected support $Z \supsetneq S_z$ carrying it. When $Z \subsetneq G$ this is proper delocalization, and when $Z = G$ it is global delocalization; in both cases the branch is exit (6).

Each of these alternatives is exactly the absorbed case in Definition 14.11. This contradicts that z is surviving. Thus $d_G(z) \neq 3$, and the already-proved inequality $d_G(z) \geq 3$ gives $d_G(z) \geq 4$. $\square$

Lemma 14.25 (High-degree separator handoff). *Let X be a Type A support, and let z be a surviving first separator for a finite family of declared response coordinates through one completion port. Then z , together with the separated connector tails from z to their first-entry data in X , produces a decorated handoff fan envelope in the sense of Definition 14.19.*

Proof. By Lemma 14.24, $d_G(z) \geq 4$, so z is an eligible high-degree decoration. Take $Y = X$ as the counted remainder core and take $H = \{z\}$. For each continuation class that leaves z through a next neighbour a , let $A_{z,a}$ be the remaining connector tail from a to its first-entry receiver in X ; if $a \in X$, this tail has length 0. These are simple paths because the receiver-entry returns are simple. Distinct first neighbours give distinct first-neighbour data. If two tails later meet before their first entry into X , the common suffix is part of the finite arm union and is recorded once. Thus the required handoff-arm data are well-defined.

It remains to verify the fan-safe condition on K_z . Suppose that two assigned first neighbours $a, b \in K_z$ are not adjacent in $\mathcal{F}_z$. The five defining failures of Definition 14.17 give, respectively, a return in $G - z$ of length $2^j - 2$, a P_{13} label collision, a target-defective quotient, a target-complete compression, or proper/global delocalization. In the first case the return together with the two edges za, zb is a power-of-two cycle in G . In the remaining cases the branch is one of exits (3)–(6) of Definition 14.16, so the separator is absorbed rather than surviving. Hence every pair in K_z is fan-safe, and K_z is a clique in $\mathcal{F}_z$. The data just constructed are exactly a decorated handoff fan envelope. $\square$

Lemma 14.26 (Decorated fan admissibility). *If exit (7) of Definition 14.16 occurs in a saturated Type A branch, the resulting decorated handoff fan envelope carries exactly the data required by the Type B fan calculation: contextual dyadic-safety, a P_{13} -free empty-3-core remainder core, hereditary target-uncompressibility of the decorated boundaried profile, and fan-return-safety at every decoration $h \in H$.*

Proof. For a Type A exit (7), the counted core Y is the Type A support X used in Lemma 14.25. Hence Y inherits contextual dyadic-safety, P_{13} -freeness, the empty internal 3-core condition, and hereditary target-uncompressibility from Definition 13.3. The only new vertices counted outside Y are the decorations $h \in H$, and their surplus appears only through the displayed term $d_G(h) - 3$ in Definition 14.19.

For a decoration h , the assigned set K_h consists of actual neighbours of h . If two vertices $a, b \in K_h$ failed the first fan-safe condition, then there would be a simple a – b return in $G - h$ of length $2^j - 2$; adding the two fan edges ha, hb would give a cycle of length 2^j in G . Failures of the other four fan-safe conditions are exactly the label, target-defect, target-compression, and delocalization exits already removed before exit (7). Thus K_h is a clique in $\mathcal{F}_h$. The handoff arms are declared coordinates in the signature of Definition 3.94, so they do not add undeclared response data. The decorated envelope therefore satisfies the Type B hypotheses later consumed by Lemma 14.98. This lemma verifies those hypotheses; it does not use the conclusion of Lemma 14.98. $\square$

Lemma 14.27 (Raw Type A thresholds). *Let $w \in \mathcal{R}_j(X)$, so $d_X(w) = 2 - j$ and $q(w) = j + 1$. The first routed load value at which w leaves the unsaturated Type A discharging branch is*

$$H_j = 4q(w) = 4(j + 1).$$

Consequently

$$H_0 \leq 4, \quad H_1 \leq 8, \quad H_2 \leq 12.$$

If the saturated branch has been eliminated, then every receiver in $\mathcal{R}_j(X)$ satisfies

$$L(w) \leq H_j - 1 = 4(j + 1) - 1.$$

Proof. The charge of a receiver is

$$\text{ch}(w) = q(w) - \frac{1}{4}.$$

Each routed cubic vertex assigned to w costs $\frac{1}{4}$. The largest load that can be paid while keeping nonnegative final charge is the largest integer L with

$$q(w) - \frac{1}{4} - \frac{1}{4}L \geq 0,$$

namely $L \leq 4q(w) - 1$. The next value, $4q(w)$, is exactly the saturated threshold. Substituting $q(w) = 1, 2, 3$ for internal degrees 2, 1, 0 gives 4, 8, 12. After the saturated cases are removed, the surviving capacities are 3, 7, 11. $\square$

Lemma 14.28 (Visible saturated receiver: exits 1–7). *Let w be a receiver in a Type A support. If some completion port of w carries four visible receiver-entry returns, then the branch realizes one of exits (1)–(7) of Definition 14.16.*

Proof. Fix a port $\vec{e} = (w, h)$ carrying four visible receiver-entry returns. Each return has the form $\Gamma \circ Q$, with the same completion port and with Q carrying a distinct canonical trace T_u . If one return has length in $\mathcal{M}$, Lemma 2.2 gives a power-of-two cycle; this is exit (1). If two returns are internally disjoint as anchored h -to- w paths and their full anchored-return lengths sum to a power of two, the union is a forbidden cycle by Lemma 14.6; this is exit (2). If two traces pass through a common P_{13} window, their labels are governed by the relations C_s of Lemma 6.1; failure of the corresponding C_s test is the stated label collision, which is exit (3).

Assume exits (1)–(3) do not occur. For each of the four visible returns record the local target-response coordinate consisting of its completion port, its first-entry receiver, its connector label, and its receiver-entry channel. These four coordinates lie in one boundary-degree fibre in the sense of Definition 3.88. If identifying two of them fails the context-universal target test of Lemma 3.102, the quotient is target-defective; this is exit (4). If the identification is target-complete on a nontrivial response state, this is exit (5); when it is realized on a proper atom, it contradicts Corollary 3.105. If the identification becomes target-complete only after adjoining a larger support, it is proper delocalization when that support is proper and global delocalization when the support is all of G ; these are excluded by Lemmas 10.6 and 10.9 and give exit (6).

It remains that the four response coordinates are pairwise distinct in every target-complete quotient and are separated in the actual outside context. Lemma 14.23, applied to this four-coordinate family, gives a surviving first separator for their outside connector germs. By Lemma 14.24, this separator has ambient degree at least 4. By Lemma 14.25, the separator and its separated connector tails produce a decorated handoff fan envelope. This is exit (7). Lemma 14.26 is used here only to record the handoff interface data; the Type B exclusion conclusion is not used in this Type A lemma. $\square$

Definition 14.29 (Excess basin). *Let w be a receiver with $L(w) \geq 4q(w)$, and assume no completion port of w carries four visible receiver-entry returns. The canonical payable set $A(w)$ is formed in the visible-first elimination order: first list the visible routed loads of w in canonical port order, breaking ties by the canonical trace order, and then list the remaining routed loads by the canonical trace order. Let $A(w)$ be the first $4q(w) - 1$ routed loads in this order, and*

$$E(w) = \{u : r(u) = w\} \setminus A(w)$$

is the excess basin. Let

$$\mathcal{U}(w) = \{u \in E(w) : u \text{ is silent}\}.$$

The set $E(w)$ is nonempty because

$$L(w) \geq 4q(w) > 4q(w) - 1 = |A(w)|.$$

Lemma 14.30 (Visible-first excess is silent and quantitative). *Let X be a Type A support. Suppose that no saturated receiver of X has a completion port carrying four visible receiver-entry returns, and form the visible-first excess basins of Definition 14.29. Define*

$$S_{\text{sil}}^{\text{exc}}(X) = \sum_w |\mathcal{U}(w)|,$$

where the sum is over the receivers of X and $\mathcal{U}(w) = \emptyset$ for unsaturated receivers. Put

$$D_A(X) = \frac{1}{4}|V(X)| - \partial_+(X).$$

Then

$$S_{\text{sil}}^{\text{exc}}(X) \geq n_3 - 3n_2 - 7n_1 - 11n_0 = 4D_A(X),$$

where $n_i = |\{v \in V(X) : d_X(v) = i\}|$.

Proof. For a receiver w , put

$$c(w) = 4q(w) - 1.$$

If $L(w) \leq c(w)$, then w contributes no unpaid routed vertex. If $L(w) > c(w)$, then w is saturated. By hypothesis no completion port of w carries four visible receiver-entry returns. Since w has exactly $q(w)$ completion ports, each port carries at most three visible loads, so

$$L_{\text{vis}}(w) \leq 3q(w).$$

For $q(w) \in \{1, 2, 3\}$ one has

$$3q(w) \leq 4q(w) - 1 = c(w).$$

The visible-first order therefore pays every visible routed load before the payable set is exhausted. Hence every unpaid routed vertex at w is silent, and

$$|\mathcal{U}(w)| = L(w) - c(w) \quad \text{whenever } L(w) > c(w).$$

For receivers with $L(w) \leq c(w)$ the equality is replaced by the inequality $|\mathcal{U}(w)| = 0 \geq L(w) - c(w)$. Summing over all receivers gives

$$S_{\text{sil}}^{\text{exc}}(X) \geq \sum_w (L(w) - c(w)).$$

By Lemma 14.4, every internal degree-3 vertex is routed to exactly one receiver, so $\sum_w L(w) = n_3$. The receivers of internal degree 2, 1, 0 have $q = 1, 2, 3$, and hence capacities 3, 7, 11, respectively. Thus

$$\sum_w c(w) = 3n_2 + 7n_1 + 11n_0.$$

This proves the first displayed inequality.

Finally,

$$\partial_+(X) = 3n_0 + 2n_1 + n_2, \quad |V(X)| = n_0 + n_1 + n_2 + n_3,$$

so

$$4D_A(X) = 4 \left(\frac{1}{4}|V(X)| - \partial_+(X) \right) = n_3 - 11n_0 - 7n_1 - 3n_2.$$

This is the same expression as above. □

Lemma 14.31 (Reduced silent-excess residual). *Let w be a saturated receiver in a Type A support, and assume no completion port of w carries four visible receiver-entry returns. Then either one of exits (4)–(7) of Definition 14.16 occurs, or every $u \in \mathcal{U}(w)$ has a target-complete-minimal trace basin B_u .*

Proof. Fix $u \in \mathcal{U}(w)$ and consider its trace basin B_u from Definition 14.14. The definition of target-complete-minimality lists four possible ways for B_u to fail to be residual-minimal. If a trace-local non-carrier quotient or identification inside $\rho_u(B_u)$ is distinguished by an outside context, then the attempted quotient is target-defective, which is exit (4). If $\rho_u(B_u)$ has a nontrivial target-complete response quotient with the same boundary degree profile, then this is exit (5). When the quotient is realized by a smaller proper atom, it is forbidden by hereditary target-uncompressibility; when it is only a trace-response quotient, it is already failure alternative (b) in Definition 14.14. If an equality among coordinates of $\rho_u(B_u)$ becomes target-complete only after adjoining a larger connected support, then the dependence is proper delocalization when the support is proper and global delocalization when the support is all of G ; this is exit (6). If two declared connector germs in $\rho_u(B_u)$ have a surviving first separator, then Lemma 14.24 gives ambient degree at least 4, and Lemma 14.25 records that separator as a decorated handoff fan envelope. This is exit (7). Therefore, if none of exits (4)–(7) occurs, none of the four failure alternatives occurs for any $u \in \mathcal{U}(w)$, and each such trace basin is target-complete-minimal. $\square$

Lemma 14.32 (Silent and mixed excess: exits 4–8). *Let w be a saturated receiver in a Type A support, and assume no completion port of w carries four visible receiver-entry returns. Then the excess basin $E(w)$ realizes one of exits (4)–(8) of Definition 14.16.*

Proof. The vertices in $E(w)$ are precisely the routed cubic vertices that cannot be paid by the nonnegative receiver charge. Put $c(w) = 4q(w) - 1$. Since no completion port of w carries four visible receiver-entry returns and w has exactly $q(w)$ completion ports,

$$L_{\text{vis}}(w) \leq 3q(w) \leq 4q(w) - 1 = c(w).$$

The payable set $A(w)$ is ordered visible-first, so all visible routed loads are included in $A(w)$. Since

$$E(w) = \{u : r(u) = w\} \setminus A(w),$$

no vertex of $E(w)$ is visible at any completion port of w ; equivalently, every excess vertex is silent. Let $B(w)$ be the boundaried subgraph consisting of the excess basin, the canonical traces from its vertices to w , and the completion-port boundary data through which those traces can be completed. Its boundary response is a coordinate in the target-response collection of Definition 3.88.

If replacing $B(w)$ by its boundary response changes the target predicate in some context, the induced quotient is target-defective; this is exit (4). If the replacement is a nontrivial target-complete response compression, this is exit (5); when supported on a proper atom it is forbidden by Corollary 3.105, and when it is only a trace-response quotient it violates target-complete-minimality. If the equality of responses becomes target-complete only after adjoining a larger support, the dependence is either proper delocalization, excluded by Lemma 10.6, or global delocalization, excluded by Lemma 10.9; this is exit (6). If the response of the basin is separated in the outside context, then Lemma 14.23 either has already routed it to exits (4)–(6), or supplies a surviving first separator. By Lemmas 14.24 and 14.25, that separator produces a decorated handoff fan envelope. The admissibility lemma Lemma 14.26 records the Type B interface data for that envelope; it does not use the Type B exclusion theorem. This is exit (7).

If none of exits (4)–(7) occurs, then [Lemma 14.31](#) applies: every unpaid silent routed vertex in $\mathcal{U}(w)$ has a target-complete-minimal trace basin. The only remaining data are exactly those listed in the silent-core residual profile of [Definition 14.15](#): the excess loads are silent after the visible-first payment, their indexed trace basins are target-complete-minimal, no target-defective quotient occurs, no target-complete compression occurs, no delocalization occurs, and no high-degree fan decoration is created. In that residual case, [Lemmas 14.9](#) and [14.12](#) still impose nonempty anchored-return and connector-band constraints on every completion port. This is exit (8). $\square$

Definition 14.33 (Exit-(4) peeling data). *Let w be a Type A receiver and put*

$$\mathcal{L}(w) = \{u : r(u) = w\}.$$

An exit-(4) witness for a routed load $u \in \mathcal{L}(w)$ consists of a quotient $q \in \mathcal{Q}_4(w)$, two realizations in the same boundary-degree fibre, and a compatible outside context distinguishing their target predicates, such that the declared support of q contains the canonical coordinate of u generated by [Definition 3.94](#). A peeling set $P_4(w) \subseteq \mathcal{L}(w)$ is a set of routed loads, each equipped with one exit-(4) witness, and with no routed load listed twice. The witness serves only as a routing record for its selected load. The receiver charge changes because the selected routed load is removed from the pure Type A receiver calculation. The residual load after peeling is

$$L_4(w) = L(w) - |P_4(w)|.$$

Lemma 14.34 (Exact charge after exit-(4) peeling). *Let w be a Type A receiver with peeling set $P_4(w)$. After the loads in $P_4(w)$ have left the Type A receiver calculation through exit (4), the remaining receiver charge is*

$$q(w) - \frac{1}{4} - \frac{1}{4}L_4(w).$$

In particular, if $L_4(w) \leq 4q(w) - 1$, then the remaining charge at w is nonnegative.

Proof. The receiver w has initial charge $q(w) - \frac{1}{4}$, and every routed cubic vertex assigned to w costs exactly $\frac{1}{4}$ in [Lemma 14.43](#). The peeling set removes exactly $|P_4(w)|$ routed loads from this receiver calculation, with no routed load removed twice. The target-defective witnesses justify routing those selected loads out of the pure Type A branch; the displayed charge counts the selected removed loads. Therefore the number of routed loads still charged to w is $L(w) - |P_4(w)| = L_4(w)$, giving the displayed formula. If $L_4(w) \leq 4q(w) - 1$, then

$$q(w) - \frac{1}{4} - \frac{1}{4}L_4(w) \geq q(w) - \frac{1}{4} - \frac{1}{4}(4q(w) - 1) = 0.$$

$\square$

Lemma 14.35 (Unpeeled visible-return routing). *Let w be a saturated Type A receiver with a peeling set $P_4(w)$. Suppose that a completion port of w carries four visible receiver-entry returns whose routed loads lie in $\mathcal{L}(w) \setminus P_4(w)$. Then the unpeeled data realize one of exits (1)–(7) of [Definition 14.16](#). If the realized exit is (4), the target-defective quotient has declared routed-load support containing at least one of those four unpeeled loads.*

Proof. Let $u_1, u_2, u_3, u_4 \in \mathcal{L}(w) \setminus P_4(w)$ be the four visible loads, and let the corresponding receiver-entry returns pass through the same completion port $\vec{e} = (w, h)$. Each return has the form $\Gamma_i \circ Q_i$, where Q_i contains the canonical trace T_{u_i} . The construction uses only these four returns and their declared response coordinates, so previously peeled loads are outside the data used below.

If one return has length in $\mathcal{M}$, [Lemma 2.2](#) gives exit (1). If two of the four anchored returns are internally vertex-disjoint and their lengths sum to a power of two, [Lemma 14.6](#) gives exit (2). If two traces pass through a common P_{13} -window and fail the legal C_s -relation of [Lemma 6.1](#), this is exit (3).

Assume exits (1)–(3) do not occur. For each i , form the declared target-response coordinate determined by the completion port, the first-entry receiver, the connector label, and the receiver-entry channel of the return for u_i . These four coordinates lie in one boundary-degree fibre. If a quotient identifying two of them is distinguished by a compatible outside context, then it is a target-defective quotient in the canonical family $\mathcal{Q}_4(w)$, of type (Q1) in [Definition 14.49](#). Its declared routed-load support contains the corresponding u_i 's, so exit (4) supports an unpeeled load.

If the identification is target-complete on a nontrivial response state, the branch realizes exit (5). If target-completeness requires adjoining a larger support, the branch realizes exit (6), proper or global according to whether the support is proper or all of G . Finally, if the four coordinates remain pairwise distinct in every target-complete quotient and are separated by the actual outside context, [Lemma 14.23](#) supplies either one of exits (4)–(6) or a surviving first separator. In the exit-(4) subcase, the quotient is of type (Q4) in [Definition 14.49](#), and its declared routed-load support is contained in $\{u_1, u_2, u_3, u_4\}$. In the separator subcase, [Lemmas 14.24](#) and [14.25](#) produce the decorated handoff fan envelope, which is exit (7). $\square$

Lemma 14.36 (Unpeeled silent-excess routing). *Let w be a saturated Type A receiver with peeling set $P_4(w)$. Suppose*

$$L_4(w) \geq 4q(w)$$

and no completion port of w carries four visible receiver-entry returns from unpeeled loads. Order $\mathcal{L}(w) \setminus P_4(w)$ by the canonical visible-first order, let $A_4(w)$ be the first $4q(w) - 1$ unpeeled loads, and put

$$E_4(w) = (\mathcal{L}(w) \setminus P_4(w)) \setminus A_4(w).$$

Then $E_4(w)$ is nonempty and realizes one of exits (4)–(8) of [Definition 14.16](#). If the realized exit is (4), the target-defective quotient has declared routed-load support containing a load of $E_4(w)$.

Proof. The inequality $L_4(w) \geq 4q(w)$ gives

$$|E_4(w)| = L_4(w) - (4q(w) - 1) \geq 1.$$

Since no completion port has four visible unpeeled returns and w has exactly $q(w)$ completion ports, the number of visible unpeeled loads is at most $3q(w)$. For $q(w) \in \{1, 2, 3\}$,

$$3q(w) \leq 4q(w) - 1.$$

The visible-first order therefore places every visible unpeeled load in $A_4(w)$. Hence every load in $E_4(w)$ is silent at every completion port of w .

Let $B_4(w)$ be the boundaried subgraph consisting of the loads in $E_4(w)$, their canonical traces to w , and the completion-port boundary data through which those traces can be completed. If the boundary response of $B_4(w)$ is target-defective in some compatible outside context, the corresponding quotient is of type (Q2) in [Definition 14.49](#); its declared routed-load support is $E_4(w)$. This is exit (4), and it supports an unpeeled residual excess load.

If the boundary response of $B_4(w)$ has a nontrivial target-complete compression, the branch realizes exit (5). If equality of the relevant responses becomes target-complete only after adjoining

a larger support, the branch realizes exit (6), proper or global according to that support. If the declared connector response data in $B_4(w)$ have a surviving first separator, then [Lemmas 14.24](#) and [14.25](#) produce a decorated handoff fan envelope, giving exit (7).

Assume none of exits (4)–(7) occurs. For each $u \in E_4(w)$, consider its trace basin B_u . The four failure alternatives in [Definition 14.14](#) are precisely: a trace-local target-defective quotient, a nontrivial target-complete response quotient, proper/global delocalization of an equality, or a surviving first separator. In the first case the quotient is canonical of type (Q3) and has declared routed-load support $\{u\}$, so it is exit (4). In the second case the branch realizes exit (5). In the third case it realizes exit (6). In the fourth case [Lemmas 14.24](#) and [14.25](#) produce the exit-(7) decorated handoff fan envelope. Since all four exits are absent, every $u \in E_4(w)$ has a target-complete-minimal trace basin. The loads in $E_4(w)$ are silent, no target-defective quotient occurs, no target-complete compression occurs, no delocalization occurs, and no high-degree fan decoration is created. This is exactly the silent-core residual profile of [Definition 14.15](#), namely exit (8). $\square$

Lemma 14.37 (Residual saturated routing after peeling). *Let w be a saturated Type A receiver with a peeling set $P_4(w)$. If*

$$L_4(w) \geq 4q(w),$$

then the unpeeled routed loads at w realize one of exits (1)–(8) of [Definition 14.16](#). If the realized exit is (4), the peeling set can be enlarged by one additional unpeeled routed load.

Proof. If a completion port contains four visible receiver-entry returns from unpeeled loads, [Lemma 14.35](#) gives one of exits (1)–(7). If the exit is (4), the same lemma gives an unpeeled supported load; add the lexicographically first such load to $P_4(w)$.

If no completion port contains four visible receiver-entry returns from unpeeled loads, [Lemma 14.36](#) gives one of exits (4)–(8). If the exit is (4), that lemma gives a supported load in the residual unpeeled excess set; add the lexicographically first such load to $P_4(w)$. $\square$

Lemma 14.38 (Exit-(4) peels one routed load). *Let w be a saturated Type A receiver with peeling set $P_4(w)$. If exit (4) occurs through a quotient $q \in \mathcal{Q}_4(w)$ whose declared support contains the canonical coordinate of an unpeeled routed load $u \in \mathcal{L}(w) \setminus P_4(w)$, then $P_4(w) \cup \{u\}$ is a valid peeling set and the remaining receiver deficit is reduced by exactly $1/4$.*

Proof. By definition of exit (4), q is boundary-degree-preserving and target-defective. Hence there are two realizations in the same boundary-degree fibre and a compatible outside context for which the target predicates differ. The quotient family $\mathcal{Q}_4(w)$ is canonical: its members are exactly the visible-coordinate, silent-basin, trace-basin, continuation/switch, and two-carrier deletion quotients listed in [Definition 14.49](#). The hypothesis says that the witness uses the declared coordinate of the unpeeled load u . Since $u \notin P_4(w)$, adjoining u preserves the no-duplicate condition in [Definition 14.33](#). Thus

$$P_4(w) \cup \{u\}$$

is a peeling set.

By [Lemma 14.34](#), replacing $P_4(w)$ by $P_4(w) \cup \{u\}$ changes $L_4(w)$ to $L_4(w) - 1$. The remaining receiver charge therefore increases by exactly $\frac{1}{4}$, equivalently the remaining receiver deficit decreases by exactly $\frac{1}{4}$. The bookkeeping records only the removal of the selected load u from the pure Type A receiver sum. $\square$

Lemma 14.39 (Closed saturated exits, exit-(4) peeling, and Type B handoff). *Let a saturated Type A receiver realize one of exits (1)–(7) of [Definition 14.16](#). If the realized exit is one of (1)–(3), (5), or (6), then that branch cannot remain as a negative-charge Type A discharging case. If the realized*




Figure 22: Visual proof of the Type A saturated-exit drain. Panel I shows a Type A receiver w with missing-port count $q(w) = 3 - d_X(w)$ and saturated routed load $L(w) \geq 4q(w)$. Panel II records the exits that close inside the Type A analysis: exits (1)–(3) create a target cycle or target-defective state, exit (5) gives target-complete compression, and exit (6) gives proper or global delocalization. Panel III isolates exit (4): one routed load using the canonical family $\mathcal{Q}_4(w)$ is peeled out of the pure Type A receiver sum, changing $L_4(w)$ to $L_4(w) - 1$ and reducing the remaining receiver deficit by $1/4$; the dashed loop indicates that the receiver is then tested again with smaller residual load. Panel IV shows the two forward exits: exit (7) is reclassified as decorated Type B handoff fan data, while exit (8) is the route-8 trace-basin residual treated by the later carrier-reduction and no-go figures.

exit is (4), one routed load is peeled in the sense of Lemma 14.38. If the realized exit is (7), the branch is reclassified as a decorated handoff fan envelope and leaves the Type A charge calculation.

Proof. Exit (1) gives an edge-rooted Mersenne return, hence a power-of-two cycle by Lemma 2.2. Exit (2) is a power-of-two common-port return cycle by Lemma 14.6. Exit (3) is precisely failure of the legal P_{13} label relation from Lemma 6.1; by definition of the relation, it creates a target event or a target-defective local state. Exit (4) has the one-load role proved in Lemma 14.38: it peels exactly one routed load whose declared coordinate is used by the canonical target-defective quotient, and it reduces the remaining receiver deficit by exactly $1/4$. Exit (5) is a nontrivial target-complete response compression. If the compression is realized by a smaller proper atom, it contradicts hereditary target-uncompressibility (Corollary 3.105); if it occurs only at the trace-basin response level, it is exactly failure alternative (b) in Definition 14.14 and therefore is not an admissible route-8 residual. Exit (6) is excluded by Lemma 10.6 in the proper-support case and by Lemma 10.9 in the whole-graph case. Exit (7) is not a Type A charge case: it is the decorated handoff fan envelope of Definition 14.19, with admissibility supplied by Lemma 14.26. Thus exits (1)–(3), (5), and (6) are closed inside the Type A analysis, exit (4) is an exact one-load peeling step, and exit (7) is transferred to the Type B branch without using Lemma 14.98. $\square$

Lemma 14.40 (Saturated receiver handoff). *Let X be a Type A support. If a receiver w satisfies $L(w) \geq 4q(w)$, then repeated application of the saturated receiver routing has one of the following outcomes: a closed exit among (1)–(3), (5), and (6) occurs; the Type B handoff exit (7) occurs; the route-8 residual profile of Definition 14.15 occurs; or exit (4) peels finitely many routed loads and the remaining receiver load becomes unsaturated. Consequently every receiver that remains in the pure Type A discharging branch after the closed exits, exit-(4) peeling, route 8, and the Type B*

handoff have been removed satisfies

$$L(w) \leq 4q(w) - 1.$$

Proof. Start with the empty peeling set $P_4(w) = \emptyset$. If the residual load $L_4(w)$ is at most $4q(w) - 1$, the receiver is already unsaturated in the remaining Type A charge calculation. Otherwise [Lemma 14.37](#) applies to the unpeeled loads. Exits (1)–(3), (5), and (6) close by [Lemma 14.39](#); exit (7) is transferred to the decorated handoff fan branch by the same lemma; and exit (8) is the route-8 residual profile. If exit (4) occurs, then [Lemma 14.38](#) enlarges $P_4(w)$ by one routed load, so $L_4(w)$ decreases by one. Since $\mathcal{L}(w)$ is finite, this exit-(4) peeling step cannot occur indefinitely. Therefore the iteration ends in one of the listed outcomes. A receiver that remains in the pure Type A discharging branch is precisely one for which the iteration ended by reaching $L_4(w) \leq 4q(w) - 1$; with all peeled loads removed from that branch, this is the displayed unsaturated bound. $\square$

Remark 14.41 (Use of exit-(4) peeling). *The peeling set is a routing record. A load in $P_4(w)$ exits through the target-defect branch before the pure Type A charge is summed. The charge calculation therefore applies only to the unpeeled loads counted by $L_4(w)$. Consequently [Lemma 14.45](#) is invoked only after the target-defect alternative has been removed; a support with a remaining exit-(4) witness is routed by alternative (iii) of [Lemma 14.44](#).*

Remark 14.42 (Type A–Type B stratification). *The saturated Type A analysis*

[Lemmas 14.28, 14.32, 14.39, 14.40 and 14.43](#)

uses no conclusion of [Lemma 14.98](#). The only Type B reference in that block is the handoff interface: [Lemmas 14.25](#) and [14.26](#) construct and validate the decorated envelope data. The logical direction then reverses: [Lemma 14.98](#) may apply the completed Type A saturated and unsaturated conclusions to the post-ledger core.

The possible alternation terminates. Type A handoffs are assigned to grouped decorated envelopes by the connected components of the core–center incidence graph in [Definition 14.20](#). Each such component uses its finite set of center tokens once in the grouped-envelope ledger. After the Type B ledger removes the corresponding fan/envelope entries, the remaining core is a proper subcore of the previous support or has fewer unused incidence component tokens. Since the support and center-token sets are finite, an $A \rightarrow B \rightarrow A$ alternation cannot continue indefinitely.

Lemma 14.43 (Unsaturated Type A discharging). *Let X be a Type A support in which every receiver satisfies $L(w) \leq 4q(w) - 1$. Then*

$$\partial_+(X) \geq \frac{1}{4}|V(X)|.$$

Equivalently, if $n_i = |\{v \in V(X) : d_X(v) = i\}|$, then

$$n_3 \leq 3n_2 + 7n_1 + 11n_0.$$

Proof. Set

$$\text{ch}(v) = (3 - d_X(v)) - \frac{1}{4}.$$

Then $\sum_v \text{ch}(v) = \partial_+(X) - |V(X)|/4$. A cubic vertex has charge $-1/4$, and every cubic vertex is assigned to exactly one receiver by [Lemma 14.4](#). A receiver w has initial charge $q(w) - 1/4$ and pays $1/4$ for each assigned cubic vertex. The unsaturated bound gives

$$q(w) - \frac{1}{4} - \frac{1}{4}L(w) \geq q(w) - \frac{1}{4} - \frac{1}{4}(4q(w) - 1) = 0.$$

All other non-cubic vertices have nonnegative charge and pay nothing. Hence the final total charge is nonnegative, so $\partial_+(X) - |V(X)|/4 \geq 0$. Written by degree classes, the receiver capacities are 3, 7, 11 for internal degrees 2, 1, 0, which is exactly $n_3 \leq 3n_2 + 7n_1 + 11n_0$. $\square$

Lemma 14.44 (Density forces discharging, target defect, Type B handoff, or route 8). *Let X be a connected admissible support (Definition 13.3) with $\sigma(X) = 0$. Thus X is contextually dyadic-safe (Definition 2.5), P_{13} -free, has empty internal 3-core, has every deficient vertex supplied from the packing W (Lemma 8.4), and is hereditarily target-uncompressible (Corollary 3.105). Then one of the following holds:*

(i)

$$\partial_+(X) \geq \frac{1}{4}|V(X)|, \quad \text{equivalently} \quad \frac{2|E(X)|}{|V(X)|} \leq 2.75.$$

(ii) X produces a decorated handoff fan envelope in the sense of Definition 14.19.

(iii) X has an exit-(4) witness for a routed load in the sense of Definition 14.33; this branch is routed to the target-defect exit and that routed load is removed from the Type A receiver calculation by Lemma 14.38.

(iv) X contains an admissible silent-core residual profile in the sense of Definition 14.15.

Proof. If a receiver is saturated, Lemma 14.40 applies. Exits (1)–(3), (5), and (6) contradict the standing target, replacement, or delocalization invariants. Exit (4) gives alternative (iii) and peels one routed load by Lemma 14.38. Exit (7) gives alternative (ii), and exit (8) gives alternative (iv). If no saturated receiver remains after excluding alternatives (ii)–(iv), every receiver in the pure Type A support is unsaturated. The charge calculation of Lemma 14.43 gives

$$\partial_+(X) \geq \frac{1}{4}|V(X)|.$$

Since $\sigma(X) = 0$, this is equivalent to $2|E(X)|/|V(X)| \leq 2.75$. $\square$

Lemma 14.45 (Type A split and charge bound). *Let X be a connected admissible support with $\sigma(X) = 0$, empty internal 3-core, P_{13} -free and dyadic-safe. If X contains no exit-(4) witness for a routed load and contains neither an admissible silent-core residual profile nor a decorated handoff fan envelope, then*

$$\partial_+(X) \geq \frac{1}{4}|V(X)|, \quad \text{equivalently} \quad \frac{2|E(X)|}{|V(X)|} \leq 2.75,$$

so that $\mathcal{N}(X) \geq 0$. Consequently a Type A support with $\mathcal{N}(X) < 0$ must carry an admissible route-8 residual profile, produce the Type B handoff, or contain an exit-(4) witness for a routed load.

Proof. The hypotheses are the admissibility hypotheses used in Lemma 14.44. Since the route-8 residual and Type B handoff alternatives and the exit-(4) witness alternative are excluded by hypothesis, that lemma gives $\partial_+(X) \geq \frac{1}{4}|V(X)|$. Since $\sigma(X) = 0$,

$$\mathcal{N}(X) = \partial_+(X) - \frac{1}{4}|V(X)| \geq 0.$$

Thus a Type A support outside the route-8 residual class, outside the Type B handoff, and outside the exit-(4) target-defect exit cannot have negative net charge. $\square$

Definition 14.46 (Large-budget Type A deficit). *Let $\mathcal{X}$ be a collection of Type A supports whose saturated receivers survive only through the route-8 residual profile of Definition 14.15. Put*

$$D_A(\mathcal{X}) = \sum_{X \in \mathcal{X}} \left(\frac{1}{4} |V(X)| - \partial_+(X) \right).$$

The collection $\mathcal{X}$ carries the large-budget Type A deficit if

$$D_A(\mathcal{X}) \geq \left(\frac{1}{4} - \tau_{\text{win}} \right) |R| - o(|R|).$$

Lemma 14.47 (Quantitative burden of the reduced route-8 residual). *Let $\mathcal{X}$ be a collection of Type A supports whose saturated receivers survive only through the route-8 residual profile of Definition 14.15; in particular, no receiver in any $X \in \mathcal{X}$ realizes exits (1)–(7). Let $D_A(\mathcal{X})$ be as in Definition 14.46. Let*

$$N_{\text{basin}}(\mathcal{X}) = \sum_{X \in \mathcal{X}} \sum_w |\mathcal{U}_X(w)|,$$

where w ranges over the saturated route-8 receivers of X and $\mathcal{U}_X(w)$ denotes the unpaid silent routed vertices at w . Thus N_{basin} counts the indexed trace-basin entries (u, B_u) ; two different vertices may contribute two entries even if their underlying boundaried subgraphs coincide. Then

$$N_{\text{basin}}(\mathcal{X}) \geq 4D_A(\mathcal{X}).$$

Consequently, if $\mathcal{X}$ carries the large-budget Type A deficit, then it contains at least

$$4 \left(\frac{1}{4} - \tau_{\text{win}} \right) |R| - o(|R|)$$

indexed trace-basin entries. In the limiting window-only cap this coefficient is

$$4 \left(\frac{1}{4} - 0.22817486846 \dots \right) = 0.08730052616 \dots$$

Proof. For each $X \in \mathcal{X}$, the absence of exits (1)–(7) implies in particular that no saturated receiver has a completion port carrying four visible receiver-entry returns. Hence Lemma 14.30 applies to X and gives

$$S_{\text{sil}}^{\text{exc}}(X) \geq 4 \left(\frac{1}{4} |V(X)| - \partial_+(X) \right).$$

By Lemma 14.31, every unpaid silent excess vertex counted by $S_{\text{sil}}^{\text{exc}}(X)$ has a target-complete-minimal trace basin unless one of exits (4)–(7) occurs. Those exits are absent by hypothesis. Counting the resulting indexed pairs (u, B_u) and summing the displayed inequality over $X \in \mathcal{X}$ proves $N_{\text{basin}}(\mathcal{X}) \geq 4D_A(\mathcal{X})$. No distinctness of the underlying subgraphs B_u is used.

If $\mathcal{X}$ carries the large-budget Type A deficit, then

$$D_A(\mathcal{X}) \geq \left(\frac{1}{4} - \tau_{\text{win}} \right) |R| - o(|R|)$$

by Definition 14.46. Multiplying by the factor 4 from the first part gives the stated lower bound. Substituting $\tau_{\text{win}} = 0.22817486846 \dots$ gives the displayed coefficient. $\square$

Definition 14.48 (Route-8 response carriers and two-carrier entries). *Let $\mathcal{X}$ be a collection of Type A supports whose saturated receivers survive only through route 8. For a Type A support X , put*

$$\partial_E X = \{(x, xy) : x \in V(X), y \notin V(X), xy \in E(G)\},$$

the set of oriented boundary incidences leaving X . Since $\sigma(X) = 0$, every vertex of X has ambient degree 3, and hence

$$|\partial_E X| = \partial_+(X).$$

For an indexed entry $\xi = (X, w, u, B_u)$, every declared u -supported coordinate of $\rho_u(B_u)$ that uses the ambient exterior of X records the oriented boundary incidence of X through which it leaves X . More precisely, the declared support of such a coordinate contains a completion-port incidence, first-entry incidence, connector-band endpoint, boundary-degree entry, or packed-window interface incidence; when that incidence is represented by an edge $xy \in E(G)$ with $x \in V(X)$ and $y \notin V(X)$, its ambient carrier is the element $c = (x, xy) \in \partial_E X$. Incidences of ∂B_u internal to X have no ambient carrier unless the declared coordinate canonically assigns them to the terminal completion port of T_u , in which case the carrier is that completion-port incidence. Thus ambient carriers are not extra data attached to B_u ; they are the $\partial_E X$ -labels already recorded by the declared coordinate signature.

An indexed route-8 entry is a tuple

$$\xi = (X, w, u, B_u),$$

where $u \in \mathcal{U}_X(w)$ and B_u is the target-complete-minimal trace basin assigned to u . The set of all indexed route-8 entries of $\mathcal{X}$ is denoted

$$\Xi(\mathcal{X}).$$

More generally, a tagged Type A entry family is any finite family $\mathfrak{E}$ of tuples $\xi = (X, w, u, B_u)$ with $\sigma(X) = 0$, $u \in \mathcal{U}_X(w)$, and B_u the trace basin assigned to u . Target-complete-minimality is required only when ξ is called a route-8 entry. The carrier constructions below apply to every $\xi \in \mathfrak{E}$. For a pure route-8 collection the tagged family is $\mathfrak{E} = \Xi(\mathcal{X})$. A boundary incidence $c \in \partial_E X$ is active for ξ if the declared support of a u -supported coordinate in $\rho_u(B_u)$ contains c as a completion-port incidence, first-entry incidence, connector-band endpoint, boundary-degree entry, or packed-window interface incidence under [Definition 3.94](#).

For a coordinate $r \in \mathcal{R}_u(B_u)$, let

$$\text{car}_\xi(r) \subseteq \partial_E X$$

be the set of active carriers used by the declared support of r . If $D \subseteq \partial_E X$, the D -restriction $\rho_u(B_u)|_D$ is the response quotient obtained from $\rho_u(B_u)$ by retaining the full boundary degree profile $\mathbf{d}_\partial(B_u)$ and retaining, apart from this boundary-degree profile, exactly those u -supported coordinates r with $\text{car}_\xi(r) \subseteq D$; all other u -supported coordinates are forgotten. Thus every carrier restriction is taken inside the original boundary-degree fibre.

A set $D \subseteq \partial_E X$ is a target-complete carrier set for ξ if the quotient

$$\rho_u(B_u) \longrightarrow \rho_u(B_u)|_D$$

is target-complete in the sense of [Definition 14.14](#): for every outside context compatible with the boundary profile, all realizations of $\rho_u(B_u)|_D$ give the same target predicate as $\rho_u(B_u)$ after gluing. Such sets exist: for $D = \partial_E X$, all declared u -supported coordinates are retained, so the completeness clause in [Definition 14.14](#) applies. Let

$$\mathcal{C}_{\text{ess}}(\xi)$$

be the lexicographically first inclusion-minimal target-complete carrier set for ξ . Its elements are the essential carriers of ξ , and the carriers in $\partial_E X \setminus \mathcal{C}_{\text{ess}}(\xi)$ are noncore carriers. Thus

$$\rho_u(B_u) \longrightarrow \rho_u(B_u)|_{\mathcal{C}_{\text{ess}}(\xi)}$$

is target-complete by definition, while for every $c \in \mathcal{C}_{\text{ess}}(\xi)$ the smaller restriction $\rho_u(B_u)|_{\mathcal{C}_{\text{ess}}(\xi) \setminus \{c\}}$ is not target-complete. Put

$$\alpha_{\mathfrak{E}}(\xi) = |\mathcal{C}_{\text{ess}}(\xi)|.$$

The value is independent of $\mathfrak{E}$; the subscript records the current branch state in which privacy is evaluated.

An essential carrier for ξ is private for ξ relative to $\mathfrak{E}$ if it is not essential for any other tagged entry in $\mathfrak{E}$. Let

$$\pi_{\mathfrak{E}}(\xi)$$

be the number of private essential carriers of ξ relative to $\mathfrak{E}$. When $\mathfrak{E} = \Xi(\mathcal{X})$, abbreviate these quantities by $\alpha_{\mathcal{X}}(\xi)$ and $\pi_{\mathcal{X}}(\xi)$.

The tagged entry ξ is two-carrier relative to $\mathfrak{E}$ if

$$\pi_{\mathfrak{E}}(\xi) \leq 2.$$

When $\mathfrak{E} = \Xi(\mathcal{X})$ and ξ is a route-8 entry, this is called a two-carrier route-8 entry. An admissible two-carrier route-8 residual is a large-budget route-8 collection containing such an entry.

Definition 14.49 (Canonical exit-(4) quotient family). Let w be a saturated receiver in a Type A support X . Write $\mathcal{Q}_4^{\text{pre}}(w)$ for the boundary-degree-preserving response quotients generated by clauses (Q1)–(Q4) below; these clauses do not use carrier or privacy data. Every bare occurrence of $\mathcal{Q}_4(w)$ before a tagged family is fixed denotes this preliminary family.

For the collection-dependent clause (Q5), freeze the entire current tagged Type A entry family $\mathfrak{E}$; in particular, $\mathfrak{E}$ contains every current entry at w . Then $\mathcal{Q}_4(w; \mathfrak{E})$ is the union of $\mathcal{Q}_4^{\text{pre}}(w)$ and the quotients generated by clause (Q5). For a pure route-8 collection $\mathcal{X}$, the notation $\mathcal{Q}_4(w; \mathcal{X})$ means $\mathcal{Q}_4(w; \Xi(\mathcal{X}))$.

(Q1) the visible receiver-entry coordinate identifications tested in [Lemma 14.28](#);

(Q2) the silent/excess boundary-response quotient of the basin $B(w)$ in [Lemma 14.32](#);

(Q3) the trace-local non-carrier quotient or identification of a trace basin in the sense of [Definition 14.14](#), as used in [Lemma 14.31](#);

(Q4) the finite continuation or cubic-switch quotient produced by [Lemmas 14.23](#) and [14.24](#);

(Q5) for an indexed route-8 entry $\xi = (X, w, u, B_u) \in \mathfrak{E}$, the carrier-deletion quotient

$$\rho_u(B_u)|_{\mathcal{C}_{\text{ess}}(\xi)} \longrightarrow \rho_u(B_u)|_{\mathcal{C}_{\text{ess}}(\xi) \setminus \{c\}}$$

when $c \in \mathcal{C}_{\text{ess}}(\xi)$, the entry is two-carrier relative to the frozen tagged family, $\pi_{\mathfrak{E}}(\xi) \leq 2$, and the distinguishing event is recorded by a declared deletion witness in the sense of [Definition 14.51](#).

Every member of $\mathcal{Q}_4^{\text{pre}}(w)$, and every member of $\mathcal{Q}_4(w; \mathfrak{E})$, retains the boundary degree profile of the response state on which it is defined. At the preliminary stage, exit (4) occurs precisely when a quotient in $\mathcal{Q}_4^{\text{pre}}(w)$ is target-defective. After $\mathfrak{E}$ is frozen, the Q5 refinement tests the additional members of $\mathcal{Q}_4(w; \mathfrak{E})$. The declared routed-load support is part of the canonical construction: in (Q1) it is the set of visible routed loads whose response coordinates are identified; in (Q2) it is the routed loads in the basin $B(w)$; in (Q3) it is the trace load u whose basin is being tested; in (Q4) it is the routed loads whose declared connector coordinates form the finite family $\mathcal{K}$; and in (Q5) it is the indexed load u of the route-8 entry $\xi = (X, w, u, B_u)$.

Definition 14.50 (True route-8 residual entries). *Let $\mathfrak{E}$ be a frozen tagged Type A entry family and let $\xi = (X, w, u, B_u) \in \mathfrak{E}$ be an indexed route-8 entry. The entry is a preliminary route-8 entry when w is saturated, its data form an admissible silent-core residual profile, B_u is target-complete-minimal, exits (1)–(3) and (5)–(7) are absent, and exit (4) is absent for clauses (Q1)–(Q4) of Definition 14.49.*

After $\mathfrak{E}$, its essential carrier sets, and the privacy counts $\pi_{\mathfrak{E}}$ have been fixed, perform the Q5 refinement:

(Q5a) if a clause-(Q5) quotient in $\mathcal{Q}_4(w; \mathfrak{E})$ is target-defective, the entry leaves through exit (4), with the indexed load u as its declared routed-load support;

(Q5b) if no clause-(Q5) quotient is target-defective, record Q5 absence.

The test changes only the branch tag. It does not alter G , the support, the trace basin, the carrier sets, the privacy counts, the receiver burden, or any previously recorded absent exit.

A preliminary route-8 entry on the Q5-absent branch is a true route-8 residual entry. Equivalently, it satisfies:

(R1) w is a saturated receiver of X ;

(R2) exits (1)–(7) of Definition 14.16 are absent at w , with exit (4) interpreted through the refined canonical family $\mathcal{Q}_4(w; \mathfrak{E})$;

(R3) the data at w form an admissible silent-core residual profile of Definition 14.15; and

(R4) B_u is the target-complete-minimal trace basin assigned to u .

For a route-8 Type A collection $\mathcal{X}$, take $\mathfrak{E} = \Xi(\mathcal{X})$. It is a true route-8 residual collection when every indexed entry lies on the Q5-absent branch. Thus, in a true route-8 residual entry, any occurrence of one of exits (1)–(7) contradicts the definition of the residual.

Definition 14.51 (Carrier-deletion witnesses). *Let*

$$\xi = (X, w, u, B_u)$$

be an indexed route-8 entry, and put

$$\mathcal{C} = \mathcal{C}_{\text{ess}}(\xi).$$

For $c \in \mathcal{C}$, the c -deletion quotient is the quotient

$$q_{\xi, c} : \rho_u(B_u)|_{\mathcal{C}} \longrightarrow \rho_u(B_u)|_{\mathcal{C} \setminus \{c\}}.$$

A c -deletion witness is a realization S of $\rho_u(B_u)|_{\mathcal{C} \setminus \{c\}}$, together with a compatible outside context Y , such that

$$S \oplus Y$$




Figure 23: Visual definition of a route-8 trace-basin entry and its carrier core. The grey rectangle is the Type A support X , and the blue dashed shaded region is the target-complete-minimal trace basin $B_u \subseteq X$ assigned to the routed vertex u . Black edges are support edges inside X , the blue arrowed path T_u is the canonical trace from u to the receiver w , and ∂B_u denotes the boundary of the basin inside X . The entry is $\xi = (X, w, u, B_u)$. The labels $c_1, \dots, c_5$ denote oriented boundary incidences in $\partial_E X$. Orange arrows mark active carrier incidences used by declared u -supported coordinates; dashed black incidences illustrate boundary incidences that remain outside the essential core. Red carrier nodes are elements of the essential carrier set $\mathcal{C}_{\text{ess}}(\xi)$, grey carrier nodes are noncore carriers, and the red crosses on c_1, c_2 mark carriers counted as private for this indexed entry inside the collection $\mathcal{X}$. The response state $\rho_u(B_u)$ records the boundary degree profile $\mathbf{d}_\partial$, the u -supported response coordinates $\mathcal{R}_u$, and the boundary target profile $\mathbf{H}_\partial$. A D -restriction keeps only the coordinates whose declared carriers lie in the chosen set D ; the displayed restriction uses $D = \mathcal{C}_{\text{ess}}(\xi)$. The number $\pi_{\mathcal{X}}(\xi)$ counts the private essential carriers of ξ , and a two-carrier route-8 entry is one with $\pi_{\mathcal{X}}(\xi) \leq 2$.

has a different target predicate from the original core state $\rho_u(B_u)|_{\mathcal{C}}$ glued to the corresponding context. A witnessing target event is chosen as the lexicographically first simple edge-rooted return or simple cycle responsible for this target difference. The declared support of the witness is the finite set of declared response coordinates from Definition 3.94 used by that event, and its carrier support is the union of the ambient carriers of those coordinates in $\partial_E X$. The witness is internal when this event uses an edge of $B_u - \partial B_u$. It is mixed when, in addition, the event uses an edge outside X .

Lemma 14.52 (Essential carriers have deletion witnesses). *Let $\xi = (X, w, u, B_u)$ be an indexed route-8 entry. For every*

$$c \in \mathcal{C}_{\text{ess}}(\xi),$$

the c -deletion quotient $q_{\xi,c}$ is not target-complete. Equivalently, there exists a c -deletion witness in the sense of Definition 14.51.

Proof. By definition, $\mathcal{C}_{\text{ess}}(\xi)$ is inclusion-minimal among target-complete carrier sets for ξ . Therefore

$$\mathcal{C}_{\text{ess}}(\xi) \setminus \{c\}$$

is not a target-complete carrier set. The full-to-core quotient

$$\rho_u(B_u) \longrightarrow \rho_u(B_u)|_{\mathcal{C}_{\text{ess}}(\xi)}$$

is target-complete. If $q_{\xi,c}$ were target-complete, then [Lemma 3.93](#) would make the composite

$$\rho_u(B_u) \longrightarrow \rho_u(B_u)|_{\mathcal{C}_{\text{ess}}(\xi) \setminus \{c\}}$$

target-complete, contradicting the preceding sentence. Hence $q_{\xi,c}$ is not target-complete. Unwinding the definition of target-completeness gives a compatible outside context and a realization of the smaller restricted response state with different target predicate. Choosing the lexicographically first simple edge-rooted return or simple cycle that accounts for the difference gives the stated deletion witness. $\square$

Lemma 14.53 (Deletion witnesses are declared). *Let $\xi = (X, w, u, B_u)$ be an indexed route-8 entry, and let $c \in \mathcal{C}_{\text{ess}}(\xi)$. Every c -deletion witness obtained from [Lemma 14.52](#) has declared support and carrier support in the sense of [Definition 14.51](#). Moreover its carrier support contains c .*

Proof. The witness distinguishes the core state $\rho_u(B_u)|_{\mathcal{C}_{\text{ess}}(\xi)}$ from the smaller restriction $\rho_u(B_u)|_{\mathcal{C}_{\text{ess}}(\xi) \setminus \{c\}}$ in a compatible outside context. If the witnessing target event used only retained coordinates, then its truth value would be the same in the two glued states, because the two states agree on every retained declared u -supported coordinate and on the full boundary degree profile. Hence the event uses at least one coordinate forgotten by the deletion quotient.

By the definition of D -restriction in [Definition 14.48](#), the coordinates forgotten when passing from $\mathcal{C}_{\text{ess}}(\xi)$ to $\mathcal{C}_{\text{ess}}(\xi) \setminus \{c\}$ are precisely declared u -supported coordinates whose ambient carrier support contains c . These coordinates are generated by the declared response-coordinate signature [Definition 3.94](#). Taking the finite set of declared coordinates used by the lexicographically first witnessing event gives the declared support, and the union of their ambient carriers gives the carrier support. This support contains c . $\square$

Lemma 14.54 (Two-carrier deletion quotients are canonical exit-(4) quotients). *Let $\mathfrak{E}$ be a frozen tagged Type A entry family, let $\xi = (X, w, u, B_u) \in \mathfrak{E}$ be an indexed route-8 entry with $\pi_{\mathfrak{E}}(\xi) \leq 2$, and let $c \in \mathcal{C}_{\text{ess}}(\xi)$. If c has a declared c -deletion witness, then $q_{\xi,c} \in \mathcal{Q}_4(w; \mathfrak{E})$.*

Proof. The quotient $q_{\xi,c}$ retains the full boundary degree profile by the definition of D -restriction in [Definition 14.48](#). The hypothesis $\pi_{\mathfrak{E}}(\xi) \leq 2$ says that ξ is two-carrier relative to the frozen tagged family. The declared witness hypothesis supplies the final datum required in clause (Q5) of [Definition 14.49](#). Therefore $q_{\xi,c}$ is one of the canonical exit-(4) quotients at w relative to $\mathfrak{E}$. $\square$

Lemma 14.55 (Carrier deletion is exit (4)). *Let $\mathfrak{E}$ be a frozen tagged Type A entry family, let $\xi = (X, w, u, B_u) \in \mathfrak{E}$ be an indexed route-8 entry with $\pi_{\mathfrak{E}}(\xi) \leq 2$, and put*

$$\mathcal{C} = \mathcal{C}_{\text{ess}}(\xi).$$

If $c \in \mathcal{C}$ has a declared c -deletion witness, then exit (4) of [Definition 14.16](#) occurs at the saturated receiver w .

Proof. The carrier-core restriction

$$\rho_u(B_u) \longrightarrow \rho_u(B_u)|_{\mathcal{C}}$$

is target-complete by the definition of $\mathcal{C}_{\text{ess}}(\xi)$. Thus, for every compatible outside context, replacing $\rho_u(B_u)$ by the core state $\rho_u(B_u)|_{\mathcal{C}}$ preserves the target predicate.

Let S and Y be the c -deletion witness. Thus S is a realization of

$$\rho_u(B_u)|_{C \setminus \{c\}},$$

the context Y is compatible with the same boundary degree profile, and

$$S \oplus Y$$

has a target predicate different from the core state $\rho_u(B_u)|_C$ glued to the corresponding context. Since the full-to-core quotient is target-complete, [Lemma 3.93](#) implies that the same realization S and context Y distinguish the composite quotient

$$\rho_u(B_u) \longrightarrow \rho_u(B_u)|_{C \setminus \{c\}}.$$

This composite quotient preserves the boundary degree profile and is a response quotient inside $\rho_u(B_u)$: it is obtained by forgetting exactly the declared u -supported coordinates whose carrier support uses c , after the noncore carriers have already been forgotten by the target-complete core restriction. Hence it is a boundary-degree-preserving response quotient distinguished by an outside context.

By [Lemma 14.54](#), $q_{\xi,c}$ belongs to $\mathcal{Q}_4(w; \mathfrak{E})$. The preceding paragraph proves that this canonical quotient is target-defective. This is precisely exit (4) of [Definition 14.16](#). $\square$

Lemma 14.56 (Carrier cut parity). *Let $\xi = (X, w, u, B_u)$ be an indexed route-8 entry, and replace its response state by the canonical core restriction $\rho_u(B_u)|_{C_{\text{ess}}(\xi)}$. If a surviving u -supported target event is realized by a simple edge-rooted return or a simple cycle which uses an internal edge of B_u and also uses an edge outside X , then its declared support contains at least two distinct carriers from $C_{\text{ess}}(\xi)$.*

Proof. For an edge-rooted return, add the root edge to the return path; for a cycle do nothing. In both cases we obtain a simple cycle C of G . Since the target event uses an internal edge of $B_u \subseteq X$ and also an edge outside X , the set

$$E(C) \cap \delta_G(V(X)), \quad \delta_G(V(X)) = \{ab \in E(G) : |\{a, b\} \cap V(X)| = 1\},$$

is nonempty. A cycle crosses every cut an even number of times. Because C is simple, it cannot use the same cut edge twice. Hence C uses at least two distinct boundary edges of X .

Each such crossing is recorded in the declared support of the corresponding u -supported coordinate: it is a completion-port incidence, a first-entry incidence, a connector endpoint, a boundary-degree entry, or a packed-window interface incidence in the signature of [Definition 3.94](#). Since the event survives in the restricted state $\rho_u(B_u)|_{C_{\text{ess}}(\xi)}$, every carrier in its declared support lies in $C_{\text{ess}}(\xi)$. The two distinct cut crossings therefore give two distinct carriers from $C_{\text{ess}}(\xi)$. $\square$

Lemma 14.57 (Internal trace quotient witnesses are mixed). *Let $\xi = (X, w, u, B_u)$ be an indexed route-8 entry, let $\mathcal{C} \subseteq \partial_E X$, and let $\rho_{\mathcal{C}}^\circ$ be the quotient of $\rho_u(B_u)|_C$ that retains the full boundary degree profile and forgets exactly the u -supported coordinates whose declared support contains an internal edge of B_u . If*

$$\rho_u(B_u)|_C \longrightarrow \rho_{\mathcal{C}}^\circ$$

is not target-complete, then some distinguishing realization contains a surviving u -supported target event which is a simple edge-rooted return or a simple cycle, whose declared support contains an internal edge of B_u , and which uses an edge outside X .

Proof. Failure of target-completeness gives a compatible outside context and two realizations of $\rho_u(B_u)|_{\mathcal{C}}$ with the same image in $\rho_{\mathcal{C}}^\circ$, but with different target predicates after gluing. Choose a target event present in one glued realization and absent in the other. If the event is an edge-rooted return, add its root edge; otherwise use the target cycle itself. The event is represented by a simple cycle.

The two realizations have the same image in $\rho_{\mathcal{C}}^\circ$, so a distinguishing event must use at least one coordinate forgotten by $\rho_u(B_u)|_{\mathcal{C}} \rightarrow \rho_{\mathcal{C}}^\circ$. The forgotten coordinates are precisely the u -supported coordinates whose declared support contains an internal edge of B_u . Hence the chosen event is u -supported and uses such an internal edge.

If the event were contained entirely in X , it would be a dyadic cycle in the contextually dyadic-safe support X , impossible in the counterexample. Thus the event also uses an edge outside X . Since the event is present in a realization of the restricted state, it survives in $\rho_u(B_u)|_{\mathcal{C}}$. $\square$

Lemma 14.58 (Zero- and one-carrier silent entries collapse). *Let $\mathfrak{E}$ be a tagged Type A entry family and let $\xi = (X, w, u, B_u) \in \mathfrak{E}$ be an indexed route-8 entry. If*

$$\alpha_{\mathfrak{E}}(\xi) \leq 1,$$

then one of exits (4)–(7) of Definition 14.16 occurs. Consequently, every indexed entry in a true route-8 residual satisfies

$$\alpha_{\mathfrak{E}}(\xi) \geq 2.$$

Proof. Let

$$\mathcal{C} = \mathcal{C}_{\text{ess}}(\xi).$$

Replacing $\rho_u(B_u)$ by the restriction $\rho_u(B_u)|_{\mathcal{C}}$ deletes exactly the noncore carriers. This quotient is target-complete by the definition of the canonical carrier core.

Assume $|\mathcal{C}| \leq 1$. Let $\rho_{\mathcal{C}}^\circ$ be the further quotient of $\rho_u(B_u)|_{\mathcal{C}}$ that retains the full boundary degree profile and forgets every u -supported coordinate whose declared support contains an internal edge of B_u . This quotient is nontrivial because Definition 14.14 includes the trace-incidence coordinate of $T_u \subseteq B_u$, and the trace contains an internal edge of B_u .

We claim that $\rho_u(B_u)|_{\mathcal{C}} \rightarrow \rho_{\mathcal{C}}^\circ$ is target-complete. Suppose not. By Lemma 14.57, there is a surviving u -supported target event in the core-restricted response state which uses an internal edge of B_u and an edge outside X . Hence Lemma 14.56 applies and forces its declared support to contain at least two distinct carriers from $\mathcal{C}$, contradicting $|\mathcal{C}| \leq 1$. Thus the quotient is target-complete.

The nontrivial target-complete quotient above is precisely failure alternative (b) in the definition of target-complete-minimality (Definition 14.14), unless one of the other alternatives occurs first: a distinguishing outside context is alternative (a), validity only after adjoining a larger support is alternative (c), and first separation at a high-degree outside vertex is alternative (d). These are exits (4), (6), and (7), respectively, while alternative (b) is exit (5). Therefore an indexed route-8 entry with at most one essential carrier cannot survive while exits (4)–(7) are absent. In a true route-8 residual those exits are absent by definition, so every surviving indexed entry has at least two essential carriers. $\square$

Proposition 14.59 (Carrier reduction for route 8). *Let $\mathcal{X}$ be a collection of route-8 Type A supports that carries the large-budget Type A deficit in the sense of Definition 14.46. Then $\mathcal{X}$ contains a two-carrier route-8 entry. Equivalently, if no two-carrier route-8 entry exists, no route-8 Type A collection carries the large-budget Type A deficit.*

Proof. Suppose that no indexed route-8 entry in $\mathcal{X}$ is two-carrier. By [Definition 14.48](#), $\Xi(\mathcal{X})$ is the set of indexed trace-basin entries. Then $|\Xi(\mathcal{X})| = N_{\text{basin}}(\mathcal{X})$, and every $\xi \in \Xi(\mathcal{X})$ has at least three private essential carriers:

$$\pi_{\mathcal{X}}(\xi) \geq 3.$$

Private essential carriers belonging to different indexed entries are disjoint by definition. They are boundary incidences of the corresponding Type A supports, so

$$3N_{\text{basin}}(\mathcal{X}) \leq \sum_{X \in \mathcal{X}} |\partial_E X| = \sum_{X \in \mathcal{X}} \partial_+(X).$$

The collection $\mathcal{X}$ is a subcollection of the canonical support decomposition. Since the canonical supports are the connected components of R , their positive deficiencies add exactly, and all summands are nonnegative. Hence

$$\sum_{X \in \mathcal{X}} \partial_+(X) \leq \partial_+(R).$$

By [Corollary 8.7](#) and the near-cubic spine estimate $\sigma(G) = o(|R|)$ from [Definition 4.1](#),

$$\partial_+(R) \leq \tau_{\text{win}}|R| + o(|R|).$$

Hence

$$3N_{\text{basin}}(\mathcal{X}) \leq \tau_{\text{win}}|R| + o(|R|).$$

On the other hand, [Lemma 14.47](#) and [Definition 14.46](#) give

$$N_{\text{basin}}(\mathcal{X}) \geq 4 \left(\frac{1}{4} - \tau_{\text{win}} \right) |R| - o(|R|).$$

Combining the two inequalities yields

$$\tau_{\text{win}}|R| + o(|R|) \geq 12 \left(\frac{1}{4} - \tau_{\text{win}} \right) |R| - o(|R|).$$

This is impossible, because

$$12 \left(\frac{1}{4} - \tau_{\text{win}} \right) = 0.26190157848 \dots > 0.22817486846 \dots = \tau_{\text{win}}.$$

Equivalently, the numerical condition is

$$\tau_{\text{win}} < \frac{3}{13}.$$

Thus a surviving large-budget route-8 collection must contain a two-carrier route-8 entry. $\square$

Definition 14.60 (Terminal two-carrier route-8 obstruction). *A terminal two-carrier route-8 obstruction is a pair*

$$(\mathfrak{E}, \xi), \quad \xi = (X, w, u, B_u) \in \mathfrak{E},$$

where $\mathfrak{E}$ is the frozen tagged entry family of the current branch, with the following entry-level properties.

(T2) *The entry ξ is a true route-8 residual entry in the sense of [Definition 14.50](#): exits (1)–(7) of [Definition 14.16](#) are absent at its saturated receiver, and B_u is target-complete-minimal.*




Figure 24: Visual proof of the route-8 carrier-reduction squeeze. Panel I records the lower side: a route-8 Type A collection $\mathcal{X}$ carrying the large-budget deficit has $N_{\text{basin}}(\mathcal{X}) \geq 4D_A(\mathcal{X})$, and D_A is linear in $(1/4 - \tau_{\text{win}})|R|$. Panel II shows the negation of the conclusion: if no indexed entry is two-carrier, then every entry has at least three private essential carriers. Panel III records the upper side: private essential carriers are disjoint boundary incidences of the canonical supports, so $3N_{\text{basin}}(\mathcal{X}) \leq \partial_+(R) \leq \tau_{\text{win}}|R| + o(|R|)$. Panel IV combines the lower and upper bounds; since $12(1/4 - \tau_{\text{win}}) > \tau_{\text{win}}$, equivalently $\tau_{\text{win}} < 3/13$, the no-two-carrier alternative is impossible.

(T3) *The canonical carrier core is nontrivial:*

$$\alpha_{\mathfrak{E}}(\xi) = |\mathcal{C}_{\text{ess}}(\xi)| \geq 2.$$

(T4) *The entry is two-carrier in the private-carrier sense:*

$$\pi_{\mathfrak{E}}(\xi) \leq 2.$$

(T5) *For every $c \in \mathcal{C}_{\text{ess}}(\xi)$, the c -deletion witness of Lemma 14.52 is part of the obstruction data, with its witnessing target event, declared support, and carrier support recorded as in Lemma 14.53.*

The definition allows essential carriers of ξ that are also essential for other tagged entries. The large-budget argument is used only to produce such an entry; it is not a field of the terminal package and is not used by the no-go theorem. The package is not restricted to the special case $|\mathcal{C}_{\text{ess}}(\xi)| = 2$.

Theorem 14.61 (Two-carrier route-8 no-go). *There is no terminal two-carrier route-8 obstruction in the sense of Definition 14.60.*

Proof. Suppose that $(\mathfrak{E}, \xi)$ is a terminal two-carrier route-8 obstruction, and write

$$\xi = (X, w, u, B_u), \quad \mathcal{C} = \mathcal{C}_{\text{ess}}(\xi).$$

By condition (T3) in Definition 14.60, the set $\mathcal{C}$ has at least two elements; choose $c \in \mathcal{C}$. Condition (T4) says that ξ is two-carrier, and condition (T5) records the declared c -deletion witness supplied by Lemmas 14.52 and 14.53. By Lemma 14.54, the corresponding deletion quotient lies in $\mathcal{Q}_4(w; \mathfrak{E})$. Therefore Lemma 14.55 gives exit (4) at the saturated receiver w .

Condition (T2) says that ξ is a true route-8 residual entry, so exits (1)–(7) are absent at w . This contradicts the occurrence of exit (4). Hence no terminal two-carrier route-8 obstruction exists. $\square$




Figure 25: Visual proof of the terminal two-carrier route-8 no-go. Panel I shows a terminal entry $\xi = (X, w, u, B_u)$ in the frozen tagged family $\mathfrak{E}$. The grey rectangle is the Type A support X , the blue dashed region is the target-complete-minimal basin B_u , the orange vertex is the saturated receiver w , and the red boundary nodes are essential carrier incidences. The displayed terminal data are $\alpha_{\mathfrak{E}}(\xi) \geq 2$, $\pi_{\mathfrak{E}}(\xi) \leq 2$, and true route-8 residuality. Panel II isolates one essential carrier c . Its declared deletion witness has declared support and carrier support, and the two-carrier deletion lemma places the corresponding quotient in the canonical exit-(4) family $\mathcal{Q}_4(w; \mathfrak{E})$. Panel III records the contradiction: the carrier-deletion exit lemma realizes exit (4) at w , while true route-8 residuality says that exits (1)–(7), including exit (4), are absent. Thus the terminal two-carrier obstruction of Definition 14.60 cannot occur.

Proposition 14.62 (Route-8 closure). *No route-8 Type A collection carries the large-budget Type A deficit.*

Proof. Let $\mathcal{X}$ be a route-8 Type A collection carrying the large-budget Type A deficit, and freeze $\mathfrak{E} = \Xi(\mathcal{X})$. By Proposition 14.59, $\mathcal{X}$ contains an indexed two-carrier entry

$$\xi = (X, w, u, B_u)$$

with $\pi_{\mathfrak{E}}(\xi) \leq 2$. Since the collection survives only through route 8, ξ is a true route-8 residual entry in the sense of Definition 14.50; in particular, exits (1)–(7) are absent. By Lemma 14.58, every indexed entry in a true route-8 residual has at least two essential carriers, so

$$\alpha_{\mathfrak{E}}(\xi) \geq 2.$$

For each essential carrier, Lemmas 14.52 and 14.53 supplies the required declared deletion witness. Therefore $(\mathfrak{E}, \xi)$ is a terminal two-carrier route-8 obstruction, contradicting Theorem 14.61. Hence no such collection $\mathcal{X}$ exists. $\square$

Remark 14.63 (Type A route-8 closure). *After Proposition 14.59, the only Type A route-8 configuration not eliminated by the private-carrier count is a target-complete-minimal silent trace basin with at least two essential response carriers and at most two private essential carriers. The zero- and one-essential-carrier cases are removed by Lemma 14.58. The remaining two-carrier entry is excluded by Lemmas 14.54 and 14.55 and Theorem 14.61: an essential carrier has a declared deletion witness, and, for the frozen tagged family $\mathfrak{E}$, the two-carrier condition places the corresponding deletion quotient in the canonical family $\mathcal{Q}_4(w; \mathfrak{E})$, and exit (4) is absent in a true route-8 residual.*

Remark 14.64 (Admissibility is essential; the bound is not intrinsic). *Subcubicity and an empty 3-core alone do not yield 2.75: a 2-degenerate subcubic graph can have average degree approaching 3. The point of Lemma 14.44 is therefore not a bare extremal statement about isolated subcubic graphs. It is a discharging statement for admissible supports, where every receiver also carries the contextual target algebra, the replacement constraints, the window-stub accounting, and the saturated-exit alternatives of Definition 14.16. The raw receiver thresholds are 4, 8, 12, not the discharging capacities. They mark the saturated handoff values H_0, H_1, H_2 of Lemma 14.27. Once exits 1–3, 5, and 6 are closed, exit 4 has peeled its target-defective routed loads, and exit 7 is handed to Type B by Lemma 14.40, the remaining Type A charge capacities are 3, 7, 11 for internal degrees 2, 1, 0. Route 8 is not a bare finite graph class: by Lemma 14.47 it is reduced to indexed target-complete-minimal silent trace-basin entries with a linear burden whenever it carries the large-budget charge deficit, and Proposition 14.59 reduces that branch further to the two-carrier residual of Definition 14.48. Thus the density bound is a statement about completable atoms carrying the full target algebra and replacement constraints, not about isolated subcubic graphs.*

Framework branch table. The Type A subsection closes its local residuals by the following framework table.

Branch	Residual condition	Demand	Payer	Bounded multiplicity	Closure
Direct exits	Saturated exits (1)–(3) occur.	A target hit.	The Mersenne-return algebra.	One local receiver.	Closed by Definition 14.16 and Lemma 14.39.
Target defect	Exit (4) supplies a routed-load witness.	One peel token.	The exit-(4) ledger.	Finite descent in the peeling set.	Definition 14.33 and Lemmas 14.34 and 14.38.
Compression or delocalization	Exits (5) and (6) occur.	A proper support.	Replacement and smearing.	The declared support.	Lemma 14.39.
Type B handoff	Exit (7) occurs.	Decorated fan data.	The Type B ledger.	One grouped handoff envelope.	Lemmas 14.25 and 14.40.
Unsaturated atom	Receiver loads remain below the saturated thresholds.	3/7/11 discharging capacity.	The deficiency budget.	The receiver partition.	Lemmas 14.43 and 14.45.
Route 8	A silent trace-basin residual remains.	Carrier burden.	Carrier-deletion witnesses.	The terminal two-carrier package.	Proposition 14.62.

14.2 Type B: high-degree fan-safe surplus obstruction

Definition 14.65 (Assigned Type B charge ledger). *Let X be a connected assigned Type B support. Write Y_X for the counted remainder core of X , and write H_X for the high-degree fan centers whose surplus units are assigned to X by [Definition 13.2](#). Thus*

$$\partial_+(X) = \partial_+(Y_X), \quad |V(X)| = |V(Y_X)|, \quad \sigma(X) = \sum_{h \in H_X} (d_G(h) - 3).$$

A high-degree vertex that appears only as boundary context for another support is not an element of that support's set H_X . If the same ambient vertex lies in the local drawing of a support and also supplies a surplus unit, it is entered once in H_X and is not available as a non-center carrier.

For $y \in Y_X$ set

$$\delta_X^+(y) = \max\{0, 3 - d_{Y_X}(y)\}, \quad \text{ch}_X(y) = \delta_X^+(y) - \frac{1}{4},$$

and for $h \in H_X$ set

$$\text{ch}_X(h) = -(d_G(h) - 3) - \frac{1}{4}.$$

The Type B augmented ledger is

$$\widehat{\text{Ch}}_B(X) = \sum_{y \in Y_X} \text{ch}_X(y) + \sum_{h \in H_X} \text{ch}_X(h).$$

It satisfies the exact accounting identity

$$\mathcal{N}(X) = \widehat{\text{Ch}}_B(X) + \frac{1}{4}|H_X|. \quad (\text{B-ledger})$$

Consequently $\widehat{\text{Ch}}_B(X) \geq 0$ implies $\mathcal{N}(X) \geq 0$, hence

$$\partial_+(X) - \sigma(X) \geq \frac{1}{4}|V(X)|.$$

In a fan envelope, a cubic-closed neighbour has $\delta_X^+ = 0$ and contributes $-\frac{1}{4}$; a fan neighbour that is not cubic-closed has internal degree at most 2 in the assigned envelope and contributes at least $\frac{3}{4}$.

For a grouped decorated Type B envelope support $\mathfrak{X}_{\mathfrak{C}}^$ of [Definition 14.20](#), the same ledger is used with $Y_X = Y_{\mathfrak{C}}^*$ and with the component center-token sum $\omega(\mathfrak{C})$ in place of the canonically assigned surplus units. The grouped ledger is never summed separately from the canonical Type A pieces that generated it; the transfer identity of [Lemma 14.21](#) is the summation rule. Hence an external handoff center, including a center inside a packed P_{13} -window, can enter the Type B bridge only through its grouped envelope component. In the global fan-mass bound, that surplus unit may also appear once through the ordinary canonical support containing the center, and this is the only allowed double use.*

Suppose X carries high-degree surplus, and let $h \in H_X$. By [Lemma 3.2](#), the vertices of H_X are independent. Write

$$d_G(h) = 3 + t_h.$$

The surplus at h is t_h , and every unordered pair of neighbours $u, v \in N(h)$ gives a fan path

$$u - h - v$$

of length 2. If $G - h$ contains a simple $u-v$ path of length

$$2^j - 2 \in \{2, 6, 14, 30, \dots\},$$

then that path, together with the two edges hu and hv , is a cycle of length $(2^j - 2) + 2 = 2^j$. Thus every neighbour pair of h must be fan-return-safe.

For every $h \in H_X$,

$$N(h) \text{ is a clique in } \mathcal{F}_h,$$

because the failure of one of the five fan-safe conditions is, respectively, a power-of-two cycle, a legal-label contradiction, a target-defective quotient, a forbidden target-complete compression, or a proper/global delocalization. Hence

$$d_G(h) \leq \omega(\mathcal{F}_h).$$

Also,

$$d_G(h) \leq \rho(A(\mathcal{F}_h)) + 1$$

by the spectral clique bound: if a clique has size r , the corresponding principal submatrix is $J_r - I_r$, whose largest eigenvalue is $r - 1$, and eigenvalue interlacing gives $r - 1 \leq \rho(A(\mathcal{F}_h))$. The Type B branch is the case in which a negative-charge support contains at least one assigned high-degree fan center. The fan-safe graph and marked fan data are those of [Definitions 14.17](#) and [14.18](#); the fan ledger below accounts for the surplus terms at those centers.

Lemma 14.66 (Fan label packing and certificate-closed charge). *In the Type B fan-certificate graph, the P_{13} -label projection with adjacency relation $C_2(S, T) = 1$ has clique number at most 8. Every clique in that projection containing a non-singleton label has size at most 7. Consequently every certificate-marked Type B fan has $d_G(h) \leq 8$, and every certificate-marked Type B fan with*

$$c - \frac{11 - k}{4} \leq 0, \quad k = d_G(h), \quad c = |\{u \in N(h) : u \text{ is cubic-closed}\}|,$$

has nonnegative closed-neighbourhood charge in the augmented ledger of [Definition 14.65](#). These are the certificate-closed fans of [Definition 14.71](#).

Proof. Let D be the graph on $\{0, \dots, 12\}$ in which two distinct indices are adjacent exactly when their difference is 4 or 12. If $S_1, \dots, S_m$ are pairwise C_2 -compatible labels, choose one representative $r_i \in S_i$ from each label. The relation $C_2(S_i, S_j) = 1$ says precisely that no cross-difference between S_i and S_j is 0, 4, or 12. Hence the representatives $r_1, \dots, r_m$ are distinct and form an independent set in D .

The graph D is the disjoint union of the component

$$0 - 4 - 8 - 12 - 0$$

and the three paths

$$1 - 5 - 9, \quad 2 - 6 - 10, \quad 3 - 7 - 11.$$

Thus

$$\alpha(D) = 2 + 2 + 2 + 2 = 8,$$

and every C_2 -compatible label family has size at most 8. In a certificate-marked fan the labelling $S_h : N(h) \rightarrow \mathcal{L}$ sends the neighbour set to such a C_2 -compatible label family. Since $C_2(S, S) = 0$, the labels of distinct neighbours are distinct. Hence

$$d_G(h) = |N(h)| = |S_h(N(h))| \leq 8.$$

Now suppose the clique contains a non-singleton label S . Every other label is contained in

$$A = \{0, \dots, 12\} \setminus N_D[S],$$

where $N_D[S]$ is the closed D -neighbourhood of S , because a cross index equal to an index of S or at distance 4 or 12 from it would violate C_2 . Choose two distinct indices $a, b \in S$. Deleting the closed D -neighbourhood of $\{a, b\}$ reduces the independence number of D by at least 2. If a and b lie in different components of D , each deletion reduces the corresponding component's independence number by at least 1. If they lie in the same component, that component is either a three-vertex path or the four-cycle above; in either case the closed neighbourhoods of two distinct vertices leave independence number 0 in that component, whereas the original component had independence number 2. Since $N_D[\{a, b\}] \subseteq N_D[S]$, it follows that

$$\alpha(D[A]) \leq 6.$$

Representatives of all labels other than S form an independent set in $D[A]$, so there are at most 6 of them. Therefore every clique containing a non-singleton label has size at most 7.

For the charge statement, write $k = d_G(h)$ and let c be the number of cubic-closed neighbours. The first packing bound gives $k \leq 8$. The certificate-closed condition used below is exactly the nonpositive-deficit condition

$$c - \frac{11 - k}{4} \leq 0.$$

Equivalently, $c \leq 1$ when $k \leq 7$, and $c = 0$ when $k = 8$. With the augmented Type B charge of [Definition 14.65](#), the center contributes

$$\text{ch}_X(h) = 3 - k - \frac{1}{4}.$$

In the worst certificate-closed case with $k \leq 7$, one neighbour is cubic-closed and contributes $-\frac{1}{4}$, while the remaining $k - 1$ neighbours contribute at least $\frac{3}{4}$ each. Therefore

$$\text{ch}_X(h) + \sum_{u \in N(h)} \text{ch}_X(u) \geq \left(3 - k - \frac{1}{4}\right) - \frac{1}{4} + \frac{3}{4}(k - 1) = \frac{7 - k}{4} \geq 0.$$

If $k = 8$, certificate-closed means $c = 0$, so all eight neighbours contribute at least $\frac{3}{4}$, and

$$\text{ch}_X(h) + \sum_{u \in N(h)} \text{ch}_X(u) \geq \left(3 - 8 - \frac{1}{4}\right) + 8 \cdot \frac{3}{4} = \frac{3}{4} > 0.$$

Thus every certificate-closed marked fan has nonnegative closed-neighbourhood charge. $\square$

Remark 14.67 (The fan-degree bound is label packing). *The clique bound $d_G(h) \leq \omega(\mathcal{F}_h)$ is an explicit structural consequence of the label algebra, not a free parameter. Two neighbours attaching to a common P_{13} window with labels S, T are fan-return-safe through that window exactly when $C_2(S, T) = 1$: the wedge $u - h - v$ is the length-two outside path, closing the window into a cycle of length $4 + |i - j| \notin \mathcal{D}$. Since $C_2(S, S) = 0$, distinct neighbours carry distinct labels. [Lemma 14.66](#) proves that every pairwise C_2 -compatible set has size at most 8, and that every such set containing a non-singleton label has size at most 7. The spectral inequality $d_G(h) \leq \rho(A(\mathcal{F}_h)) + 1$ is a secondary quantitative form of this same packing constraint. A reader who treats the fan as a bare subcubic graph, ignoring this label algebra, will not see the bound—just as ignoring admissibility hides the Type A bound ([Remark 14.1](#)).*

Definition 14.68 (Type B window-incidence profile). *Let W be the union of the fixed maximal packed induced P_{13} windows. A Type B fan-window profile records the assigned fan neighbours that lie in the remainder side of a decorated envelope; thus each such neighbour lies in $R = G - W$. A window incidence of a cubic-closed neighbour u is a triple*

$$(u, P, i)$$

where $u \notin W$, $P = p_0 \cdots p_{12}$ is a packed window, and $up_i \in E(G)$. A Type B fan-window profile is fully window-supported when every cubic-closed neighbour has both non- h neighbours in W . It is same-window supported at u when the two non- h neighbours of u lie in one packed window.

A non-window fan incidence of a cubic-closed neighbour u is an edge uz with $z \notin W \cup \{h\}$. A cubic-closed neighbour is window-supported, mixed-supported, or internal-supported according as it has two, one, or zero window incidences among its two non- h incidences. Thus a mixed-supported neighbour has one window incidence and one non-window fan incidence, while an internal-supported neighbour has two non-window fan incidences.

For a surplus port $p = (h, x)$ whose port vertex x is recorded on the remainder side of the fan-window profile, the neighbour $x \in N(h)$ is treated as a cubic-closed fan neighbour once the two non- h incidences xa_p and xb_p are assigned to the fan envelope. Its support type is read from the positions of a_p and b_p : it is window-supported, mixed-supported, or internal-supported according as two, one, or neither of those endpoints lies in W . This convention does not depend on whether a_pb_p is an edge of G .

Definition 14.69 (Fan-closed surplus port). *Let h be a high-degree fan center, let $p = (h, x)$ be a surplus port at h , and write*

$$N_G(x) = \{h, a_p, b_p\}.$$

In an assigned Type B fan-window profile $\mathfrak{F}_h$, the port p is fan-closed if the following three conditions hold.

- (a) *The port vertex x is recorded as a remainder-side fan neighbour of h .*
- (b) *The two non- h incidences xa_p and xb_p are assigned to the fan envelope of $\mathfrak{F}_h$.*
- (c) *Each assigned incidence is recorded as either a window incidence or a non-window fan incidence in the sense of [Definition 14.68](#).*

Thus a fan-closed port contributes one cubic-closed neighbour to the count $c(\mathfrak{F}_h)$ of [Definition 14.71](#). A fan-closed port may be triangular, when $a_pb_p \in E(G)$, or open, when $a_pb_p \notin E(G)$.

Lemma 14.70 (Fan-compatible open pairs give fan-closed local data). *Let h be a high center, and let $p = (h, x)$ and $q = (h, y)$ be fan-compatible open ports at h . In any assigned Type B fan-window profile that records x and y as remainder-side fan neighbours of h and assigns the four incidences*

$$xa_p, \quad xb_p, \quad ya_q, \quad yb_q$$

to the fan envelope, the ports p and q are two distinct fan-closed ports. Moreover, these four incidences are pairwise distinct as local incidence carriers.

Proof. The ports are distinct, so $x \neq y$. Since both ports have center h , both x and y lie in $N_G(h)$; by [Lemma 3.9](#), they are cubic. Hence

$$N_G(x) = \{h, a_p, b_p\}, \quad N_G(y) = \{h, a_q, b_q\}.$$

The hypotheses of the assigned profile are exactly the three conditions in [Definition 14.69](#) for p and for q . Thus both ports are fan-closed.

It remains to check that the four displayed incidences are locally distinct. The two incidences at x are distinct because $a_p \neq b_p$, and the two incidences at y are distinct because $a_q \neq b_q$. No incidence incident with x can equal an incidence incident with y , since $x \neq y$. [Definition 3.10](#) gives $x \notin s(q)$, $y \notin s(p)$, and $s(p) \cap s(q) = \emptyset$. Therefore no endpoint of one displayed incidence can equal the endpoint of another displayed incidence except at its own port vertex, and the four incidences are pairwise distinct as local carriers. □

Definition 14.71 (Positive-deficit Type B fan-window residual). *Let $\mathfrak{F}$ be a marked Type B fan centered at h , and write*

$$k = d_G(h), \quad c(\mathfrak{F}) = |\{u \in N(h) : u \text{ is cubic-closed in } \mathfrak{F}\}|.$$

The closed-neighbour deficit of $\mathfrak{F}$ is

$$D_B(\mathfrak{F}) = c(\mathfrak{F}) - \frac{11 - k}{4}.$$

The fan is certificate-closed when

$$D_B(\mathfrak{F}) \leq 0.$$

Equivalently, certificate-closed means $c(\mathfrak{F}) \leq 1$ for $k \leq 7$ and $c(\mathfrak{F}) = 0$ for $k = 8$.

A positive-deficit Type B fan-window residual is an admissible marked Type B fan-window profile with $D_B(\mathfrak{F}) > 0$ whose recorded window labels are pairwise C_2 -compatible whenever they lie in a common packed window, and whose fan-window data do not realize a Mersenne return, a target-defective quotient, a target-complete compression, or a proper/global delocalization. In this residual every cubic-closed neighbour is recorded with its two non- h incidences, marked as internal-supported, mixed-supported, or window-supported according as zero, one, or two of those incidences lie in W . Thus for $k \leq 7$ the residual has at least two cubic-closed neighbours; for $k = 8$ the single cubic-closed case is also residual, because its local charge deficit is $D_B = 1/4$.

Definition 14.72 (Hybrid Type B incidence accounting). *Let $\mathfrak{F}$ be a positive-deficit Type B fan-window residual. Write*

$$c_W(\mathfrak{F}), \quad c_M(\mathfrak{F}), \quad c_I(\mathfrak{F})$$

for the numbers of window-supported, mixed-supported, and internal-supported cubic-closed neighbours, respectively. Thus

$$c = c_W + c_M + c_I.$$

The window-incidence count and non-window incidence count are


$$I_W(\mathfrak{F}) = 2c_W + c_M, \quad I_N(\mathfrak{F}) = c_M + 2c_I.$$

The hybrid non-window demand is

$$D_N(\mathfrak{F}) = \max \left\{ 0, D_B(\mathfrak{F}) - \frac{1}{2} I_W(\mathfrak{F}) \right\}.$$

Equivalently,

$$D_N(\mathfrak{F}) = \max \left\{ 0, \frac{1}{2} I_N(\mathfrak{F}) - \frac{11 - k}{4} \right\}.$$



If the labelling S_h is absent, h is a fan-certificate residual center and is charged to the Type B fan-mass ledger rather than to the certificate-closed local fan ledger.

Figure 26: Visual definition of the Type B fan-window and hybrid ledger data. The left panel is the local marked Type B fan, with high-degree center h and neighbours in $N(h)$; the right panel is the hybrid incidence ledger. The horizontal black path $P = p_0 \cdots p_{12}$ is a packed induced P_{13} window, with the displayed indices naming its vertices. The labelling $S_h : N(h) \rightarrow \mathcal{L}$ assigns fan certificates, and the condition $C_2(S_h(u), S_h(v)) = 1$ states pairwise fan-safety for distinct neighbours. The neighbour u_W is window-supported because both of its non- h incidences go to window vertices; u_M is mixed-supported because one incidence goes to the window and one to a non-window vertex z_M ; u_I is internal-supported because its two non- h incidences go to non-window vertices z_1, z_2 ; and u_0 is an open neighbour, so it is not counted as cubic-closed in this local ledger. Orange edges are window incidences, green edges are non-window incidences, and each displayed $\frac{1}{2}$ is the half-credit contributed by one such incidence. In the ledger, $k = d_G(h)$, $c = c_W + c_M + c_I$ counts cubic-closed neighbours by support type, $D_B(\mathfrak{F}) = c - (11 - k)/4$ is the closed-neighbour deficit, $I_W = 2c_W + c_M$ is the window-incidence count, $I_N = c_M + 2c_I$ is the non-window-incidence count, and $D_N = \max\{0, D_B - \frac{1}{2}I_W\}$ is the remaining non-window demand. B1 denotes the local payment of one positive-deficit fan, while B2 denotes the cross-center disjointness ledger needed to sum those payments globally. The bottom boxed note records the alternate case: if the labelling S_h is absent, then h is charged as a fan-certificate residual center in the Type B fan-mass ledger.

Definition 14.73 (Closed fan-window pair). *Let $P = p_0 p_1 \cdots p_{12}$ be a packed induced P_{13} window. Suppose u is a cubic-closed neighbour of a Type B fan center h and is same-window supported at P . Its closed label*

$$S(u) = \{a_u, b_u\} \subseteq \{0, \dots, 12\}, \quad a_u < b_u,$$

meaning that u is adjacent to h, p_{a_u}, p_{b_u} and has no other incident edge in the assigned fan-window envelope. For two cubic-closed neighbours u, v same-window supported at P , put $S(v) = \{a_v, b_v\}$ and define

$$L_h(x, y) = 4 + |x - y| \quad (x \in S(u), y \in S(v)).$$

If the four endpoints interlace, say $a_u < a_v < b_u < b_v$ after possibly swapping u and v , define

$$L_\times(u, v) = 4 + (a_v - a_u) + (b_v - b_u).$$

If the endpoints do not interlace, $L_\times(u, v)$ is undefined.

Definition 14.74 (Direct-cycle-free closed pair). *A closed fan-window pair (u, v) is direct-cycle-free when all of the following arithmetic conditions hold:*

(a)

$$b_u - a_u, b_v - a_v \notin \{2, 6\};$$

(b)

$$|x - y| \notin \{0, 4, 12\} \quad (x \in S(u), y \in S(v));$$

(c) if $L_\times(u, v)$ is defined, then

$$L_\times(u, v) \notin \{8, 16\}.$$

Lemma 14.75 (Direct fan-window cycle eliminations). *Let u, v be distinct cubic-closed neighbours of a Type B fan center h that are same-window supported at the same packed window P . If the closed fan-window pair (u, v) is not direct-cycle-free and $h \notin V(P)$, then the envelope contains a power-of-two cycle. The condition $h \notin V(P)$ holds for an ordinary remainder-side fan center. A handoff center lying in a packed window is handled only through its grouped decorated envelope and is not an instance of this local lemma.*

Proof. Write $S(u) = \{a_u, b_u\}$ and $S(v) = \{a_v, b_v\}$ as in Definition 14.73. If $b_u - a_u \in \{2, 6\}$, then

$$up_{a_u} P p_{b_u} u$$

is a simple cycle of length $(b_u - a_u) + 2 \in \{4, 8\}$, because $u \notin W$ and P is a simple induced path. If $b_v - a_v \in \{2, 6\}$, then

$$vp_{a_v} P p_{b_v} v$$

is a simple cycle of length $(b_v - a_v) + 2 \in \{4, 8\}$ for the same reason.

If $|x - y| \in \{0, 4, 12\}$ for some $x \in S(u)$ and $y \in S(v)$, then

$$uhvp_y P p_x u$$

is a simple cycle: the vertices u and v are distinct and lie outside W , whereas the segment $p_y P p_x$ lies inside the packed window. Its length is $2 + 1 + |x - y| + 1 = 4 + |x - y| \in \{4, 8, 16\}$.

Finally suppose the endpoints interlace. After possibly swapping u and v , write

$$a_u < a_v < b_u < b_v.$$

The two path segments $p_{a_u} P p_{a_v}$ and $p_{b_u} P p_{b_v}$ are internally vertex-disjoint, and $u, v \notin W$. Hence

$$up_{a_u} P p_{a_v} vp_{b_v} P p_{b_u} u$$

is a simple cycle of length

$$4 + (a_v - a_u) + (b_v - b_u) = L_\times(u, v).$$

If this length is 8 or 16, the cycle is a power-of-two cycle. These cases are exactly the negations of the direct-cycle-free conditions. $\square$

Lemma 14.76 (Two-window direct cycle eliminations). *Let u, v be distinct cubic-closed neighbours of a Type B fan center h . Suppose u and v have incidences to two distinct packed windows*

$$P = p_0 \cdots p_{12}, \quad Q = q_0 \cdots q_{12},$$

with

$$up_i, vp_j, uq_a, vq_b \in E(G).$$

If

$$|i - j| + |a - b| \in \{0, 4, 12\},$$

then the fan-window envelope contains a power-of-two cycle.

Proof. By the definition of a window incidence, u and v lie outside W . The packed windows are vertex-disjoint, so the path segments $p_i P p_j$ and $q_a Q q_b$ are vertex-disjoint and avoid u and v . Hence

$$u p_i P p_j v q_b Q q_a u$$

is a simple cycle. Its length is

$$1 + |i - j| + 1 + 1 + |a - b| + 1 = 4 + |i - j| + |a - b|.$$

If $|i - j| + |a - b| \in \{0, 4, 12\}$, this length is 4, 8, or 16. $\square$

Lemma 14.77 (Positive-deficit fan-window budget). *Let $\mathfrak{F}$ be an admissible positive-deficit Type B fan-window residual centered at h , with $k = d_G(h)$ and $c = c(\mathfrak{F})$. Then*

$$4 \leq k \leq 8, \quad D_B(\mathfrak{F}) > 0,$$

and more explicitly either $k \leq 7$ and $c \geq 2$, or $k = 8$ and $c \geq 1$. If $k = 8$, no cubic-closed neighbour is same-window supported with a non-singleton label. For every k , every same-window supported cubic-closed neighbour has internal label gap outside $\{2, 6\}$, every same-window closed fan-window pair in $\mathfrak{F}$ is direct-cycle-free, and in any assigned Type B support containing $\mathfrak{F}$ the closed fan neighbourhood has augmented charge lower bound

$$\text{ch}_X(h) + \sum_{u \in N(h)} \text{ch}_X(u) \geq \frac{11 - k}{4} - c.$$

Equivalently, the remaining Type B charge deficit is exactly

$$D_B(\mathfrak{F}) = c - \frac{11 - k}{4} > 0.$$

Assume in addition that $\mathfrak{F}$ is fully window-supported and contains no direct cycle of the forms in Lemmas 14.75 and 14.76. Then its c cubic-closed neighbours use exactly $2c$ distinct packed-window incidences, and

$$D_B(\mathfrak{F}) \leq \frac{1}{2} \cdot 2c - \frac{3}{4}.$$

Proof. The residual is a marked Type B fan, so its certificate labelling $S_h : N(h) \rightarrow \mathcal{L}$ is part of the data. Therefore the fan degree bound gives $k \leq 8$ by Lemma 14.66. If $k = 8$, then a same-window supported cubic-closed neighbour would carry a non-singleton label, and the non-singleton clique cap in Lemma 14.66 would force $k \leq 7$, a contradiction. Hence no cubic-closed neighbour of an actual 8-fan residual is same-window supported with a non-singleton label. Since the fan is Type B, $k \geq 4$, so $4 \leq k \leq 8$. The positive-deficit hypothesis is $D_B(\mathfrak{F}) > 0$, i.e.

$$c > \frac{11 - k}{4}.$$

As c is integral, this gives $c \geq 2$ for $k \leq 7$ and $c \geq 1$ for $k = 8$.

If a cubic-closed neighbour u is same-window supported at $P = p_0 \cdots p_{12}$ and has closed label $S(u) = \{a_u, b_u\}$ with $b_u - a_u \in \{2, 6\}$, then

$$u p_{a_u} P p_{b_u} u$$

is a simple cycle of length 4 or 8, contradicting admissibility. If some same-window closed pair were not direct-cycle-free, Lemma 14.75 would give a power-of-two cycle, contradicting admissibility.

Hence every single same-window label has gap outside $\{2, 6\}$, and every same-window closed pair is direct-cycle-free.

For the augmented Type B charge, the fan center contributes

$$\text{ch}_X(h) = (3 - k) - \frac{1}{4}.$$

Each cubic-closed neighbour contributes $-\frac{1}{4}$, and each of the remaining $k - c$ neighbours has internal degree at most 2 in the assigned fan support and therefore contributes at least $\frac{3}{4}$. Hence

$$\text{ch}_X(h) + \sum_{u \in N(h)} \text{ch}_X(u) \geq \left(3 - k - \frac{1}{4}\right) - \frac{c}{4} + \frac{3}{4}(k - c) = \frac{11 - k}{4} - c.$$

Because $D_B(\mathfrak{F}) > 0$, this lower bound is strictly negative, and the deficit is $D_B(\mathfrak{F}) = c - \frac{11-k}{4}$.

Now assume $\mathfrak{F}$ is fully window-supported and contains no direct cycle of the two displayed forms. Each cubic-closed neighbour has exactly two non- h neighbours, both in packed windows, so the number of packed-window incidences counted with multiplicity is $2c$. These incidences are distinct. Indeed, two different closed neighbours cannot use the same packed-window vertex: if up_i and vp_i were both present, then the fan-cross cycle

$$u - h - v - p_i - u$$

has length 4. One closed neighbour cannot use the same packed-window vertex twice because the graph is simple. Incidences in different packed windows are distinct because the packing is vertex-disjoint. Thus the residual uses exactly $2c$ distinct packed-window incidences.

Since $k \leq 8$,

$$\frac{11 - k}{4} \geq \frac{3}{4}.$$

Therefore

$$D_B(\mathfrak{F}) = c - \frac{11 - k}{4} \leq c - \frac{3}{4} = \frac{1}{2} \cdot 2c - \frac{3}{4}.$$

Equivalently, half a unit of reserved credit per distinct packed-window incidence pays the local Type B deficit with at least three quarters of a unit of slack. $\square$

Lemma 14.78 (Hybrid incidence budget for positive-deficit fans). *Let $\mathfrak{F}$ be an admissible positive-deficit Type B fan-window residual centered at h . Suppose that none of the direct-cycle conclusions of Lemmas 14.75 and 14.76 occurs. Then the $2c(\mathfrak{F})$ non- h incidences of the cubic-closed neighbours of h are pairwise disjoint as incidence carriers. The window incidences supply half-credit $\frac{1}{2}I_W(\mathfrak{F})$, and the non-window incidences supply a disjoint local reserve of size at least*

$$D_N(\mathfrak{F}) = \max \left\{ 0, D_B(\mathfrak{F}) - \frac{1}{2}I_W(\mathfrak{F}) \right\}.$$

Consequently the hybrid fan entry supplies at least $D_B(\mathfrak{F})$.

Proof. Let u and v be distinct cubic-closed neighbours of h . If a non- h neighbour z were incident with both u and v , then

$$u - h - v - z - u$$

would be a simple 4-cycle, because u and v are distinct, $z \neq h$, and the fan neighbours have ambient degree 3. This is forbidden by dyadic-safety. One cubic-closed neighbour cannot use the same non- h

endpoint twice, since G is simple. Hence the two non- h incidences of all cubic-closed neighbours are pairwise disjoint as incidence carriers.

The window incidences among them are counted by $I_W(\mathfrak{F})$ and the non-window incidences by $I_N(\mathfrak{F})$. By the local accounting rule in [Definition 14.72](#), each such incidence carries half-credit; the window/non-window classification determines which of the two local pools contains it. No global carrier choice is used at this point. Thus the total local incidence capacity is

$$\frac{1}{2}I_W(\mathfrak{F}) + \frac{1}{2}I_N(\mathfrak{F}) = \frac{1}{2} \cdot 2c(\mathfrak{F}) = c(\mathfrak{F}).$$

Since

$$D_B(\mathfrak{F}) = c(\mathfrak{F}) - \frac{11-k}{4}$$

and $k \leq 8$, this total capacity pays $D_B(\mathfrak{F})$ with slack at least $\frac{11-k}{4} \geq \frac{3}{4}$.

After the window half-credit has been applied, the remaining demand is

$$\max \left\{ 0, D_B(\mathfrak{F}) - \frac{1}{2}I_W(\mathfrak{F}) \right\} = D_N(\mathfrak{F}).$$

The available non-window half-credit is

$$\frac{1}{2}I_N(\mathfrak{F}) = c(\mathfrak{F}) - \frac{1}{2}I_W(\mathfrak{F}) = D_B(\mathfrak{F}) + \frac{11-k}{4} - \frac{1}{2}I_W(\mathfrak{F}),$$

which is at least $D_N(\mathfrak{F})$. The disjointness proved in the first paragraph permits these non-window incidence carriers to be chosen disjointly inside the single fan entry. $\square$

Lemma 14.79 (Local hybrid fan entry). *Let $\mathfrak{F}$ be an admissible positive-deficit Type B fan-window residual. After the direct-cycle, target-defect, compression, and delocalization alternatives are excluded, the packed-window incidences and the local internal/mixed incidence reserve form a disjoint local fan entry whose total capacity is at least $D_B(\mathfrak{F})$.*

Proof. The excluded target-defect, compression, and delocalization alternatives are part of the definition of a positive-deficit residual. The direct-cycle case has also been excluded. Therefore [Lemma 14.78](#) supplies the asserted disjoint local hybrid incidence entry of capacity at least $D_B(\mathfrak{F})$. $\square$

Proposition 14.80 (Fan-closed ports route to the Type B fan ledger). *Let h be a high-degree fan center of degree $k \geq 4$, and let $\mathfrak{F}_h$ be an admissible Type B fan-window profile that records h as a fan center. Suppose that $\mathfrak{F}_h$ contains a set $\mathcal{Q}$ of $r \geq 2$ fan-closed surplus ports at h . Then the following hold.*

- (a) *The ports in $\mathcal{Q}$ contribute at least r cubic-closed neighbours of h .*
- (b) *If the fan at h is certificate-marked and survives the direct-cycle, target-defect, compression, and delocalization alternatives in B1, then it is a positive-deficit Type B fan-window residual with*

$$D_B(\mathfrak{F}_h) \geq r - \frac{11-k}{4} \geq \frac{k-3}{4} > 0.$$

- (c) *The hybrid B1 ledger supplies local incidence capacity at least $D_B(\mathfrak{F}_h)$. The paying incidences include the two non- h incidences xa_p and xb_p of every fan-closed port $p = (h, x) \in \mathcal{Q}$, classified by [Definition 14.68](#) as window, mixed, or non-window incidences according to the locations of a_p and b_p .*

Proof. Let $p = (h, x) \in \mathcal{Q}$. Since p is fan-closed, [Definition 14.69](#) records x as a remainder-side fan neighbour of h and assigns the two non- h incidences xa_p and xb_p to the fan envelope. By [Lemma 3.9](#), x is cubic. Hence x has internal fan degree 3 in the assigned envelope, so it is cubic-closed in the sense of [Definition 14.18](#). Distinct ports at h have distinct port vertices, and therefore the r ports in $\mathcal{Q}$ give at least r cubic-closed neighbours. This proves (a).

Assume now that the fan is certificate-marked and survives the direct-cycle, target-defect, compression, and delocalization alternatives in the B1 statement. Let $c = c(\mathfrak{F}_h)$ be the number of cubic-closed neighbours recorded in the fan envelope. By (a), $c \geq r$. The closed-neighbour deficit is

$$D_B(\mathfrak{F}_h) = c - \frac{11 - k}{4}.$$

Therefore

$$D_B(\mathfrak{F}_h) \geq r - \frac{11 - k}{4} \geq 2 - \frac{11 - k}{4} = \frac{k - 3}{4}.$$

Since $k \geq 4$, the last quantity is positive. Thus the fan satisfies the positive-deficit condition of [Definition 14.71](#). The profile is admissible by hypothesis; the certificate marking gives the required C_2 -compatibility of the recorded labels; and the absence of the direct-cycle, target-defect, compression, and delocalization alternatives is exactly the surviving B1 hypothesis. Hence all clauses of [Definition 14.71](#) are satisfied. This proves (b).

Under the hypotheses of (b), [Lemma 14.79](#) applies. It invokes [Lemma 14.78](#), which gives disjoint local half-credit on the $2c$ non- h incidences of the cubic-closed fan neighbours and pays at least $D_B(\mathfrak{F}_h)$. For every fan-closed port $p = (h, x) \in \mathcal{Q}$, the two relevant incidences are exactly xa_p and xb_p . [Definition 14.68](#) classifies them according to whether their non- h endpoints lie in the packed-window union W . This proves (c). □

Corollary 14.81 (Fan-compatible open pairs route to the Type B fan ledger). *Let h be a light center of degree $k \geq 4$, and let $p = (h, x)$ and $q = (h, y)$ be fan-compatible open ports at h . In any admissible Type B fan-window profile that records h as a fan center, records x and y as remainder-side fan neighbours, and assigns the four non- h incidences of p and q , the fan at h has two fan-closed ports. Consequently it routes through [Proposition 14.80](#). If the direct-cycle, target-defect, compression, or delocalization alternative occurs, the branch is already closed. If the fan is certificate-marked and those closed alternatives are absent, then*

$$D_B(\mathfrak{F}_h) \geq \frac{k - 3}{4} > 0,$$

so the B1 ledger pays the local positive-deficit fan. If the fan is not certificate-marked, it is passed to the later Type B bridge-residual routing.

Proof. [Lemma 14.70](#) shows that p and q are two distinct fan-closed ports in the assigned profile, with locally distinct non- h incidence carriers. Applying [Proposition 14.80](#) with $\mathcal{Q} = \{p, q\}$ gives the displayed deficit bound and the local B1 alternative. □

Corollary 14.82 (Triangular ports route to the Type B fan ledger). *Let h be a heavy center of degree $k \geq 5$, and suppose that the triangular alternative in [Corollary 3.13](#) holds at h . Let $\mathcal{T}_h$ be a set of $k - 2$ triangular ports at h . In any admissible Type B fan-window profile that records h as a fan center, records the port vertices x_i of $\mathcal{T}_h$ as remainder-side fan neighbours, and assigns the two triangle incidences of every port in $\mathcal{T}_h$, the fan at h has at least $k - 2$ fan-closed ports. Hence it routes through [Proposition 14.80](#); in the certificate-marked B1 survivor,*

$$D_B(\mathfrak{F}_h) \geq (k - 2) - \frac{11 - k}{4} = \frac{5k - 19}{4} > 0.$$

Proof. For a triangular port $p_i = (h, x_i)$, the vertex x_i is a neighbour of h and is cubic by [Lemma 3.9](#). Its two non- h neighbours are precisely a_i and b_i , and both incidences $x_i a_i$, $x_i b_i$ are present in G . Once those two incidences are assigned to the fan envelope, [Definition 14.69](#) makes p_i fan-closed. The $k - 2$ ports in $\mathcal{T}_h$ have distinct port vertices, so they contribute at least $k - 2$ fan-closed, hence cubic-closed, neighbours. Applying [Proposition 14.80](#) with $\mathcal{Q} = \mathcal{T}_h$ gives the stated alternatives and the stronger displayed deficit bound. $\square$

Definition 14.83 (Type B ledger carriers). *Let X be a connected assigned Type B support. A Type B carrier in X is one of the following pieces of support-charge data.*

- (a) *a non-center vertex $v \in Y_X$ which is not entered as an assigned high-degree center, carrying the charge*

$$\text{ch}_X(v) = \delta_X^+(v) - \frac{1}{4}$$

from [Definition 14.65](#);

- (b) *an assigned high-degree center $h \in H_X$, carrying the augmented center charge*

$$\text{ch}_X(h) = -(d_G(h) - 3) - \frac{1}{4};$$

- (c) *a packed-window incidence (u, P, i) of [Definition 14.68](#), carrying capacity $1/2$ when it is used as the half-incidence credit of [Lemma 14.77](#);*

- (d) *a non-window fan incidence uz of [Definition 14.68](#), carrying capacity $1/2$ when it is used as the local internal/mixed half-credit of [Lemma 14.78](#);*

- (e) *a local internal/mixed reserve block Q , contained in the assigned non-window part of a positive-deficit fan, whose capacity is the sum of the support charges of its vertices and incidences in the same support-charge ledger.*

The carrier support of a block or incidence is its underlying vertex/incidence set. Two carriers are disjoint when their underlying supports are disjoint and neither is part of the ordinary deficiency reserve. The ordinary deficiency reserve is the unrefined pool of positive-deficiency units $\delta_X^+(y)$ and packed-window boundary-incidence units counted by $\partial_+(X)$ and by the stub supply in [Lemma 8.4](#). A Type B carrier is disjoint from this reserve precisely when the same vertex or incidence unit is not also used as one of those unrefined reserve units. The incidence carriers in (c) and (d) are the global support-ledger realizations of the two local half-credit pools constructed in [Lemma 14.78](#); this identification is made only after the global carrier universe has been defined.

Definition 14.84 (Candidate Type B ledger entries). *Let X be a connected assigned Type B support. A Type B demand is attached to a marked fan center $h \in H_X$.*

- (a) *If the marked fan at h is certificate-closed, a candidate ledger entry is a set $A_h \subseteq N(h) \setminus H_X$ such that*

$$\text{ch}_X(h) + \sum_{v \in A_h} \text{ch}_X(v) \geq 0.$$

For the finite ledger this condition is checked in quarter units from the assigned incidences themselves. If $u \in N(h)$, let

$$a_h(u) = |\{z \in N_G(u) \setminus \{h\} : uz \text{ is assigned to the fan support}\}| \quad \text{and} \quad q_h(u) = 11 - 4(1 + a_h(u)).$$

Deletion criticality gives $d_G(u) = 3$, so $a_h(u) \in \{0, 1, 2\}$. The assigned fan support may, however, contain decorative or packed-window incidences whose remote endpoint is outside the counted core Y_X . Put

$$a_X(u) = |N_G(u) \cap (Y_X \setminus \{h\})|, \quad e_h(u) = |\{z \notin Y_X : uz \text{ is assigned to the fan support}\}|.$$

Every core edge is assigned, and the assignment contains only graph edges; hence $a_h(u) = a_X(u) + e_h(u)$. Consequently the correct transfer identity is

$$q_h(u) + 4e_h(u) = 11 - 4(1 + a_X(u)) = 4\text{ch}_X(u), \quad (\text{B-port-transfer})$$

and in particular $q_h(u) \leq 4\text{ch}_X(u)$. Thus $a_h(u) = 2$ gives the assigned cubic-closed weight $q_h(u) = -1$, but it need not say that u has three neighbours in the counted core: the missing amount is exactly the external correction $4e_h(u)$. If $a_h(u) \leq 1$, then $q_h(u) \geq 3$.

The finite certificate-candidate test is

$$(11 - 4d_G(h)) + \sum_{u \in A_h} q_h(u) \geq 0. \quad (\text{B-candidate-charge})$$

The test is a lower-bound calculation in the actual induced-core ledger, not an equality obtained by forgetting external incidences. Summing (B-port-transfer) transfers it to the displayed augmented-charge inequality. For a positive-deficit entry, every packed-window incidence is an external assigned shoulder and is paid from the corresponding $4e_h(u)$ correction. A selected non-window half-credit is allowed only when its remote endpoint z is an ordinary available core vertex: $z \in Y_X \setminus H_X$, z is adjacent to no assigned center, and $4\text{ch}_X(z) \geq 2$. These ordinary vertices, the complete fan-neighbour support, and the certificate-selected vertices are included in the candidate carrier support, so B2 disjointness prevents their charge from being used twice. Its carriers are the center h together with the vertices of A_h .

- (b) If the marked fan at h is positive-deficit, a candidate ledger entry is the fan envelope at h , together with all packed-window incidences and enough non-window fan incidences used by its cubic-closed neighbours to pay the remaining demand $D_N(\mathfrak{F}_h)$ of [Definition 14.72](#). Each chosen incidence has capacity $1/2$. By [Lemma 14.78](#), this entry pays $D_B(\mathfrak{F}_h)$ with slack at least $(11 - k)/4$ before unused incidence credits are discarded. Its carriers are the center h , the cubic-closed fan neighbours whose closed-neighbour deficit is being paid, and the chosen incidence carriers.

A disjoint refined Type B ledger is a choice of one candidate entry for every Type B demand in X such that all chosen carriers are pairwise disjoint. It is maximal when, after the chosen entries are removed, the remaining core has no high-degree center and no decorated Type B handoff whose center belongs to X .

Definition 14.85 (Minimal Type B overlap obstruction). *Let X be a connected assigned Type B support. A Type B overlap obstruction is a pair*

$$\mathcal{O} = (\mathcal{D}, \{\mathcal{E}(d)\}_{d \in \mathcal{D}})$$

where $\mathcal{D}$ is a nonempty finite set of Type B demands in X and $\mathcal{E}(d)$ is the set of candidate ledger entries for d , such that no choice of entries

$$E_d \in \mathcal{E}(d) \quad (d \in \mathcal{D})$$

has pairwise disjoint carriers. The obstruction is minimal when every proper nonempty subfamily of $\mathcal{D}$ does admit such a disjoint choice.

For each demand, its declared carrier universe is the center together with every vertex or incidence in its finite local item universe, whether or not reserve conflicts leave a valid candidate entry. The overlap support $Z(\mathcal{O})$ is the union of the declared carrier universes of the demands in $\mathcal{D}$, together with the paths or window segments witnessing the corresponding fan-safe relations. Thus a demand with an empty valid candidate set is not erased from the obstruction. When a connected support is required below, connected components of this union are considered explicitly; connectivity is not part of the finite-choice theorem itself.

Lemma 14.86 (Finite choice-or-overlap for a declared Type B family). *Let $\mathcal{D}_X$ be a finite duplicate-free family of declared Type B demands in X , and let every demand carry one of the two local entry types of Definition 14.84. Then exactly the following exhaustive alternative is available:*

- (i) *the whole family $\mathcal{D}_X$ admits a choice of pairwise disjoint candidate entries; or*
- (ii) *a nonempty subfamily is a minimal Type B overlap obstruction.*

The conclusion also holds when $\mathcal{D}_X = \emptyset$, in which case (i) is the vacuous choice. If $\mathcal{D}_X$ is the complete family of centers required by a support assignment, conclusion (i) is the full disjoint-carrier part of B2 for that assignment.

Proof. If the full family has a pairwise-disjoint choice, this is (i). If it does not, the full family is a nonempty overlap obstruction: it cannot be empty, because the empty family has the vacuous choice. Among the nonempty obstructing subfamilies choose one of least cardinality. Every proper nonempty subfamily of this least obstruction has a disjoint choice, so it is minimal and gives (ii). This is a proof-level finite-choice argument; it does not require enumerating products of candidate sets or the powerset of $\mathcal{D}_X$. $\square$

For an ordinary vertex support X , the complete high-center schedule is obtained by filtering $V(X)$ at ambient degree at least four. There is then a separate exhaustive local-resolution split: either some actual high center has neither a certificate-closed entry nor a verified positive-B1 entry, or every high center has one of those entries and hence forms the complete declared family to which the preceding lemma applies. This split is distinct from the later residual-core assertion involving decorated handoffs.

Lemma 14.87 (Bridge failure produces a minimal overlap obstruction). *Let X be a connected assigned Type B support with no fan-certificate residual center. If the disjoint-carrier part of the B2 refined support ledger of Definition 14.91 fails for X , then X contains a minimal Type B overlap obstruction in the sense of Definition 14.85.*

Proof. For each high-degree fan center of X , form the finite set of candidate ledger entries in Definition 14.84. The sets are finite because X is finite, every fan neighbour has degree 3 by Lemma 3.2, and a packed-window incidence uses one of the finitely many labels $0, \dots, 12$ in a fixed packed window.

Assume that no choice of one candidate entry for every high-degree fan demand has pairwise disjoint carriers. Apply Lemma 14.86 to the complete declared family. Its first alternative fails, so the second supplies a least nonempty obstructing subfamily $\mathcal{O}$. Its overlap support is the declared carrier union of Definition 14.85; no connectivity conclusion is used in this finite-choice step.

The hybrid B1 fan entry is supplied locally by Lemma 14.79. Thus a disjoint-carrier failure after B1 is exactly the overlap obstruction constructed above. A failure of the maximal residual

core condition is a different alternative: the remaining core still contains a high-degree center or a decorated Type B handoff, and is routed again through [Lemma 14.40](#) and the Type B support assignment rather than through the overlap-obstruction construction.

Applied to the four-carrier local entry supplied by [Lemma 14.70](#), this records failure to choose that entry disjointly from the remaining Type B entries as a minimal Type B overlap obstruction. $\square$

Corollary 14.88 (Global bridge routing for fan-closed local entries). *Under the hypotheses of [Proposition 14.80](#), an unmarked fan center is a fan-certificate residual center. If the fan is certificate-marked and its local B1 entry cannot be made disjoint from the other Type B entries in the assigned support, then that support contains a minimal Type B overlap obstruction. These are the two Type B bridge residual alternatives of [Definition 14.91](#).*

Proof. The first assertion is the fan-certificate residual clause in [Definition 14.18](#). In the second case the global disjoint-carrier clause B2 fails, so [Lemma 14.87](#) supplies the minimal overlap obstruction. $\square$

Lemma 14.89 (Global-to-local reflection of Type B overlap). *Let $\mathcal{O}$ be a minimal Type B overlap obstruction contained in a connected admissible support X . Then its overlap support $Z = Z(\mathcal{O})$ inherits the following global constraints.*

- (a) Contextual dyadic safety: *no cycle of G meeting $V(Z)$ has power-of-two length. Equivalently, no oriented edge in the overlap support has a Mersenne return through the actual outside context.*
- (b) High-degree separation: *the high-degree fan centers in Z are independent, and every neighbour of such a center has ambient degree 3.*
- (c) Window compatibility: *every literal same-window or two-window datum carried by a demand in Z satisfies the direct-cycle exclusions of [Lemmas 14.75](#) and [14.76](#).*
- (d) Replacement obstruction: *every concrete normalized nonempty-boundary proper atom supported on Z admits no target-complete smaller representative with the same boundary-degree profile. Equivalently, any literal CT3 compression datum for such an atom is contradictory.*
- (e) Minimal overlap: *every proper connected sub-obstruction of $\mathcal{O}$ has a refined disjoint ledger. Moreover, the demand schedule cannot be partitioned into two nonempty blocks whose complete declared carrier universes are disjoint across the partition.*

Proof. The support X is admissible by construction of the local analysis. Hence it is embedded in the fixed minimal counterexample G and carries all global constraints listed in [Definition 13.3](#) and [Remark 14.1](#). Since $Z \subseteq X$, any power-of-two cycle of G meeting $V(Z)$ would also meet $V(X)$, contradicting contextual dyadic safety of X ; this proves (a).

For (b), apply [Lemma 3.2](#). No edge joins two vertices of ambient degree at least 4, so the fan centers are independent. If $u \in N(h)$ for a high-degree center h , then the edge hu has one endpoint of degree 3; since $d_G(h) \geq 4$, the endpoint forced to have degree 3 is u .

For (c), if same-window or two-window labels violated the displayed arithmetic exclusions, the corresponding literal direct-cycle lemma would produce a power-of-two cycle in G , contradicting (a).

For (d), supply the literal normalized boundaried atom, its nonempty boundary enumeration, and an alleged smaller target-complete representative. The CT3 gluing theorem derives the whole-graph rank decrease, baseline preservation, and target transport, and therefore contradicts minimality by [Corollary 3.105](#). No abstract response flag or unconstructed quotient is used in this assertion.

Finally, (e) is the minimality clause in [Definition 14.85](#): every proper nonempty subfamily of demands admits disjoint candidate entries, hence every proper connected sub-obstruction has a refined disjoint ledger. If the full demand family split into two nonempty blocks with disjoint declared carrier universes, take a disjoint choice on each proper block and append the choices. Candidate carrier supports lie inside their declared universes, so the appended entries are disjoint across the split, contradicting that $\mathcal{O}$ is an obstruction. $\square$

Proposition 14.90 (Type B bridge residual as a global-to-local obstruction). *Every disjoint-carrier Type B bridge residual is represented by a minimal Type B overlap obstruction satisfying all constraints of [Lemma 14.89](#). Conversely, if an admissible support outside the route-8 residual has no fan-certificate residual center and no such overlap obstruction, then the disjoint-carrier clauses B2(a)–(c) hold there. Clause B2(d) completes the maximal support assignment. Together with the hybrid B1 ledger of [Lemma 14.79](#), this gives the quantitative alternative of [Proposition 14.97](#): either the Type B charge is nonnegative or its negative part is bounded by the assigned surplus.*

Proof. The first assertion is [Lemma 14.87](#) followed by [Lemma 14.89](#). For the converse, [Lemma 14.79](#) gives the local B1 entry for every positive-deficit fan. The absence of fan-certificate residual centers means that every high-degree center in the Type B support is certificate-marked. Since no multi-demand overlap obstruction is present, [Lemma 14.86](#) supplies a maximal disjoint refined Type B ledger. This proves the disjoint-carrier portion of B2, but does not by itself pay the deleted-boundary term in (B-boundary). The additional absence of an overlap obstruction supplies the full disjoint choice, and B2(d) makes the support assignment maximal. The boundary term is then bounded from the literal size and ambient degree of the selected support, as in [Proposition 14.97](#). Thus

$$\max\{0, -\mathcal{N}(X)\} \leq 200 \sum_{h \in H_X} (d_G(h) - 3)$$

for every connected admissible Type B support outside the saturated and route-8 residual branches. When the processed quarter-charge slots admit the boundary injection, the same proposition gives $\mathcal{N}(X) \geq 0$. $\square$

Definition 14.91 (Type B bridge statements). *The Type B bridge consists of the local hybrid fan ledger B1 and the global disjoint support ledger B2. The local ledger B1 is proved in [Lemma 14.79](#); B2 is the global disjointness and maximal-assignment statement used in [Proposition 14.97](#). Deleted-boundary accounting is derived from the selected support and is not a clause of B2.*

(B1) Hybrid fan ledger. *Every positive-deficit Type B fan-window residual is one of the following: it contains a direct power-of-two cycle of the forms in [Lemmas 14.75](#) and [14.76](#); it realizes a target-defective quotient, a target-complete compression, or a proper/global delocalization; or its packed-window incidences together with its local internal/mixed non-window incidences contain a disjoint hybrid fan entry of capacity at least $D_B(\mathfrak{F})$. The window part contributes $\frac{1}{2}I_W(\mathfrak{F})$, and the non-window part contributes the remaining demand $D_N(\mathfrak{F})$ of [Definition 14.72](#).*

(B2) Refined support ledger. *For any connected assigned Type B support X with no fan-certificate residual center, the high-degree fan centers of X admit a canonical refined ledger in the sense of [Definition 14.84](#). Clauses (a)–(c) are the disjoint-carrier part of B2, and clause (d) is the maximal support-assignment part.*




Figure 27: Visual proof of the Type B overlap-obstruction routing. Panel I shows the disjoint-carrier B2 failure mechanism: two Type B fan demands may have local B1 entries, but their candidate carrier sets can compete for a shared carrier x , so the chosen entries are not disjoint. Panel II records the finite minimality step of [Lemma 14.87](#): among all failed demand families, a minimal one gives a Type B overlap obstruction $\mathcal{O} = (\mathcal{D}, \{\mathcal{E}(d)\})$, and its involved fan centers, carriers, paths, and window segments form the connected overlap support $Z(\mathcal{O})$. Panel III displays the global constraints inherited by $Z(\mathcal{O})$ in [Lemma 14.89](#): contextual dyadic safety, high-degree separation, window compatibility, replacement/delocalization routing, and minimality of the overlap. Panel IV gives the routing consequence of [Proposition 14.90](#): if no such obstruction remains, the disjoint local choice exists. Its deleted boundary is subsequently bounded by the ambient degree sum of the selected support.

- (a) Each certificate-closed fan center is assigned a set of incident non-center vertices whose augmented charge, together with the center charge, is nonnegative. The center and the assigned non-center vertices are carriers of that entry, and no carrier is assigned to two different fan centers.
- (b) For any maximal canonical family $\mathcal{F}$ of positive-deficit Type B residual fans in X , each fan center has its hybrid B1 ledger entry: packed window incidences carry credit $1/2$, and non-window fan incidences carry the local internal/mixed half-credit needed to pay D_N . These incidence carriers are disjoint from the ordinary deficiency reserve of [Definition 14.83](#).
- (c) The certificate-closed entries in (a), the local reserve entries supplied by B1 for positive-deficit fans, and the half-incidence entries in (b) are mutually disjoint refinements of the augmented ledger $\widehat{\text{Ch}}_B(X)$ of [Definition 14.65](#). After the support assignment is forgotten, this ledger gives $\mathcal{N}(X)$ by the identity (B-ledger).
- (d) The ledger is maximal for the Type B support assignment: after the ledger entries are removed, the remaining core has no high-degree center, no decorated Type B handoff whose center belongs to X , and no external decorated handoff that is not already covered by a grouped decorated envelope processed in the same canonical step.

For a full B2 choice, let U be the used non-center vertices, put $P = U \cup H_X$ and $Y_0 = Y_X \setminus P$, and define

$$B_{\partial} = \sum_{v \in Y_0} |N_X(v) \cap P|.$$

A certificate entry uses at most eight non-center vertices. A positive-deficit entry uses at most eight fan neighbours and at most sixteen remote incidence endpoints. Hence

$$|U| \leq 24|H_X|, \quad |P| \leq 25|H_X|. \quad (\text{B-support-size})$$

Every vertex of P has ambient degree at most eight, and every $h \in H_X$ has $d_G(h) - 3 \geq 1$. The finite cut bound therefore gives

$$B_\partial \leq \sum_{x \in P} d_G(x) \leq 200|H_X| \leq 200 \sum_{h \in H_X} (d_G(h) - 3). \quad (\text{B-boundary-size})$$

The boundary demand consists of $4B_\partial$ quarter-charge units. The canonical finite transfer either injects these units into the processed quarter-charge slots or returns the strict numerical overload. This alternative is a consequence of the two finite cardinalities; it requires no enumeration of transfer functions.

A Type B bridge residual is either:

- (i) a connected assigned Type B support containing a fan-certificate residual center of [Definition 14.18](#); or
- (ii) a connected assigned Type B support containing a certificate-marked center for which no certificate-tied certificate-closed or positive-deficit local entry is available; or
- (iii) a connected assigned Type B support in which the local B1 entries have been supplied, the support assignment has been made maximal, and the disjoint-carrier part of B2 fails; or
- (iv) a grouped decorated Type B envelope support whose envelope ledger cannot be closed by the center tokens of its core-center incidence component.
- (v) a connected assigned Type B support for which the disjoint local choices exist but the boundary demand strictly exceeds the available processed quarter-charge slots.

In cases (iii) and (iv), the support contains a minimal Type B overlap obstruction in the sense of [Definition 14.85](#); by [Proposition 14.90](#), that obstruction carries the global-to-local constraints of [Lemma 14.89](#). Case (v) is a separate arithmetic residual. Its deficit is bounded by (B-boundary-size) and is not identified with an overlap obstruction.

Corollary 14.92 (The local hybrid entry realizes B1). *The local entry of [Lemma 14.79](#) is the hybrid fan-ledger alternative B1 in [Definition 14.91](#).*

Proof. The alternatives excluded in [Lemma 14.79](#) are exactly the first two alternatives in B1. In the remaining case that lemma supplies the disjoint hybrid entry with window credit $I_W/2$, remaining non-window demand D_N , and total capacity at least D_B , which is precisely the third alternative in the definition of B1. $\square$

Lemma 14.93 (Exact post-B2 charge decomposition). *Let X be an assigned Type B support and suppose that a full disjoint B2 choice has been made. Let U be the union of the literal counted-core vertices whose charges are used by its local entries. Thus $U \subseteq Y_X$, $U \cap H_X = \emptyset$, and the per-center subsets forming U are pairwise disjoint. Put*

$$Y_0 = Y_X \setminus (U \cup H_X).$$




Figure 28: Visual definition of the Type B bridge and fan-mass summation. Panel I shows the local B1 entry for a positive-deficit Type B fan centered at h_1 . Window incidences and non-window incidences contribute half-credit, and the B1 ledger pays the local deficit $D_B(\mathfrak{F}_{h_1})$. Panel II shows the global B2 disjointness requirement. Centers h_2 and h_3 may use their own carrier sets C_2 and C_3 , but a shared carrier x is forbidden by B2; a minimal unresolved sharing pattern is the minimal Type B overlap obstruction of Definition 14.85. Panel III shows the fan-mass summation used for Type B bridge residuals. A center without the required certificate enters the fan-certificate residual class, a route-8 non-window core is extracted into the Type A route-8 deficit ledger D_A , and the remaining Type B bridge entries are charged to assigned surplus units. After this extraction, the Type B bridge mass is sublinear in the near-cubic spine, so it cannot carry the linear large-budget deficit.

In quarter units, the selected fan ledger is nonnegative:

$$0 \leq \sum_{h \in H_X} (-4(d_G(h) - 3) - 1) + \sum_{u \in U} (4\delta_X^+(u) - 1). \quad (\text{B-selected})$$

Moreover the net charge has the exact decomposition

$$\begin{aligned} 4\mathcal{N}(X) = & \left[\sum_{h \in H_X} (-4(d_G(h) - 3) - 1) + \sum_{u \in U} (4\delta_X^+(u) - 1) \right] \\ & + \left[\sum_{h \in H_X} (4\delta_X^+(h) - 1) + |H_X| \right] + \sum_{v \in Y_0} (4\delta_X^+(v) - 1). \end{aligned} \quad (\text{B-postledger})$$

The second bracket equals $4 \sum_{h \in H_X} \delta_X^+(h)$ and is therefore nonnegative. Consequently either $\mathcal{N}(X) \geq 0$, or the literal remaining core has negative charge:

$$\sum_{v \in Y_0} (4\delta_X^+(v) - 1) < 0. \quad (\text{B-to-A})$$

Proof. For a certificate entry, (B-port-transfer) bounds every selected candidate weight by the actual charge of its selected core vertex. For a positive-deficit entry, the external-shoulder correction

pays the window credit, and each selected non-window half-credit is bounded by the charge of its ordinary available remote endpoint. Candidate carrier support contains all these vertices and every fan neighbour whose charge occurs in the local calculation. B2 therefore makes the corresponding subsets pairwise disjoint, and summing the local nonnegative balances gives (B-selected).

The sets U , H_X , and Y_0 are disjoint and partition Y_X . Split the ordinary core-charge sum over this partition, add the augmented center sum, and then use the exact identity

$$4\mathcal{N}(X) = 4\widehat{\text{Ch}}_B(X) + |H_X|.$$

This gives (B-postledger) by rearrangement. The asserted form of its second bracket is pointwise, since $(4\delta_X^+(h) - 1) + 1 = 4\delta_X^+(h)$. If the final sum is nonnegative, all three terms on the right of (B-postledger) are nonnegative; otherwise (B-to-A) holds. $\square$

Lemma 14.94 (Post-ledger Type A boundary identity). *Let X be a Type B support and let $P = U \cup H_X$ be the vertices removed by a disjoint local ledger. Put $Y_0 = Y_X \setminus P$, and let Z be a connected component of the graph induced by Y_0 . Every vertex of Y_0 has ambient degree 3. For $v \in Z$, let*

$$b_Z(v) = |N_X(v) \cap P|$$

be the number of original-core incidences deleted at v . Then

$$d_Z(v) = d_X(v) - b_Z(v), \quad \delta_Z^+(v) = \delta_X^+(v) + b_Z(v),$$

and hence, in quarter units,

$$\sum_{v \in Z} (4\delta_Z^+(v) - 1) = \sum_{v \in Z} (4\delta_X^+(v) - 1) + 4 \sum_{v \in Z} b_Z(v). \quad (\text{B-boundary})$$

The component Z is P_{13} -free and has no internal 3-core. Thus unsaturated Type A receiver discharging proves nonnegativity of the left-hand side of (B-boundary). It does not, by itself, prove nonnegativity of the first sum on the right.

Proof. The remaining graph is induced on Y_0 , so the only neighbors counted by $d_X(v)$ and not by $d_Z(v)$ are the neighbors in P . Since $v \notin H_X$, the minimum-degree-three baseline and the definition of H_X give $d_G(v) = 3$. Therefore $d_X(v) \leq 3$, no truncation occurs in $\delta^+ = 3 - d$, and the two pointwise identities follow. Summing the charge identity gives (B-boundary).

The graph on Z is an induced subgraph of the verified P_{13} -free remainder, so it is P_{13} -free. Any internal subgraph of minimum degree at least 3 in Z would also be such a subgraph of the remainder, contrary to the Hegde–Sandeep–Shashank consequence. The last assertion is then exactly the unsaturated 3/7/11 receiver calculation applied to Z . The term $4 \sum b_Z(v)$ is retained explicitly; dropping it would reverse the valid implication. $\square$

Definition 14.95 (Center-deleted Type A overload). *Let X be an assigned Type B support, let H_X be its complete set of assigned vertices of ambient degree at least 4, and put*

$$Y_X^\circ = V(X) \setminus H_X.$$

The induced graph on Y_X° is subcubic. For each of its degree-3 vertices choose a reachable receiver of induced degree at most 2, as in Definition 14.3. If $L_X(w)$ is the resulting receiver load, put

$$\Omega_A(X) = \sum_{w \in Y_X^\circ : d_{Y_X^\circ}(w) \leq 2} \max\{0, L_X(w) - (4(3 - d_{Y_X^\circ}(w)) - 1)\}.$$

Thus $\Omega_A(X)$ is the exact excess above the 3/7/11 receiver capacities for this fixed routing. It vanishes when every receiver is unsaturated.

Proposition 14.96 (Unconditional Type B deficit bound). *For every assigned Type B support X contained in the verified P_{13} -free remainder,*

$$\mathcal{N}_-(X) \leq \frac{21}{4} \sum_{h \in H_X} (d_G(h) - 3) + \frac{1}{4} \Omega_A(X). \quad (\text{B-unconditional})$$

No fan-certificate marking, local candidate, or B2 disjointness hypothesis is required.

Proof. Write

$$S_X = \sum_{h \in H_X} (d_G(h) - 3), \quad Q_X = 4\mathcal{N}(X).$$

Every vertex of Y_X° has ambient degree 3: the baseline gives the lower bound, while the definition of H_X excludes ambient degree at least 4. The graph induced by Y_X° is P_{13} -free and dyadic-safe. The Hegde–Sandeep–Shashank consequence therefore excludes an internal 3-core, so the receiver routing in [Definition 14.95](#) exists.

Let R_X be the raw quarter-charge of Y_X° , and let B_X be the number of edges of the induced support cut between Y_X° and H_X . The exact boundary identity gives induced quarter-charge $R_X + 4B_X$. The receiver-fibre calculation, without an unsaturated assumption, gives

$$-(R_X + 4B_X) \leq \Omega_A(X). \quad (1)$$

Deleting only the high centers in the augmented Type B identity gives

$$Q_X = (-4S_X - |H_X|) + \left(\sum_{h \in H_X} (4\delta_X^+(h) - 1) + |H_X| \right) + R_X.$$

The middle bracket is $4 \sum_{h \in H_X} \delta_X^+(h) \geq 0$. Moreover,

$$|H_X| \leq S_X, \quad B_X \leq \sum_{h \in H_X} d_G(h) \leq 4S_X,$$

because $d_G(h) \geq 4$ implies $d_G(h) \leq 4(d_G(h) - 3)$. Combining these inequalities with (1) yields

$$-Q_X \leq 4S_X + |H_X| + 4B_X + \Omega_A(X) \leq 21S_X + \Omega_A(X).$$

Taking the positive part and dividing by 4 proves (B-unconditional). $\square$

Proposition 14.97 (Quantitative Type B reduction under the refined ledger). *Let X be a connected admissible Type B support outside the route-8 Type A residual. Assume that X contains no fan-certificate residual center and that the refined support ledger B2 of [Definition 14.91](#) holds. Then either the remaining induced core contains a saturated Type A receiver, or X satisfies*

$$\max\{0, -\mathcal{N}(X)\} \leq 200 \sum_{h \in H_X} (d_G(h) - 3). \quad (\text{B-deficit})$$

More precisely, the unsaturated branch gives either $\mathcal{N}(X) \geq 0$, or a strict boundary-transfer overload together with (B-deficit). A successful boundary transfer gives $\mathcal{N}(X) \geq 0$.

Proof. Since X contains no fan-certificate residual center, every high-degree center assigned to X is certificate-marked. Use the canonical refined ledger supplied by bridge statement (B2). Its certificate-closed entries are nonnegative by Lemma 14.66, and B2(a) makes the corresponding non-center vertices disjoint across fan centers. A positive-deficit fan that realizes a direct cycle, quotient, compression, or delocalization is forbidden by the standing admissibility and minimality invariants. For every remaining positive-deficit fan $\mathfrak{F}$, the hybrid B1 ledger of Lemma 14.79 supplies a fan entry of capacity at least $D_B(\mathfrak{F})$: packed-window incidences pay their half-credit and non-window incidences pay the residual demand $D_N(\mathfrak{F})$. Bridge statement B2(b) makes these incidence carriers disjoint from the ordinary deficiency reserve of Definition 14.83, and B2(c) makes the resulting fan entries mutually disjoint from the certificate-closed entries and from the other positive-deficit fan entries. Clause (c) records these credits as entries of the same augmented support-charge ledger. Summing over the positive-deficit fans therefore supplies at least the total positive-deficit Type B deficit.

After the fan centers are processed, the remaining core has no high-degree surplus, no decorated Type B handoff whose center belongs to X , and no external decorated handoff left uncovered by the grouped envelope step, by the maximality clause B2(d). Apply Lemma 14.94 to its connected components. A saturated receiver is the first alternative in the statement. Assume that no such receiver exists. Then every component is in the unsaturated branch, and Lemma 14.43 gives nonnegative *induced* charge on each remaining component. Write S for the selected-fan bracket, C for the center-correction bracket, R for the raw remaining charge in (B-postledger), and $4B_\partial$ for the boundary term in (B-boundary). The Type A conclusion is

$$R + 4B_\partial \geq 0.$$

The selected and center-correction brackets are nonnegative, so

$$-4\mathcal{N}(X) = -(S + C + R) \leq -R \leq 4B_\partial.$$

By (B-boundary-size),

$$4 \max\{0, -\mathcal{N}(X)\} \leq 800 \sum_{h \in H_X} (d_G(h) - 3),$$

which is (B-deficit). The finite boundary-transfer alternative is obtained by comparing its required and available unit cardinalities. If the transfer exists, then $S + C \geq 4B_\partial$, and

$$4\mathcal{N}(X) = (S + C - 4B_\partial) + (R + 4B_\partial) \geq 0.$$

If it does not exist, the required-unit cardinality is strictly larger than the available-unit cardinality; this is the boundary-transfer overload. $\square$

Lemma 14.98 (Type B exclusion under the refined ledger). *Let X be a connected admissible support carrying high-degree surplus. Assume that X contains no fan-certificate residual center and admits the refined support ledger B2 of Definition 14.91. If X contains neither an admissible route-8 residual profile nor an admissible positive-deficit Type B fan-window residual, then one of the following holds: the post-ledger core contains a saturated Type A receiver; $\mathcal{N}(X) \geq 0$; or the boundary transfer is strictly overloaded and*

$$\max\{0, -\mathcal{N}(X)\} \leq 200 \sum_{h \in H_X} (d_G(h) - 3). \quad (\text{B-exclusion-deficit})$$

Consequently, a Type B support with $\mathcal{N}(X) < 0$ and with no saturated Type A, route-8, or positive-deficit fan-window residual is a Type B bridge residual: it witnesses either failure of the disjoint B2 choice or a strict boundary-transfer overload.

Proof. Use the assigned Type B notation of [Definition 14.65](#). The set H_X consists of the high-degree fan centers whose surplus units are assigned to X ; it is independent by [Lemma 3.2](#), and

$$\sigma(X) = \sum_{h \in H_X} (d_G(h) - 3).$$

The local calculation below proves that the selected fan bracket is nonnegative. Step 2 retains and bounds the induced-boundary correction.

Step 1 (fan degree is bounded; the closed fan neighbourhood carries nonnegative charge). Fix $h \in H_X$ and write $k = d_G(h) = |N(h)|$. Since X contains no fan-certificate residual center, the fan at h is certificate-marked in the sense of [Definition 14.18](#). By the fan-return-safety analysis above, $N(h)$ is a clique in the fan-safe graph $\mathcal{F}_h$, built from the five local safety relations (i)–(v).

Degree bound. The certificate labelling $S_h : N(h) \rightarrow \mathcal{L}$ satisfies $C_2(S_h(u), S_h(v)) = 1$ for every distinct $u, v \in N(h)$. Hence $S_h(N(h))$ is a clique in the fan-certificate graph of [Lemma 14.66](#). That lemma gives

$$k = d_G(h) \leq 8,$$

and the spectral inequality $d_G(h) \leq \rho(A(\mathcal{F}_h)) + 1$ is the quantitative form of this bound.

Closed-neighbourhood charge for the certificate-closed subcase. We use the augmented charge ch_X of [Definition 14.65](#). The center has

$$\text{ch}_X(h) = 3 - k - \frac{1}{4}.$$

Cubic-closed neighbours contribute $-\frac{1}{4}$, while non-cubic-closed fan neighbours have internal degree at most 2 in the assigned fan envelope and contribute at least $\frac{3}{4}$.

Let $\mathfrak{F}_h$ be the marked fan centered at h , let $c = c(\mathfrak{F}_h)$ be its number of cubic-closed neighbours, and put

$$D_B(\mathfrak{F}_h) = c - \frac{11 - k}{4}.$$

If $D_B(\mathfrak{F}_h) > 0$, then, because X is admissible, the fan-window data cannot realize a target cycle, target-defective quotient, target-complete compression, or delocalization. Moreover, if a cubic-closed neighbour u is same-window supported at $P = p_0 \cdots p_{12}$ with closed label $\{a, b\}$ and $b - a \in \{2, 6\}$, then it gives the cycle

$$up_a P p_b u$$

of length 4 or 8, and any two-neighbour same-window incompatibility gives a direct fan-window cycle by [Lemma 14.75](#). Hence the recorded same-window labels satisfy the compatibility requirements of [Definition 14.71](#). Thus [Definition 14.71](#) would make $\mathfrak{F}_h$ a positive-deficit Type B residual, contrary to the hypothesis of this lemma. Therefore

$$D_B(\mathfrak{F}_h) \leq 0.$$

The charge calculation is then explicit:

$$\text{ch}_X(h) + \sum_{u \in N(h)} \text{ch}_X(u) \geq \left(3 - k - \frac{1}{4}\right) - \frac{c}{4} + \frac{3}{4}(k - c) = \frac{11 - k}{4} - c = -D_B(\mathfrak{F}_h) \geq 0.$$

This proves nonnegative charge for every non-residual fan neighbourhood. In particular, when $k = 8$, the exclusion of the positive-deficit residual forces $c = 0$.

Step 2 (B2 ledger and residual core). Step 1 is a local calculation at a single fan center. It does not identify disjoint paying carriers for different centers. The required disjointness is exactly the

refined support ledger B2. By B2(a), every certificate-closed fan center is assigned a non-center vertex set whose augmented charge together with the center charge is nonnegative, and no carrier is assigned to two fan centers. Since the hypothesis excludes positive-deficit fan-window residuals, every marked fan centered in H_X is certificate-closed by Step 1. Thus all high-degree centers of X are covered by pairwise disjoint nonnegative ledger entries.

Remove these ledger entries and call the remaining support X_0 . By B2(d), X_0 has no high-degree center, no decorated Type B handoff whose center belongs to X , and no external decorated handoff left uncovered by the grouped envelope step. Apply Lemma 14.94 and decompose X_0 into connected admissible Type A components. A saturated receiver is the first alternative in the statement. Assume that no component has one. Then every receiver in every component of X_0 is unsaturated, and Lemma 14.43 gives nonnegative induced charge on those components. With the notation of Lemma 14.94, this is $R + 4B_\partial \geq 0$, not $R \geq 0$. The cardinal cut bound (B-boundary-size) gives $B_\partial \leq 200 \sum_{h \in H_X} (d_G(h) - 3)$. Since $S + C \geq 0$,

$$-4\mathcal{N}(X) = -(S + C + R) \leq -R \leq 4B_\partial.$$

This proves (B-exclusion-deficit). If the finite boundary transfer exists, then $S + C \geq 4B_\partial$ and the preceding identity instead gives $\mathcal{N}(X) \geq 0$. Hence a negative branch forces the strict cardinal overload. $\square$

Remark 14.99. *The fan-certificate labelling bounds the degree of a certificate-marked fan by the label-packing constant 8. A high-degree center without such a labelling is not placed in the certificate-closed local discharging calculation; it is a Type B bridge residual center and is charged by the fan-mass estimate below. For certificate-marked fans, the charge proof closes exactly the nonpositive-deficit cases, i.e. $D_B = c - (11 - k)/4 \leq 0$. Thus a fan with $k \leq 7$ becomes residual only when it has at least two cubic-closed neighbours, while an 8-fan becomes residual already with one cubic-closed neighbour. If that 8-fan neighbour were same-window supported with a non-singleton label, the non-singleton label cap would force $k \leq 7$; therefore any actual $k = 8, c = 1$ residual must be internal, mixed, or supported on two windows. Outside the positive-deficit residual of Definition 14.71, outside fan-certificate residual centers, and outside route 8, the exclusion requires the refined support ledger B2. That ledger supplies exactly the disjointness needed to sum the certificate-closed fan charges without counting any carrier twice.*

Degree-four Type B ledger. The degree-four remainder is the no-branch of node [68] in Fig. 6, expanded as node [78] in Fig. 7. Here

$$k = d_G(h) = 4, \quad d_G(h) - 3 = 1, \quad D_B(\mathfrak{F}_h) = c - \frac{7}{4},$$

where c is the number of cubic-closed neighbours of the fan center h . Thus the exact local possibilities are

c	0	1	2	3	4
$D_B = c - \frac{7}{4}$	$-\frac{7}{4}$	$-\frac{3}{4}$	$\frac{1}{4}$	$\frac{5}{4}$	$\frac{9}{4}$
closed-neighbour charge $-D_B$	$\frac{7}{4}$	$\frac{3}{4}$	$-\frac{1}{4}$	$-\frac{5}{4}$	$-\frac{9}{4}$
local B1 incidence capacity c	0	1	2	3	4

For $c \geq 2$, the hybrid incidence capacity pays the exact local deficit with slack

$$c - D_B = \frac{7}{4}.$$

The table below records what a degree-four Type B support must satisfy if it is still present after all named local exits have been tested. The word *survivor* means a connected assigned Type B support with negative net charge outside route 8 which has not been routed to the fan-mass residual node.

Constraint	Quantitative consequence when $k = 4$	What a survivor must have	What a survivor must not have	Route
High-degree bookkeeping	The center supplies exactly one surplus unit: $\sigma(h) = d_G(h) - 3 = 1$. By deletion-criticality, every neighbour of h has ambient degree 3, and high-degree centers are independent.	A single assigned degree-four fan center, with cubic fan neighbours in the assigned Type B support.	An edge from h to another high-degree center, or a non-cubic fan neighbour.	[68]→[78]→[79]
Local closed-neighbour charge	$D_B = c - \frac{7}{4}$. Hence $c = 0, 1$ gives nonnegative local charge, while $c = 2, 3, 4$ gives deficits $\frac{1}{4}, \frac{5}{4}, \frac{9}{4}$, respectively.	At least two cubic-closed neighbours if it is locally positive-deficit.	A negative local fan charge with $c \leq 1$; those cases are certificate-closed.	[79]→[81]
Certificate marking	If the fan has the certificate labelling $S_h : N(h) \rightarrow \mathcal{L}$, the degree cap $k \leq 8$ is automatic and the fan is eligible for the certificate/B1 ledger. If the labelling is absent, the center is charged as a bridge residual.	The labelling S_h if it is to proceed to the local fan ledger.	An unlabeled fan center remaining in the certificate-closed local calculation.	[80]→[84]
Exact local-entry resolution	For each marked center, the local entry must use the certificate fixed at node [80]. The finite resolution either supplies such an entry for every center or returns an actual center whose exact certificate has no verified certificate-closed or positive-deficit entry.	An exact certificate-tied local entry at every marked center before the B2 completion is attempted.	Replacing the node-[80] certificate by another witness, or assuming that an empty local-entry fibre is inhabited. Either failure is retained as a bridge residual and routed to [84].	[80]→[81]→[84]
Direct dyadic-cycle exclusions	Any same-window single label gap in $\{2, 6\}$ gives a 4- or 8-cycle. For two same-window closed neighbours, endpoint distances $\{0, 4, 12\}$ or interlace length 8 or 16 give a dyadic cycle. For two windows, $ i-j + a-b \in \{0, 4, 12\}$ gives a 4-, 8-, or 16-cycle.	All recorded same-window and two-window labels must satisfy exactly these arithmetic exclusions.	Any violating label pair; the displayed endpoint arithmetic constructs an explicit power-of-two cycle.	[81]

Constraint	Quantitative consequence when $k = 4$	What a survivor must have	What a survivor must not have	Route
Target-algebra exits	A positive-deficit fan is tested for target-defect, target-complete compression, and proper/global delocalization before the B1 ledger is used.	Target-compatible fan-window data after those exits fail.	A target-defective quotient, a target-complete smaller representative, or a proper/global delocalization.	[81]
Local B1 incidence budget	The c cubic-closed neighbours use $2c$ pairwise distinct non- h incidence carriers. Half-credit gives total capacity c , pays $D_B = c - \frac{7}{4}$, and leaves slack $\frac{7}{4}$. Window credit is $\frac{1}{2}I_W$; the remaining non-window demand is $D_N = \max\{0, D_B - \frac{1}{2}I_W\}$.	A local hybrid B1 entry for every positive-deficit fan $c = 2, 3, 4$.	A claimed local deficit with no assigned incidence carriers paying D_B .	[81]
Global B2 disjointness	B1 pays one fan at a time. To sum several degree-four fan entries, their candidate entries and half-incidence carriers must be pairwise disjoint and disjoint from the ordinary deficiency reserve.	A refined disjoint Type B ledger for all degree-four fan demands in the support. If it exists, exact post-ledger accounting gives either nonnegative assigned charge or a literal negative remaining core routed to route 8.	Double-use of a vertex, incidence carrier, or ordinary deficiency-reserve unit inside the asserted summed ledger.	[81]→[82]
B2 failure	If B1 exists locally but B2 cannot be chosen globally, the support contains a minimal Type B overlap obstruction. Every proper connected sub-obstruction has a disjoint ledger.	A minimal family of conflicting degree-four fan demands, with all proper subfamilies disjointly payable.	A nonminimal conflict, or a conflict removable by a quotient, compression, or delocalization already excluded by the target algebra.	[81]→[83]
Global-to-local reflection	The overlap support inherits contextual dyadic safety, high-degree separation, window-stub capacity, direct-cycle-free labels, replacement obstruction, and minimal overlap.	All those inherited constraints simultaneously.	A local picture that violates any inherited constraint; such a violation exits through a dyadic cycle, a target defect, compression, delocalization, or a smaller disjoint ledger.	[83]

Constraint	Quantitative consequence when $k = 4$	What a survivor must have	What a survivor must not have	Route
Fan-mass bound	For degree-four centers, $\sum_h (d_G(h) - 3)$ is just the number of assigned centers. If $m_4(X)$ is that number, then after any route-8 non-window core is extracted the bridge residual deficit satisfies $\mathcal{N}_-(X) \leq 208m_4(X)$. Summed in the near-cubic spine, allowing one canonical use and one grouped-envelope use of a center surplus, $M_B \leq 416\sigma(G) = O(\sqrt{n}) = o(R)$.	Only sublinear total Type B bridge mass, unless the support hands a non-window core to the Type A route-8 ledger.	A linear large-budget Type B deficit outside route 8 carried solely by degree-four bridge residuals.	[84]→[85]

Lemma 14.100 (Local selected-center fan-mass bound). *Fix one Type B support scope supplied by the exact local ledger, and let $S \subseteq H_X$ be either the singleton unresolved center or the literal set of centers in its minimal-overlap obstruction. Then*

$$|S| \leq \sum_{h \in S} (d_G(h) - 3) \leq \sum_{h \in H_X} (d_G(h) - 3).$$

The same local decision retains the nonnegative and route-(8) outcomes without changing their payloads.

Proof. The center list is the graph-computed list of actual high centers. Hence $d_G(h) \geq 4$ for every $h \in S$, so each selected center has unit lower mass at most its capacity $d_G(h) - 3$. Summing this pointwise inequality on the literal selected list proves the first bound. That list is duplicate free: in the overlap branch it is a sublist of the fixed center enumeration, and in the unresolved branch it is a singleton. The second bound is the inclusion $S \subseteq H_X$. This computation inspects only the fixed local center list; it does not construct a family of supports or envelopes. $\square$

The global ordinary/grouped aggregation is a separate obligation. In particular, [Lemma 14.100](#) does not construct the canonical support family and does not imply either coefficient-208 incidence bound used below. These are the open node-[84] producer contract.

Definition 14.101 (Type B residual fan mass). *Let $\mathcal{X}_B$ be a subcollection of the canonical admissible supports which are Type B bridge residuals in the sense of [Definition 14.91](#), together with any grouped decorated Type B envelope supports produced from Type A exit-(7) handoffs. Window-center envelopes are included in the same family; by [Lemma 14.22](#) they use the surplus $d_G(h) - 3$ of their high-degree window center, not the packed-window stub reserve. For an ordinary assigned Type B support X , let*

$$H_X = \{h : h \text{ is a high-degree Type B fan center assigned to } X\}$$

using the assigned-center convention of [Definition 14.91](#). For a grouped decorated envelope support $X = \mathfrak{X}_{\mathfrak{C}}^$, put $H_X = H_{\mathfrak{C}}$ and $\omega(h) = d_G(h) - 3$ for every $h \in H_{\mathfrak{C}}$. Let*

$$\mathcal{N}_-(X) = \max\{0, -\mathcal{N}(X)\}.$$

The Type B residual fan mass of $\mathcal{X}_B$ is

$$M_B(\mathcal{X}_B) = \sum_{X \in \mathcal{X}_B} \mathcal{N}_-(X), \quad H_B(\mathcal{X}_B) = \bigsqcup_{X \in \mathcal{X}_B} H_X, \quad S_B(\mathcal{X}_B) = \sum_{X \in \mathcal{X}_B} \sum_{h \in H_X} s_X(h),$$

where $s_X(h) = d_G(h) - 3$ for ordinary assigned Type B supports and $s_X(h) = \omega(h) = d_G(h) - 3$ for grouped decorated envelope supports. The symbol $\bigsqcup$ is a tagged union of center occurrences. Within the ordinary role, high-degree centers are disjoint by the canonical assignment in [Definition 13.2](#). Within the grouped-envelope role, centers are disjoint because the core-center incidence components of [Definition 14.20](#) partition the handoff-center set. The same vertex may occur once in each role, and this is the only allowed double occurrence. Hence each surplus unit is counted at most twice in S_B , and

$$S_B(\mathcal{X}_B) \leq 2\sigma(G).$$

Lemma 14.102 (Boundary size of processed Type B envelopes). *For each $h \in H_X$, let P_h consist of the center, its processed fan neighbours, and the remote endpoints of the processed non-center fan incidences. Put $P_X = \bigcup_{h \in H_X} P_h$. If $Y_0 = Y_X \setminus P_X$, then*

$$B_\partial(X) := \sum_{v \in Y_0} |N_X(v) \cap P_X| \leq 200 \sum_{h \in H_X} (d_G(h) - 3). \quad (\text{B-envelope-boundary})$$

The assertion remains valid when the sets P_h overlap.

Proof. Fix $h \in H_X$ and write $k = d_G(h) \geq 4$. The envelope uses at most k fan neighbours and at most two remote endpoints for each of at most k cubic-closed fan neighbours. Thus $P_h \setminus \{h\}$ has at most $3k$ vertices. Every such vertex is an ordinary processed core vertex and has ambient degree three; the center has degree k . Consequently

$$\sum_{x \in P_h} d_G(x) \leq k + 3(3k) = 10k \leq 200(k - 3).$$

The last inequality holds for every $k \geq 4$. Every edge counted by $B_\partial(X)$ has exactly one endpoint in P_X , so the finite cut bound and subadditivity over the possibly overlapping family give

$$B_\partial(X) \leq \sum_{x \in P_X} d_G(x) \leq \sum_{h \in H_X} \sum_{x \in P_h} d_G(x) \leq 200 \sum_{h \in H_X} (d_G(h) - 3).$$

□

Lemma 14.103 (Bounded deficit for grouped decorated envelopes). *Let $X = \mathfrak{X}_\mathfrak{C}^*$ be a grouped decorated Type B envelope support with center set $H_\mathfrak{C}$. Assume that the non-window core left after deleting the grouped fan envelopes of $\mathfrak{C}$ contains neither a saturated Type A receiver nor an admissible route-8 Type A residual profile. Then*

$$\mathcal{N}_-(X) \leq 208 \sum_{h \in H_\mathfrak{C}} (d_G(h) - 3).$$

Proof. All Type A cores and all centers in the incidence component $\mathfrak{C}$ are grouped before the Type B ledger is evaluated. Thus each center $h \in H_\mathfrak{C}$ contributes its negative center term once in the grouped envelope, and the fan neighbours and half-incidence carriers are taken in the same component ledger. The local calculation at a fixed center is the same as the one used for an

ordinary Type B bridge center: with $k = d_G(h)$ and $c \leq k$ cubic-closed neighbours, the negative part contributed by that center envelope is at most

$$\left(k - 3 + \frac{1}{4}\right) + \frac{c}{4} \leq \frac{5}{4}k - \frac{11}{4} \leq 8(k - 3),$$

the last inequality being equivalent to $27k \geq 85$. Summing this estimate over $h \in H_{\mathfrak{C}}$ bounds the envelope contribution below by $-8 \sum_h (d_G(h) - 3)$. Grouping by incidence components prevents either a core with several handoff centers or a center serving several cores from being counted more than once inside the envelope family.

After the grouped fan envelopes of $\mathfrak{C}$ are removed, no handoff incident with a core of $\mathcal{Y}_{\mathfrak{C}}$ remains uncovered; any such handoff would be an edge of the same incidence component. By the hypothesis, the remainder has no route-8 Type A residual profile. [Lemma 14.94](#) decomposes the remainder into admissible Type A components, and [Lemma 14.43](#) gives nonnegative *induced* charge on those components. If R denotes their raw charge, (B-boundary) gives

$$R \geq -B_{\partial}(X).$$

By [Lemma 14.102](#), $B_{\partial}(X) \leq 200 \sum_h (d_G(h) - 3)$. Adding this boundary loss to the eight-unit envelope bound gives

$$\mathcal{N}(X) \geq -208 \sum_{h \in H_{\mathfrak{C}}} (d_G(h) - 3),$$

which proves the assertion. $\square$

Lemma 14.104 (Grouped decorated envelope bound with route-8 cores extracted). *Let $X = \mathfrak{X}_{\mathfrak{C}}^*$ be a grouped decorated Type B envelope support with center set $H_{\mathfrak{C}}$. After deleting the grouped fan envelopes of $\mathfrak{C}$, let $\mathcal{A}_X$ be the canonical collection of Type A route-8 residual supports contained in the remaining non-window core. Assume that every remaining component outside $\mathcal{A}_X$ is unsaturated. Then*

$$\mathcal{N}(X) \geq -D_A(\mathcal{A}_X) - 208 \sum_{h \in H_{\mathfrak{C}}} (d_G(h) - 3).$$

Proof. The center-envelope estimate in the proof of [Lemma 14.103](#) gives total unpaid center-envelope contribution at least

$$-8 \sum_{h \in H_{\mathfrak{C}}} (d_G(h) - 3).$$

The processed-envelope cut contributes at worst

$$-B_{\partial}(X) \geq -200 \sum_{h \in H_{\mathfrak{C}}} (d_G(h) - 3)$$

by [Lemma 14.102](#) and (B-boundary). After the grouped fan envelopes of $\mathfrak{C}$ are removed, no decorated handoff incident with a core of the incidence component remains uncovered. [Lemma 14.94](#) decomposes the remaining non-window core into admissible Type A components. The components in $\mathcal{A}_X$ are exactly those carrying route-8 residual profiles. Every other component has no route-8 residual profile and is unsaturated by hypothesis, so [Lemma 14.43](#) gives nonnegative induced charge on it.

For $Y \in \mathcal{A}_X$, one has $\sigma(Y) = 0$, and hence

$$\mathcal{N}(Y) = \partial_+(Y) - \frac{1}{4}|V(Y)| = -\left(\frac{1}{4}|V(Y)| - \partial_+(Y)\right).$$

Summing over $Y \in \mathcal{A}_X$ gives the route-8 core contribution $-D_A(\mathcal{A}_X)$. Combining this with the center-envelope and boundary estimates gives the displayed inequality. $\square$

Lemma 14.105 (Bounded deficit per Type B bridge center). *Let X be a connected assigned Type B bridge residual. Assume that the non-window core left after deleting the Type B fan envelopes contains neither a saturated Type A receiver nor an admissible route-8 Type A residual profile. Then*

$$\mathcal{N}_-(X) \leq 208 \sum_{h \in H_X} (d_G(h) - 3).$$

Proof. The estimate uses the augmented ledger of [Definition 14.65](#). It does not assume that the B2 disjoint-carrier ledger exists. It bounds the unpaid negative part left by a failure of B2.

Fix a center $h \in H_X$, and write

$$k = d_G(h), \quad t_h = k - 3 \geq 1.$$

Let c be the number of cubic-closed neighbours in the fan envelope at h ; if h is not certificate-marked, c is still the number of fan neighbours whose two non- h incidences are assigned to the envelope. In all cases $c \leq k$. Let E_h be the fan envelope consisting of h , its assigned fan neighbours in Y_X , and the non- h incidences of the cubic-closed neighbours that appear in the local B1 candidate entry when such an entry exists. In the augmented ledger the center contributes

$$-(k - 3) - \frac{1}{4}.$$

Each cubic-closed neighbour contributes $-\frac{1}{4}$. Every other fan neighbour has internal degree at most 2 in the assigned fan envelope and contributes at least $\frac{3}{4}$, and every incidence carrier in the B1 entry carries nonnegative capacity. Therefore the negative part of E_h is bounded by

$$\left(k - 3 + \frac{1}{4}\right) + \frac{c}{4} \leq \left(k - 3 + \frac{1}{4}\right) + \frac{k}{4} = \frac{5}{4}k - \frac{11}{4} \leq 8(k - 3) = 8t_h. \quad (1)$$

The last inequality is equivalent to $27k \geq 85$, and holds for every $k \geq 4$.

If h is not certificate-marked, there is no local certificate-closed entry to choose; by definition it is already a Type B bridge residual center. The bound (1) is the only estimate used for that center. If h is certificate-marked and certificate-closed, the closed-neighbourhood calculation ([Lemmas 14.66](#) and [14.98](#)) supplies a nonnegative candidate entry. If h is certificate-marked and positive-deficit, [Lemma 14.79](#) supplies a local hybrid incidence entry whose capacity is at least $D_B(\mathfrak{F}_h)$. Thus every center is either bounded directly by (1), or has a local candidate envelope E_h whose unpaid negative part is bounded by (1). A bridge residual is a certificate failure or a failure to choose the local candidate entries with disjoint carriers after the support assignment is maximal. Omitting a contested carrier deletes only one of the nonnegative carrier capacities listed in [Definition 14.83](#); the negative terms that remain unpaid are contained in the corresponding envelopes E_h . Hence, after all overlaps are ignored rather than resolved,

$$\widehat{\text{Ch}}_B(X) \geq -8 \sum_{h \in H_X} (d_G(h) - 3). \quad (2)$$

After the fan-envelope terms are removed, the residual core has no high-degree center, no decorated Type B handoff whose center belongs to X , and no external decorated handoff left uncovered by the grouped envelope step, by the maximality clause in [Definition 14.91](#). By the hypothesis of the lemma it also has no route-8 Type A residual profile. [Lemma 14.94](#) decomposes it into admissible Type A components. Therefore the Type A saturated handoff [Lemma 14.40](#) and

the unsaturated discharge [Lemma 14.43](#) give nonnegative induced charge on each component. Their raw contribution to the original Type B ledger is therefore at least $-B_\partial(X)$. By [Lemma 14.102](#),

$$B_\partial(X) \leq 200 \sum_{h \in H_X} (d_G(h) - 3).$$

Combining this with (2) and the nonnegative center correction in (B-ledger) gives

$$\mathcal{N}(X) \geq -208 \sum_{h \in H_X} (d_G(h) - 3).$$

This is the asserted bound on $\mathcal{N}_-(X)$. □

Lemma 14.106 (Type B bridge bound with route-8 cores extracted). *Let X be a connected assigned Type B bridge residual. After deleting the Type B fan envelopes, let $\mathcal{A}_X$ be the canonical collection of Type A route-8 residual supports contained in the remaining non-window core. Assume that every remaining component outside $\mathcal{A}_X$ is unsaturated. Then*

$$\mathcal{N}(X) \geq -D_A(\mathcal{A}_X) - 208 \sum_{h \in H_X} (d_G(h) - 3).$$

Proof. Fix a high-degree center $h \in H_X$ and put

$$k = d_G(h), \quad t_h = k - 3.$$

Then $k \geq 4$ and $t_h \geq 1$. Let c_h be the number of cubic-closed fan neighbours whose two non- h incidences are assigned to the fan envelope at h . Since all such neighbours are neighbours of h , one has $0 \leq c_h \leq k$. In the augmented Type B ledger the center contributes negative charge

$$-(k - 3) - \frac{1}{4},$$

and the c_h cubic-closed neighbours contribute total negative charge $-c_h/4$. All other fan neighbours have internal degree at most 2 in the assigned fan envelope and contribute at least $3/4$, while the B1 incidence carriers have nonnegative capacity. Therefore the negative part of the fan envelope assigned to h is at most

$$\left(k - 3 + \frac{1}{4}\right) + \frac{c_h}{4} \leq \left(k - 3 + \frac{1}{4}\right) + \frac{k}{4} = \frac{5}{4}k - \frac{11}{4}.$$

For $k \geq 4$,

$$\frac{5}{4}k - \frac{11}{4} \leq 8(k - 3),$$

because this inequality is equivalent to $27k \geq 85$. Summing over $h \in H_X$ gives the fan-envelope bound

$$\widehat{\text{Ch}}_B(X) \geq -8 \sum_{h \in H_X} (d_G(h) - 3)$$

before the remaining non-window core is discharged.

The raw charge of that core differs from its induced Type A charge by the processed-envelope boundary. By [Lemma 14.102](#), this costs at most

$$200 \sum_{h \in H_X} (d_G(h) - 3).$$

The maximality clause in [Definition 14.91](#) leaves the remaining core with no high-degree center, no decorated Type B handoff whose center belongs to X , and no external decorated handoff left uncovered by the grouped envelope step. By [Lemma 14.94](#), decompose the core into admissible Type A components. The pieces in $\mathcal{A}_X$ are exactly the components carrying route-8 residual profiles. Every other Type A component is unsaturated by hypothesis, so [Lemma 14.43](#) gives nonnegative induced charge on it.

For a piece $Y \in \mathcal{A}_X$, $\sigma(Y) = 0$, hence

$$\mathcal{N}(Y) = \partial_+(Y) - \frac{1}{4}|V(Y)| = -\left(\frac{1}{4}|V(Y)| - \partial_+(Y)\right).$$

Summing over $Y \in \mathcal{A}_X$ gives induced core contribution $-D_A(\mathcal{A}_X)$. Combining this contribution with the fan-envelope bound, the processed-boundary cost, and the identity (B-ledger) gives the displayed inequality. $\square$

Proposition 14.107 (Type B bridge residuals are sublinear). *Assume [Definition 4.1](#). Let $\mathcal{X}_B$ be any canonical subcollection of Type B bridge residual supports whose non-window cores contain neither a saturated Type A receiver nor an admissible route-8 Type A residual profile. Then*

$$M_B(\mathcal{X}_B) \leq 208 S_B(\mathcal{X}_B) \leq 416 \sigma(G) = O(\sqrt{n}) = o(|R|).$$

Proof. For an ordinary assigned Type B bridge residual, use [Lemma 14.105](#). For a grouped decorated envelope residual, use [Lemma 14.103](#). Summing these estimates over $X \in \mathcal{X}_B$ gives $M_B(\mathcal{X}_B) \leq 208 S_B(\mathcal{X}_B)$. The at-most-twice convention in [Definition 14.101](#) gives

$$S_B(\mathcal{X}_B) \leq 2\sigma(G) = 2 \sum_{v \in V(G)} (d_G(v) - 3).$$

The near-cubic spine hypothesis gives $\sigma(G) = O(\sqrt{n})$ ([Definition 4.1](#)), while [Remark 7.67](#) gives $|R| \geq 0.83489768623 \dots n - o(n)$. Hence $\sigma(G) = o(|R|)$, and the displayed estimate follows. $\square$

Framework branch table. The Type B fan-safe surplus obstruction is discharged by the following framework branch table.

Branch	Residual condition	Demand	Payer	Bounded multiplicity	Closure
Fan-compatible pair	An open pair closes two ports at a high center.	Positive-deficit fan data.	The Type B local ledger.	One fan center.	Corollary 14.81 and Proposition 14.80 .
Triangular ports	The heavy center has triangular alternatives.	Fan-closed ports.	The Type B ledger.	Triangular shoulder cases.	Corollary 14.82 .
Certificate-closed fan	The local fan ledger has nonpositive deficit.	Local deficit.	The fan-certificate calculation.	One marked fan.	Lemmas 14.66 and 14.98 .

Positive-deficit fan	The certificate-closed calculation does not pay the fan.	Hybrid incidence demand.	B1 half-credit incidences.	Window and non-window incidence carriers.	Lemmas 14.78 and 14.79.
B2 failure	The disjoint ledger cannot be selected globally.	A minimal overlap obstruction.	The fan-mass ledger.	Assigned surplus units.	Definition 14.85 and Proposition 14.90.
Bridge residual	Non-window cores carry no route-8 profile.	Bridge mass.	Global surplus.	At most twice assigned surplus.	Proposition 14.107.

15 The branch-closure theorem

Theorem 15.1 (Entropy-admissible closure outside explicit residuals). *Assume Definition 4.1. The two local obstruction classes outside the explicit residuals isolated in Section 14 are empty:*

- (a) **Type A exclusion** (Lemma 14.45). *Every admissible subcubic P_{13} -free dyadic-safe atom C with $\sigma(C) = 0$ and empty internal 3-core, no exit-(4) witness for a routed load, no admissible route-8 residual profile, and no decorated Type B handoff satisfies*

$$\partial_+(C) \geq \frac{1}{4}|V(C)|.$$

- (b) **Type B bridge reduction** (Proposition 14.97, with Lemma 14.98 as the certificate-closed specialization). *Every connected admissible high-degree fan support X outside the Type B bridge residual and containing no admissible route-8 residual profile satisfies*

$$\partial_+(X) - \sigma(X) \geq \frac{1}{4}|V(X)|.$$

Then every large-budget residual outside the target-defect exit, outside the Type A route-8 residual branch, and outside the Type B bridge residual branch is closed.

Proof. If the high-entropy remaining budget is small, Proposition 12.2 closes that high-entropy branch. In every branch that survives to the large-budget analysis, including the low-entropy branches of Proposition 11.9, the near-cubic window-only cap gives

$$\frac{\partial_+(R) - \sigma_R}{|R|} = \frac{\partial_+(R) - \sigma(R)}{|R|} \leq \tau_{\text{win}} + o(1).$$

Since $\tau_{\text{win}} < 1/4$, the local lower bound

$$\partial_+(R) - \sigma(R) \geq \frac{1}{4}|R| - o(|R|)$$

would contradict the cap. More explicitly, for all sufficiently large n ,

$$\partial_+(R) - \sigma(R) \leq \left(\frac{1}{4} - \varepsilon\right)|R|, \quad \varepsilon = \frac{1}{2}\left(\frac{1}{4} - \tau_{\text{win}}\right) > 0.$$

Thus the hypothesis of Proposition 13.8 holds, and if the branch survives there is a connected admissible support X with $\mathcal{N}(X) < 0$.

By the Type A/Type B analysis, such an X is either a low-deficiency atom or a high-degree fan-safe surplus support. If X is Type A and contains an exit-(4) witness for a routed load, then it leaves the pure Type A net-charge calculation through the target-defect exit of [Definition 14.33](#) and [Lemma 14.38](#); this exit is excluded in the present branch. If X is Type A and contains the route-8 residual profile, then it lies in the explicit Type A route-8 residual branch, which is excluded by the hypothesis of this theorem. If X is Type A and produces a decorated Type B handoff, the handoff is transferred to the grouped decorated Type B envelope ledger by [Lemma 14.21](#); if the handoff center lies in a packed window, [Lemma 14.22](#) charges it to the center's own surplus token or gives exit (3). Outside the target-defect exit, the Type A route-8 branch, and the Type B bridge residual, [Lemma 14.45](#) excludes the pure Type A case. In the Type B case, the local B1 ledger is [Lemma 14.79](#), and being outside the bridge residual means that B2 holds. Hence [Proposition 14.97](#) gives either $\mathcal{N}(X) \geq 0$ or a strict boundary-transfer overload. The latter is itself a Type B bridge residual, so only the first alternative remains; in the certificate-closed specialization this is exactly [Lemma 14.98](#). This contradicts $\mathcal{N}(X) < 0$. Hence every large-budget residual outside the target-defect exit and the explicit residual branches is closed. $\square$

Framework branch table. The branch-closure theorem uses the following branch table.

Branch	Residual condition	Demand	Payer	Bounded multiplicity	Closure
High entropy	The admissible entropy cap is active.	Window, remainder, and curvature bits.	The skeleton budget.	One coordinate family.	Proposition 12.2 .
Negative support	The large-budget branch survives.	Negative net charge.	Local support analysis.	The canonical support.	Proposition 13.8 .
Pure Type A	The support avoids exit (4), route 8, and handoff.	Low-deficiency atom charge.	Type A discharging.	The receiver partition.	Lemma 14.45 .
Type B	The support avoids the bridge residual.	Fan-safe surplus support charge.	B1 and B2 ledger entries.	Assigned fan centers.	Proposition 14.97 .
Explicit residuals	Target defect, Type A route 8, or Type B bridge mass remains.	Residual mass.	Pressure and bridge ledgers.	Global accounting.	Routed to Theorem 15.48 and Proposition 14.107 .

15.1 External pressure accounting for the target-defect residual

Throughout assume [Definition 4.1](#). The purpose of this section is to account, in the large-budget net-charge ledger, for the supports that leave the pure Type A branch through the target-defect exit of [Definition 14.33](#). Such a support has strictly negative net charge (a peeled routed load

is a cubic vertex still carrying $-\frac{1}{4}$ in $\sum_v \text{ch}(v)$), so it is neither a nonnegative support nor, in general, a route-8 residual nor a Type B bridge residual. Any partition that counts only those three classes omits the target-defect negative mass. The large-budget ledger therefore sums over a single collection that contains *every* negative Type A support, and the carrier reduction is applied to that unified collection.

15.1.1 The unified negative Type A collection

Definition 15.2 (Unified negative collection). *Let R be the P_{13} -free remainder of a minimal counterexample satisfying [Definition 4.1](#), decomposed into the canonical connected admissible supports of [Definition 13.2](#). Let*

$$\tilde{\mathcal{X}} = \{ X : \sigma(X) = 0, \mathcal{N}(X) < 0, X \text{ produces no decorated Type B handoff} \}.$$

For $X \in \tilde{\mathcal{X}}$ write $\delta(X) = -\mathcal{N}(X) = \frac{1}{4}|V(X)| - \partial_+(X) > 0$ and set

$$\tilde{D}_A = \sum_{X \in \tilde{\mathcal{X}}} \delta(X).$$

Remark 15.3. $\tilde{\mathcal{X}}$ contains, simultaneously, the Type A supports carrying an admissible route-8 residual profile and the Type A supports that leave through the target-defect exit; the two are not distinguished at the level of net charge. A support that hands off to Type B is excluded from $\tilde{\mathcal{X}}$ and is treated in the Type B bridge ledger of [Definition 14.101](#).

Lemma 15.4 (The unified collection carries the whole large-budget deficit). *In the large-budget branch,*

$$\tilde{D}_A \geq \left(\frac{1}{4} - \tau_{\text{win}}\right)|R| - o(|R|).$$

Proof. By [Lemma 13.5](#),

$$\sum_i \mathcal{N}(X_i) = \partial_+(R) - \sigma(R) - \frac{1}{4}|R| \leq -\left(\frac{1}{4} - \tau_{\text{win}}\right)|R| + o(|R|),$$

the inequality being the large-budget net-deficiency cap obtained from [Remark 7.67](#) and [Lemma 8.4](#). Split the canonical supports into three groups: the members of $\tilde{\mathcal{X}}$; the Type B bridge residual ledger entries $\mathcal{X}_B$; and all remaining supports. A remaining support has $\sigma(X) = 0$ and $\mathcal{N}(X) \geq 0$, or is a nonnegative Type B support; in either case it contributes nonnegatively. Hence

$$\sum_i \mathcal{N}(X_i) \geq -\tilde{D}_A - \sum_{X \in \mathcal{X}_B} \mathcal{N}_-(X) \geq -\tilde{D}_A - 208S_B(\mathcal{X}_B) - o(|R|) \geq -\tilde{D}_A - o(|R|),$$

using $\mathcal{N}(X) = -\delta(X)$ on $\tilde{\mathcal{X}}$, [Proposition 14.107](#) for the middle inequality, and the at-most-twice bound $S_B(\mathcal{X}_B) \leq 2\sigma(G) = o(|R|)$ of [Definition 14.101](#) for the last. Combining the two displays gives the claim. $\square$

15.1.2 Every unit of deficit is a silent-excess entry with a nonempty carrier core

The next lemma is the point of the reformulation: it shows that the deficit of *every* $X \in \tilde{\mathcal{X}}$ – route-8 or target-defect – is carried by silent-excess routed loads, each of which possesses a trace basin, and hence a canonical essential-carrier core. It uses only counting and the dyadic-safety of admissible supports; it does not use target-complete-minimality.

Definition 15.5 (Unified indexed entries). For $X \in \tilde{\mathcal{X}}$ and a saturated receiver w of X , let $\mathcal{U}_X(w)$ be its silent unpaid routed loads (Definition 14.29). For $u \in \mathcal{U}_X(w)$ let B_u be its trace basin (Definition 14.14). The set of unified indexed entries is

$$\tilde{\Xi} = \{ \xi = (X, w, u, B_u) : X \in \tilde{\mathcal{X}}, u \in \mathcal{U}_X(w) \}, \quad \tilde{N} = |\tilde{\Xi}|.$$

An entry is a route-8 entry if B_u is target-complete-minimal (Definition 14.14); otherwise it is a target-defect entry, and one of the four failure alternatives of Definition 14.14 holds for B_u . The set $\tilde{\Xi}$ is a tagged Type A entry family in the sense of Definition 14.48. All carrier cores and privacy counts in the unified ledger are therefore evaluated relative to the full tagged family:

$$\alpha_{\tilde{\Xi}}(\xi) = |\mathcal{C}_{\text{ess}}(\xi)|, \quad \pi_{\tilde{\Xi}}(\xi) = |\{c \in \mathcal{C}_{\text{ess}}(\xi) : c \text{ is private relative to } \tilde{\Xi}\}|.$$

Lemma 15.6 (Unified burden).

$$\tilde{N} \geq 4\tilde{D}_A.$$

Proof. Fix $X \in \tilde{\mathcal{X}}$. Because X produces no Type B handoff, no saturated receiver of X realizes exit (7); because X lies in a minimal counterexample, none realizes exits (1)–(3), (5), or (6), since each of those is a standing-invariant contradiction (Lemma 14.39). Thus every saturated receiver of X realizes only exit (4) or exit (8); in particular no completion port carries four visible receiver-entry returns. Hence Lemma 14.30 applies and gives

$$S_{\text{sil}}^{\text{exc}}(X) \geq 4\left(\frac{1}{4}|V(X)| - \partial_+(X)\right) = 4\delta(X).$$

By Lemma 14.31, every vertex counted by $S_{\text{sil}}^{\text{exc}}(X)$ lies in some $\mathcal{U}_X(w)$ and thus indexes an entry of $\tilde{\Xi}$ (its basin is target-complete-minimal, giving a route-8 entry, or realizes one of exits (4)–(7); exits (5),(6) are contradictions and exit (7) is excluded, so the only alternative to route 8 is a target-defect entry via exit (4)). Summing over $X \in \tilde{\mathcal{X}}$,

$$\tilde{N} = \sum_{X \in \tilde{\mathcal{X}}} \sum_w |\mathcal{U}_X(w)| \geq \sum_{X \in \tilde{\mathcal{X}}} S_{\text{sil}}^{\text{exc}}(X) \geq 4 \sum_{X \in \tilde{\mathcal{X}}} \delta(X) = 4\tilde{D}_A.$$

□

Lemma 15.7 (Minimality-free carrier lower bound). Every entry $\xi = (X, w, u, B_u) \in \tilde{\Xi}$, route-8 or target-defect, has

$$\alpha_{\tilde{\Xi}}(\xi) = |\mathcal{C}_{\text{ess}}(\xi)| \geq 2.$$

Proof. For a route-8 entry this is Lemma 14.58. Let ξ be a target-defect entry and put $\mathcal{C} = \mathcal{C}_{\text{ess}}(\xi)$. Suppose for contradiction that $|\mathcal{C}| \leq 1$.

Because ξ is a target-defect entry, one of the four alternatives of Definition 14.14 holds for B_u . Alternatives (b)–(d) are exits (5),(6),(7): the first two are standing-invariant contradictions (Corollary 3.105 and Lemmas 10.6 and 10.9) and the third is excluded because $X \in \tilde{\mathcal{X}}$ produces no handoff. Hence alternative (a) holds: a trace-local non-carrier quotient q_0 of $\rho_u(B_u)$ is target-defective, i.e. distinguished by a compatible outside context. By definition of a trace-local non-carrier quotient (Definition 14.14), q_0 preserves the boundary degree profile and forgets an internal coordinate whose declared support contains an internal edge of B_u .

Apply Lemma 14.57 to ξ with the carrier set $\mathcal{C}$ and the quotient $\rho_{\mathcal{C}}^\circ$ that forgets exactly the u -supported coordinates whose declared support contains an internal edge of B_u . Since q_0 is target-defective and is one such forgetting quotient, the core-to- ρ° quotient $\rho_u(B_u)|_{\mathcal{C}} \rightarrow \rho_{\mathcal{C}}^\circ$ is not target-complete. The lemma then supplies a surviving u -supported target event realized by a simple

edge-rooted return or simple cycle that uses an internal edge of B_u and an edge outside X . Add the root edge if the event is edge-rooted. The result is a simple cycle which uses an edge of X and an edge outside X . It therefore crosses the cut $\delta_G(V(X))$ a positive even number of times. Because the cycle is simple, two distinct cut edges occur. The event survives in the carrier-restricted state, so both cut crossings are recorded by declared coordinates retained by $\mathcal{C}$. Thus its declared support contains at least two distinct carriers of $\mathcal{C}$, so $|\mathcal{C}| \geq 2$, contradicting $|\mathcal{C}| \leq 1$. Therefore $\alpha_{\tilde{\Xi}}(\xi) \geq 2$. $\square$

15.1.3 Unified carrier reduction

Proposition 15.8 (Unified two-carrier reduction). *In the large-budget branch, $\tilde{\Xi}$ contains a two-carrier entry, that is, an entry ξ with $\pi_{\tilde{\Xi}}(\xi) \leq 2$.*

Proof. Suppose not; then every $\xi \in \tilde{\Xi}$ has $\pi_{\tilde{\Xi}}(\xi) \geq 3$. By Lemma 15.7 the essential-carrier cores are well defined and, by definition of privacy (Definition 14.48), private essential carriers of distinct entries are disjoint oriented boundary incidences of their supports. Hence

$$3\tilde{N} \leq \sum_{X \in \tilde{\mathcal{X}}} |\partial_E X| = \sum_{X \in \tilde{\mathcal{X}}} \partial_+(X) \leq \partial_+(R) \leq \tau_{\text{win}}|R| + o(|R|),$$

using $|\partial_E X| = \partial_+(X)$ for ambient-cubic supports and the near-cubic window-only cap $\partial_+(R) \leq \tau_{\text{win}}|R| + o(|R|)$ of Corollary 8.7. On the other hand, Lemmas 15.4 and 15.6 give

$$\tilde{N} \geq 4\tilde{D}_A \geq 4\left(\frac{1}{4} - \tau_{\text{win}}\right)|R| - o(|R|).$$

Combining,

$$12\left(\frac{1}{4} - \tau_{\text{win}}\right)|R| - o(|R|) \leq \tau_{\text{win}}|R| + o(|R|),$$

which forces $12\left(\frac{1}{4} - \tau_{\text{win}}\right) \leq \tau_{\text{win}}$, i.e. $\tau_{\text{win}} \geq \frac{3}{13}$. This contradicts $\tau_{\text{win}} = 0.22817\dots < 0.23076\dots = \frac{3}{13}$ (Remark B.1). Hence a two-carrier entry exists. $\square$

Remark 15.9 (The margin is unchanged). *Proposition 15.8 spends the boundary-incidence budget $\partial_+(R) \leq \tau_{\text{win}}|R|$ exactly once, at three carriers per entry, and draws no additional boundary or window-stub capacity. The tight inequality is still $\tau_{\text{win}} < \frac{3}{13}$ with the 1.12% relative margin of Remark B.1; no new capacity is consumed, so that margin is preserved verbatim.*

15.1.4 External pressure accounting for target-defect entries

The route-8 two-carrier case is already closed in the manuscript: if B_u is target-complete-minimal and $\pi_{\tilde{\Xi}}(\xi) \leq 2$, the declared deletion witnesses make $(\tilde{\Xi}, \xi)$ a terminal two-carrier route-8 obstruction, which is excluded by Theorem 14.61. The pressure accounting below treats only the other two-carrier outcome, namely an entry whose trace basin fails target-complete-minimality by an exit-(4) target-defective quotient.

Definition 15.10 (Exit-(4) pressure token). *An exit-(4) pressure token is a tuple*

$$\mathbf{t} = (\xi, q, S_0, S_1, Y, E), \quad \xi = (X, w, u, B_u) \in \tilde{\Xi},$$

with the following data.

(P1) ξ is a target-defect entry: B_u is not target-complete-minimal, and the surviving failure alternative is alternative (a) of Definition 14.14.

(P2) q is the corresponding trace-local non-carrier quotient. It preserves the boundary degree profile of $\rho_u(B_u)$, forgets at least one declared u -supported coordinate whose support contains an internal edge of B_u , and is distinguished by the compatible outside context Y .

(P3) S_0 and S_1 are two realizations in the same boundary-degree fibre with $q(S_0) = q(S_1)$, but

$$S_0 \oplus Y \quad \text{and} \quad S_1 \oplus Y$$

have different target predicates.

(P4) E is the lexicographically first simple edge-rooted return or simple cycle which witnesses this target difference. Its declared support is the set of declared coordinates of [Definition 3.94](#) used by E .

The carrier set of $\mathfrak{t}$ is

$$K(\mathfrak{t}) = \{c \in \partial_E X : c \text{ occurs in the declared support of } E\}.$$

The token is cheap if $|K(\mathfrak{t})| = 2$. It is actual-context if the context Y is the actual exterior of X in G , with the recorded boundary degree profile. Otherwise it is profile-only.

Lemma 15.11 (Pressure tokens have at least two carriers). *Every exit-(4) pressure token $\mathfrak{t}$ satisfies*

$$|K(\mathfrak{t})| \geq 2.$$

Proof. The quotient q forgets a u -supported coordinate whose declared support contains an internal edge of B_u , and the witnessing event E uses one of the forgotten coordinates. Hence E uses an internal edge of B_u . If E were contained in X , adding the root edge in the edge-rooted case would give a dyadic cycle contained in the contextually dyadic-safe support X , impossible in the counterexample. Thus E also uses an edge outside X .

Add the root edge if E is edge-rooted. We obtain a simple cycle of G which meets both sides of the cut $\delta_G(V(X))$. A cycle crosses a cut an even number of times, and a simple cycle cannot use one cut edge twice. Thus there are at least two distinct boundary edges of X in E . Each such crossing is recorded in the declared support of E as a completion-port incidence, first-entry incidence, connector endpoint, boundary-degree entry, or packed-window interface incidence. Therefore the set $K(\mathfrak{t})$ of ambient carriers in the declared support has size at least two. $\square$

Definition 15.12 (Pressure ledger). *The ledger is used after exits (5)–(7) have been removed or transferred. Thus every surviving target-defect entry fails target-complete-minimality through alternative (a) and has a pressure token. Let $\mathcal{T}_4$ be the finite set of those tokens, obtained from the target-defect entries of $\tilde{\Xi}$ by choosing the lexicographically first witness data in [Definition 15.10](#). A 2/3-pressure ledger on $\tilde{\Xi}$ consists of the following partition:*

$$\tilde{\Xi} = \Xi_3 \sqcup \Xi_2 \sqcup \Xi_{\text{res}},$$

together with an assignment $A(\xi) \subseteq \partial_E X$ for every $\xi = (X, w, u, B_u) \in \Xi_3 \cup \Xi_2$, subject to:

(L1) if ξ is a route-8 entry, then either the branch is already closed by [Theorem 14.61](#), or $\xi \in \Xi_3$ and $A(\xi)$ consists of three private essential carriers of ξ , relative to $\tilde{\Xi}$. Thus no surviving route-8 entry is placed in Ξ_2 or Ξ_{res} ;

(L2) if ξ is a target-defect entry and its token is $\mathfrak{t}_\xi$, then $A(\xi) \subseteq K(\mathfrak{t}_\xi)$;

(L3) $|A(\xi)| = 3$ for $\xi \in \Xi_3$, and $|A(\xi)| = 2$ for $\xi \in \Xi_2$;

(L4) the incidence sets $A(\xi)$ are pairwise disjoint over $\Xi_3 \cup \Xi_2$;

(L5) Ξ_{res} contains exactly the entries for which no such assignment was made.

Among all such ledgers, choose once and for all the lexicographically first ledger maximizing first $|\Xi_3|$, then $|\Xi_2|$. Define

$$N_3 = |\Xi_3|, \quad N_2 = |\Xi_2|, \quad N_{\text{res}} = |\Xi_{\text{res}}|.$$

The external-pressure defect of the ledger is

$$P_{\text{ext}} = N_2 + 3N_{\text{res}}.$$

Lemma 15.13 (Pressure ledger does not overcount). *For the pressure ledger of Definition 15.12,*

$$3N_3 + 2N_2 \leq \sum_{X \in \tilde{\mathcal{X}}} |\partial_E X| = \sum_{X \in \tilde{\mathcal{X}}} \partial_+(X) \leq \partial_+(R).$$

Equivalently,

$$3\tilde{N} - P_{\text{ext}} \leq \partial_+(R).$$

Proof. Every assigned carrier is an oriented boundary incidence of the Type A support which contains the corresponding entry. Clause (L4) says that no assigned incidence is used by two entries. Hence the total number of assigned incidences is at most $\sum_{X \in \tilde{\mathcal{X}}} |\partial_E X|$, while clauses (L3) and (L4) make this total exactly $3N_3 + 2N_2$.

For $X \in \tilde{\mathcal{X}}$ one has $\sigma(X) = 0$; hence every vertex of X has ambient degree 3, and the identity in Definition 14.48 gives $|\partial_E X| = \partial_+(X)$. The canonical supports are disjoint components of the remainder decomposition, so their positive deficiencies add into $\partial_+(R)$, with all omitted terms nonnegative. This proves the first display. Since

$$\tilde{N} = N_3 + N_2 + N_{\text{res}},$$

we have

$$3\tilde{N} - P_{\text{ext}} = 3(N_3 + N_2 + N_{\text{res}}) - (N_2 + 3N_{\text{res}}) = 3N_3 + 2N_2,$$

which gives the equivalent form. □

Definition 15.14 (Two-terminal pressure records). *Let $\xi \in \Xi_2 \cup \Xi_{\text{res}}$ be a target-defect entry, and let $\mathbf{t}_\xi$ be its token. If $\mathbf{t}_\xi$ is actual-context, choose the lexicographically first ordered pair (c_1, c_2) of distinct carriers in $K(\mathbf{t}_\xi)$ that occur as consecutive crossings of $\partial_E X$ along the witnessing event E . The corresponding actual outside corridor is the subpath of E between those crossings whose internal vertices lie outside X . The record*

$$\mathbf{a}(\xi) = (X, w, u, B_u, c_1, c_2, E, \text{outside corridor})$$

is the actual two-terminal pressure record of ξ .

If $\mathbf{t}_\xi$ is profile-only, the two realizations S_0, S_1 differ in at least one declared coordinate forgotten by q . Let $r(\xi)$ be the lexicographically first such coordinate. The record

$$\mathbf{p}(\xi) = (X, w, u, B_u, q, r(\xi), S_0, S_1, Y)$$

is the profile pressure record of ξ .

Lemma 15.15 (Pressure records are canonical). *Every target-defect entry in $\Xi_2 \cup \Xi_{\text{res}}$ has exactly one of the two records in Definition 15.14. Actual records and profile records are disjoint classes of data.*

Proof. The token is either actual-context or profile-only by definition, and not both. In the actual-context case, the mixed event E uses an edge outside X . A simple cycle crosses the cut $\delta_G(V(X))$ an even positive number of times; hence there is at least one pair of consecutive boundary crossings bounding an outside subpath. The imposed lexicographic order chooses one such pair and one outside corridor.

In the profile-only case, S_0 and S_1 have the same image under q but different target predicates after gluing to Y . They cannot be the same exact response state. Since q preserves the boundary degree profile, their difference is recorded by at least one declared coordinate forgotten by q . The lexicographic rule chooses $r(\xi)$. Thus exactly one record is assigned in either case, and the record type is determined by whether the context is actual or profile-only. $\square$

Proposition 15.16 (External-pressure reduction). *After the route-8 two-carrier no-go has been applied, every surviving unified entry is accounted for in exactly one of the following ways:*

class	charge used now	remaining data
Ξ_3	3 disjoint boundary incidences	none
Ξ_2	2 disjoint boundary incidences	actual or profile pressure record
Ξ_{res}	0 assigned incidences	actual or profile pressure record

No boundary incidence is counted twice in the first two rows, and no entry is discarded.

Proof. The route-8 case with at most two private essential carriers is closed by Theorem 14.61; otherwise the ledger definition puts the route-8 entry in Ξ_3 with three private essential carriers relative to $\tilde{\Xi}$. A target-defect entry has the pressure token of Definition 15.10, and Lemma 15.11 gives at least two available carriers for that token. The pressure ledger chooses the maximal disjoint assignment into Ξ_3 and Ξ_2 , leaving the unassigned entries in Ξ_{res} . Clause (L4) gives the no-overcount assertion for assigned boundary incidences. Finally, Lemma 15.15 attaches an actual or profile pressure record to every target-defect entry not paid at the three-incidence level. Since surviving route-8 entries never lie in $\Xi_2 \cup \Xi_{\text{res}}$, the displayed table is exactly this partition. $\square$

15.1.5 Pressure-defect absorption ledger

Definition 15.17 (Actual and profile pressure defects). *For $\xi \in \Xi_2 \cup \Xi_{\text{res}}$, define its pressure weight by*

$$\omega(\xi) = \begin{cases} 1, & \xi \in \Xi_2, \\ 3, & \xi \in \Xi_{\text{res}}. \end{cases}$$

Let $\Xi_{\text{act}} \subseteq \Xi_2 \cup \Xi_{\text{res}}$ be the entries whose canonical record in Definition 15.14 is actual-context, and let Ξ_{prof} be the entries whose canonical record is profile-only. Put

$$P_{\text{act}} = \sum_{\xi \in \Xi_{\text{act}}} \omega(\xi), \quad P_{\text{prof}} = \sum_{\xi \in \Xi_{\text{prof}}} \omega(\xi).$$

Lemma 15.18 (Exact pressure-defect split). *The pressure defect splits without overlap:*

$$P_{\text{ext}} = P_{\text{act}} + P_{\text{prof}}.$$

Proof. By Lemma 15.15, each target-defect entry in $\Xi_2 \cup \Xi_{\text{res}}$ has exactly one canonical record, and that record is actual-context or profile-only, but not both. Since surviving route-8 entries do not lie in $\Xi_2 \cup \Xi_{\text{res}}$ by Definition 15.12, the sets Ξ_{act} and Ξ_{prof} form a partition of $\Xi_2 \cup \Xi_{\text{res}}$. Therefore

$$P_{\text{act}} + P_{\text{prof}} = \sum_{\xi \in \Xi_2 \cup \Xi_{\text{res}}} \omega(\xi) = |\Xi_2| + 3|\Xi_{\text{res}}| = N_2 + 3N_{\text{res}} = P_{\text{ext}}.$$

□

Lemma 15.19 (Profile-pressure dependence routing). *Let*

$$\mathcal{A}_{\text{prof}} = \{ r(\xi) : \xi \in \Xi_{\text{prof}} \}$$

be the family of profile pressure coordinates of Definition 15.14, with repeated coordinates listed once. If a functional admissible rank quotient is rank-reducing on $\mathcal{A}_{\text{prof}}$, then one of the following occurs:

- (i) *a target-defective quotient;*
- (ii) *an admissible nontrivial target-complete compression of a proper support;*
- (iii) *a target-dependence whose validity requires adjoining a connected support strictly larger than the support carrying the dependent coordinates.*

Proof. Let $\mathcal{I} \subseteq \mathcal{A}_{\text{prof}}$ be maximal among subfamilies which survive every functional admissible rank quotient, and choose $a \in \mathcal{A}_{\text{prof}} \setminus \mathcal{I}$. By Definition 3.99, the quotient that kills a 's label gives a finite subfamily $\mathcal{B} \subseteq \mathcal{I}$ whose quotient-image vector determines the quotient-image of a . Choose such a pair $(a, \mathcal{B})$ with inclusion-minimal connected declared support. This is the target-dependence extracted in Lemma 10.3, but the argument uses only the definitions of exact response profile, functional quotient, and declared support.

Let C be the connected support of this minimal determination certificate. The determination is tested against T -boundaried outside contexts, where T is the boundary of C . If some compatible context distinguishes two realizations with the same quotient image, the quotient is target-defective, giving (i). If no compatible context distinguishes them and the support is already the proper support C , then the quotient is target-complete on a proper support. Since it is rank-reducing and admissible, it has a strictly smaller proper representative by Definition 3.98; this is the target-complete compression in (ii). The only remaining case is that the determination becomes context-universal only after adjoining further connected support. Minimality of the certificate gives a connected support strictly larger than C , which is (iii). □

Definition 15.20 (Pressure absorbers). *A pressure unit is a pair (ξ, j) , where $\xi \in \Xi_2 \cup \Xi_{\text{res}}$ and $1 \leq j \leq \omega(\xi)$. Let $\mathcal{U}_{\text{press}}$ be the set of all pressure units, so that*

$$|\mathcal{U}_{\text{press}}| = P_{\text{ext}}.$$

A pressure unit is absorbed if it is assigned one of the following single-use certificates.

- (A1) *an oriented boundary incidence $c \in \partial_E X$, for the same support X as ξ , which is not already used in any set $A(\xi')$ of the 2/3-pressure ledger and is not assigned to another pressure unit;*

(A2) a profile-dependence certificate of [Lemma 15.19](#) whose conclusion is an admissible proper-support compression, or whose larger-support dependence is one of the delocalization alternatives closed by the branch split. A certificate whose conclusion is another target-defective quotient is not a type (A2) absorber; it re-enters the target-defect pressure ledger and is counted there unless it is absorbed by type (A1).

Among all assignments satisfying the single-use conditions in (A1)–(A2), choose the lexicographically first assignment which maximizes, in order, the number of absorbed units and then the number of type (A1) boundary-incidence absorptions. Let $\mathcal{U}_{\text{abs}} \subseteq \mathcal{U}_{\text{press}}$ be its absorbed units, and put

$$P_{\text{open}} = |\mathcal{U}_{\text{press}} \setminus \mathcal{U}_{\text{abs}}|.$$

Lemma 15.21 (Absorbed units improve the boundary count). *Let B_{abs} be the number of pressure units absorbed by boundary incidences of type (A1) in [Definition 15.20](#). Then*

$$3\tilde{N} - P_{\text{open}} \leq \partial_+(R) + B_{\text{dep}},$$

where B_{dep} is the number of units absorbed by dependence certificates of type (A2). On any branch where the type (A2) dependence certificates have routed to their proper-compression or delocalization alternatives, $B_{\text{dep}} = 0$, and hence

$$3\tilde{N} - P_{\text{open}} \leq \partial_+(R).$$

Proof. The original 2/3-pressure ledger assigns $3N_3 + 2N_2$ boundary incidences, and [Lemma 15.13](#) proves that these incidences are pairwise disjoint. Each type (A1) absorber is, by definition, an additional oriented boundary incidence not used by the 2/3-pressure ledger and not assigned to another absorbed unit. Therefore the 2/3-ledger together with all type (A1) absorbers uses

$$3N_3 + 2N_2 + B_{\text{abs}}$$

distinct boundary incidences of the supports in $\tilde{\mathcal{X}}$. Since these supports are ambient-cubic and disjoint in the remainder decomposition, each counted external incidence is incident with exactly one deficient boundary slot of R , and no deficient boundary slot is used twice. Hence the number of counted incidences is bounded by the total positive deficiency of R :

$$3N_3 + 2N_2 + B_{\text{abs}} \leq \partial_+(R).$$

The total pressure defect decomposes as

$$P_{\text{ext}} = B_{\text{abs}} + B_{\text{dep}} + P_{\text{open}}.$$

Using

$$3\tilde{N} - P_{\text{ext}} = 3N_3 + 2N_2$$

from [Lemma 15.13](#), we get

$$3\tilde{N} - P_{\text{open}} - B_{\text{dep}} = 3N_3 + 2N_2 + B_{\text{abs}} \leq \partial_+(R).$$

This is the first displayed inequality. If the type (A2) absorbers route to their proper-compression or delocalization alternatives, no surviving pressure residual contains those units; equivalently the surviving residual has $B_{\text{dep}} = 0$, giving the second display. $\square$

Proposition 15.22 (Exit-(4) closure from open pressure). *Assume Definition 4.1. On a branch where all dependence absorbers of type (A2) have routed to their proper-compression or delocalization alternatives, the target-defect exit-(4) residual is impossible whenever*

$$\limsup_{|R| \rightarrow \infty} \frac{P_{\text{open}}}{|R|} < \varepsilon_{\text{press}}, \quad \varepsilon_{\text{press}} = 12 \left(\frac{1}{4} - \tau_{\text{win}} \right) - \tau_{\text{win}}.$$

Proof. By Lemmas 15.4 and 15.6,

$$\tilde{N} \geq 4 \left(\frac{1}{4} - \tau_{\text{win}} \right) |R| - o(|R|).$$

By Lemma 15.21 and the hypothesis that the type (A2) units have routed to their proper-compression or delocalization alternatives,

$$3\tilde{N} - P_{\text{open}} \leq \partial_+(R).$$

The window-only deficiency cap gives

$$\partial_+(R) \leq \tau_{\text{win}}|R| + o(|R|).$$

Combining the three inequalities gives

$$\left[12 \left(\frac{1}{4} - \tau_{\text{win}} \right) - \frac{P_{\text{open}}}{|R|} \right] |R| \leq \tau_{\text{win}}|R| + o(|R|).$$

The displayed limsup assumption makes the coefficient on the left eventually strictly larger than τ_{win} , a contradiction. $\square$

15.1.6 Window-blocker accounting for the open pressure

Definition 15.23 (Canonical window blocker of an open unit). *Let $v = (\xi, j) \in \mathcal{U}_{\text{press}} \setminus \mathcal{U}_{\text{abs}}$ be an open pressure unit, with $\xi = (X, w, u, B_u)$, and let $\mathbf{t}_\xi$ be the pressure token of ξ . Since X is a component of the P_{13} -free remainder R , every ambient boundary incidence of X leaves X through an edge from R to the packed-window set W . Choose the lexicographically first carrier $c(v) \in K(\mathbf{t}_\xi)$. If*

$$c(v) = (x, xy), \quad x \in V(X), \quad y \in W,$$

let $P(v) \in \mathcal{P}$ be the unique packed induced P_{13} window containing y . The pair

$$\mathbf{b}(v) = (P(v), c(v))$$

is the canonical window blocker of v . For a packed window P , put

$$B_{\text{open}}(P) = |\{v \in \mathcal{U}_{\text{press}} \setminus \mathcal{U}_{\text{abs}} : P(v) = P\}|.$$

Lemma 15.24 (Open pressure is exactly the window-blocker load). *The canonical window-blocker loads satisfy*

$$P_{\text{open}} = \sum_{P \in \mathcal{P}} B_{\text{open}}(P).$$

Proof. Every open pressure unit has a pressure token. By Lemma 15.11, its carrier set is nonempty. The lexicographic rule in Definition 15.23 therefore assigns exactly one carrier $c(v)$. Because X is a connected component of R , the other endpoint of the corresponding boundary edge lies in W , and because the packed windows are vertex-disjoint, that endpoint lies in a unique window $P(v)$. Thus the sets counted by $B_{\text{open}}(P)$, $P \in \mathcal{P}$, partition $\mathcal{U}_{\text{press}} \setminus \mathcal{U}_{\text{abs}}$. Their total size is P_{open} by the definition of P_{open} . $\square$




Figure 29: Exit-(4) attachment to a packed P_{13} window and the associated accounting. A target-defect pressure token for the entry $\xi = (X, w, u, B_u)$ forgets a declared u -supported coordinate using an internal edge of B_u . The witnessing event must also use the exterior of X , so cut parity gives at least two boundary carriers. Carriers assigned by the 2/3-pressure ledger are pairwise disjoint and are counted once in $\partial_+(R)$. An unabsorbed pressure unit chooses the lexicographically first carrier $c(v)$; because X is a component of the P_{13} -free remainder, the other endpoint lies in the packed-window set, and the disjointness of the packing gives a unique window $P(v)$. Thus $P_{\text{open}} = \sum_P B_{\text{open}}(P)$. On each window, the first two open units form the base term $2p_{13}$; any third open unit forms an overload triple. Routed triples leave through the direct dyadic certificate, an unused boundary absorber, a proper-compression or delocalization absorber, or the Type B fan-window handoff. If two open units have a recorded P -avoiding corridor Q with $b \in \text{Sh}_{s(Q)}(a)$, the displayed cycle has length $s(Q) + 2 + |a - b| \in \mathcal{D}$, so the triple is routed. The only surviving excess is the zero-shadow term in $P_{\text{open}} \leq 2p_{13} + P_{\text{zero}}^+$. Independently of that auxiliary window audit, every surviving target-defect two-carrier entry is a canonical (Q3) exit-(4) quotient, so [Lemmas 15.44](#) and [15.47](#) peel the routed load u , decrease the remaining deficit by exactly $1/4$, and strictly decrease Λ_4 .

Definition 15.25 (Same-window overload triples). *For a packed window $P \in \mathcal{P}$, put*

$$\mathcal{U}(P) = \{v \in \mathcal{U}_{\text{press}} \setminus \mathcal{U}_{\text{abs}} : P(v) = P\}.$$

A same-window overload triple at P is a three-element subset $\mathcal{T} \subseteq \mathcal{U}(P)$. If $v \in \mathcal{T}$, write

$$c(v) = (x_v, x_v p_{a(v)}), \quad P = p_0 p_1 \cdots p_{12},$$

for its canonical window carrier and its window index.

The triple $\mathcal{T}$ is routed if at least one of the following certificates is present.

- (O1) Direct dyadic window certificate. *The pressure records of one or two units in $\mathcal{T}$, or an explicitly recorded P -avoiding corridor between their canonical window carriers, together with one or two subpaths of P , contain a simple cycle of length in $\mathcal{D}$. This includes the direct same-window and two-window cycle configurations covered by [Lemmas 14.75](#) and [14.76](#).*

- (O2) Unused boundary-incidence certificate. *Some $v = (\xi, j) \in \mathcal{T}$, with $\xi = (X, w, u, B_u)$, has an oriented boundary incidence $c \in \partial_E X$ which is not used by any set $A(\xi')$ in the 2/3-pressure ledger and is not assigned to any absorbed pressure unit.*
- (O3) Profile-dependence certificate. *A subfamily of profile coordinates belonging to units of $\mathcal{T}$ is rank-reducing under a functional admissible quotient, and the routing supplied by [Lemma 15.19](#) is a proper-support compression or one of the delocalization alternatives already removed by the branch split.*
- (O4) Type B fan-window certificate. *Two units of $\mathcal{T}$ determine fan-closed ports or cubic-closed fan neighbours in an assigned Type B fan-window profile, with the window incidences and non-window incidences required by [Lemma 14.70](#) and [Definition 14.91](#). Thus the local B1 entry is available, and any failure of cross-center disjointness is a B2 overlap obstruction in the sense of [Lemma 14.87](#).*

A same-window overload triple is primitive if it is not routed.

Lemma 15.26 (Routed overload triples do not remain open). *Work on a branch where the direct dyadic cycle alternatives, the type (A1) boundary absorbers, the type (A2) proper-compression or delocalization absorbers, and the Type B fan-window handoff have all been applied. Then every surviving same-window overload triple is primitive.*

Proof. Let $\mathcal{T}$ be a same-window overload triple. If (O1) holds, the displayed certificate is a power-of-two cycle in G , contradicting the standing counterexample hypothesis. If (O2) holds, the indicated boundary incidence is an admissible absorber of type (A1) in [Definition 15.20](#). Since the absorber assignment is maximal, the corresponding pressure unit cannot still be open. If (O3) holds, [Lemma 15.19](#) supplies an absorber of type (A2) whose conclusion is proper compression or delocalization; on the stated branch such a unit has routed away from the open pressure ledger. If (O4) holds, the data are exactly Type B fan-window data. The local B1 entry is supplied by [Definition 14.91](#); if disjointness fails, [Lemma 14.87](#) records the B2 residual. In either case the triple is no longer an open Type A pressure overload. Therefore a surviving same-window overload triple has none of (O1)–(O4), and is primitive by definition. $\square$

Definition 15.27 (Primitive same-window overload excess). *Assume the branch of [Lemma 15.26](#). A packed window P is primitive-overloaded if*

$$B_{\text{open}}(P) = |\mathcal{U}(P)| \geq 3$$

and the lexicographically first three-element subset of $\mathcal{U}(P)$ is a primitive same-window overload triple. The primitive overload excess is

$$P_{\text{prim}}^+ = \sum_{\substack{P \in \mathcal{P} \\ P \text{ primitive-overloaded}}} (B_{\text{open}}(P) - 2).$$

Definition 15.28 (Window attachment shadows). *Let $P = p_0 p_1 \cdots p_{12}$ be a packed window. For an integer $s \geq 0$, define the forbidden distance set*

$$D_s = \{d \in \{0, \dots, 12\} : s + 2 + d \in \mathcal{D}\}.$$

For an attachment index $a \in \{0, \dots, 12\}$, define its s -shadow on P by

$$\text{Sh}_s(a) = \{b \in \{0, \dots, 12\} : |a - b| \in D_s\}.$$

Equivalently,

$$b \in \text{Sh}_s(a) \iff C_s(\{a\}, \{b\}) = 0$$

for the P_{13} -label relation C_s of the main window algebra.

Lemma 15.29 (Exact singleton-shadow table). *For every $s \geq 0$,*

$$D_s = \{2^q - s - 2 : q \geq 2, 0 \leq 2^q - s - 2 \leq 12\}.$$

In particular,

s	D_s
0	$\{2, 6\}$
1	$\{1, 5\}$
2	$\{0, 4, 12\}$
3	$\{3, 11\}$
4	$\{2, 10\}$
5	$\{1, 9\}$
6	$\{0, 8\}$
7	$\{7\}$
8	$\{6\}$
9	$\{5\}$
10	$\{4\}$
11	$\{3\}$
12	$\{2\}$
13	$\{1\}$
14	$\{0\}$
15, 16, 17	$\emptyset$
18	$\{12\}$
19	$\{11\}$
20	$\{10\}$
21	$\{9\}$
22	$\{8\}$
23	$\{7\}$
24	$\{6\}$
25	$\{5\}$
26	$\{4\}$
27	$\{3\}$
28	$\{2\}$
29	$\{1\}$
30	$\{0\}$.

For $s \geq 18$, the set D_s has at most one element.

Proof. By [Definition 15.28](#), $d \in D_s$ exactly when $s + 2 + d = 2^q$ for some $q \geq 2$. This is equivalent to

$$d = 2^q - s - 2$$

with $0 \leq d \leq 12$, giving the displayed formula. The table is obtained by substituting $s = 0, \dots, 30$ into this formula. If $s \geq 18$, then $[s + 2, s + 14] \subseteq [20, \infty)$. Consecutive powers of two in $[16, \infty)$ differ by at least $16 > 12$, so the interval $[s + 2, s + 14]$ contains at most one power of two. Hence D_s has at most one element. $\square$

Definition 15.30 (Recorded window-shadow hit). *Let $v, v' \in \mathcal{U}(P)$ be distinct open pressure units on the same packed window $P = p_0 \cdots p_{12}$, with canonical carriers*

$$c(v) = (x, xp_a), \quad c(v') = (y, yp_b).$$

A recorded P -avoiding corridor from v to v' is a simple x - y path Q in $G - V(P)$, recorded in the actual pressure data of one or both units or in the Type B fan-window handoff data, whose internal vertices avoid P and whose first and last edges are not the canonical carrier edges xp_a, yp_b . Its length is denoted $s(Q) = |E(Q)|$.

The ordered pair (v, v') has a recorded window-shadow hit if the two canonical carrier edges xp_a and yp_b are distinct, such a corridor Q exists, and

$$b \in \text{Sh}_{s(Q)}(a).$$

When $x = y$, the corridor Q may be the trivial length-zero path; the distinct-carrier condition then says $a \neq b$.

Lemma 15.31 (Window-shadow hits are direct dyadic certificates). *If two open pressure units on the same packed window have a recorded window-shadow hit, then every same-window overload triple containing them is routed by certificate (O1) of Definition 15.25.*

Proof. Use the notation of Definition 15.30. Since $b \in \text{Sh}_{s(Q)}(a)$, the definition of the shadow gives

$$s(Q) + 2 + |a - b| \in \mathcal{D}.$$

The path

$$x Q y p_b P p_a x$$

is a simple cycle. Indeed, Q lies in $G - V(P)$, the segment $p_b P p_a$ lies in the induced path P , and the only edges between the two pieces used in the display are the two distinct canonical carrier edges yp_b and xp_a . If $x = y$, then Q is the trivial path and distinctness of the carrier edges gives $a \neq b$, so the displayed cycle is $xp_b P p_a x$. In all cases the length is exactly $s(Q) + 2 + |a - b|$, hence is a power of two. This is a direct dyadic window certificate, which is certificate (O1). $\square$

Definition 15.32 (Zero-shadow primitive excess). *Assume the branch of Lemma 15.26. A primitive-overloaded packed window P is zero-shadow if the set $\mathcal{U}(P)$ has no pair of distinct open units with a recorded window-shadow hit. The zero-shadow primitive excess is*

$$P_{\text{zero}}^+ = \sum_{\substack{P \in \mathcal{P} \\ P \text{ zero-shadow}}} (B_{\text{open}}(P) - 2).$$

Lemma 15.33 (Primitive excess equals zero-shadow excess after window-shadow routing). *Work on a branch where the routed overload triples of Definition 15.25 have been removed. Then*

$$P_{\text{prim}}^+ = P_{\text{zero}}^+.$$

Proof. Let P be primitive-overloaded. Its lexicographically first three open units form a surviving same-window overload triple. If some pair in that triple had a recorded window-shadow hit, then Lemma 15.31 would route the triple by certificate (O1), contrary to the branch hypothesis. More generally, if any distinct pair $v, v' \in \mathcal{U}(P)$ had a recorded window-shadow hit, then, since $B_{\text{open}}(P) \geq 3$, choosing any third unit of $\mathcal{U}(P)$ would give a same-window overload triple routed by (O1). Hence no pair in $\mathcal{U}(P)$ has a recorded window-shadow hit, so P is zero-shadow. Therefore every summand counted by P_{prim}^+ is counted by P_{zero}^+ . Conversely, a zero-shadow window is primitive-overloaded by definition, so every summand counted by P_{zero}^+ is counted by P_{prim}^+ . $\square$

Lemma 15.34 (Open pressure after routed overloads). *On a branch where the routed overload triples of Definition 15.25 have been removed,*

$$P_{\text{open}} \leq 2p_{13} + P_{\text{zero}}^+.$$

Proof. By Lemma 15.24,

$$P_{\text{open}} = \sum_{P \in \mathcal{P}} B_{\text{open}}(P).$$

If $B_{\text{open}}(P) \leq 2$, then its contribution is at most 2. If $B_{\text{open}}(P) \geq 3$, then the lexicographically first three-element subset of $\mathcal{U}(P)$ is a surviving same-window overload triple. By Lemma 15.26, it is primitive, so P is primitive-overloaded. If its first triple had a recorded window-shadow hit, Lemma 15.31 would route it by certificate (O1), contrary to the branch hypothesis. If any other pair in $\mathcal{U}(P)$ had a recorded window-shadow hit, that pair together with any third open unit would give a routed same-window overload triple, again contrary to the branch hypothesis. Hence P is zero-shadow, and its excess $B_{\text{open}}(P) - 2$ is one summand of P_{zero}^+ . Therefore, for every packed window P ,

$$B_{\text{open}}(P) \leq 2 + \mathbf{1}_{\{P \text{ zero-shadow}\}}(B_{\text{open}}(P) - 2)$$

with no overlap between windows. Summing this pointwise inequality over $P \in \mathcal{P}$ gives the result. $\square$

Lemma 15.35 (Window-blocker accounting audit). *On a branch where the routed overload triples of Definition 15.25 have been removed, every open pressure unit is counted in exactly one of the following two places:*

- (a) among the first two open units assigned to its packed window, contributing to the base term $2p_{13}$;
- (b) among the excess units of a zero-shadow primitive-overloaded packed window, contributing to P_{zero}^+ .

No open unit is counted for two packed windows, and every same-window overload not counted in P_{zero}^+ has routed through one of (O1)–(O4).

Proof. Lemma 15.24 partitions the open pressure units by the unique packed window $P(v)$. Inside a fixed packed window, choose the first two units in the global lexicographic order as the base units; if fewer than two units are present, all units are base units. Any remaining unit exists only when $B_{\text{open}}(P) \geq 3$. On the stated branch, the first three units then form a surviving same-window overload triple. If that triple were routed by one of (O1)–(O4), the branch would have removed that overload before P_{open} is counted. Hence, for an open residual unit, the triple is not routed. Lemma 15.26 makes it primitive, and Lemma 15.31 removes every same-window overload triple containing a pair with a recorded window-shadow hit. Since $B_{\text{open}}(P) \geq 3$, any recorded-hit pair in $\mathcal{U}(P)$ would lie in such a triple. Thus the window is zero-shadow, and all units beyond the first two are counted in the summand $B_{\text{open}}(P) - 2$ of P_{zero}^+ . The packed windows are vertex-disjoint and the assignment $P(v)$ is single-valued, so no open unit can appear in two window classes. These two alternatives exhaust the open units because they are applied after the exact partition of Lemma 15.24. $\square$

Lemma 15.36 (Exhaustion of the final open-pressure charge). *Assume the type (A2) absorbers of Definition 15.20 have routed to their proper-compression or delocalization alternatives, and assume*

the routed overload triples of [Definition 15.25](#) have been removed. Then the surviving open pressure admits a disjoint partition

$$\mathcal{U}_{\text{press}} \setminus \mathcal{U}_{\text{abs}} = \mathcal{U}_{\text{base}} \sqcup \mathcal{U}_{\text{zero}}$$

such that

$$|\mathcal{U}_{\text{base}}| \leq 2p_{13}, \quad |\mathcal{U}_{\text{zero}}| = P_{\text{zero}}^+,$$

and hence

$$3\tilde{N} - P_{\text{open}} \leq \partial_+(R), \quad P_{\text{open}} \leq 2p_{13} + P_{\text{zero}}^+.$$

Proof. The boundary-incidence inequality

$$3\tilde{N} - P_{\text{open}} \leq \partial_+(R)$$

is [Lemma 15.21](#) after the type (A2) alternatives have routed away.

It remains to identify the open units. By [Lemma 15.24](#), every open pressure unit has a unique packed window $P(v)$. For a fixed window P , order $\mathcal{U}(P)$ by the global lexicographic order. Put the first two units, if present, in $\mathcal{U}_{\text{base}}$. Since there are p_{13} packed windows, this gives $|\mathcal{U}_{\text{base}}| \leq 2p_{13}$.

Let v be any remaining open unit on P . Then $B_{\text{open}}(P) \geq 3$. The first three units of $\mathcal{U}(P)$ form a same-window overload triple. If that triple were routed by one of (O1)–(O4), it would have been removed before the present open-pressure ledger was formed. Hence the triple is primitive and P is primitive-overloaded. If any pair in $\mathcal{U}(P)$ had a recorded window-shadow hit, that pair and any third open unit of $\mathcal{U}(P)$ would form a same-window overload triple routed by (O1) via [Lemma 15.31](#); this is again excluded on the present branch. Therefore P is zero-shadow. Put every unit of $\mathcal{U}(P)$ after the first two in $\mathcal{U}_{\text{zero}}$.

The preceding paragraph applies to every open unit not in $\mathcal{U}_{\text{base}}$, and the construction is window-by-window, so the two sets are disjoint and cover $\mathcal{U}_{\text{press}} \setminus \mathcal{U}_{\text{abs}}$. For a zero-shadow window P , exactly $B_{\text{open}}(P) - 2$ units are placed in $\mathcal{U}_{\text{zero}}$, and no non-zero-shadow window contributes to $\mathcal{U}_{\text{zero}}$. Summing over P gives

$$|\mathcal{U}_{\text{zero}}| = \sum_{\substack{P \in \mathcal{P} \\ P \text{ zero-shadow}}} (B_{\text{open}}(P) - 2) = P_{\text{zero}}^+.$$

The last displayed inequality follows from

$$P_{\text{open}} = |\mathcal{U}_{\text{press}} \setminus \mathcal{U}_{\text{abs}}| = |\mathcal{U}_{\text{base}}| + |\mathcal{U}_{\text{zero}}|.$$

□

Proposition 15.37 (Exit-(4) closure from zero-shadow slack). *Assume [Definition 4.1](#). Work on a branch where the type (A2) absorbers of [Definition 15.20](#) have routed to their proper-compression or delocalization alternatives, and where the routed overload triples of [Definition 15.25](#) have been removed. If*

$$\limsup_{|R| \rightarrow \infty} \frac{P_{\text{zero}}^+}{|R|} < \varepsilon_{\text{prim}},$$

where

$$\varepsilon_{\text{press}} = 12 \left(\frac{1}{4} - \tau_{\text{win}} \right) - \tau_{\text{win}}, \quad \varepsilon_{\text{prim}} = \varepsilon_{\text{press}} - \frac{2\theta_{\text{win}}}{1 - 13\theta_{\text{win}}} = 0.00330339423 \dots$$

then the target-defect exit-(4) residual is impossible.

Proof. By Lemma 15.36,

$$P_{\text{open}} \leq 2p_{13} + P_{\text{zero}}^+.$$

The near-cubic P_{13} -density bound gives

$$p_{13} \leq \theta_{\text{win}} n + o(n).$$

Writing $\theta = p_{13}/n$, the remainder size is

$$|R| = (1 - 13\theta)n + o(n),$$

and $2\theta/(1 - 13\theta)$ is increasing for $0 \leq \theta < 1/13$. Therefore

$$\limsup_{|R| \rightarrow \infty} \frac{2p_{13}}{|R|} \leq \frac{2\theta_{\text{win}}}{1 - 13\theta_{\text{win}}}.$$

Hence the displayed hypothesis implies

$$\limsup_{|R| \rightarrow \infty} \frac{P_{\text{open}}}{|R|} < \frac{2\theta_{\text{win}}}{1 - 13\theta_{\text{win}}} + \varepsilon_{\text{prim}} = \varepsilon_{\text{press}}.$$

Proposition 15.22 now closes the target-defect exit-(4) residual. □

Remark 15.38 (Numerical size of the final shadow residual). *Since*

$$\frac{p_{13}}{|R|} \leq \frac{\theta_{\text{win}}}{1 - 13\theta_{\text{win}}} + o(1) = 0.01521165789 \dots + o(1),$$

the available slack

$$\varepsilon_{\text{prim}} = 0.00330339423 \dots$$

is exactly

$$\frac{\varepsilon_{\text{prim}}}{\theta_{\text{win}}/(1 - 13\theta_{\text{win}})} = 0.2171620118 \dots$$

zero-shadow excess units per packed window on average. Thus the zero-shadow criterion alone leaves only residuals satisfying

$$P_{\text{zero}}^+ \geq 0.2171620118 \dots p_{13} - o(|R|).$$

Every packed window contributing to this sum carries at least a third open blocker, no pair of open blockers on that window has a recorded C_s -shadow hit, and the corresponding open units have no unused boundary absorber, no proper-compression or delocalization absorber, and no Type B fan-window handoff.

Definition 15.39 (Same-window open-blocker cap). *The open-pressure ledger satisfies the same-window two-blocker cap if there is an exceptional set*

$$\mathcal{E}_{\text{win}} \subseteq \mathcal{U}_{\text{press}} \setminus \mathcal{U}_{\text{abs}}$$

with $|\mathcal{E}_{\text{win}}| = o(|R|)$ such that, after deleting $\mathcal{E}_{\text{win}}$, every packed window $P \in \mathcal{P}$ receives at most two canonical open blockers:

$$|\{v \in (\mathcal{U}_{\text{press}} \setminus \mathcal{U}_{\text{abs}}) \setminus \mathcal{E}_{\text{win}} : P(v) = P\}| \leq 2.$$

Lemma 15.40 (The two-blocker cap has sublinear overload excess). *If the same-window two-blocker cap of Definition 15.39 holds, then*

$$P_{\text{prim}}^+ = o(|R|) \quad \text{and} \quad P_{\text{zero}}^+ = o(|R|).$$

Proof. Let $\mathcal{E}_{\text{win}}$ be the exceptional set in Definition 15.39. For a fixed packed window P , write $e(P) = |\mathcal{E}_{\text{win}} \cap \mathcal{U}(P)|$. After deleting $\mathcal{E}_{\text{win}}$, at most two units remain on P , so

$$\max\{0, B_{\text{open}}(P) - 2\} \leq e(P).$$

The primitive overload excess is bounded by the total overload excess:

$$P_{\text{prim}}^+ \leq \sum_{P \in \mathcal{P}} \max\{0, B_{\text{open}}(P) - 2\} \leq \sum_{P \in \mathcal{P}} e(P) = |\mathcal{E}_{\text{win}}| = o(|R|).$$

The same displayed total overload excess also bounds P_{zero}^+ , since zero-shadow windows are a subfamily of the packed windows. $\square$

Lemma 15.41 (Numerical window-blocker squeeze). *Assume Definition 4.1. If the same-window two-blocker cap of Definition 15.39 holds, then*

$$\limsup_{|R| \rightarrow \infty} \frac{P_{\text{open}}}{|R|} \leq \frac{2\theta_{\text{win}}}{1 - 13\theta_{\text{win}}} = 0.03042331579\dots < 0.03372671002\dots = \varepsilon_{\text{press}}.$$

Proof. By Lemma 15.24 and the same-window cap,

$$P_{\text{open}} \leq 2p_{13} + o(|R|).$$

The near-cubic P_{13} -density bound gives

$$p_{13} \leq \theta_{\text{win}}n + o(n).$$

Writing $\theta = p_{13}/n$, the remainder size is

$$|R| = (1 - 13\theta)n + o(n),$$

and the function $2\theta/(1 - 13\theta)$ is increasing on the relevant interval $0 \leq \theta < 1/13$. Therefore

$$\limsup_{|R| \rightarrow \infty} \frac{P_{\text{open}}}{|R|} \leq \frac{2\theta_{\text{win}}}{1 - 13\theta_{\text{win}}}.$$

Substituting $\theta_{\text{win}} = 0.01270017798\dots$ gives

$$\frac{2\theta_{\text{win}}}{1 - 13\theta_{\text{win}}} = 0.03042331579\dots$$

The open-pressure threshold from Proposition 15.22 is

$$\varepsilon_{\text{press}} = 12 \left(\frac{1}{4} - \tau_{\text{win}} \right) - \tau_{\text{win}} = 0.03372671002\dots,$$

so the displayed bound lies strictly below the threshold. $\square$

Proposition 15.42 (Exit-(4) closure from the same-window cap). *Assume Definition 4.1. On a branch where all dependence absorbers of type (A2) have routed to their proper-compression or delocalization alternatives, the target-defect exit-(4) residual is impossible whenever the same-window two-blocker cap of Definition 15.39 holds.*

Proof. By Lemma 15.41, the same-window cap gives

$$\limsup_{|R| \rightarrow \infty} \frac{P_{\text{open}}}{|R|} < \varepsilon_{\text{press}}.$$

The conclusion is exactly Proposition 15.22. $\square$

Proposition 15.43 (Sharp external-pressure closure criterion). *Assume Definition 4.1. In the large-budget branch, suppose the pressure ledger of Definition 15.12 satisfies*

$$\limsup_{|R| \rightarrow \infty} \frac{P_{\text{ext}}}{|R|} < \varepsilon_{\text{press}}, \quad \varepsilon_{\text{press}} = 12 \left(\frac{1}{4} - \tau_{\text{win}} \right) - \tau_{\text{win}}.$$

Then no minimal counterexample survives. Numerically,

$$\varepsilon_{\text{press}} = 12(0.02182513154 \dots) - 0.22817486846 \dots = 0.03372671002 \dots$$

In particular, a sublinear external-pressure defect $P_{\text{ext}} = o(|R|)$ closes the branch.

Proof. By Lemmas 15.4 and 15.6,

$$\tilde{N} \geq 4 \left(\frac{1}{4} - \tau_{\text{win}} \right) |R| - o(|R|).$$

By Lemma 15.13 and the window-only deficiency cap,

$$3\tilde{N} - P_{\text{ext}} \leq \partial_+(R) \leq \tau_{\text{win}}|R| + o(|R|).$$

Combining the two displays gives

$$\left[12 \left(\frac{1}{4} - \tau_{\text{win}} \right) - \frac{P_{\text{ext}}}{|R|} \right] |R| \leq \tau_{\text{win}}|R| + o(|R|).$$

If the normalized limsup of P_{ext} is strictly below $\varepsilon_{\text{press}}$, the coefficient on the left is eventually larger than τ_{win} by a fixed positive amount, contradicting the displayed inequality. The numerical value follows from the stated value of τ_{win} . The sublinear case has normalized limsup 0, so it is a special case. $\square$

Lemma 15.44 (Target-defect pressure is canonical exit-(4) peel data). *Let $\xi = (X, w, u, B_u) \in \tilde{\Xi}$ be a target-defect entry, and let $\mathbf{t}_\xi = (\xi, q, S_0, S_1, Y, E)$ be its lexicographically chosen pressure token. Then $q \in \mathcal{Q}_4(w)$, the declared routed-load support of q contains the load u , and q is an exit-(4) witness for u in the sense of Definition 14.33. Consequently, if $u \notin P_4(w)$, then Lemma 14.38 enlarges the peeling set by u and reduces the remaining receiver deficit by exactly $1/4$.*

Proof. Because ξ is target-defect, the surviving failure alternative in Definition 14.14 is alternative (a): a trace-local non-carrier quotient or identification inside $\rho_u(B_u)$ is target-defective. The token definition chooses exactly this quotient q , two same-boundary-degree realizations S_0, S_1 , and a compatible context Y which distinguishes their target predicates.

Clause (Q3) of [Definition 14.49](#) places every such trace-local non-carrier quotient of the trace basin B_u in the canonical family $\mathcal{Q}_4(w)$. The same definition records the declared routed-load support of a (Q3) quotient as the trace load u whose basin is being tested. Thus $q \in \mathcal{Q}_4(w)$, it is boundary-degree-preserving, it is target-defective in the compatible context Y , and its declared support contains the canonical coordinate of u . This is precisely an exit-(4) witness for u by [Definition 14.33](#). If u is unpeeled, the last assertion is [Lemma 14.38](#). $\square$

Definition 15.45 (Peeling-reduced unified ledger). *Let $P_4 = (P_4(w))_w$ be a family of valid exit-(4) peeling sets in the sense of [Definition 14.33](#). The P_4 -reduced unified ledger is obtained from [Definitions 15.2](#) and [15.5](#) by deleting every routed load listed in its receiver's peeling set and appearing in the unified silent-entry ledger. Put*

$$\mathcal{U}^{\text{un}}(w) = \bigcup_{X \in \tilde{\mathcal{X}}} \mathcal{U}_X(w), \quad P_4^{\text{un}}(w) = P_4(w) \cap \mathcal{U}^{\text{un}}(w).$$

Thus

$$\mathcal{U}_X^{P_4}(w) = \mathcal{U}_X(w) \setminus P_4^{\text{un}}(w),$$

and

$$\tilde{\Xi}^{P_4} = \{ (X, w, u, B_u) : X \in \tilde{\mathcal{X}}, u \in \mathcal{U}_X^{P_4}(w) \}, \quad \tilde{N}^{P_4} = |\tilde{\Xi}^{P_4}|.$$

The current entry set $\tilde{\Xi}^{P_4}$ is the frozen tagged family for this stage. Essential carrier sets are unchanged by peeling, while privacy is recomputed relative to the surviving entries; thus the current two-carrier condition is $\pi_{\tilde{\Xi}^{P_4}}(\xi) \leq 2$. The remaining Type A deficit represented by this ledger is

$$\tilde{D}_A^{P_4} = \tilde{D}_A - \frac{1}{4} \sum_w |P_4^{\text{un}}(w)|.$$

If $\tilde{D}_A^{P_4} \leq 0$, the Type A target-defect/route-8 ledger has no negative mass left to carry.

Lemma 15.46 (Peeling-reduced carrier reduction). *For every valid family P_4 of exit-(4) peeling sets, the reduced ledger satisfies*

$$\tilde{N}^{P_4} \geq 4\tilde{D}_A^{P_4}.$$

Moreover, if

$$\tilde{D}_A^{P_4} \geq \left(\frac{1}{4} - \tau_{\text{win}} \right) |R| - o(|R|),$$

then $\tilde{\Xi}^{P_4}$ contains an entry ξ with $\pi_{\tilde{\Xi}^{P_4}}(\xi) \leq 2$.

Proof. Each load in $P_4^{\text{un}}(w)$ is equipped with one exit-(4) witness and is removed exactly once from the unified Type A receiver calculation. By [Lemmas 14.34](#) and [14.38](#), each such removal decreases the remaining receiver deficit by exactly 1/4 and does not alter the graph, the support boundary, or any unpeeled load. Hence the silent excess count of [Lemma 15.6](#) loses exactly one indexed possible entry for each load in P_4^{un} , while the deficit side loses exactly 1/4. The proof of [Lemma 15.6](#) therefore gives

$$\tilde{N}^{P_4} \geq 4 \left(\tilde{D}_A - \frac{1}{4} \sum_w |P_4^{\text{un}}(w)| \right) = 4\tilde{D}_A^{P_4}.$$

Assume now that the displayed lower bound on $\tilde{D}_A^{P_4}$ holds and that every entry of $\tilde{\Xi}^{P_4}$ satisfies $\pi_{\tilde{\Xi}^{P_4}}(\xi) \geq 3$. Then every remaining entry has at least three private essential carriers relative to the current tagged family. These private carriers are oriented boundary incidences of the supports in

$\tilde{\mathcal{X}}$, and private carriers belonging to different entries are disjoint by the definition of $\pi_{\tilde{\Xi}^{P_4}}$ in [Definition 14.48](#). Deleting peeled loads has not changed G , the supports, or their boundary incidences, so selecting three private carriers for each entry gives

$$3\tilde{N}^{P_4} \leq \sum_{X \in \tilde{\mathcal{X}}} |\partial_E X| = \sum_{X \in \tilde{\mathcal{X}}} \partial_+(X) \leq \partial_+(R) \leq \tau_{\text{win}}|R| + o(|R|),$$

where the equality uses $\sigma(X) = 0$ for $X \in \tilde{\mathcal{X}}$, and the last inequality is [Corollary 8.7](#). The first part of the lemma and the assumed reduced-deficit lower bound give

$$\tilde{N}^{P_4} \geq 4\left(\frac{1}{4} - \tau_{\text{win}}\right)|R| - o(|R|).$$

Combining the two displays yields

$$12\left(\frac{1}{4} - \tau_{\text{win}}\right)|R| - o(|R|) \leq \tau_{\text{win}}|R| + o(|R|),$$

and hence $12(\frac{1}{4} - \tau_{\text{win}}) \leq \tau_{\text{win}}$. Equivalently $\tau_{\text{win}} \geq 3/13$, contradicting $\tau_{\text{win}} = 0.22817\dots < 3/13$. Therefore $\tilde{\Xi}^{P_4}$ contains a two-carrier entry. $\square$

Lemma 15.47 (Finite exit-(4) descent). *Run the saturated Type A receiver analysis with the peeling sets $P_4(w)$ of [Definition 14.33](#). Whenever the current large-budget residual contains a two-carrier target-defect entry, one exit-(4) peeling step is performed, the integer*

$$\Lambda_4 = \sum_w |\mathcal{L}(w) \setminus P_4(w)|$$

decreases by one, and no standing invariant of the current branch is weakened. Hence an infinite target-defect peeling loop is impossible.

Proof. Let $\xi = (X, w, u, B_u)$ be a two-carrier target-defect entry in the current residual. Entries of the current residual are indexed by unpeeled routed loads; a load already in $P_4(w)$ has left the Type A receiver calculation by [Remark 14.41](#). Hence $u \notin P_4(w)$. By [Lemma 15.44](#), the pressure token of ξ is an exit-(4) witness for u . Applying [Lemma 14.38](#) replaces $P_4(w)$ by $P_4(w) \cup \{u\}$. Thus $L_4(w)$ decreases by one and every other residual load count is unchanged, so Λ_4 decreases by one.

The operation changes only the bookkeeping ledger: the graph G , the packed P_{13} -windows, the boundary degree profiles, the dyadic-safety constraints, and all target-response quotients remain the same. The selected load is removed from the pure Type A receiver sum, and [Lemma 14.34](#) gives the exact 1/4-unit charge update. Thus the next stage satisfies the same standing invariants; if its large-budget deficit has disappeared, the branch is closed, and otherwise the same large-budget reduction applies to the updated unpeeled ledger. Since Λ_4 is a nonnegative integer, the target-defect peeling step cannot occur indefinitely. $\square$

15.1.7 Large-budget closure from pressure accounting and finite descent

With the target-defect class accounted for, the large-budget branch closes by a finite exit-(4) descent. The pressure ledger proves that no target-defect unit is missing from the carrier count; the descent lemma then identifies each surviving target-defect two-carrier unit with one of the canonical peelable loads already present in the saturated Type A analysis.

Theorem 15.48 (Large-budget closure with target-defect pressure accounting). *Assume [Definition 4.1](#). No minimal counterexample survives in the large-budget branch.*

Proof. Start with the peeling sets $P_4(w) = \emptyset$ and run the following deterministic procedure. At any stage, let $\tilde{D}_A^{P_4}$ be the remaining Type A deficit of [Definition 15.45](#). If

$$\tilde{D}_A^{P_4} < \left(\frac{1}{4} - \tau_{\text{win}}\right)|R| - o(|R|),$$

then the still-unresolved negative mass is too small to realize the large-budget branch: all supports outside the P_4 -reduced Type A ledger are nonnegative or belong to the Type B bridge ledger, and the latter has total negative mass $o(|R|)$ by [Proposition 14.107](#). Hence the large-budget net-deficiency cap and [Theorem 15.1](#) close that stage. Thus a surviving stage must satisfy

$$\tilde{D}_A^{P_4} \geq \left(\frac{1}{4} - \tau_{\text{win}}\right)|R| - o(|R|).$$

[Lemma 15.46](#) then produces a two-carrier unified entry in the current unpeeled receiver ledger.

If that two-carrier entry is a route-8 entry, then [Proposition 14.62](#) applies: the entry gives the terminal two-carrier route-8 obstruction and [Theorem 14.61](#) excludes it. If instead the entry is target-defect, then [Lemma 15.44](#) makes its pressure token a canonical exit-(4) witness for its routed load, and [Lemma 15.47](#) performs one peeling step.

The second alternative cannot occur infinitely many times, because the integer Λ_4 of [Lemma 15.47](#) strictly decreases at each target-defect peel. Hence the procedure terminates. It cannot terminate with a large-budget survivor. Indeed, a terminal surviving stage still satisfies the displayed lower bound on $\tilde{D}_A^{P_4}$, so the reduced carrier reduction again supplies a two-carrier entry. The target-defect case would perform another peel and lower Λ_4 , contrary to terminality. The two-carrier entry is therefore route 8, and the previous paragraph excludes it. This contradiction proves that the large-budget branch has no minimal counterexample. $\square$

Corollary 15.49 (Large-budget closure with sublinear open pressure). *Assume [Definition 4.1](#). After the type (A2) proper-compression and delocalization alternatives of [Definition 15.20](#) have been removed, if*

$$P_{\text{open}} = o(|R|),$$

then no minimal counterexample survives in the large-budget branch.

Proof. The normalized limsup of $P_{\text{open}}/|R|$ is 0. The threshold equals

$$12 \left(\frac{1}{4} - \tau_{\text{win}} \right) - \tau_{\text{win}} = 0.03372671002\dots,$$

which is positive. Therefore [Proposition 15.22](#) closes the target-defect exit-(4) residual. The route-8 two-carrier residual is closed by [Theorem 14.61](#), and the remaining unified entries are already paid by the pressure ledger or by the Type B bridge ledger used in [Lemma 15.4](#). Hence no large-budget survivor remains. $\square$

Remark 15.50 (Unified pressure accounting). *A lower bound for $\sum_i \mathcal{N}(X_i)$ that uses only the route-8 collection and the Type B bridge ledger is not sufficient. Such a bound has the form*

$$\sum_i \mathcal{N}(X_i) \geq -D_A(\mathcal{X}_A^0) - 208S_B,$$

where the right-hand side sums only over the route-8 collection $\mathcal{X}_A^0$ and the Type B bridge ledger. It is insufficient to combine this estimate with [Theorem 15.1](#) by treating every other support as nonnegative: [Theorem 15.1\(a\)](#) proves $\mathcal{N}(X) \geq 0$ only for supports with no exit-(4) witness, and

a support that peels routed loads through the target-defect exit has $\mathcal{N}(X) \geq -\frac{1}{4}|P_4| < 0$. Such a support is neither nonnegative nor (in general) a route-8 residual nor a Type B bridge residual, so its negative mass was omitted from the lower bound, and the deduction “ $\mathcal{X}_A^0$ carries the large-budget deficit” did not follow: the deficit could reside entirely in the target-defect class. [Definition 15.2](#) removes the omission by summing over all negative Type A supports at once; [Lemma 15.7](#) shows the extra (target-defect) entries obey the same essential-carrier lower bound as route-8 entries – this is where the minimality-free [Lemmas 14.56](#) and [14.57](#) are essential. The target-defect mass is assigned to the 2/3-pressure ledger of [Definition 15.12](#). The no-overcount lemma [Lemma 15.13](#) proves the exact raw boundary-incidence inequality

$$3\tilde{N} - P_{\text{ext}} \leq \partial_+(R), \quad P_{\text{ext}} = N_2 + 3N_{\text{res}},$$

and [Proposition 15.43](#) gives the raw closure threshold

$$P_{\text{ext}} < \left(12 \left(\frac{1}{4} - \tau_{\text{win}}\right) - \tau_{\text{win}}\right) |R| - o(|R|).$$

The absorption ledger of [Definition 15.20](#) refines this by charging additional single-use units before the final squeeze is applied. After those absorptions, the same numerical threshold applies to the explicitly defined open remainder P_{open} by [Proposition 15.22](#). The window-blocker ledger then assigns every open unit to a unique packed P_{13} -window. Same-window overload triples that realize a direct dyadic cycle, an unused boundary absorber, a proper-compression or delocalization absorber, or Type B fan-window data are removed by the listed mechanisms in [Definition 15.25](#). After those routed triples are removed, [Lemma 15.31](#) removes every overload triple containing a pair of open blockers with a recorded window-shadow hit. The exact residual quantity beyond the two-blocker-per-window base is therefore the zero-shadow primitive excess P_{zero}^+ , where no open pair on the overloaded window has such a hit, and

$$P_{\text{open}} \leq 2p_{13} + P_{\text{zero}}^+.$$

The remaining numerical slack is

$$\varepsilon_{\text{prim}} = 0.03372671002 \dots - 0.03042331579 \dots = 0.00330339423 \dots$$

Thus [Proposition 15.37](#) closes exit (4) whenever $P_{\text{zero}}^+ < \varepsilon_{\text{prim}}|R| - o(|R|)$. The same-window two-blocker cap is the special case $P_{\text{zero}}^+ = o(|R|)$, closed by [Lemma 15.41](#). No target-defect unit is discarded or counted twice: every unit is charged to the original 2/3-pressure ledger, to an unused boundary incidence, to a routed proper-compression or delocalization certificate, to a Type B handoff, to a recorded dyadic window-shadow route, or to a unique zero-shadow window-overload slot.

Framework branch table. The target-defect pressure section closes the remaining large-budget residual by the following table.

Branch	Residual condition	Demand	Payer	Bounded multiplicity	Closure
--------	--------------------	--------	-------	----------------------	---------

Unified negative collection	A negative Type A support remains.	$\delta(X) = -\mathcal{N}(X)$.	The unified Type A ledger.	The canonical decomposition.	Definition 15.2 and Lemma 15.4 .
Pressure ledger	Target-defect or route-8 entries remain.	Pressure units.	Boundary incidences.	The no-overcount pressure ledger.	Lemma 15.13 .
Open pressure	Boundary incidences do not pay all target-defect units.	Open pressure.	P_{13} -attachment accounting.	Same-window and primitive excess caps.	Propositions 15.22 , 15.37 and 15.42 .
Finite descent	An exit-(4) peel is available.	One routed load.	The integer Λ_4 .	Strict decrease by one.	Lemmas 15.44 and 15.47 .
Terminal route 8	No target-defect peel remains.	A two-carrier route-8 entry.	Deletion witnesses.	The canonical carrier core.	Proposition 14.62 and Theorem 14.61 .

A Finite curvature enumeration code

This appendix reproduces the finite enumerations that certify the constants used above. Each listing is self-contained and was run under CPython 3 (standard library, with `numpy` for the second block); the displayed output immediately below each listing is the certificate. The first block certifies the legal-label count $|\mathcal{L}| = 399$ of [Lemma 6.1](#) together with the two-step curvature counts 543958, 432672, 111286 of [Lemma 6.3](#), and hence the value of c_Ω .

```

from itertools import combinations
from collections import Counter
import math

labels=[]
for mask in range(1,1<<13):
    S=[i for i in range(13) if (mask>>i)&1]
    ok=True
    for a,b in combinations(S,2):
        if abs(a-b) in (2,6):
            ok=False
            break
    if ok:
        labels.append(tuple(S))

print(len(labels))
print(Counter(map(len,labels)))

powers={2**k for k in range(2,20)}

def C(s,S,T):
    return all(s+2+abs(i-j) not in powers for i in S for j in T)

W=0

```

```

pos=0
flat=0
for S in labels:
    for A in labels:
        if not C(1,S,A):
            continue
    for T in labels:
        if not C(1,A,T):
            continue
    W+=1
    if not C(2,S,T):
        pos+=1
    else:
        flat+=1

print(W,pos,flat)
print(pos/W, flat/W, math.log2(W/flat))

```

The output is:

```

399
Counter({3: 122, 4: 122, 5: 63, 2: 60, 6: 17, 1: 13, 7: 2})
543958 432672 111286
0.7954143518433409 0.20458564815665917 2.28922315244071

```

The second block certifies the constant

$$c_{13} = 118.108581006\dots$$

used in [Remarks 7.1](#) and [7.67](#), as the sum of the 91 finite curvature barriers; it uses matrix multiplication to avoid an explicit $399^3 \cdot 91$ loop.

```

from itertools import combinations
import numpy as np, math

L=[]
for mask in range(1,1<<13):
    inds=[i for i in range(13) if (mask>>i)&1]
    ok=True
    for a,b in combinations(inds,2):
        if abs(a-b) in (2,6):
            ok=False
            break
    if ok:
        L.append(mask)
assert len(L)==399

def indices(mask):
    return [i for i in range(13) if (mask>>i)&1]

Ms=[]
for s in range(15):
    forb={2**r-s-2 for r in range(2,20) if 0 <= 2**r-s-2 <= 12}
    M=np.ones((len(L),len(L)), dtype=np.int8)
    for i,S in enumerate(L):

```

```

    Si=indices(S)
    for j,T in enumerate(L):
        Tj=indices(T)
        bad=any(abs(a-b) in forb for a in Si for b in Tj)
        if bad:
            M[i,j]=0
    Ms.append(M)

total=0.0
vals=[]
for a in range(1,15):
    for b in range(1,15-a): # a,b >= 1 and a+b <= 14
        middle_counts=Ms[a].astype(np.int32) @ Ms[b].astype(np.int32)
        W=int(middle_counts.sum())
        bad=int((middle_counts*(1-Ms[a+b])).sum())
        flat=W-bad
        val=math.log2(W/flat)
        vals.append((a,b,W,bad,flat,val))
        total += val

print(len(vals))
print(total)
print(min(vals, key=lambda x: x[-1]))

```

The output is:

```

91
118.10858100609198
(4, 4, 6648206, 1190878, 5457328, 0.284770329720...)

```

The formal companion stores the same fifteen 399×399 Boolean matrices as fixed bit rows. Fifteen independent audit shards recheck every entry against the power-of-two attachment predicate and recheck every accepted safe and flat count against the generic finite-relation formula. In particular, it recovers the first-block counts at $(a,b) = (1,1)$, proves that exactly 91 pairs are present, and certifies the rounding-free bound

$$2^{118} \prod_{a,b} F_{a,b} < \prod_{a,b} W_{a,b}.$$

This is the finite node [21] certificate. The construction and routing of independent window coordinates belong to the subsequent hot/cold branch at nodes [22]–[24]; they are not premises of the finite computation.

The third block reproduces the Type B fan constants used in [Remark 14.67](#) and [Lemma 14.98](#). The proof of those constants is the structural label-packing argument in [Lemma 14.66](#). The listing builds the same graph on the 399 labels with adjacency relation $C_2(S, T) = 1$, computes its clique number, and computes the marked variant corresponding to the non-singleton-label cap. The residual class in [Definition 14.71](#) is determined by the positive closed-neighbour deficit $D_B > 0$: for $k \leq 7$ this means at least two cubic-closed neighbours, while for $k = 8$ it also includes the single cubic-closed border case.

```

from fractions import Fraction
from itertools import combinations

labels=[]

```

```

for mask in range(1,1<<13):
    S=tuple(i for i in range(13) if (mask>>i)&1)
    if all(abs(a-b) not in (2,6) for a,b in combinations(S,2)):
        labels.append(S)

powers={2**k for k in range(2,20)}

def C(s,S,T):
    return all(s+2*abs(i-j) not in powers for i in S for j in T)

n=len(labels)
adj=[0]*n
edges=0
for i,S in enumerate(labels):
    for j in range(i+1,n):
        if C(2,S,labels[j]):
            adj[i] |= 1<<j
            adj[j] |= 1<<i
            edges += 1

def maximum_clique(mask):
    best_size=0
    best_bits=0
    def color_sort(P):
        verts=[]
        colors=[]
        U=P
        color=0
        while U:
            color += 1
            Q=U
            while Q:
                bit=Q & -Q
                v=bit.bit_length()-1
                verts.append(v)
                colors.append(color)
                U &= ~bit
                Q &= ~bit
                Q &= ~adj[v]
            return verts, colors
    def expand(size,bits,P):
        nonlocal best_size,best_bits
        if not P:
            if size > best_size:
                best_size=size
                best_bits=bits
            return
        verts,colors=color_sort(P)
        for idx in range(len(verts)-1,-1,-1):
            if size + colors[idx] <= best_size:
                return
            v=verts[idx]
            bit=1<<v
            if P & bit:

```

```

        expand(size+1,bits|bit,P & adj[v])
        P &= ~bit
        if size + P.bit_count() <= best_size:
            return
    expand(0,0,mask)
    return best_size,best_bits

omega,wit=maximum_clique((1<<n)-1)
omega_ns=0
wit_ns=0
for i,S in enumerate(labels):
    if len(S) >= 2:
        size,bits=maximum_clique(adj[i] | (1<<i))
        if size > omega_ns:
            omega_ns=size
            wit_ns=bits

def bit_labels(bits):
    return [labels[i] for i in range(n) if (bits>>i)&1]

simple_charge_k_le_7=min(
    (Fraction(3-k,1)-Fraction(1,4)-Fraction(1,4)
     + Fraction(3,4)*(k-1), k)
    for k in range(4,8)
)
charge_k8_no_closed=Fraction(3-8,1)-Fraction(1,4)+8*Fraction(3,4)

print('labels', n)
print('fan_safe_edges', edges)
print('omega_C2', omega)
print('omega_C2_witness', bit_labels(wit))
print('omega_with_nonsingleton', omega_ns)
print('nonsingleton_witness', bit_labels(wit_ns))
print('simple_charge_k_le_7', simple_charge_k_le_7[0], 'at_k',
      simple_charge_k_le_7[1])
print('charge_k8_no_closed', charge_k8_no_closed)

```

The output is:

```

labels 399
fan_safe_edges 18301
omega_C2 8
omega_C2_witness [(1,), (2,), (3,), (4,), (9,), (10,), (11,), (12,)]
omega_with_nonsingleton 7
nonsingleton_witness [(0, 1), (2,), (3,), (8,), (9,), (10,), (11,)]
simple_charge_k_le_7 0 at_k 7
charge_k8_no_closed 3/4

```

Remark A.1 (Fan-packing reproducibility check). *The fan-packing bound used in the Type B proof is established structurally in [Lemma 14.66](#). The listing above is retained only as a reproducibility check of the same constants for the 399-label graph; it is not an additional premise in the proof.*

Framework branch table. The enumeration appendix is a computational-certificate register. Each row states which finite datum is exported and where the proof consumes it.

Branch	Residual condition	Demand	Payer	Bounded multiplicity	Closure
Legal-label certificate	All subsets of $\{0, \dots, 12\}$ are tested for forbidden gaps.	$ \mathcal{L} = 399$ and the size distribution.	The finite subset enumeration.	The predicate in the code is the predicate in Lemma 6.1 .	Consumed by Lemma 6.1 .
Curvature certificate	All ordered legal-label triples are tested.	543958, 432672, 111286	The finite curvature tensor.	The code applies C_1, C_2 exactly as defined in the text.	Consumed by Lemma 6.3 .
Window-package rate	All 91 pairs $a, b \geq 1$, $a + b \leq 14$, are summed.	c_{13} .	The matrix-computed barrier list.	The printed count 91 verifies that every scale pair is included once.	Consumed by Remarks 7.1 and 7.67 .
Type B reproducibility	The 399-label fan graph is checked independently.	Clique and marked-clique constants.	The structural proof of Lemma 14.66 .	The code is a reproducibility check, not a separate premise.	Checks Remark 14.67 and Lemma 14.98 .




Figure 30: External-pressure accounting for exit-(4) target-defect tokens. The three-carrier class is paid directly by disjoint boundary incidences. The two-carrier paid class receives two disjoint boundary incidences and leaves one unit of pressure defect; this class includes genuinely cheap two-carrier tokens and tokens whose additional carriers collide with earlier ledger choices. The unassigned class leaves three units of pressure defect. Each two-carrier paid or unassigned target-defect token is represented by either an actual outside corridor in $G - X$ or by a profile-only exact response coordinate. These give the exact split $P_{\text{ext}} = P_{\text{act}} + P_{\text{prof}}$. The absorption ledger then assigns pressure units either to unused boundary incidences, which improve the same $\partial_+(R)$ count, or to profile-dependence certificates, which route to proper-compression or delocalization alternatives. Target-defective dependence outputs re-enter the pressure ledger instead of being treated as absorbed. The only pressure left in the surviving branch is P_{open} . The window-blocker ledger assigns every open unit to a unique packed P_{13} -window. A same-window overload triple is routed when it realizes a direct dyadic cycle, an unused boundary absorber, a proper-compression or delocalization absorber, or a Type B fan-window handoff. After these routed triples and their recorded window-shadow dyadic hits have been removed, the only extra term beyond two blockers per window is the zero-shadow primitive excess P_{zero}^+ , with $P_{\text{open}} \leq 2p_{13} + P_{\text{zero}}^+$. Thus [Proposition 15.37](#) closes the branch whenever $P_{\text{zero}}^+ < \varepsilon_{\text{prim}}|R| - o(|R|)$. The same-window two-blocker cap is the special case $P_{\text{zero}}^+ = o(|R|)$, where the strict numerical inequality is $2\theta_{\text{win}}/(1 - 13\theta_{\text{win}}) = 0.03042331579\dots < 0.03372671002\dots = \varepsilon_{\text{press}}$.




Figure 31: Visual proof of the branch-closure mass ledger. Panel I shows the large-budget cap localizing negative net charge to a connected admissible support X with $\mathcal{N}(X) < 0$. Panel II shows the local branch drain used in [Theorem 15.1](#): Type A supports close through the non-route-8 exits or target-defect peeling, while Type B supports close through the B1 ledger together with B2, giving $\mathcal{N}(X) \geq 0$ outside the explicit residuals. Panel III records the remaining global ledger. Type B bridge and fan-mass residuals are charged to assigned surplus units and become $o(|R|)$ after route-8 non-window cores are extracted into the Type A ledger. The unified pressure ledger counts both target-defect and route-8 Type A entries. Target-defect two-carrier entries are canonical exit-(4) peel data, so finite descent removes them. The terminal survivor would have to be a two-carrier route-8 obstruction, which is excluded by the carrier-reduction and no-go figures.

B Proof resilience: residual counterexamples under lemma failure

For every *closing cell* (terminal node) of the proof-dependency diagram, this appendix records what survives if that cell’s closing lemma is *false* and no redundancy covers it. The entry is the forced structure of a hypothetical minimal counterexample confined to that cell—that is, the reduction theorem the proof degrades into at that leaf. The logic is that a closing lemma does not merely delete a possibility when it fails; its negation is a structural hypothesis. A surviving counterexample must satisfy *all standing invariants accumulated up to that node* (it does not lose them) *together with* the negation of the failed lemma’s conclusion. Thus every degradation cell is a specified residual problem, not an unconstrained failure mode.

Notation is as in the manuscript: G is the minimal counterexample, $\delta(G) \geq 3$, R is the P_{13} -free remainder, $\theta = p_{13}/n$, $\partial_+(X) = \sum_v \max(0, 3 - d_X(v))$, $\sigma(X) = \sum_v \max(0, d_G(v) - 3)$, $\mathcal{N}(X) = \partial_+(X) - \sigma(X) - \frac{1}{4}|V(X)|$, W_2 is the two-step curvature supply, c_Ω the flatness entropy cost, and $\beta = m - n + 1$. “Standing invariants” are the invariants established at earlier-or-equal nodes; they are retained in every residual cell below.

Framework branch table. The resilience appendix already gives the detailed degradation rows below; the framework branch table for the appendix is the following compact index.

Branch	Residual condition	Demand	Payer	Bounded multiplicity	Closure
Covered leaf	A listed closure is removed but a redundant route remains.	The residual cell.	The redundant closure.	The standing invariant set.	Recorded as covered in the detailed tables.
Dense-window auxiliary	The window entropy branch degrades to dense-window blocks.	Quiet-block demand.	The target-compatible external input.	Canonical block boundaries.	Lemma B.2.
Global smearing	The rank-drop branch degrades to a whole-graph dependence.	Closed exact-profile data.	Minimality or label-injectivity.	The empty-boundary profile.	Lemma B.3.
Type A failure	The route-8 or exit-(4) closure is weakened.	Terminal carrier package.	The carrier ledger.	Declared witnesses.	Lemma B.4.
Type B failure	The fan ledger or B2 closure is weakened.	Bridge or overlap residual.	The fan-mass ledger.	Assigned surplus.	Remark B.5.

How to read a row.

Column	Meaning
Node	Terminal node id $[k]$ from the diagram, with tikz id.

Column	Meaning
Closing lemma	The single labeled result that discharges this leaf.
Closing condition	The inequality or identity whose truth closes the leaf.
Slack	Margin by which it currently closes (how wrong the constant can be).
Redundant cover?	Whether another standing route re-closes this cell if the lemma fails.
Residual counterexample	Forced structure of a surviving G if the lemma is false and no cover applies.

A. Trivial / external-input leaves. These do not degrade to nontrivial residual reductions: their negation is either a direct cycle (so G is not a counterexample at all) or a failure of an external theorem, not of this proof.

Node	Closing lemma	Closing condition	Slack	Redundant cover?	Residual counterexample (if false)
[3] notbad	def. of counterexample	$\delta(G) \geq 3$ and no C_{2^j}	exact (definitional)	—	None — if it fails, G already has a power-of-two cycle (not a counterexample).
[7] cycle	Lemma 2.2	a Mersenne return $\iff$ a 2^k cycle	exact (arithmetic identity)	—	None — failure means the return/cycle dictionary is misstated; G would have a 2^k cycle. Pure bookkeeping.
[16] hss	Theorem 5.1 (HSS, external)	P_{13} -free + $\delta \geq 3 \Rightarrow 2^k$ cycle	external	—	Not this proof’s leaf. The proof uses the Hegde–Sandeep–Shashank theorem as its black-box external input; if that theorem is unavailable, the window/remainder spine has no entry point.

B. Density / supply leaves (close by counting; degrade to “dense-window” reductions).

Node	Closing lemma	Closing condition	Slack	Redundant cover?	Residual counterexample (if false, no cover)
[23] overflow	Remark 7.67 via Eq. (1)	$\theta \leq 0.011985$ (else window entropy > skeleton budget)	feasibility gap $0.1 + 116.8\theta \leq 1.5$; binds only at $\theta \approx 0.012$	Target-compatible auxiliary closure. The dense correlated-window residual is closed in Lemma B.2 by the stated quiet-block input.	G has a <i>dense induced-P_{13} packing</i> : $\theta > 0.011985$, so the maximal disjoint packing covers > 15.6% of vertices, yet the window entropy still fits under the skeleton budget. The target-compatible auxiliary closure shows that such entropy-cheap density must either create close cycle lengths in a dyadic-safe atom or pay quadratic boundary, in both cases restoring a contradiction.
[29] $\rightarrow$ [28] stub	Lemma 8.4	each induced P_{13} exports $\leq 15 + \sigma(P)$ incidences	exact (12-edge identity); only θ feeds the normalized form	Framework-independent. Uses only $\delta \geq 3$, “induced P_{13} has 12 edges”, and disjointness.	If the constant 15 were miscounted, the residual G would be one whose remainder borrows more than 15 incidences per window, contradicting the induced-path edge count. Thus failure here is a numerical misstatement of the incidence identity, not a new graph-theoretic residual.

C. Curvature / rank leaves (close by supply-vs-demand; degrade to “flat-remainder” reductions).

Node	Closing lemma	Closing condition	Slack	Redundant cover?	Residual counterexample (if false, no cover)
[30] wedge (demand floor)	Lemma 9.1	under Definition 4.1, $W_2(R) \geq 2.543 \cdot R $; high entropy gives $2.574 \cdot R $	$4.1 \times$ over the 0.611 supply cap	Yes inside the near-cubic spine. The net-charge collision (diagram node [60], 0.25 vs 0.2282) reaches a contradiction without the curvature magnitude; Remark 12.4 states closure already holds from the window-only cap.	G has a <i>curvature-poor remainder</i> : $W_2(R) < 2.543 \cdot R $, i.e. the P_{13} -free remainder admits fewer than 2.543 two-step curvature tests per vertex. With the standing invariants (empty 3-core, deficiency ≤ 0.2282) this forces an unusually <i>flat</i> atom remainder—locally tree-like wedge structure avoiding curvature-positive triples. <i>Reduction</i> : “a counterexample’s remainder is curvature-flat below density 2.543; classify flat P_{13} -free atoms.” Recoverable via the net-charge route, which does not consume the wedge magnitude.
[32]/[47] fullrank	Lemma 10.10	$r_\Omega(R) \geq W_2 - o(R)$ (no rank drop)	absorbs $o(R)$; rank-drop branch D separately closes at [46]	Yes. Branch D (Lemma 10.9, inv 34/35) independently closes the rank-drop case; the two-budget split also re-routes.	G has a <i>rank-deficient curvature operator</i> : $r_\Omega(R) < W_2 - \Omega(R)$, so a linear fraction of curvature tests are dependent. The proper-support cases are routed to target defect, compression, or proper smearing. The whole-graph case is closed by the exact closed response profile: empty context does not identify labelled curvature coordinates, and a non-injective closed quotient must produce a smaller closed counterexample.
[45] globalterm	Lemma 10.9	no silent global delocalization of dyadic support	qualitative (closed exactness is label-injective on raw curvature tests)	Closed in the main text. Lemma B.3 records the corresponding degradation-cell closure.	A surviving row [45] residual would require a target-complete whole-graph quotient of $\rho_{\mathcal{D}}^{\text{ex}}(G)$ that is not label-injective on raw curvature tests and has no smaller closed representative. This contradicts Definition 3.98 and Lemma 10.9.
[50]/[54] entcap	Proposition 11.2 / Proposition 11.9	high- θ branch: entropy cap $c_\Omega \cdot \text{rank}$ exceeds budget when the rank input is available	routing is exhaustive; only the high-entropy and repetitive low-entropy subcases add new rank/entropy information	Partial routing. The nonrepetitive low-entropy subcase is deliberately passed to the large-budget net-charge analysis, not closed here.	G sits in a <i>nonrepetitive low-entropy regime</i> : $\eta(R) < \frac{1}{10} \log_2 n$ per vertex, the radius-2 type coordinate has linear bit-content, and the local-type dominance lemma does not apply. <i>Reduction</i> : this branch carries no extra curvature closure claim and is handled only by the later net-charge/local-residual analysis.

D. Net-charge terminus leaves (close by discharging; degrade to “charged-atom” reductions).

Node	Closing lemma	Closing condition	Slack	Redundant cover?	Residual counterexample (if false, no cover)
[60] contrad	Proposition 13.45 + Lemma 13.5	$\frac{1}{4} R \leq \partial_+(R) - \sigma(R) = \partial_+(R) - \sigma_R \leq \tau_{\text{win}} R + o(R)$ with $\tau_{\text{win}} < \frac{1}{4}$	$\sim 9.6\%$ (0.25 vs 0.2282 per $ R $)	Yes. The second independent collision form; the curvature route ([30]) reaches a contradiction on its own with large slack.	G has a <i>net-charge-negative residual that does not localize</i> : $\partial_+(R) - \sigma(R) < \frac{1}{4} R $ globally, yet net charge fails to concentrate on a connected support (superadditivity localization fails). <i>Reduction</i> : “a counterexample distributes a $\frac{1}{4}$ -deficit across R with $\mathcal{N}(X) \geq 0$ on every connected piece”—forbids the very cell-wise accounting the discharging needs; isolates a precise superadditivity gap to repair.
[91] atomthm (Type A, $\sigma = 0$)	Lemma 14.45 via Lemmas 14.27, 14.34 to 14.38, 14.40 and 14.43	outside target-defect peeling, route 8, and the Type B handoff, $\partial_+(X) \geq \frac{1}{4} V(X) \iff 2 E(X) / V(X) \leq 2.75$, so $\mathcal{N}(X) \geq 0$	raw thresholds $H_0 \leq 4$, $H_1 \leq 8$, $H_2 \leq 12$; exits 1–3, 5, and 6 close; exit 4 peels one target-defective routed load with exact 1/4 receiver-charge accounting; exit 7 hands off to Type B; then unsaturated discharging $11n_0 + 7n_1 + 3n_2 \geq n_3$	Local split. The 4/8/12 saturated cases in exits 1–3, 5, and 6 are removed by target, compression, and delocalization, while exit 4 removes its supported target-defective load before the 3/7/11 charge calculation. Exit 7 is represented by Type B, and route 8 is the reduced silent trace-basin branch of Definition 14.15 .	G would contain a <i>dense low-deficiency admissible atom</i> : a connected, $\sigma = 0$, empty-3-core, P_{13} -free support with average degree > 2.75 . The Type A split sends exits 1–3, 5, and 6 to closed contradictions, exit 4 to target-defect peeling, exit 7 to Type B, and the remaining unsaturated case to discharging; if the route-8 branch carries the large-budget Type A deficit, the declared u -supported response algebra and canonical minimal target-complete carrier cores reduce it to a two-carrier route-8 residual, with quiet-block boundedness recorded in Lemma B.4 .

Node	Closing lemma	Closing condition	Slack	Redundant cover?	Residual counterexample (if false, no cover)
[74] fan thm (Type B, $\sigma > 0$)	Lemma 14.98 and Proposition 14.107 via Definition 14.18, Definition 14.68, Definition 14.72, Definition 14.73, Definition 14.74, Lemma 14.75, Lemma 14.76, Lemma 14.77, Lemma 14.78, Lemma 14.79, Definition 14.83, Definition 14.84, Definition 14.85, Lemma 14.87, Lemma 14.89, Proposition 14.90, Definition 14.91, Proposition 14.97, Definition 14.71, Definition 14.101, Lemma 14.102, Lemma 14.104, Lemma 14.105, Lemma 14.106, Remark 14.67, and Lemma 14.66	outside route 8, outside fan-certificate residual centers, and outside B2 failure, every certificate-marked fan has $d_G(h) \leq 8$, every nonpositive-deficit marked fan has nonnegative closed- neighbourhood charge, every positive-deficit fan has the local hybrid B1 incidence ledger, and B2 supplies the disjoint ledger needed to sum the fans; in the near-cubic spine, residual centers and B2 failures satisfy $M_B \leq 416\sigma(G) = o(R)$ after any route-8 non-window core is extracted into D_A and grouped-envelope centers are counted at most twice	certificate-closed iff $D_B \leq 0$; every positive-deficit residual has $D_B = c - \frac{11-k}{4}$ paid by half-credit on the $2c$ disjoint non- h incidences, split into window credit $I_W/2$ and non-window demand D_N ; each bridge residual is charged by at most 208 times its assigned surplus plus the extracted route-8 Type A deficit	Structural reduction to B2 plus fan-mass closure. Direct fan-window and two-window cycles close by explicit path construction; local B1 is the hybrid incidence budget; every disjoint-carrier B2 failure is represented by a minimal overlap obstruction carrying contextual dyadic safety, window-stub capacity, replacement/delocalization routing, and a sublinear fan-mass budget in the near-cubic spine after route-8 cores are charged to D_A ; fan-certificate residual centers and grouped envelopes enter this fan-mass bound directly.	In the near-cubic spine, a surviving Type B bridge obstruction has only sublinear bridge mass outside the Type A pressure/route-8 ledger, so it cannot carry the linear large-budget obstruction without entering node [123].

Node	Closing lemma	Closing condition	Slack	Redundant cover?	Residual counterexample (if false, no cover)
[76]/[85] branch (local closures)	Theorem 15.1	combines [30]+[60]+entropy: outside explicit Type A target-defect/route-8 and Type B bridge residuals, the entropy-admissible branch has no survivor	inherits the available component slacks: about $4.1\times$ in the curvature form and about 9.6% in the net-charge form; Type B bridge mass is $o(R)$ after route-8 cores are extracted into the Type A ledger	Local join. It leaves only the explicit Type A pressure/route-8 ledger for node [123].	<i>G survives the branch squeeze</i> only if the negative local support contributes to the explicit Type A pressure/route-8 ledger; Type B fan-certificate residual centers and B2 failures are too small in total to carry the required linear deficit.
[123] closed (pressure descent)	Theorem 15.48	after the branch-kill reduction, the remaining large-budget Type A target-defect/route-8 ledger either peels a target-defect load by exit (4) or forces a two-carrier route-8 entry	inherits the unified burden lower bound, the pressure no-overcount ledger, finite descent of Λ_4 , the route-8 burden lower bound, and the private-carrier upper bound; Type B bridge mass has already been removed at [76]/[85]	Join node feeding the terminal no-go. Target-defect two-carrier entries decrease Λ_4 ; the terminal non-peeling case enters node [124], which is discharged by carrier deletion.	<i>G can survive node [123]</i> only if the finite exit-(4) descent has terminated and the remaining entry is the terminal two-carrier route-8 obstruction through the canonical carrier-core ledger.

Node	Closing lemma	Closing condition	Slack	Redundant cover?	Residual counterexample (if false, no cover)
[124] nogo (two-carrier route 8)	Definition 14.50 , Lemma 14.53 to 14.55 , and Theorem 14.61 , with Definition 14.60 and Proposition 14.62	terminal two-carrier route-8 obstruction exists	inherits the route-8 lower bound, the private-carrier upper bound, cut parity, declared deletion witnesses, and all absent-exit constraints	Terminal local no-go. In a two-carrier entry, declared essential deletion witnesses first give canonical exit-(4) quotients and then realize exit (4), which is absent in a true route-8 residual.	A false instance is exactly a terminal two-carrier route-8 obstruction: a true route-8 entry with at least two essential carriers, at most two private essential carriers, and declared carrier-deletion witnesses for every essential carrier.

Remark B.1 (Tight margin in the route-8 carrier reduction). *The private-carrier reduction at node [123] uses the numerical inequality*

$$12 \left(\frac{1}{4} - \tau_{\text{win}} \right) > \tau_{\text{win}}, \quad \text{equivalently} \quad \tau_{\text{win}} < \frac{3}{13}.$$

With the printed constants, $\tau_{\text{win}} = 0.22817486846\dots$, while $3/13 = 0.23076923076\dots$. The relative margin is therefore about

$$\frac{3/13 - \tau_{\text{win}}}{3/13} = 0.01124\dots$$

This is the tight carrier-reduction margin. Any later argument that spends additional boundary-incidence or stub capacity before [Proposition 14.59](#) must recheck this inequality.

External-input closures for partial degradation cells. The next three lemmas use target-compatible forms of the auxiliary external inputs listed in the constraint ledger. The named external theorems supply the graph-theoretic alternatives; the target-compatible clauses below state exactly how those alternatives interact with the Mersenne-return algebra and the canonical support ledger of this paper. Thus no closure below depends on an unstated propagation step. For a connected support Y , write $b(Y) = |\partial Y|$ for its boundary size in the ambient counterexample.

Lemma B.2 (Dense-window residual closes by quiet-block input). *Assume the following target-compatible Bondy–Vince/Gao–Ma input for the canonical blocks generated by a row [23] dense-window residual.*

- (a) *Each canonical block Y is almost cubic, and its sub-cubic vertices are boundary vertices charged to the packed-window incidence ledger of [Lemma 8.4](#).*
- (b) *If Y exports two cycle lengths differing by 1 or 2, then the corresponding two edge-rooted returns lie in one boundary fibre and force either a return length in $\mathcal{M}$ or a target-complete replacement in the sense of [Lemma 3.103](#).*
- (c) *If Y is quiet, then $|Y| < 5b(Y)^2$, and the canonical quiet blocks are boundary-disjoint up to packed-window interfaces; summing their boundary charges restores the feasibility inequality [Eq. \(1\)](#).*

Then the dense-window residual in row [23] is empty.

Proof. Suppose row [23] survives. Then the maximal induced- P_{13} packing has $\theta > 0.011985$, while the corresponding window states remain below the skeleton entropy bound. Apply the target-compatible input to the canonical blocks of this residual. In the exported-length alternative, clause (b) gives a Mersenne return or a target-complete replacement. The first contradicts Lemma 2.2; the second contradicts minimality through Lemma 3.103. In the quiet alternative, clause (c) restores Eq. (1), contradicting $\theta > 0.011985$ in the entropy-cheap dense-window branch. The two alternatives exhaust the input, so row [23] is empty. $\square$

Lemma B.3 (Global smearing closure). *The rank-one global-mode residual in row [45] is empty.*

Proof. Let a row [45] residual survive. By Lemma 10.9, the corresponding global dependence is target-defective, has a strictly smaller admissible closed representative, or contributes exact label-injective global response data. The first alternative cannot witness a target-dependence by Definition 10.2; the second contradicts the lexicographic minimality of G ; and the third supplies no rank reduction of $r_\Omega(R)$ by Definitions 3.95 and 3.98. This contradicts the definition of a row [45] rank-one global-mode residual. Hence row [45] is empty. $\square$

Lemma B.4 (Route-8 Type A residual is bounded by quiet-block input). *Assume the following target-compatible Bondy–Vince/Gao–Ma input for route-8 Type A supports.*

- (a) *Every maximal almost-cubic block of a route-8 Type A support either exports two cycle lengths differing by 1 or 2, or is quiet and satisfies $|Y| < 5b(Y)^2$.*
- (b) *In the exported-length alternative, completion through the ambient boundary gives a Mersenne return or a target-complete replacement.*
- (c) *If all maximal blocks are quiet, gluing them along the canonical boundary of the Type A support gives $|X| < 5b(X)^2$.*

Then every admissible silent-core residual profile of Definition 14.15 lies in a finite boundary-labelled residual class determined by

$$|X| \leq 6142, \quad |X| < 5b(X)^2,$$

with all boundary return profiles avoiding $\mathcal{M}$.

Proof. Let X be the Type A support underlying an admissible silent-core residual profile. It is connected, subcubic, P_{13} -free, has empty internal 3-core, satisfies $\sigma(X) = 0$, and has $\partial_+(X) < |X|/4$. The P_{13} -free diameter argument in the Type A analysis already gives $|X| \leq 6142$.

Apply the Bondy–Vince/Gao–Ma input to every maximal almost-cubic block of X . If one of them realizes the exported-length alternative, clause (b) creates a Mersenne return or a target-complete replacement, both forbidden in an admissible support. Hence every maximal block is quiet. Clause (c) gives $|X| < 5b(X)^2$. Since $|X| \leq 6142$ and the boundary labels are drawn from the finite 399-label P_{13} algebra, only finitely many boundary-labelled return profiles remain. Thus the auxiliary input bounds route 8 by a finite boundary-labelled residual class with explicit Mersenne-avoiding return profiles; it does not by itself close that residual. $\square$

Remark B.5 (Type B residual after the structural fan-window reduction). *The unused external inputs do not close node [74]; the available closure is the structural label-packing lemma recorded in Lemma 14.66. Bondy–Vince/Gao–Ma controls quiet almost-cubic blocks, and Kleitman–West*

controls global tree-cotree profiles; the Type B obstruction is a local fan-safe clique around a high-degree surplus vertex. The packing lemma proves the C_2 clique cap, the marked non-singleton cap, and the spectral bound $d_G(h) \leq \rho(A(\mathcal{F}_h)) + 1$ for certificate-marked fans, and it discharges certificate-closed marked fans. Fan-certificate residual centers are charged directly to their assigned surplus units. The structural fan-window reduction of Lemmas 14.75 to 14.79 then removes every positive-deficit pair that already contains a direct dyadic cycle and pays every remaining positive-deficit fan by hybrid half-credit on its window and non-window incidences. The refined support ledger B2 is the local disjointness statement needed to sum those fan entries. By Lemmas 14.87 and 14.89, any disjoint-carrier failure of that bridge ledger is represented by a minimal overlap obstruction carrying contextual dyadic safety, high-degree separation, window compatibility, replacement routing, and minimal-overlap constraints. The fan-mass closure Definition 14.101, Lemmas 14.105 and 14.106, and Proposition 14.107 charges fan-certificate residual centers and disjoint-carrier failures to their assigned surplus units and gives total negative Type B bridge mass $O(\sigma(G)) = o(|R|)$. Thus neither Type B residual form is a terminal large-budget residual.

Summary: the degradation ladder. Ordered by how much the proof loses if that single leaf fails with no redundant cover:

1. **Definitional or exact-counting leaves:** [29] stub — framework-independent counting; [3]/[7] definitional.
2. **Redundantly covered or routed forward:** [30] curvature demand (covered by [60] net-charge, about $4.1 \times$ slack); [60] net-charge (covered by [30], about 9.6% slack); [32]/[47] fullrank (covered by Branch D); [50]/[54] entropy routing (high-entropy and repetitive low-entropy sub-cases add information, while nonrepetitive low entropy is routed to net charge); [123] branch-kill outside explicit residuals.
3. **Closed or bounded by auxiliary bookkeeping:** [23] overflow via the quiet-block accounting in Lemma B.2; [45] globalterm by the main global-smearing lemma, recorded again in Lemma B.3. The route-8 Type A branch is reduced to the terminal two-carrier obstruction by the canonical carrier-core ledger of Proposition 14.59, supplied with declared deletion witnesses by Lemmas 14.52 and 14.53, and closed at node [124] because two-carrier declared deletion quotients are canonical exit-(4) quotients by Lemma 14.54 and realize exit (4) by Lemma 14.55; it remains quiet-block bounded by Lemma B.4.
4. **Closed by structural Type B bookkeeping:** [74] fanthm uses the certificate-marked fan-safe clique bound for the local certificate-closed fan charge, the direct fan-window cycle lemmas for positive-deficit pairs, and the hybrid incidence budget of Lemmas 14.78 and 14.79. The degree-greater-than-four side leaf enters this ledger through fan-closed ports: a fan-compatible open pair routes by Corollary 14.81, and the triangular alternative routes by Corollary 14.82. In both cases Proposition 14.80 supplies $D_B \geq (k-3)/4 > 0$, with the triangular alternative giving the stronger $D_B \geq (5k-19)/4 > 0$. The local closure is reduced to B2 in Definition 14.91; fan-certificate residual centers and disjoint-carrier failures are then bounded globally by the fan-mass estimate Proposition 14.107, while Lemmas 14.104 and 14.106 extract route-8 non-window cores into the Type A deficit ledger. Disjoint-carrier failures are also reflected as minimal overlap obstructions of Definition 14.85.
5. **Reduced by exact surplus-token bookkeeping and geometric routing:** [125]–[144] use the active surplus family, the free/blocked pair split, the exact token-fibre identity, the classwise

token supplies, and the window-join pressure identity. The exact partition [Lemma 3.56](#) assigns every active-demand pair to exactly one total token–role fibre. The matching–star bound [Lemma 3.53](#) and the classwise budget theorem [Theorem 3.57](#) give the sharp load alternative, and [Lemma 3.84](#) and [Theorem 3.85](#) discharge any large homogeneous same-token fibre: it realizes a sparse surplus exit, produces ordinary Type B fan data, or is bounded by fixed homogeneous caps. In the bounded case [Corollary 3.50](#) gives $\sigma(G) = O(\sqrt{n})$.

No leaf degrades to nothing. Every failed leaf, stripped of redundancy, leaves a hypothetical counterexample carrying all standing invariants plus one explicit extra constraint. The quiet-block auxiliary input closes the dense-window reduction, and the main global-smearing lemma closes the global-mode reduction. In Type A, exits 1–3, 5, and 6 close, exit 4 peels one supported target-defective routed load, exit 7 hands off to Type B, the unsaturated branch is discharged, and route 8 is reduced to the terminal two-carrier trace-basin obstruction by [Proposition 14.59](#): each entry first passes to its canonical minimal target-complete carrier core inside the declared u -supported response algebra, zero/one-essential-core entries realize exits (4)–(7), each essential carrier has a deletion witness, each witness has declared support by [Lemma 14.53](#), and in a two-carrier entry each such witness gives a canonical exit-(4) quotient by [Lemma 14.54](#) and realizes exit (4) by [Lemma 14.55](#). The indexed linear burden is [Lemma 14.47](#). In Type B, certificate-closed fans have nonnegative local charge and are globally summed through B2, direct fan-window and two-window cycles are structurally removed, and every degree-greater-than-four side leaf supplies two fan-closed ports by [Corollaries 14.81](#) and [14.82](#); the hybrid B1 incidence budget pays every resulting positive-deficit fan locally. Disjoint-carrier failure of B2 is the minimal overlap obstruction of [Definition 14.85](#) with the global constraints of [Lemma 14.89](#); fan-certificate residual centers, grouped handoff envelopes, and disjoint-carrier failures are bounded by [Proposition 14.107](#) after route-8 non-window cores are extracted by [Lemmas 14.104](#) and [14.106](#). In the non-near-cubic branch, the exact token ledger and [Theorem 3.73](#) give the explicit lower bound $\psi(N_*(G)/(36(4n + 15p_{13} + 3\sigma(G))))$ for a homogeneous subtype load whenever bounded caps fail; [Lemma 3.84](#) then turns that load into a sparse surplus exit, ordinary Type B fan data, or fixed caps implying the near-cubic estimate. Hence, inside the near-cubic spine or after the surplus-token bottleneck discharge, no linear large-budget deficit remains in the Type A route-8 carrier ledger.

Cross-check against the demand > supply collision. Outside the explicit residuals, the proof’s contradiction is reached in two arithmetically independent forms, so no *single* constant kills both.

Form	Demand	Supply	Slack	Feeder leaves
Curvature	$W_2 \geq 2.543 \cdot R $	$\leq 0.611 \cdot R $ independent tests	$4.1 \times$	[30], [47], [48]
Net charge	$\partial_+(R) - \sigma(R) \geq 0.25 \cdot R $	$\partial_+(R) - \sigma_R =$ $\partial_+(R) - \sigma(R) \leq 0.2282 \cdot R $ (τ_{win})	$\sim 9.6\%$	[60], [91], [74]

The shared ancestor of both forms is invariant 24 ([Lemma 8.4](#), diagram node [29]), the exact incidence-counting leaf.

References

- [1] A. S. Hegde, R. B. Sandeep, and P. Shashank, *Erdős–Gyárfás conjecture on graphs without long induced paths*, arXiv:2410.22842v2, 2025. <https://arxiv.org/pdf/2410.22842v2>

- [2] J. A. Bondy and A. Vince, *Cycles in a graph whose lengths differ by one or two*, J. Graph Theory 27 (1998), 11–15. <https://people.clas.ufl.edu/avince/files/Cycles.pdf>
- [3] J. Gao and J. Ma, *On a conjecture of Bondy and Vince*, arXiv:1902.05701, 2019. <https://arxiv.org/abs/1902.05701>
- [4] D. J. Kleitman and D. B. West, *Spanning trees with many leaves*, SIAM J. Discrete Math. 4 (1991), 99–106. <https://dwest.web.illinois.edu/pubs/manyleaf.pdf>